

Development of a beam-based phase feed-forward demonstration at the CLIC Test Facility (CTF3).

Jack Roberts
New College, Oxford

Thesis submitted in fulfilment of the requirements for the degree of Doctor of Philosophy at the University of Oxford

Trinity Term, 2016

Abstract

This is the abstract TeX for the thesis and the stand-alone abstract.

Acknowledgements

Acknowledgements.

Contents

1	Introduction	1
1.1	Particle Physics	1
1.2	Colliders	2
1.3	Motivation for Future Linear Colliders	3
1.4	The Compact Linear Collider (CLIC)	4
1.4.1	Drive Beam Recombination	7
1.5	Phase Feedforward for CLIC	9
1.6	The CLIC Test Facility CTF3	11
1.7	The PFF Prototype at CTF3	12
1.7.1	Hardware	14
1.7.2	Differences Between PFF at CTF3 and CLIC	16
1.8	Definitions of Different Phase Statistics	17
1.9	Thesis Overview	18
2	Design of the PFF Chicane	19
2.1	Introduction to Optics	19
2.2	Kicker Design	21
2.3	TL2	22
2.3.1	Integration of PFF Hardware	24
2.4	TL2 Optics Constraints	26
2.4.1	Nominal Optics Constraints	26
2.4.2	PFF Optics Constraints	28
2.5	TL2 Optics Measurements	30
2.5.1	Method	30
2.5.2	Results with Original MADX Model	31
2.5.3	Sources of Errors in MADX Model	31
2.5.4	Corrections to MADX Model	34
2.6	Matched TL2 Optics	36
2.6.1	Matching Process	39
2.6.2	PFF Optics	40
3	Phase Monitor Performance	46
3.1	Phase Monitor Design	46
3.2	Phase Monitor Electronics	50
3.3	Resolution Definition	57
3.4	Digitisation of Phase Monitor Signals	59

3.5	Fitting Method	61
3.6	Signal Generator Measurements	63
3.6.1	Experimental Setup	63
3.6.2	Results	64
3.6.3	Mixer Performance	64
3.6.4	Diode Performance	72
3.6.5	Consequences for Normal Operation	78
3.7	Calibrations	79
3.7.1	Calibration on SiS Digitisers	80
3.7.2	Calibration on FONT5a Board	83
3.7.3	Multi-Sample Results	83
3.7.4	Zero Crossing	85
3.8	Phase Shifter Noise	87
3.9	Resolution Measurements	93
3.9.1	Best Resolution	93
3.9.2	With Sample Averaging	94
3.9.3	Dependence of Resolution on LO Phase	94
3.10	Bandwidth	95
3.11	Comparison of Measured Phase Along Pulse	95
3.12	Effect of Variations in Calibration Constant	100
3.13	Dependence on Position	100
4	Phase Propagation	107
4.1	Feedforward Algorithm	107
4.2	Characteristics of Uncorrected Phase Jitter	108
4.3	First Order Energy Dependencies	111
4.3.1	Beam Energy Variations	111
4.3.2	Correlation between Phase and Energy	115
4.3.3	R56	116
4.3.4	Effect of R56 in TL2	118
4.4	Mitigation of First Order Energy Dependence	121
4.4.1	Matched Optics for TL1	124
4.4.2	Scans of R56 in TL1	128
4.5	Higher Order Energy Dependencies	131
4.5.1	Simulated Effect of T_{566} on the Downstream Phase	133
4.5.2	R56 Scans whilst Varying Beam Energy	138
4.5.3	Mitigation of Higher Order Dependencies	148
4.6	Best Phase Propagation	150
4.7	Possible Other Sources of Phase Jitter	154
5	Setup and Commissioning of the PFF System	160
5.1	Feedforward Controller (FONT5a Board)	160
5.1.1	Implementation of PFF Correction	164
5.1.2	ADC Droop Correction	169
5.1.3	Gain Calculation	172

5.1.4	Effect of Using Small Angle Approximation	175
5.2	Amplifier	176
5.2.1	Design	177
5.2.2	Linearity	178
5.2.3	Shape	183
5.3	Data Acquisition	186
5.3.1	Acquisition Tools	186
5.3.2	Monitoring Tools	186
5.4	Kicker and Optics Performance Verification	187
5.4.1	Correction Range	187
5.4.2	Variations Along Pulse	190
5.4.3	Shape	191
5.4.4	Orbit Closure	191
5.5	Correction Output Timing	195
5.5.1	Kicker Cable Lengths	195
5.5.2	Absolute Timing	199
5.5.3	Relative Kicker Timing	204
5.6	Early Phase Feedforward Attempts and Simulations	208
5.6.1	Gain Scans	208
5.6.2	Effect of Limited Correction Range	208
6	Best Achieved Phase Stabilisation	209
6.1	Lowest Achieved Phase Jitter	209
6.1.1	Mean Phase Jitter	210
6.1.2	Correction of Pulse Shape	212
6.1.3	Phase Jitter Along the Pulse	213
6.2	Correction on Longer Time Scales	214
6.2.1	Upstream Phase Drifts	216
6.2.2	Gain Stability	220
6.2.3	Results	221
6.3	Possible Areas for Future Improvement	230
7	Conclusions	231
7.1	Summary	231
7.2	Future Work	231
Bibliography		232

Chapter 1

Introduction

1.1 Particle Physics

Particle physics is the study of the building blocks of the universe. At its heart is the standard model, which defines a small number of fundamental particles, spin half fermions, that form the constituents of all visible matter in the universe, as well integer spin bosons that mediate the forces between the fermions. Tables 1.1 and 1.2 list all the standard model fermions and bosons, respectively.

The fermions are split in to two types – leptons with an integer charge of -1 (charged leptons) or 0 (neutrinos), and quarks with fractional charges of +2/3 or -1/3. The fermions are also split in to three generations, each with a charged lepton, neutrino, positively charged quark and negatively charged quark. The mass of the particles increases with the generation, and all stable matter is comprised of the first generation particles. Every particle in the standard model has an associated anti-particle, with the same mass but opposite charge.

The fermions in the standard model interact via three fundamental forces – the electromagnetic force, mediated by the photon, the weak force, mediated by the W and Z bosons, and the strong force, mediated by the gluon. The strong force binds quarks together to form hadrons, which are either a quark—anti-quark pair (mesons), or a bound state of three quarks (hadrons). Protons are hadrons comprising of two up quarks and one down quark. Quarks also interact via the electromagnetic and weak forces. Leptons do not interact with the strong force – neutrinos only interact via the weak force, and the charged leptons interact via the weak and electromagnetic forces. The final component of the standard model is the Higgs field, which is responsible for giving particles their masses, and its associated Higgs boson [REF].

The standard model has been the subject of extensive theoretical and experimental research over the last 50 years. This culminated in 2012 with the discovery of the final standard model particle yet to be observed experimentally, the Higgs Boson, at the Large Hadron Collider (LHC), CERN, Switzerland [REF]. Despite the incredible success of the standard model there are several known phenomena it cannot explain. For example, it does not include the gravitational force nor have viable particles or mechanisms to describe dark matter and dark

Generation	Leptons		Quarks	
1	Electron (e^-)	Electron neutrino (ν_e)	Up (u)	Down (d)
2	Muon (μ^-)	Muon neutrino (ν_μ)	Charm (c)	Strange (s)
3	Tau (τ^-)	Tau neutrino (ν_τ)	Top (t)	Bottom (b)
Charge	-1		0	+2/3 -1/3

Table 1.1: Standard model fermions.

Force	Bosons
Electromagnetic	Photon (γ)
Weak	W^\pm and Z^0 Bosons
Strong	Gluon (g)
Higgs Field	Higgs Boson (H^0)

Table 1.2: Standard model bosons.

energy, which are required to explain observations in cosmology such as the increasing rate of expansion of the universe [REF]. The LHC and other experiments around the world aim to discover new particles or find discrepancies in the standard model to explain these effects.

1.2 Colliders

The driving force behind recent discoveries in particle physics has been colliders, in which two beams of particles are accelerated to high energy and then brought in to collision with one another. Longitudinal electric fields are used to accelerate the two beams, with the fields created by injecting RF power in to cavities placed along the beam line. The source of the RF power is usually klystrons [REF], in which a low power RF input is amplified using a low energy electron beam (contained within the klystron and independent of the colliding beams). Due to the use of RF accelerating fields, the colliding beams are bunched with a frequency related to the RF frequency, rather than being continuous. Each bunch of particles can then experience the same accelerating field.

The interaction of the two beams when they are brought in to collision produces new particles that are observed in large detectors surrounding the interaction point (IP). The types of interaction that can take place and the particles that can be produced depends on the centre of mass energy of the collision. Colliding a single beam in to a fixed target reduces the available energy for the interaction as the final state must have high kinetic energy to conserve momentum. Colliding two opposing beams head on with zero net momentum therefore maximises the centre of mass energy available to produce new particles.

The rate at which a given interaction X occurs when the beams collide can be defined as:

$$R(X) = \mathcal{L}\sigma(X) \quad (1.1)$$

Where $\sigma(X)$ is the cross-section for the interaction, defined from the standard model and including dependencies on the collision energy, for example. The luminosity \mathcal{L} is a property

of the beam and can be defined as:

$$\mathcal{L} = H \frac{f N^2}{4\pi\sigma_x\sigma_y} \quad (1.2)$$

Where f is the frequency at which bunches collide, N is the number of particles in each bunch, σ_x and σ_y are the horizontal and vertical beam sizes respectively and H is a factor dependent on the electromagnetic interaction of the two beams close to the collision point. Small, dense beams and a high bunch crossing frequency are desirable to maximise the interaction rate.

1.3 Motivation for Future Linear Colliders

Colliders are typically either circular (synchrotrons) or linear (linacs) and use either electron or proton beams (and their associated anti-particles). The choice of collider shape and particle has many consequences for the properties of the resulting experiment.

In synchrotrons the two beams are bent around a path of fixed radius using magnetic fields (dipoles). The beams circulate the collider many times, being brought in to collision at one or several interaction points around the ring where detectors are placed. A large fraction of the ring can be filled with bunches, and with these bunches circulating at close to the speed of light synchrotrons therefore benefit from high luminosities thanks to their high bunch crossing frequency. For proton machines the highest achievable energy in a synchrotron is predominantly defined by the radius of the ring and the maximum sustainable field in the dipoles. Electron beams have other limitations, as described below. The LHC is a 27 km proton synchrotron with 8.3 T dipoles and a bunch crossing frequency of around 30 MHz that has reached a world record collision energy of 13 TeV [REF].

Proton collisions present a number of challenges for the particle detectors and data analysis, however. Protons are not fundamental particles, but rather contain quarks and gluons. Therefore the interactions that occur in proton colliders are in reality between the constituents of the protons, rather than the protons themselves. The precise energy of each quark or gluon is not known, which leads to increased uncertainties in the measurements. In addition, strong interactions between the quarks and gluons lead to high background noise in the collision events, making particle identification in the detectors more difficult. As electrons are (in our current knowledge) fundamental and do not partake in the strong interaction the resulting collisions in an electron (or electron–positron) collider are much cleaner and the uncertainties smaller. This motivates research into a future high energy electron collider, where the properties of recently discovered heavy particles such as the Higgs boson and top quark, or any new particles discovered by the LHC in the coming years, could be studied with high precision.

Electron machines present a different challenge to protons due to their approximately 2000 times lower mass. Charged particles bent in a magnetic field emit radiation, known as synchrotron radiation, that depends on the bending radius, r , the particle's energy, E , and

the particle's rest mass, m_0 , as follows [REF]:

$$P \propto \frac{1}{r^2} \frac{E^4}{m_0^4} \quad (1.3)$$

Where P is the radiated power due to synchrotron radiation. For electrons the power radiated is roughly a factor 2000^4 larger than for a proton beam of the same energy, which limits the beam energies achievable in electron synchrotrons. For example, LEP, an electron–positron collider previously installed in the same 27 km tunnel as the LHC [REF], achieved energies up to around 200 GeV, almost two orders of magnitude lower than is now achievable with protons in the LHC.

The only feasible way to achieve an electron collider with centre of mass energies significantly larger than what was achieved at LEP, for example at the TeV scale, is to use a linear collider (linac). In a linac particles are injected at one end of the beam line, accelerated in cavities placed along its whole length and then brought in to collision with an opposing beam from a second linac. The maximum achievable energy in a linac is defined by the length of the facility and the rate of acceleration (the peak field) in the cavities. As the particles only traverse the linac once the collision frequency is normally much lower than for synchrotrons, and thus very small, nanometre scale, beam sizes are required to provide adequate luminosity. The requirements for the accelerating cavities are also much more challenging compared to a synchrotron where the particles circulate through the cavities many times and a lower accelerating field can be accepted.

Currently there are two separate proposals for the design of a future electron–positron collider that addresses these challenges and could achieve centre of mass energies at the TeV scale – the International Linear Collider (ILC) [REF] and the Compact Linear Collider (CLIC) [REF]. The ILC uses superconducting RF cavities with an accelerating gradient up to 35 MV/m to achieve a 500 GeV collision energy with a facility approximately 30 km in length, with a possible future upgrade to 1 TeV with a longer facility. CLIC uses normal conducting cavities and a novel two beam acceleration concept to achieve an accelerating gradient of 100 MV/m and collision energies up to 3 TeV for a 50 km facility, similar in length to the ILC 1 TeV upgrade. Both the ILC and CLIC have large international collaborations and test facilities, with the ILC design at a more advanced stage having published its technical design report in 2013 [REF]. This thesis presents a contribution towards proving the feasibility of the CLIC concept.

1.4 The Compact Linear Collider (CLIC)

The parameters of CLIC, summarised in Table 1.3, are chosen to optimise the cost of an electron–positron collider with a centre of mass energy of 3 TeV whilst providing similar or higher luminosity than the LHC. Superconducting RF cavities are limited in the peak fields they can achieve, and so room temperature normal conducting cavities are used to achieve an accelerating gradient of 100 MV/m. This reduces the site length, and associated civil engineering costs, by roughly a factor 3 compared to what would be necessary with a superconducting 3 TeV machine. Studies and cost optimisations of the cavities necessary to

Parameter	CLIC 3 TeV
Peak Luminosity [$\text{cm}^{-2}\text{s}^{-1}$]	5.9×10^{34}
Site Length [km]	48.4
Accelerating Gradient [MV/m]	100
RF Frequency [GHz]	12
Number of Bunches	312
Number of Particles per Bunch	3.72×10^9
Bunch Separation [ns]	0.5
Bunch Length [μm]	44
Beam Pulse Length [ns]	156
Repetition Rate [Hz]	50
Horizontal Beam Size at IP [nm]	40
Vertical Beam Size at IP [nm]	1
AC Power Consumption [MW]	582

Table 1.3: Parameters of the CLIC main beam for a collision energy of 3 TeV [REF].

support the high peak fields, taking in to account the RF field breakdown rate for example, lead to the choice of a 12 GHz RF frequency with a beam pulse length of 156 ns [REF]. Each accelerating cavity requires an RF input power of around 65 MW [REF].

Conventional klystrons cannot produce the short pulse lengths and high peak powers required for CLIC [REF]. To achieve the necessary specifications with klystrons RF pulse compressors would have to be used to increase their output power and decrease their pulse length by roughly a factor 5 [REF]. This would require a prohibitively large number of klystrons, approximately 35000, and yeild a low RF power efficiency of about 40% [REF]. To solve this problem CLIC proposes to use a novel two beam acceleration concept, in which the high energy main beam is accelerated using RF power extracted from a secondary, high intensity but low energy, drive beam.

Figure 1.1 shows the layout of the CLIC complex. The bottom half of the figure shows the generation of the main colliding beams. The electron and positron beams are generated in dedicated injectors at the bottom of the figure, where they are accelerated to an energy of 2.86 GeV. They are then transported to a series of damping rings, in which the effects synchrotron radiation are used to decrease the beam size [REF]. After leaving the damping rings they are accelerated further to 9 GeV before being transported to each end of the main linac. Here both beams are accelerated in 21 km linacs containing the 100 MV/m accelerating structures, reaching an energy of 1.5 TeV. Finally, the beam size is squeezed further in the beam delivery system (BDS), down to the nanometre sizes necesarry to achieve high peak luminosity, before being brought in to collision at the interaction point.

The top half of the figure shows the beam lines involved in the generation of the high intensity drive beams. The drive beam accelerators are conventional linacs with an RF frequency of 1 GHz, a beam pulse length of 142 μs , a beam current of 4.2 A and a final energy of 2.4 GeV. The longer beam pulse length and lower power requirements means klystrons can be used with high efficiencies, unlike the main beams. Following the linac the




Figure 1.1: Layout of the CLIC complex.

drive beams enter a series of “delay loops” and “combiner rings” in which the drive beam is interleaved with itself to increase its intensity by a factor 24 (to 100 A) in a series of shorter 240 ns beam pulses. This drive beam recombination process is described in Section 1.4.1.

The 100 A drive beam is then directed into 24 decelerator sections. Here a series of specially designed cavities, Power Extraction and Transfer Structures (PETS), are installed which decelerate, and extract power from, the drive beam. The PETS extract 90% of the energy of the drive beam [REF], creating an RF pulse with the necessary properties to accelerate the main beam at 100 MV/m. The extracted RF power from the PETS is transferred to the main beam accelerating cavities via short waveguides.

1.4.1 Drive Beam Recombination

The combination process used to produce the high intensity drive beam pulses is a unique feature of CLIC. The intensity is increased by a factor 2 in the delay loop, a factor three in the first combiner ring (CR1) and a factor 4 in the second combiner ring (CR2) to give the overall increase of $2 \times 3 \times 4 = 24$ in beam current.

Figure 1.2 shows the recombination process in the delay loop. Although the drive beam acceleration frequency is 1 GHz it is bunched at 0.5 GHz. At 240 ns intervals along the 142 μ s drive beam pulse the phase of the bunching is switched by 180° at 0.5 GHz. This phase switching is performed by sub-harmonic bunchers (SHBs) powered by travelling wave tubes (TWTs). As 180° at 0.5 GHz is a full 360° at the acceleration frequency of 1 GHz the phase switching has no effect on the acceleration of the drive beam.

At the entrance of the delay loop a 0.5 GHz RF deflector is used. The 240 ns sub-trains alternate between having bunches that arrive when the field in the RF deflector is positive or negative. Sub-trains arriving when the field is positive are deflected into the delay loop, whereas the remaining sub-trains bypass it. The length of the delay loop is also precisely tuned to be 240 ns, so that the sub-trains leaving the delay loop merge with the following sub-train bypassing the delay loop. The bunch frequency and beam current following the delay loop are therefore doubled to 1 GHz and 8.4 A respectively. The resulting beam structure consists of 240 ns pulses separated by 240 ns gaps.

A similar process takes place in the two combiner rings to give the final 12 GHz bunch frequency and 100 A beam current. Figure 1.3 shows the recombination of a factor 4 in CR2. RF deflectors are again used to inject the beam into the ring, this time with a higher frequency 3 GHz deflecting field to match the bunch frequency of the drive beam following CR1. The first drive beam sub-train entering CR2 circulates the ring for four turns. The length of the ring is chosen so that the second sub-train arrives just as the first sub-train has completed one turn in the ring, with a delay of one quarter of the 3 GHz wave length. Repeating this for a further two turns, interleaving a new sub-train with the beam in the ring after each turn, yields the factor 4 increase in bunch frequency and therefore beam current. CR1 operates in the same way, but the beam circulates for three turns and the required path length difference is one third of the 1 GHz wavelength.

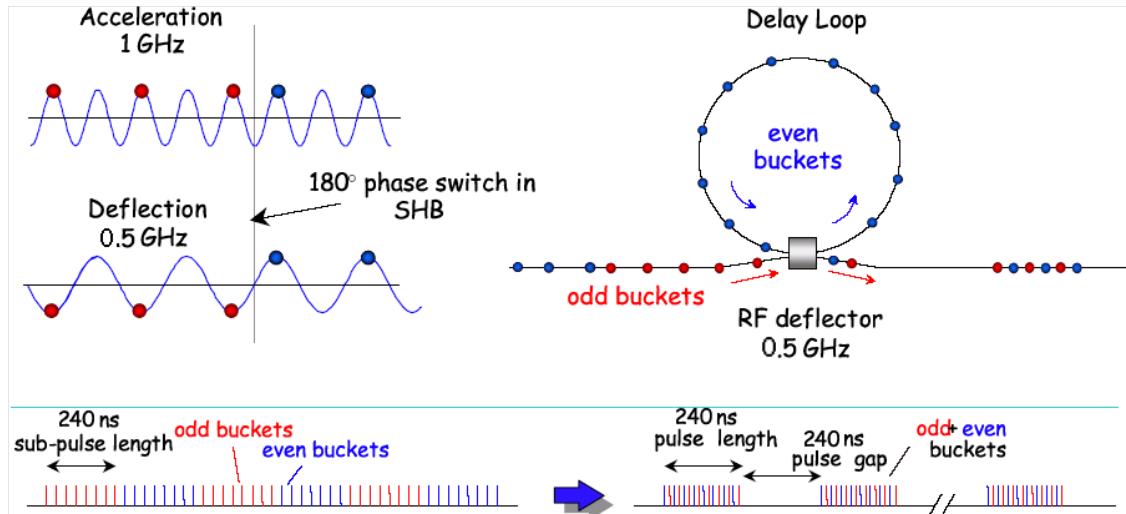


Figure 1.2: Drive beam recombination in the delay loop [REF].

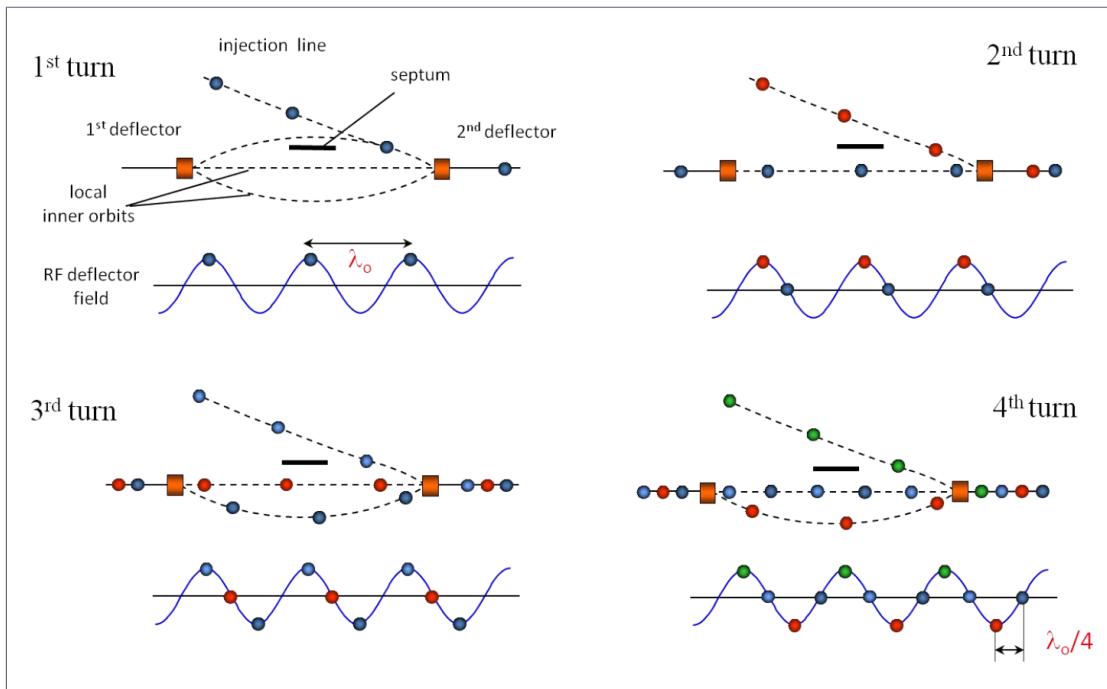


Figure 1.3: Drive beam recombination in the second combiner ring [REF].

1.5 Phase Feedforward for CLIC

In the CLIC BDS at the end of the main linacs the beam size is reduced to the nanometre sizes needed to achieve high peak luminosity. The beam is focused using magnets (mostly quadrupoles) whose focal length depends on the beam energy. Variations in beam energy therefore lead to imperfections in the focusing, which subsequently causes a growth in the beam size and a reduction in luminosity at the collision point. To keep this luminosity loss below 1% CLIC can tolerate a main beam energy spread (standard deviation or “jitter”) of 0.3% [REF]. The term jitter is used to describe the standard deviation in the remainder of the thesis.

As the CLIC main beams are accelerated using RF power extracted from the drive beams, this energy stability requirement also places strict tolerances on the parameters of the drive beams. To limit the luminosity loss to below 1% a drive beam phase jitter of 0.2° at 12 GHz, a bunch charge jitter of 0.08% and a bunch length jitter of 1% or below are required [REF], for example. To achieve this level of stability a number of correction systems will be required in the drive beam complex. This thesis presents work on the correction system required to achieve the necessary 0.2° at 12 GHz (~ 45 fs) drive beam phase stability — the Phase Feedforward (PFF) system.

The drive beam linac klystrons are expected to provide an RF power stability of 0.2% and an RF phase stability of 0.05° at 1 GHz [REF]. After being accelerated in the linac and recombined in the delay loop and combiner rings this leads to a final beam phase stability of around 2° at 12 GHz. The PFF system must therefore reduce the drive beam phase jitter by an order of magnitude prior to the power extraction and transfer to the main beam in the decelerator sections. Each drive beam complex has 24 decelerator sectors (see Figure 1.1), and CLIC proposes to install a PFF system in each sector, for a total of 48 PFF systems across both drive beams [REF]. Placing the PFF systems in the decelerator sectors minimises the possibility of additional phase errors being introduced between the correction and the power extraction.

Figure 1.4 shows a schematic of the layout of the decelerator sectors and the implementation of the PFF system inside them. At the entrance to each decelerator sector the beam is bent and redirected through a turnaround. Following the turnaround each decelerator contains a chicane in which the beam is directed on to a “C-shaped” trajectory using four dipole magnets. Finally, the drive beam is transported to the PETs, where its power is extracted and transferred to the main beam.

For the purposes of the PFF system sixteen electromagnetic kickers will be installed in the chicane, with two before and two after each dipole magnet. A voltage is applied to conducting strips inside the kickers in order to produce transversely deflecting electric and magnetic fields. The combination of four kickers around each dipole can horizontally deflect the beam by up to $\pm 375 \mu\text{rad}$. This assumes the use of one 500 kW amplifier to power each metre long, 50 mm aperture kicker [REF]. The length of the beam’s trajectory through the chicane depends on the voltage applied to the kickers — if the polarity of the applied voltage is chosen to deflect the beam towards the outside of the chicane the trajectory is




Figure 1.4: Schematic of the CLIC PFF system. The positions of the phase monitor (ϕ) before the turnaround, and the kickers (K) and dipoles (rectangles) in the chicane are shown. The box PFF represents the feedforward electronics, controls and kicker amplifiers. For the purposes of the diagram only four kickers are shown – the actual system would have 16 kickers, with two before and two after each dipole. Beam pulses arriving early at the phase monitor are directed on to longer trajectories in the chicane (blue), whereas pulses arriving late are directed on to shorter paths (red).

longer (blue), whereas the trajectory is shorter with the opposite polarity voltage applied and the beam directed towards the inside of the chicane (red).

This dependence of path length on the kicker voltage is used to correct the phase (with the phase being analogous to the arrival time or longitudinal position of the drive beam bunches). The drive beam phase is measured prior to the turnaround using a high precision phase monitor with a resolution of around 0.1° at 12 GHz (~ 23 fs). The measured phase is compared to a reference phase derived from the main beam, and the voltage to the kickers is varied depending on the relative phase offset between the two. If the drive beam arrives early it is directed on to a longer path and vice versa so that after the chicane the corrected phase is on reference.

Measuring the phase upstream of the correction location (feedforward rather than feedback) combined with the use of the turnaround gives the benefit of being able to apply the correction to the same beam pulse that was initially measured. This is possible despite the beam travelling at close to the speed of light as the path length of the signals between the phase monitors and the kickers is shorter than the trajectory of the beam. In other words the correction signals bypass the turnaround. The difference between the beam path length and the correction signal path length defines the available latency for the PFF hardware (signal processing, correction calculation and voltage generation).

The 0.2° phase stability must be maintained along the full combined beam pulse length, and thus the PFF correction will also be high bandwidth, able to correct variations along the pulse as well as any overall offsets in the mean pulse phase. The drive beam recombination and the properties of the main beam accelerating structures result in a filtering of the very high frequency phase jitter components [REF]. The PFF system must have a bandwidth in excess of 17.5 MHz in order to remove the lower frequency components and achieve the required RF phase stability. The measurement from the phase monitors and the output of the kicker amplifiers must have bandwidths above 17.5 MHz to be able to achieve this.

1.6 The CLIC Test Facility CTF3

The CLIC design requires the use of many new concepts and technologies. Therefore, to prove the feasibility of CLIC the test facility CTF3 at CERN has been in operation since 2001 [REF]. It aims to demonstrate the generation of a high intensity drive beam using the bunch recombination process, as well as the extraction of power from this drive beam and the use of the extracted power to accelerate another beam. CTF3 also hosts many related activities in areas such as the development of 12 GHz accelerating cavities and beam instrumentation, for example [REF]. The main goals of CTF3 have been achieved and 2016 will be its last year of operation [REF].

A diagram of the CTF3 facility is found in Figure 1.5, showing the layout of buildings (purple) and beam lines (red) with the names and extent of the main sections labelled in black. The source of the CTF3 beam, at the top left of the figure in the linac, is a thermionic electron gun [REF] that produces a $1.4 \mu\text{s}$ beam pulse with an intensity of 4 A

and a repetition rate of 0.8 Hz. The first 200 ns of the pulse contains a sharp energy transient and is eventually lost in the first bending sections of the facility, creating a usable pulse length of around 1.2 μ s. The continuous beam pulse from the gun is bunched at either 1.5 GHz or 3 GHz depending on the mode of operation [REF].

The beam is accelerated along the linac in 3 GHz RF cavities powered by conventional klystrons combined with RF pulse compressors which double their output power [REF]. The accelerating cavities are operated in the fully loaded mode, in which almost all the RF power sent to the cavity is absorbed by the beam for high power efficiency [REF]. At the end of the linac the beam reaches an energy of approximately 135 MeV.

Following the linac the beam intensity can be increased by up to a factor 8 using the delay loop and combiner ring, which both function in the same way as described for CLIC in Section 1.4. Experiments at CTF3 are usually performed with either a 3 GHz factor 1 (4 A) or factor 4 (16 A) beam, or a 1.5 GHz factor 8 (28 A) beam. For the setups with 3 GHz bunching the delay loop is bypassed. The number of times the beam circulates around the combiner ring can be varied between 0.5 and 3.5 turns to give an increase in beam intensity by a factor 1, 2, 3 or 4. With 1.5 GHz beam a 1.5 GHz RF deflector [REF] is used at the delay loop entrance to inject alternating 140 ns sub-trains in to the loop. The sub-trains exiting the loop merge with the 140 ns trains bypassing the loop to give a factor 2 increase in beam intensity. The intensity of this factor 2 beam can then also be increased by additional factor 4 in the combiner ring, to give the total increase of a factor 8.

Following the combiner ring the beam enters the transfer line TL2, which transports the beam to the CLIC experimental area (CLEX). In CLEX the CTF3 drive beam can be directed to two different beam lines - TBL (Test Beam Line) and TBM (Two Beam Module). In TBL a series of PETs are installed to test the extraction of power from the drive beam and to measure the properties of the produced RF power [REF]. In TBM a prototype CLIC accelerating module is installed, in which RF power is extracted from the drive beam and used to accelerate a second beam, called CALIFES. The CALIFES beam also hosts many experiments independent from the CTF3 drive beam [REF].

1.7 The PFF Prototype at CTF3

The phase feedforward system proposed for CLIC (Section ??) presents many challenges in terms of the required hardware latencies, resolutions, power and bandwidth. As a result, one of the key activities at CTF3 since 2013 has been the design, installation and operation of a prototype PFF system. The primary goal of the PFF prototype is to demonstrate the feasibility of the PFF concept, with the ultimate aim of reducing the CTF3 phase jitter to close to the CLIC requirement of 0.2° at 12 GHz, with a correction bandwidth above 17.5 MHz. The pursuit of this goal has required the development and installation of new hardware, as well as modifications and improvements to the setup and stability of the whole CTF3 drive beam complex.

The layout of the PFF prototype is shown in Figure 1.6. The overall concept is the same

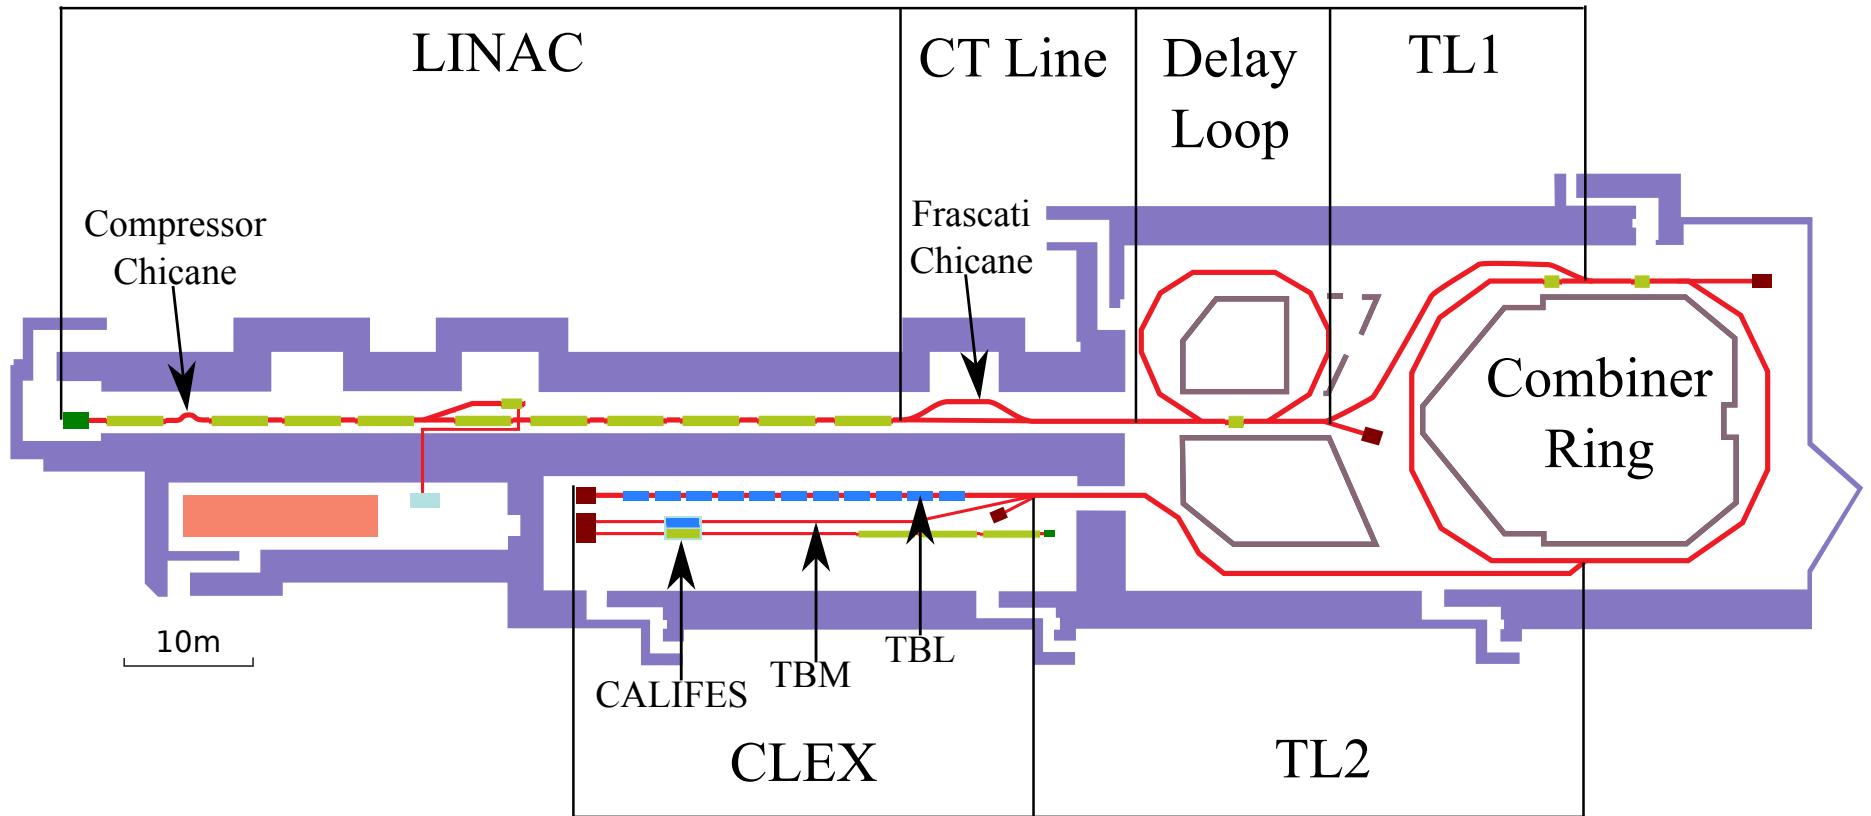


Figure 1.5: Layout of CTF3 with the main sections labelled.

as the CLIC proposal – the beam phase is measured prior to a turnaround and then corrected by changing the path length through a chicane using kickers. At CTF3 the PFF input is the phase measured in the CT line (ϕ_1 in the figure). The phase is then corrected using two kickers installed in the “dog-leg” shaped chicane in the TL2 transfer line. A 3 GHz, uncombined beam is used, bypassing the delay loop (DL) and completing only half a turn in the combiner ring (CR). With this setup the time of flight between the phase monitor (ϕ_1) and the first kicker (K1) is approximately 380 ns. Like the proposed CLIC system, the PFF prototype aims to apply the correction downstream (in TL2) to exactly the same pulse that was initially measured upstream (in the CT line), which is possible as the correction signals travel a shorter distance than the beam. The latency of the whole system, including the signal transit times in cables and hardware latencies, must therefore be less than 380 ns.

1.7.1 Hardware

The major hardware components of the PFF prototype are three phase monitors, two kickers, three sets of electronics for the phase monitors, a digitiser/feedforward controller and amplifiers to power the kickers (referred to as the kicker amplifiers). All the hardware components were designed, built and newly installed at CTF3 for the PFF prototype. The same components could be used at CLIC, with some important differences as described in Section 1.7.2. Each piece of hardware and its role is briefly introduced here, with more detail provided for each component in the remainder of the thesis.

The kickers and phase monitors are installed in the beam line at the locations shown in Figure 1.6, whereas the phase monitor electronics, feedforward controller and kicker amplifiers are installed in the CTF3 “klystron gallery”, on the floor above the accelerator hall, for easier access. The cable lengths between the phase monitors and their electronics, and between the amplifiers and the kickers, are therefore longer than they appear in Figure 1.6.

As already described, the first phase monitor (ϕ_1), in the CT line, provides the PFF input. The neighbouring second phase monitor (ϕ_2), also in the CT line, is used to cross-check and verify the performance of the phase monitors. The final phase monitor (ϕ_3), in the TBL line, measures the corrected phase jitter following the chicane in TL2. The phase monitors are designed by INFN, Italy [REF] to give a resolution below 0.2° at 12 GHz with a bandwidth above 30MHz, also taking in to account the design of the phase monitor electronics [REF].

The processed signals from the first phase monitor are sent to the “FONT5a board”, the feedforward controller designed and built by the Oxford University group of the John Adams Institute (JAI) [REF]. With a low latency of around 60 ns [REF], the FONT5a board digitises the phase monitor signals, calculates the appropriate correction to apply and provides the drive signal for the kicker amplifiers.

The kicker amplifiers have also been designed and built by Oxford University/JAI, and provide the voltage that produces the electric and magnetic fields that deflect the beam when applied to the kickers. Each amplifier module provides a power of around 20 kW with a bandwidth close to 50 MHz for small variations [REF].

Finally, the two kickers that provide the correction were designed and built by INFN,

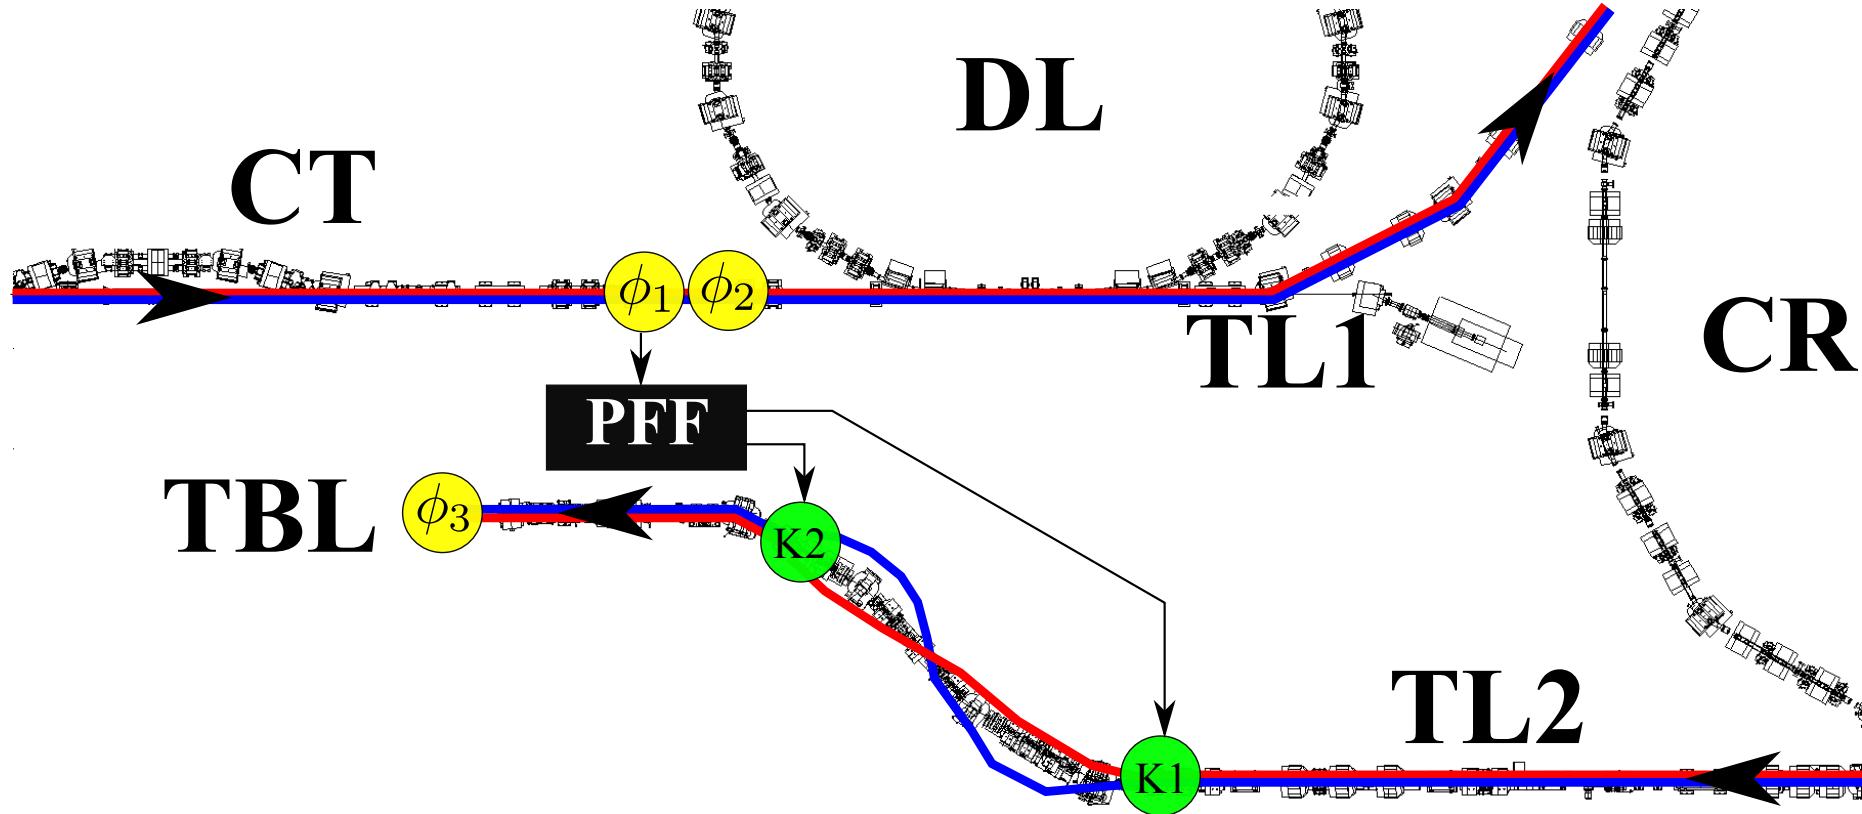


Figure 1.6: Schematic of the PFF prototype at CTF3, showing the approximate location of the phase monitors (ϕ_1 , ϕ_2 and ϕ_3) and the kickers (K1 and K2). The black box “PFF” represents the calculation and output of the correction, including the phase monitor electronics, feedforward controller and kicker amplifiers. A bunch arriving early at ϕ_1 is directed on to a longer path in the TL2 chicane using the kickers (blue trajectory), whereas a bunch arriving late will be directed on to a shorter path (red trajectory)

Italy. They are installed prior to the first and last dipoles in the TL2 dog-leg chicane, and deflect the CTF3 beam through an angle of 1 mrad for an applied voltage of around 1.3 kV [REF].

1.7.2 Differences Between PFF at CTF3 and CLIC

The goal of the prototype is to prove the general PFF concept and the feasibility of using it to achieve 0.2° drive beam phase stability. It is neither necessary nor possible for the proposed CLIC and CTF3 systems to be identical, and there are a number of differences between the two that are summarised here.

The most obvious difference between the applications at CTF3 and CLIC are the different beam energies and scale of the two complexes. CLIC will have a drive beam energy of 2.4 GeV, compared to the 135 MeV CTF3 drive beam. The CLIC PFF system therefore requires much higher power from the kicker amplifiers to deflect the beam by the same amount as the CTF3 prototype. CLIC requires up to 500 kW peak power from the amplifiers, compared to 20 kW at CTF3. Multiple amplifier modules with similar design to the CTF3 amplifiers could be combined to meet the CLIC power requirements [REF].

The CLIC proposal requires 48 separate PFF systems, with one in each of the 24 decelerator sections for each of the two drive beams. Aside from the vast difference in the quantity of required hardware components, the CLIC application also presents the challenge of synchronising the reference timing of the 48 systems along the 50 km facility with femtosecond stability. This is not addressed in the CTF3 prototype but feasible solutions have been proposed [REF].

Although both the CTF3 and CLIC schemes use chicanes to vary the beam trajectory the design of the chicanes used are different – with a four bend C-shaped chicane with sixteen PFF kickers proposed in the CLIC CDR, compared to the four bend dog leg chicane at CTF3 with two PFF kickers installed for the correction. The freedom to design a purpose built chicane with additional kickers in the CLIC design somewhat simplifies the challenges of obtaining a suitable layout for the prototype chicane at CTF3 using pre-existing beam lines (Chapter 2).

Another key difference is that at CTF3 the PFF prototype is operated on uncombined beam, bypassing the delay loop and completing only half a turn in the combiner ring. At CLIC the complete PFF system (including the PFF input) is placed after the drive beam recombination, and therefore would operate on combined beam. At CTF3 as the PFF input is placed prior to the delay loop any attempt to operate the PFF prototype with combined beam would be complicated by having to use the measured uncombined beam pulse to correct the combined pulse following the combiner ring. Nevertheless, operation of the prototype with combined beam could be possible and may be attempted in future tests.

The main effect of using the uncombined pulse at CTF3 rather than the combined pulse as foreseen for CLIC is that the beam pulse lengths are different. At CLIC 0.2° phase stability is needed across the 240 ns combined beam pulse. At CTF3 the uncombined pulse is much longer, up to 1.2 μ s. It is therefore not necessary to demonstrate 0.2° phase stability across




Figure 1.7: Phase sag along the $1.2 \mu\text{s}$ CTF3 beam pulse. The dashed black lines mark show the 240 ns combined CLIC pulse length centred around the region where the phase sag is flattest.

the full CTF3 pulse length to fulfill the CLIC requirements.

In fact, it is in any case impossible to demonstrate 0.2° phase stability across the full CTF3 pulse length with the PFF prototype due to the large phase sag that is present along the beam pulse. The RF pulse compression system at CTF3 [REF] results in an approximately parabolic variation of roughly 40° along the pulse which would not be present at CLIC. The phase sag is much larger than the correction range of the PFF prototype, which is designed to remove smaller, fast offsets. The PFF correction is therefore focused on the flatter, central part of the pulse around the minimum in the phase sag, where the phase variations across the pulse length relevant for CLIC are within the correction range. This is shown in Figure 1.7.

1.8 Definitions of Different Phase Statistics

Throughout the thesis several terms will be used to describe different ways of measuring the phase, as well as other parameters. These terms are briefly summarised here for reference. All quoted phase values throughout the thesis are in degrees at 12 GHz.

CTF3 provides an uncombined beam pulse length of up to $1.2 \mu\text{s}$. It is useful to compare results both along the pulse and for the mean of the pulse. To calculate “mean” statistics,

the average of each beam pulse is taken. Usually this is not taken across the full pulse length, but rather across a region of several hundred nanoseconds near the mid-portion of the pulse where the beam is most stable and the phase sag is flattest. Mean statistics are usually plotted against time in units of the pulse number, with CTF3 operating at a repetition rate of 0.8 Hz, or one beam pulse every 1.25 s. The mean phase jitter represents the standard deviation of these mean values across the duration of a dataset.

Any statistic instead described as being “along the pulse” represents the measured values point by point along the beam pulse, typically sampled at a rate of a few hundred MHz. The time axes for plots of statistics along the pulse are either in units of nanoseconds, or simply the point number along the pulse (sample number). The phase jitter along the pulse represents the standard deviation of the measured phases at each individual sample point taken across the duration of a dataset.

Many of the discussions in the thesis also quote correlation coefficients, which in all cases are Pearson product-moment correlation coefficients [REF].

1.9 Thesis Overview

This thesis documents the design, commissioning, operation and results of the PFF prototype at CTF3. In Chapter 2 the design of the TL2 chicane and the modifications to it that were necessary to achieve the desired phase shifting behaviour are described in more detail.

The performance of the PFF correction depends on the ability to precisely measure the beam phase, and so Chapter 3 presents the extensive work that has been completed to understand and maximise the precision of the new purpose built phase monitors that were installed at CTF3 for the PFF system.

As well as excellent precision in the measured phase, it is crucial to have high correlation between the phase at the end of the linac (the PFF input) and the phase at the correction location (TL2). With no correlation between the phases at these two points no improvement in phase jitter would be possible with the PFF system. Chapter 4 describes the process of understanding and improving this correlation.

Chapter 5 focuses on the setup of the remaining PFF hardware components and the commissioning of the complete PFF system. This includes the implementation of the PFF correction on the feedforward controller (the FONT5a board), the design and performance of the kicker amplifiers, as well as verifying the correction timing and correction range.

Chapter 6 then presents the best results that have been achieved with the PFF prototype to date following all the optimisations described in the rest of the thesis. An analysis of the current limitations of the system and possible future improvements to the PFF setup are also discussed. Finally, the conclusions from each chapter and suggestions for future work are summarised in Chapter 7.

Chapter 2

Design of the PFF Chicane

A basic beam line consists of focusing magnets (quadrupoles, as well as sextupoles and higher order magnets) and bending magnets (dipoles) connected by straight sections. A practical “real world” beam line must also include many diagnostic devices (such as beam position monitors, or BPMs) and additional elements (such as magnetic correctors) to be able to measure and remove the effects of small misalignments and imperfections in the beam line. The arrangement of devices along the line is referred to as the lattice. The collective settings (strengths) of each focusing element and the beam conditions they produce are referred to as the machine optics. The performance of the PFF system depends heavily on the lattice and optics of the correction chicane in the TL2 line at CTF3. This chapter describes the design of TL2, the modifications that have been made to its lattice for the PFF system and the derivation of suitable optics for the line taking in to account new constraints for the PFF system.

2.1 Introduction to Optics

This section presents basic aspects of lattice design and optics to introduce the terms used in the remainder of the thesis. Each element of a beam line can be expressed as a transfer matrix \mathbf{R} (sometimes also called a response matrix) that defines how it transforms the initial coordinates of a particle in the beam [REF]:

$$\vec{x}_f = \mathbf{R}\vec{x}_0 \quad (2.1)$$

Where \vec{x}_0 and \vec{x}_f are vectors describing the initial and final state of the particle. They are six dimensional vectors and the above equation can be expanded to become:

$$\begin{pmatrix} x_f \\ x'_f \\ y_f \\ y'_f \\ t_f \\ \Delta p_f/p_0 \end{pmatrix} = \begin{pmatrix} R_{11} & R_{12} & R_{13} & R_{14} & R_{15} & R_{16} \\ R_{21} & R_{22} & R_{23} & R_{24} & R_{25} & R_{26} \\ R_{31} & R_{32} & R_{33} & R_{34} & R_{35} & R_{36} \\ R_{41} & R_{42} & R_{43} & R_{44} & R_{45} & R_{46} \\ R_{51} & R_{52} & R_{53} & R_{54} & R_{55} & R_{56} \\ R_{61} & R_{62} & R_{63} & R_{64} & R_{65} & R_{66} \end{pmatrix} \begin{pmatrix} x_0 \\ x'_0 \\ y_0 \\ y'_0 \\ t_0 \\ \Delta p_0/p_{ref} \end{pmatrix} \quad (2.2)$$

In the transverse plane the vectors \vec{x} contain the horizontal and vertical offsets (x, y) and divergences ($x' = dx/ds, y' = dy/ds$, where s is the longitudinal position along the beam line). The parameters (x, y, s) define a curvilinear set of coordinates that measure the position of the particle with respect to the nominal or reference orbit, following the trajectory of the beam through bending magnets, for example [REF]. The final two longitudinal coordinates are the time offset (t) and momentum offset ($\Delta p_0/p_{ref}$) of the particle with respect to the reference or ideal particle. The time t is analogous to the phase of interest for the PFF system. The coefficients R_{ij} of the 6×6 transfer matrix \mathbf{R} define how the final value of the i^{th} coordinate after passing through the element is influenced by the initial value of the j^{th} coordinate prior to the element.

The simplest and most widely used type of accelerator lattice is a FODO cell, consisting of an equally spaced focusing and defocusing quadrupole with equal strength, as shown in Figure[REF]. For small horizontal or vertical offsets from the quadrupole centre the magnetic field linearly increases with the offset [REF]. The effect of a particle travelling through a quadrupolar field is analogous to a focusing lens with a focal length $1/kl$ where l is the length of the quadrupole and k is the strength of the quadrupole dependent on its design [REF]. Using the thin lens approximation the transfer matrix for a quadrupole is defined as [REF]:

$$\mathbf{R}_{\text{quad}} = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ \pm kl & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & \mp kl & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix} \quad (2.3)$$

The final horizontal divergences of a particle after traversing a quadrupole, using the above matrix, are $x'_f = x'_i \pm klx_i$ and $y'_f = y'_i \mp kly_i$. A quadrupole that focuses the beam in one plane therefore defocuses the beam in the other transverse plane. In a FODO cell a horizontally focusing quadrupole and horizontally defocusing quadrupole are used together to give a net focusing effect in both planes [REF]. The complete effect of a FODO cell on a particle can be determined by multiplying the transfer matrices of each element:

$$\mathbf{R}_{\text{FODO}} = \mathbf{R}_F \times \mathbf{R}_{\text{drift}} \times \mathbf{R}_D \quad (2.4)$$

$$\vec{x}_f = \mathbf{R}_{\text{FODO}} \vec{x}_0 \quad (2.5)$$

Where R_F and R_D are the transfer matrices of the focusing and defocusing quadrupoles respectively and R_{drift} is the transfer matrix for the drift space between the quadrupoles.

The same approach can be used to construct the transfer matrix for any complete beam line, and several of the transfer matrix coefficients are of particular interest for the PFF system both in the TL2 and TL1 transfer lines at CTF3. These will be explained in more detail later in this chapter but include mostly the coefficients related to horizontally deflecting (or “kicking”) the beam, so the R_{2j} and R_{i2} terms including the horizontal divergence, and the coefficients related to the final beam phase, so the R_{5j} terms. At CTF3 optics and transfer matrices are calculated using a MADX model of the machine [REF]. MADX is one of the leading tools available for the design and simulation of particle accelerators [REF]. All

the optics terms presented in this thesis use MADX coordinates and units [REF]. In some cases these are slightly modified from the coordinates defined above, and these differences are explained later when relevant.

The previous discussion shows how the propagation of a single particle through a beam line can be modelled. The matrix formalism above can be adjusted to describe the trajectories of many particles by replacing the column vectors \vec{x} with matrices of many column vectors describing each particle. However, to understand the properties of a complete beam it is also useful to introduce the general solution to the transverse equations of motion (Hill's Equation) [REF]:

$$x_i(s) = \sqrt{\beta_x(s)\epsilon_x} \cos[\mu_x(s) + \delta_{xi}] \quad (2.6)$$

Replacing x with y gives the equivalent solution in the vertical plane. The subscript i refers to the i^{th} particle. The transverse motion follows a modified harmonic oscillation with amplitude $\sqrt{\beta_x(s)\epsilon_x}$. The betatron (or beta) function $\beta_x(s)$ varies along the beam line and depends on the lattice and optics, whilst the beam emittance ϵ_x is a preserved quantity [REF]. The phase advance $\mu_x(s)$ defines the phase of the oscillation at each point along the lattice, with each particle having an initial phase offset δ_{xi} .

The solution has a constant of motion known as the Courant-Snyder invariant [REF]:

$$\gamma_x x^2 + 2\alpha_x x x' + \beta_x x'^2 = \epsilon_x \quad (2.7)$$

Where the explicit dependence on s of all the parameters apart from the emittance has been dropped for readability. β_x , α_x and γ_x are collectively known as the Twiss parameters [REF], where the α_x and γ_x functions relate to the beta function as follows [REF]:

$$\alpha_x = -\frac{1}{2} \frac{d\beta_x}{ds} \quad (2.8)$$

$$\gamma_x = \frac{1 + \alpha_x^2}{\beta_x} \quad (2.9)$$

The Courant-Snyder invariant defines an ellipse with area $\pi\epsilon_x$ in (x, x') phase space. All particles therefore follow elliptical trajectories in phase space as they progress through the beam line. At any point along the lattice one standard deviation of particles in a gaussian beam are contained within an envelope of $x(s) \leq \sqrt{\beta_x(s)\epsilon_x}$ [REF]. The beta function therefore defines the beam size at any point in the lattice (considering only transverse first order effects). In a FODO cell, for example, the optics is such that the beta function is minimum at the centre of the defocusing quadrupole and maximum at the centre of the focusing quadrupole [REF].

[TODO: fodo/ellipse figures? Or just text ok?]

2.2 Kicker Design

The two electromagnetic kickers provide the phase correction in the PFF system by deflecting the beam on to longer or shorter paths in the TL2 chicane. They have been designed and

built by INFN, Italy [REF], based on a similar design used at the DAΦNE collider [REF]. A schematic of the kicker design is shown in Figure 2.1. It consists of two parallel conducting strips placed along the left and right side of the beam pipe. Each strip is approximately one metre in length and the horizontal separation between the strips is 40 mm. The strips are tapered at their ends to reduce coupling impedance (to reduce the voltage induced on the strips by the beam) [REF].

At each end of each strip there is a transition to a 50Ω HN-type connector. A voltage is applied to the downstream end of each kicker strip, with opposite polarity on each side, for example $+V$ to the left strip and $-V$ to the right strip. The voltage is produced by the amplifier discussed in Section 5.2, and the voltage leaving the upstream ends of the kicker strips is also terminated back at the amplifier. The applied voltage $V(t)$ creates a horizontal, position independent, electric field and vertical magnetic field between the strips with related amplitudes as follows [REF]:

$$E_x \sim V(t) \quad (2.10)$$

$$B_y \sim \frac{V(t)}{c} \quad (2.11)$$

Where c is the speed of light. By the Lorentz force an electron in the beam propagating with speed v from the upstream end of the kicker to the downstream end (in the opposite direction to the voltage applied to the strips) experiences the following horizontal force [REF]:

$$F_x = e(E_x + vB_y) \sim e(1 + \beta)V(t) \sim 2eV(t) \quad (2.12)$$

Where e is the charge of an electron and $\beta = v/c$. The final expression holds for an ultra-relativistic particle where $\beta \simeq 1$, which is true for the CTF3 beam. In this case the forces resulting from the electric and magnetic fields have the same magnitude and direction. If the voltage were applied to the upstream end of the strip rather than the downstream end, the magnetic field would be in the opposite direction and the resulting electric and magnetic forces would cancel.

With the voltage correctly applied to the downstream end of the strips the force is as above and the kicker imparts a horizontal deflection to the beam. The kicker design gives a horizontal deflection of 1 mrad for an applied voltage of $\pm 1.26 \text{ kV}$ to each strip [REF], assuming the CTF3 beam energy of around 135 MeV [REF]. This value together with the peak voltage output from the amplifier and the optics of the TL2 chicane (as described below) defines the maximum phase offset that can be corrected by the PFF system (Section 5.4.1).

2.3 TL2

The transfer line TL2 at CTF3 transports the beam from the exit of the combiner ring to the experimental area CLEX (see Figure 1.5). The whole line is approximately 45 m long and contains both vertical and horizontal chicanes to align the outgoing combiner ring beam line to the CLEX entrance. The PFF system attempts to correct the beam phase using the horizontal chicane at the end of TL2, where the two kickers are installed. Further details of the design of TL2 can be found in [REF].

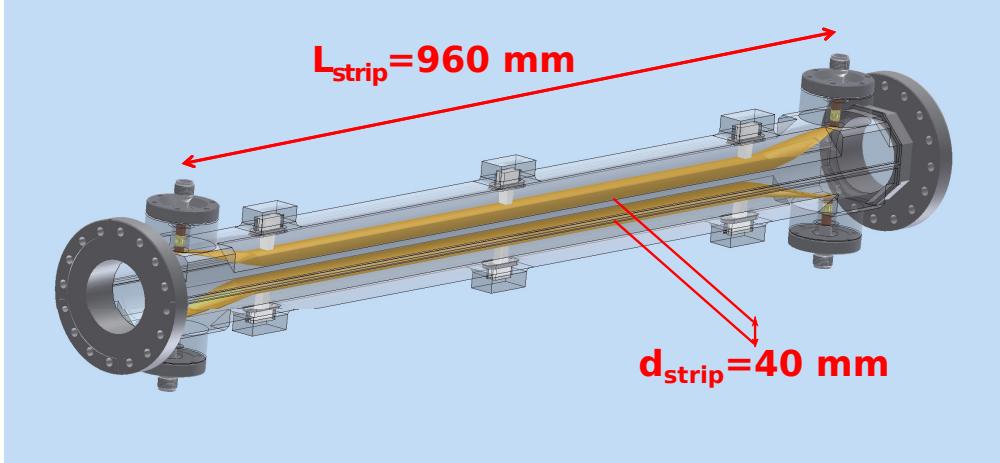


Figure 2.1: Technical drawing of the kicker design. The kicker is shown in a vertical orientation with the strips on the top and bottom. When installed in the beam line the kicker is oriented with the strips on the left and right, in order to create a horizontal electric field between the strips.

The diagram in Figure 2.2 shows a birds-eye view of the TL2 line and the lattice of the line. To interpret the diagram it is useful to introduce the device naming convention at CTF3. Devices have names of the form:

[CC].[QF][D][0840]

The first two letters refer to the section of the machine, with the prefix CC used for TL2. These are not included in the diagram to improve readability. In Chapters 3 and 4 the prefix CT is used to refer to the CT-line after the linac and transfer line TL1 prior to the combiner ring. The second group of letters refer to the type of device, the main ones being QF and QD for horizontally focusing and defocusing quadrupoles, BH and BV for horizontal and vertical dipoles, BP for beam position monitors (BPMs) and DH and DV for horizontal and vertical corrector magnets. The last letter indicates the type of that device, with four different designs of quadrupole used along the TL2 line (G-type, H-type, L-type and D-type), for example. The four final numbers indicate the position of that device along the line, in ascending order from the beginning to the end of the line.

The horizontal chicane of interest for the PFF system starts at the dipole CC.BHG0500 and ends at the dipole CC.BHG0800. The first (500) and last (800) dipoles bend the beam through $+31^\circ$ and -31° respectively. Inside the chicane there are two further dipoles of a different type – CC.BHH0600 and CC.BHH0700, which deflect the beam through $+17^\circ$ and -17° respectively. Some details on the differences between the two types of dipole are given in Section 2.5.3. The resulting overall chicane has a “dog leg” shape around 12 m in length, with three straight sections around 4 m in length between the bending magnets. Each straight section contains a triplet of quadrupoles and either one (in the first and last sections) or two (in the middle section) BPMs (of the BPI type [REF]). Although the quadrupoles are labelled as horizontally focusing or defocusing the polarity of the current sent to each can be reversed so that it focuses in the opposite plane. The F or D labels refer to whether

the magnet is horizontally focusing or defocusing when a positive current is sent to the quadrupole.

Other features along the TL2 line that are important for the derivation of optics seen later in this chapter include the vertical chicane and two long drift spaces without focusing elements. The vertical chicane starts and ends at CC.BVA0300 and CC.BVB0400 respectively, and contains a triplet of quadrupoles. Between the quadrupole CC.QFD0840 (the last shown in the diagram) and CC.QFL0910, there is a long drift space of around 4 m with no focusing elements as the beam pipe passes through into the neighbouring building where the CLEX area is located. Between the quadrupole CC.QFH0230 and CC.QFL0270 there is another long drift space, around 7 m. The Twiss beta and alpha functions entering these long drifts must be carefully chosen to avoid unrecoverable growth in the beam size.

[TODO: picture of TL2/chicane]

2.3.1 Integration of PFF Hardware

Due to building and cost constraints the PFF prototype had to make use of the pre-existing layout of the TL2 horizontal chicane, with only minor modifications possible to accommodate the PFF hardware. These changes are highlighted in yellow in Figure 2.2.

As the chicane was already densely packed with quadrupoles and other devices the integration of the two kickers was not straightforward. To maintain the functionality of the lattice quadrupoles could not be removed, and thus instead the kickers have been installed inside wide aperture ‘H-type’ quadrupoles [REF]. Two ‘L-type’ quadrupoles (now CC.QDL0330 and CC.QDL0370) from the horizontal chicane were swapped with two ‘H-type’ quadrupoles (now CC.QDH0490 and CC.QDH0790) from the vertical chicane. The two PFF kickers, CC.KHD0480 and CC.KHD0780, are then installed inside the aperture of these quadrupoles, prior to the first and last dipole of the horizontal chicane. In addition, two magnetic correctors (now CC.DHF0465 and CC.DHF0765) were installed around the PFF kickers to facilitate a complementary, large range but low bandwidth, slow phase correction [REF]. A schematic of the installation of one of the kickers inside the quadrupole and corrector is shown in Figure [REF]. The kicker CC.KHD0480 will also be referred to as the first kicker (or K1), and CC.KHD0780 as the second kicker (or K2).

Apart from the quadrupoles and correctors two BPMs (now CC.BPI0435 and CC.BPI0735) had to be moved slightly to vacate the area now occupied by the kickers. Finally, a slot near the start of TL2 (CC.PHM0245) was reserved for the installation of an additional phase monitor to verify the beam phase prior to the correction. Eventually this was not necessary and has not been pursued as the 2 cm aperture of the phase monitors (Section 3.1) compared to the 4 cm aperture of the neighbouring beam pipe would have created beam setup difficulties for normal operation at CTF3 [REF].

[TODO: pictures/diagram of kickers in quadrupole]

2.4 TL2 Optics Constraints

25

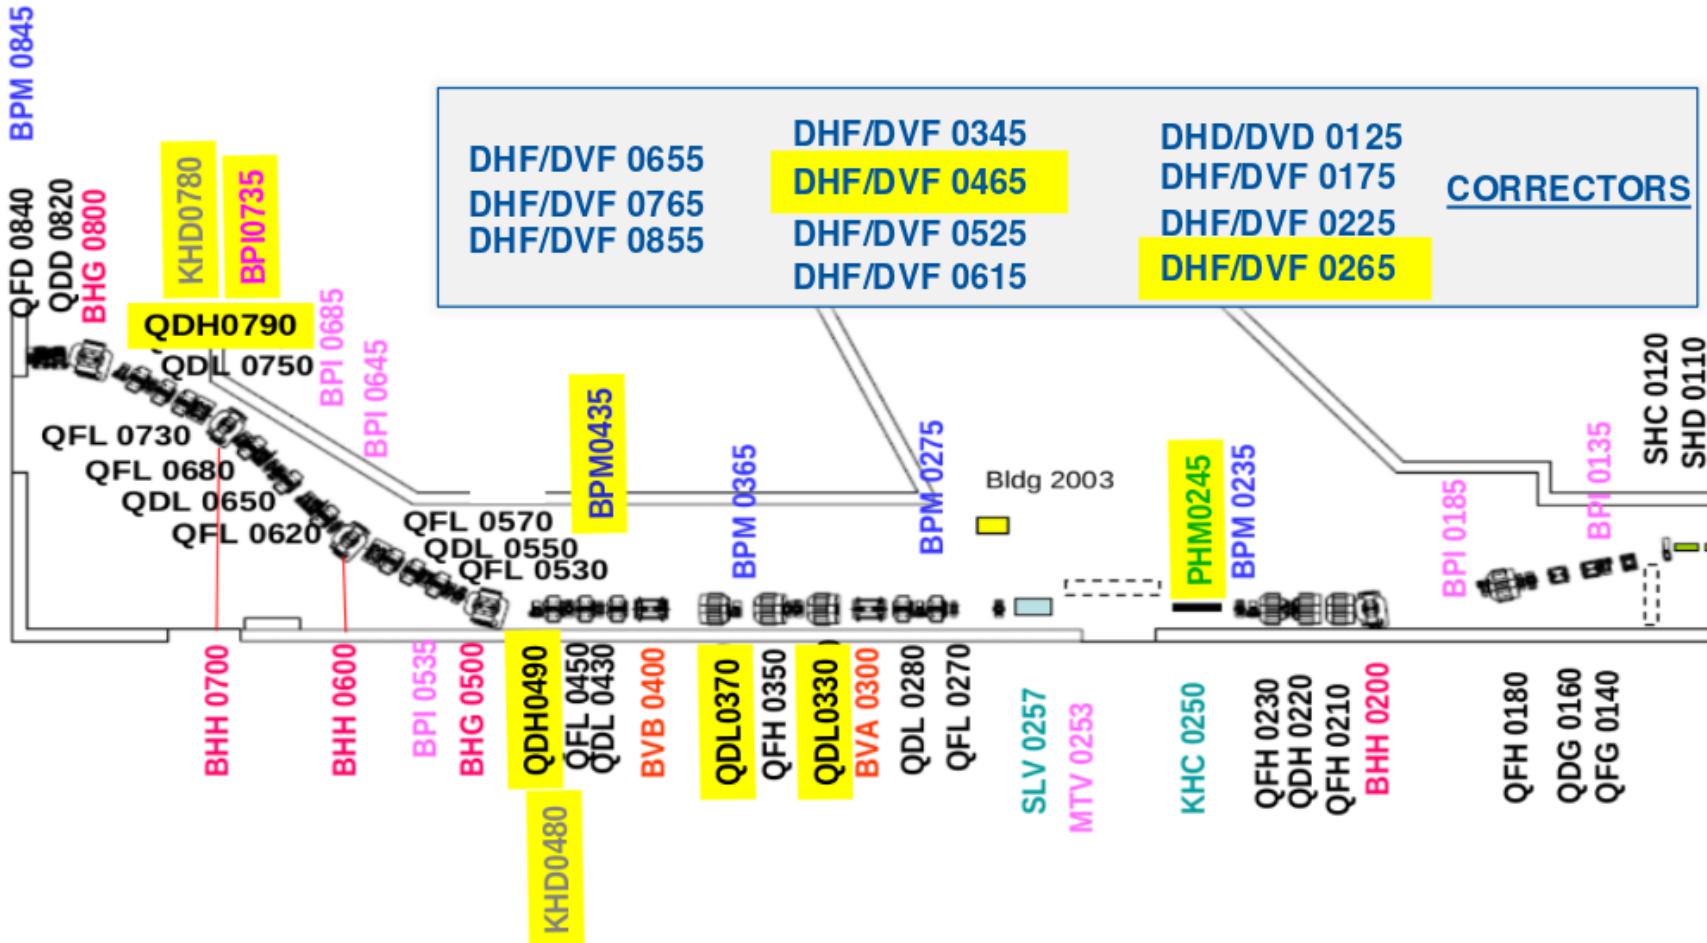


Figure 2.2: New TL2 lattice for PFF. Changes highlighted yellow.

2.4 TL2 Optics Constraints

To take in to account the changes made to the TL2 lattice new optics were needed. This section summarises the various optics constraints that must be met in TL2. These can be split in to two types – the nominal optics constraints, required to recover the same (or similar) beam conditions as before the changes, and the new optics constraints for operation of the PFF system, required to create the desired phase shifting behaviour in the chicane.

2.4.1 Nominal Optics Constraints

The nominal optics constraints are mostly put in place to minimise the transverse beam size along the line by restricting the magnitude of the dispersion and twiss functions. The final constraints, on the R_{56} transfer matrix coefficient, relates to the longitudinal stability of the beam.

Dispersion

The trajectory of a particle through a dipole depends on its energy, with lower energy particles following a smaller radius of curvature and high energy particles a larger radius. After a dipole the orbit of a particle though the following elements therefore depends on its energy. This effect is characterised using the horizontal and vertical dispersion, D_x and D_y , which are defined as follows (considering only the energy component of the beam orbit):

$$x(s) = D_x(s) \left(\frac{\Delta p}{p} \right) \quad (2.13)$$

$$y(s) = D_y(s) \left(\frac{\Delta p}{p} \right) \quad (2.14)$$

The dispersions $D_x(s)$ and $D_y(s)$ vary along the lattice and are equivalent to the transfer matrix coefficients $R_{16} = D_x$ and $R_{36} = D_y$.

Optics are usually created so that there is no dispersion ($D_x = D_y = 0$) in straight sections. However, inside chicanes and rings the dispersion can never always be zero. To give zero dispersion in the straight sections the dispersion must therefore be closed at the exit of all bending sections. Dispersion closure means that both the dispersion and its derivative ($D'_x = dD_x/ds$, and $D'_y = dD_y/ds$) are zero at the exit from the chicane or ring. In TL2 this condition applies after the bend CC.BHH0200 at the start of the line and at the exit of the horizontal chicane (CC.BHG0800) for the horizontal dispersion, and at the exit of the vertical chicane (CC.BVB0400) for the vertical dispersion.

Within the bending sections the magnitude of the dispersion should be kept as small as possible whilst meeting the other optics constraints. At CTF3 there can be peak-to-peak energy offsets at around the $\pm 1\%$ level [REF]. Dispersion is then usually the largest contribution to the beam size, with a dispersion of 1 m giving excursions up to ± 1 cm in individual particle orbits, for example. The diameter of the beam pipe in bending sections

Parameter	Value
β_x	7.26 m
β_y	5.90 m
α_x	-4.84
α_y	-1.27

Table 2.1: Initial twiss parameters for the TL2 line, taken at the entrance to CC.QFH0210.

at CTF3 is 10 cm in most cases, as opposed to 4 cm in straight sections, in order to minimise the effects of dispersion dependent beam size growth on the beam transport [REF]. However, the second kicker installed in the chicane for the PFF system (CC.KHD0780) only has the normal 4 cm aperture (2 cm radius). Dispersion around the second kicker must therefore be kept well below 2 m to avoid losing a fraction of off-energy particles on the kicker strips.

Twiss Functions

The twiss beta functions define the energy independent component of the transverse beam size and the alpha functions define the rate of change of the beta function along the line, as discussed in Section 2.1. The transverse beam size at any point in the lattice is related to the beta function and emittance via the relationship $\sqrt{\beta_x(s)\epsilon_x}$ (and the vertical equivalent). At CTF3 the beam emittance is usually around $0.5 \mu\text{m}$, but may be up to a factor two larger than this in the horizontal plane depending on the beam setup [REF]. With an emittance of $0.5 \mu\text{m}$ a beta value of 50 m at one point in the lattice corresponds to a transverse beam size of 0.5 cm at that location, for example. The beta and alpha functions should be kept as small as possible to minimise the beam size along the line. Beta is usually kept below 50 m, but values up to 100 m can be accepted if necessary [REF]. As mentioned previously there are long drift spaces with no focusing elements at the start and end of TL2, and it is most difficult to maintain low beta and alpha values following these regions.

At the start of TL2 the initial values of the twiss parameters are also constrained by the optics of the combiner ring. In other words, the beta and alpha values at the start of TL2 should be the same as the beta and alpha values at the exit of the combiner ring. This avoids discontinuities between sections in the complete CTF3 optics. The TL2 optics have usually been matched (Section 2.6) starting from CC.QFH0210, and the required initial conditions at this location are summarised in Table 2.1.

R56

The dispersion describes how the transverse orbit of a particle is changed by its energy in bending sections as already discussed. These differences can also change the longitudinal path length of the particle's trajectory, thereby shifting the particle's phase (described by the time t , the fifth coordinate in the matrix formalism). This effect is described by the

transfer matrix coefficient R_{56} :

$$t_f = t_i + R_{56} \left(\frac{\Delta p}{p} \right) \quad (2.15)$$

The R_{56} value between the entrance and exit of all bending sections at CTF3 is nominally zero so that there is no transformation of energy jitter in to phase jitter. In TL2 this places the constraints for R_{56} to be zero between the entrance and exit of the vertical chicane (CC.BVA0300 to CC.BVB0400) and between the entrance and exit of the horizontal chicane (CC.BHG0500 to CC.BHG0800).

2.4.2 PFF Optics Constraints

All the additional PFF optics constraints place requirements on the transfer matrix coefficients between the two kickers, from the exit of the first kicker to the entrance of the second kicker. There are two sets of constraints, one to maximise the correction range of the PFF system and the other to ensure the PFF system does not degrade the orbit stability of the beam after the chicane.

Correction Range

The PFF system clearly requires the path length between the two kickers to depend on the applied kick. This difference in path length dependent on the voltage applied to the kickers is what shifts the time or phase of the beam to form the correction. The transfer matrix coefficient that relates the time variable to the deflection induced by the kickers is R_{52} :

$$t_{K2} = t_{K1} + R_{52} x'_{K1} \quad (2.16)$$

Where t_{K1} and t_{K2} are the time offset of the particle at the exit of the first kicker and at the entrance to the second kicker respectively. x_{K1} is the divergence at the exit of the first kicker resulting from the applied kick. MADX uses units of metres for its ‘time’ variable [REF]. To convert these distances in to 12 GHz degrees they must be multiplied by the constant factor $360/\lambda_{12\text{GHz}}$, where $\lambda_{12\text{GHz}} = 2.5$ cm is the 12 GHz wavelength. Directly in terms of phase (in degrees) the equation above therefore becomes:

$$\phi_{K2} = \phi_{K1} + R_{52} \left(\frac{360}{\lambda_{12\text{GHz}}} \right) x'_{K1} \quad (2.17)$$

Where ϕ_{K2} is the phase at the entrance to the second kicker (the corrected phase) and ϕ_{K1} is the initial uncorrected phase at the exit of the first kicker. The maximum value of x'_{K1} is fixed by the peak voltage output from the kicker amplifiers and the design of the kickers themselves. To obtain the largest possible correction range for the PFF system given the parameters of the hardware, the R_{52} transfer matrix coefficient should be as large as possible. For example, with $R_{52} = 1$ m and a maximal kick of $x'_{K1} = \pm 1$ mrad, the correction range of the PFF system would be ± 14.4 degrees.

The path length difference in the chicane largely results from differing trajectories in the dipoles. In this way it is somewhat analogous to the dispersion, which describes the energy

dependent difference in beam orbit after dipoles. This has the unfortunate consequence of leading to optics with high R_{52} values also tending to have high peak dispersion values in the chicane. The PFF optics must therefore be a compromise that achieves a reasonable correction range whilst keeping the dispersion small enough to avoid beam losses in the chicane.

Orbit Closure

The PFF system should not change the beam orbit after the chicane, which means the beam position and divergence after the second kicker must be independent of the applied kicks. In other words, the second kicker must close the horizontal orbit bump created by the first kicker. To understand the further constraints this places on the optics the position, x_{K2} , and divergence x'_{K2} of the beam at the entrance to the second kicker will be considered first. These can be expressed as:

$$x_{K2} = R_{11}x_{K1} + R_{12}x'_{K1} \quad (2.18)$$

$$x'_{K2} = R_{21}x_{K1} + R_{22}x'_{K1} \quad (2.19)$$

Where x_{K1} and x'_{K1} are the position and divergence at the exit of the first kicker, and R_{11} , R_{12} , R_{21} and R_{22} are transfer matrix coefficients for the optics between the exit of the first kicker and the entrance to the second kicker. The beam position at the exit of the first kicker is proportional to the applied kick:

$$x_{K1} = mx'_{K1} \quad (2.20)$$

Here m is a constant that depends on the properties of the kicker and also on the strength of the quadrupole CC.QDH0490 within which the kicker is installed (Section 2.3.1). Substituting this expression in to the equations for x_{K2} and x'_{K2} gives:

$$x_{K2} = \left(R_{11} + \frac{R_{12}}{m} \right) x_{K1} \quad (2.21)$$

$$x'_{K2} = (mR_{21} + R_{22})x'_{K1} \quad (2.22)$$

As stated x_{K2} and x'_{K2} are defined above at the entrance to the second kicker. The requirement for the PFF chicane optics is that the position and divergence at the exit of the second kicker are zero independent of the applied kicks. However, a derivation of the exact expression for the optics requirements between the kickers in order to close the orbit at the exit of the second kicker is complicated by the fact that the quadrupole around the second kicker, CC.QDH0790, can have a different strength to the quadrupole around the first kicker.

For the purpose of the discussion here the simplified case where CC.QDH0790 has the same strength but opposite polarity (focuses in the opposite plane) as CC.QDH0490 will be considered. The ideal case where the two kickers can be powered with the same magnitude voltage but opposite polarity is also assumed. With these conditions, the second quadrupole/kicker effectively have precisely the opposite effect on the beam as the first quadrupole/kicker. To close the orbit after the second kicker the position and divergence at the entrance to the second kicker must therefore meet the following criteria:

$$x_{K2} = -x_{K1} \quad (2.23)$$

$$x'_{K2} = x'_{K1} \quad (2.24)$$

Comparing these two expressions to the previously derived equations for x_{K2} and x'_{K2} then yields the following optics constraints:

$$R_{11} + \frac{R_{12}}{m} = -1 \quad (2.25)$$

$$mR_{21} + R_{22} = +1 \quad (2.26)$$

There are many possible solutions to these expressions, with the simplest example being $R_{11} = -1$, $R_{12} = 0$, $R_{21} = 0$ and $R_{22} = 1$. However, the optics matching (Section 2.6) allows the quadrupoles CC.QDH0490 and CC.QDH0790 to have different strengths. MADX is then used to model the actual beam orbit in the chicane and the figure of merit is for the simulated orbit to be closed after the second kicker, rather than for the above constraints to be met. Nevertheless, the optics eventually created do satisfy Equations 2.25 and 2.26 within several percent (Section 2.6.2).

2.5 TL2 Optics Measurements

As seen above there are many optics constraints in TL2 that must be met both to ensure that the beam can be transported efficiently in to the CLEX area as well as to obtain the desired behaviour in the horizontal chicane for the PFF system. MADX can be used to create optics that meet these criteria (Section 2.6), but they will be of no use if the model of TL2 does not accurately describe the actual characteristics of the line. As such, a series of measurements has been taken to determine and improve the accuracy of the TL2 MADX model. Measurements of this type had not previously been completed for the TL2 line at CTF3, thus any errors identified in the model are not limited to only the hardware changes made for the PFF system.

2.5.1 Method

The TL2 line includes 12 magnetic correctors, as shown in Figure 2.2 (plus one at the end of TL2 in the CLEX area, which is not shown in the figure). An electric current can be applied to horizontal or vertical coils on the corrector, creating independently adjustable horizontal and vertical fields that deflect the beam. The primary purpose of the correctors is to compensate differences in beam orbit resulting from small misalignments of devices along the line. However, they can also be used to test the accuracy of the TL2 model.

Measurements were taken in which the current applied to one of the correctors was changed, causing the beam to be deflected on to a new trajectory along the rest of the line. The new orbit is observed in the BPMs downstream of the corrector, with a total of 12 BPMs in TL2 (also shown in Figure 2.2). The position offset in each BPM depends on the transfer matrix between the corrector and that BPM, and therefore on the focusing properties of all the magnetic elements between the corrector and the BPM. By deflecting the beam with each corrector along the line, and in both planes, the response of the whole line as well as individual parts of the line can be determined.

The same process can then be repeated in the MADX model, applying a current to one of the correctors and creating a simulated deflected orbit. All the BPMs are included in the model, allowing the real measured position in each BPM to be compared to the simulated position from the MADX model. Any difference between the two highlights inaccuracies in the modelled properties of the TL2 lattice.

2.5.2 Results with Original MADX Model

Figures 2.3 and 2.4 show an example of the results obtained with the original version of the TL2 MADX model, in the horizontal and vertical planes respectively. One of the first correctors in TL2, CC.DHF0175 (or the equivalent CC.DVF0175 in the vertical plane), is used, so the results are sensitive to errors in the model along the full length of the line. Three lines are shown in each figure – the red line labelled “Measurement” corresponds to the actual measured position in the BPMs, the dashed blue line “Model” shows the simulated position in the BPMs from MADX, and the final black line “Full” shows the simulated MADX orbit propagated through all elements along the line (not restricted to only the BPM positions). In a perfect model of the line the blue “Model” and red “Measurement” lines would be identical.

In the horizontal plane there is good agreement between the model and the measurement in the three BPMs following the corrector (up until CC.BPM0275). After this point the response in the model is clearly completely different to the measurement. For example, the measured position shifts by 10 mm between CC.BPM0275 and CC.BPM0365 in the measurement, but only 6 mm in the model. Towards the end of the line the model is close to being the inverse of the measurement. There is also poor agreement between the model and the measurement in the vertical plane, with a large difference already visible at CC.BPM0275 in this case. The peak-to-peak vertical orbit offset in the model is roughly a factor two larger than the measurement.

2.5.3 Sources of Errors in MADX Model

The results described above immediately demonstrated large discrepancies between the model and the actual response of TL2. Based on previous experience from corrections made to the MADX model for other sections of CTF3 [REF] two key areas were identified to investigate to try to improve the model — the properties of the “L-type” quadrupoles in TL2, and the focusing effects from the dipoles in TL2.

Quadrupole Strengths

16 out of the 27 quadrupoles in TL2 are of the “L-type”, with labels of the form CC.QDLxxxx or CC.QFLxxxx. This includes all the quadrupoles in the horizontal chicane apart from the two wide aperture “H-type” quadrupoles within which the PFF kickers are installed. These quadrupoles were reclaimed from the CELSIUS project in Uppsala, Sweden [REF]. In the

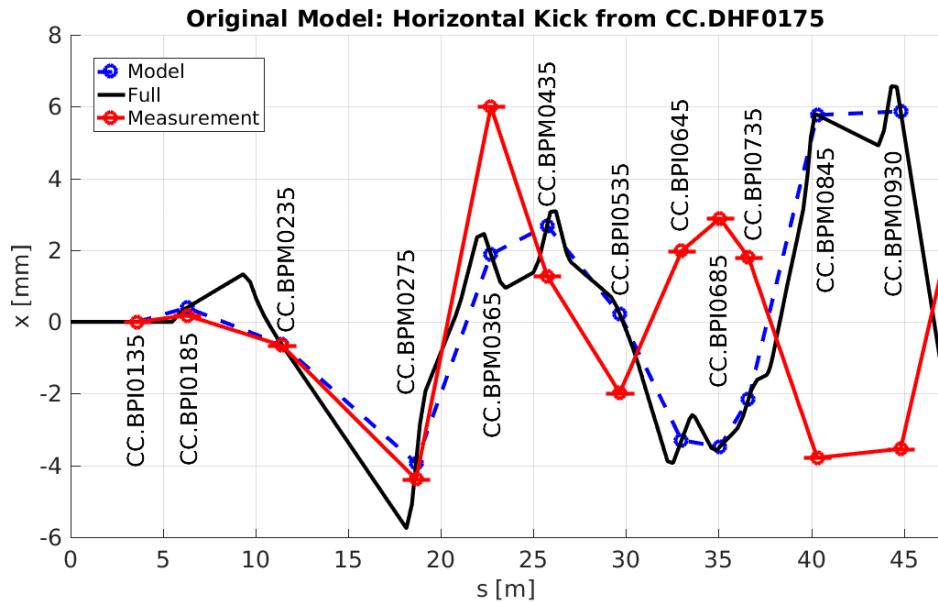


Figure 2.3: Horizontal orbit due to a kick from the corrector CC.DHF0175 compared to the original MADX model of TL2.

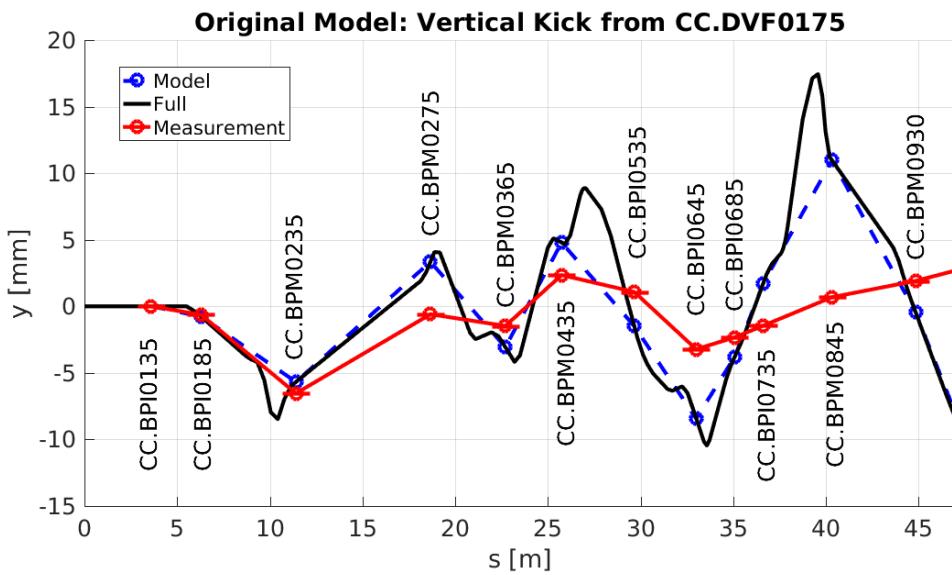


Figure 2.4: Vertical orbit due to a kick from the corrector CC.DVF0175 compared to the original MADX model of TL2.

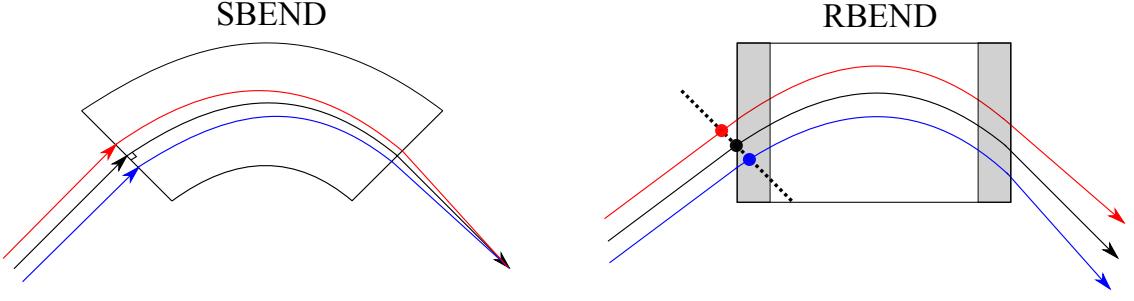


Figure 2.5: Geometry of SBEND and RBEND dipoles.

MADX model of TL2 the focusing strength of these magnets, $K1$ is defined by the following parameters:

$$K1 = \frac{FQL \times I}{E} \quad (2.27)$$

Where $FQL = 31.78$ is a constant defined by the properties of the quadrupole, I is the current delivered to the quadrupole from its power supply (a larger focusing effect) and E is the beam energy. The $K1$ value is analogous to the constant k in the quadrupole transfer matrix previously seen in Equation 2.3. This must be multiplied by the 30 cm active magnetic length of the quadrupole to determine the equivalent focal length of the quadrupole.

As the properties of the “L-type” quadrupoles were not measured in place at CERN prior to their use in CTF3 there was a large uncertainty on the correct value of FQL to use. Changing the FQL value was therefore a good candidate to try to reduce errors in the MADX model.

Dipole Focusing

Although the primary purpose of dipole magnets is to bend the beam they also give focusing effects that depend on the design of the magnet, and in particular the orientation of the pole faces. The seven dipoles in TL2 can be roughly split into two types in terms of the focusing effects they are expected to produce — sector magnets (SBENDs) and rectangular magnets (RBENDs). Figure 2.5 compares the geometry of the pole faces for SBEND and RBEND dipoles. In SBEND magnets the ends of the pole faces are oriented such that the reference trajectory of the beam (black) enters and leaves the magnet perpendicular to the pole face. Alternatively, in RBEND dipoles the pole faces at the entrance and exit of the magnet are parallel to each other. In this case the reference trajectory forms an angle $\theta/2$ with the pole faces, where θ is the angle through which the beam is deflected by the magnet [REF]. In TL2 the CC.BHH0200, CC.BHG0500, CC.BHH0600, CC.BHH0700 and CC.BHG0800 dipoles are RBENDs, and the CC.BVA0300 and CC.BVB0400 dipoles are SBENDs [REF].

The first focusing effect from dipoles relates to the path length of the beam through the magnet. In SBEND magnets this length depends on the position offset in the (horizontal or vertical) bending plane. Particles entering the magnet experience the dipole field for a longer length on one side of the reference trajectory, and a shorter length on the other side of the reference trajectory, producing a focusing effect in the bending plane. This is also

shown in Figure 2.5. In RBEND magnets the length of the trajectory is the same for all incoming position offsets, so this focusing effect is not present.

Another way in which dipoles produce focusing effects is via their fringe fields. The fringe fields describe the increase in magnetic field from zero to the peak bending field as the beam approaches the entrance to the dipole, and from the peak field back to zero as the beam leaves the dipole [REF]. The extent and rate of increase of the fringe fields depends on the design of the dipole, in particular the pole separation [REF]. In MADX the nature of the fringe fields is defined by the variables HGAP, the pole separation, and FINT, which describes the shape of the increase up to the peak field (e.g. linear) [REF].

In RBEND magnets, where the reference orbit is not perpendicular to the pole faces, the longitudinal position along the trajectory where a particle enters the fringe field linearly depends on its transverse position offset in the bending plane. This is shown in Figure 2.5, where the grey shaded areas mark the region of the fringe fields. The red trajectory enters the fringe fields later than the blue trajectory. This creates a position dependent focusing effect which is defocusing in the bending plane and focusing in the other transverse plane, similar to a quadrupole [REF]. In the ideal case this type of focusing is not present in SBEND magnets [REF].

Finally, higher order field errors in the dipole can lead to further focusing terms. MADX allows a quadrupole field coefficient (K1) to be added to a dipole to model this effect. In other sections of CTF3 it was found that the default MADX parameters did not give a good approximation of the dipole focusing terms [REF]. The FINT, HGAP and K1 parameters for each of the 7 dipoles in TL2 are therefore also good candidates to try to correct the MADX model of TL2.

2.5.4 Corrections to MADX Model

By varying the parameters of the dipole focusing and the L-type quadrupole strengths described above it has been possible to find a solution that gives much better agreement between the corrector kick measurements and the expected response in the MADX model. The new values for the parameters do not necessarily represent their true real world characteristics (for example the dipole pole separation), but are rather a set of values that adjusts the simulated MADX focusing in a way that matches the actual behaviour. An iterative process was used, with a mixture of automatic optimisations via algorithms in MATLAB and MADX [REF] and manual interventions to drive the calculated adjustments towards a reasonable solution and away from any local minima.

The first adjustment made was to the modelled strength of the L-type quadrupoles. As the majority of quadrupoles along TL2 are of this type the focusing and calculated transfer matrices in the MADX model are very sensitive to any change in their strength. It was quickly determined that the focusing strength of the L-type quadrupoles needed to be increased to improve the agreement between the measurement and the model. Figure 2.6 shows an example of a horizontal kick from the corrector CC.DHF0175 with an optimal increase of 7% in the L-type quadrupole strength. This can be directly compared to the original model in

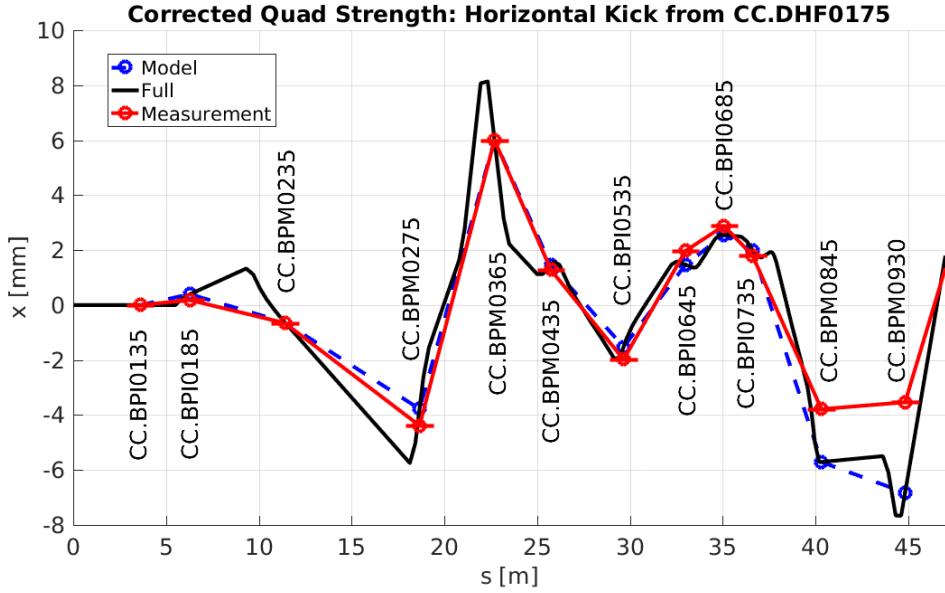


Figure 2.6: Horizontal orbit due to a kick from the corrector CC.DHF0175 with the L-type quadrupole strengths increased by 7%.

Figure 2.3. By increasing the focusing from the L-type quadrupoles the largest discrepancies between the MADX model and the measurement are removed. In particular the region from inside the vertical chicane (CC.BPM0365) to inside the horizontal chicane (CC.BPI0735) originally showed large differences including sign flips in the horizontal plane. With the adjusted L-type strength these are removed and the agreement is excellent.

Figures 2.7 and 2.8 then show an example of adjusting the focusing of the four dipoles in the horizontal chicane (the outer 500/800 pair and inner 600/700 pair) to remove remaining differences between the measurement and the model. In this case the beam is kicked from the CC.DVF0525 corrector just after the first dipole in the horizontal chicane, making the measurement insensitive to any optics errors prior to CC.DVF0525 in the line. In both figures the 7% increase in the L-type quadrupole strength has been applied in the MADX model. With the quadrupole correction in place but no adjustment to the dipole focusing, in Figure 2.7, the simulated orbit has the same overall shape as the measured orbit inside the chicane (up until CC.BPI0735). However, there are still offsets between the two and these eventually lead to a large discrepancy between the model and the measurement still being present in the two BPMs following the chicane (CC.BPIM0845 and CC.BPM0930), which is also seen in Figure 2.6.

By adjusting the HGAP, FINT and quadrupole K1 component of the dipoles in the horizontal chicane a solution is found that reduces the difference between the model and the measurement in the chicane, and then also gives much better agreement between the two after the chicane. This is shown in Figure 2.8. Repeating this process with correctors prior to the vertical chicane, and prior to the CC.BHH0200 bend at the start of TL2, yields the focusing parameters for the seven dipoles in the line summarised in Table 2.2. The largest quadrupolar component of $K1 = 0.425 \text{ m}^{-2}$ for the CC.BHH0600 and CC.BHH0700 dipoles corresponds to a focusing strength of roughly 2% of a typical quadrupole magnet in TL2

Device	Parameter	Original Value	Corrected Value
CC.BHH0200	HGAP	0 m	0.084 m
	FINT	0	0.79
	K1	0 m ⁻²	0 m ⁻²
CC.BVA0300 & CC.BVB0400	HGAP	0 m	0 m
	FINT	0	0
	K1	0 m ⁻²	-0.125 m ⁻²
CC.BHG0500 & CC.BHG0800	HGAP	0 m	0.06 m
	FINT	0	0.4
	K1	0 m ⁻²	0.15 m ⁻²
CC.BHH0600 & CC.BHH0700	HGAP	0 m	0.06 m
	FINT	0	0.2
	K1	0 m ⁻²	0.425 m ⁻²
L-type Quadrupoles	FQL	31.78	34.03

Table 2.2: Changes made to the dipole focusing and quadrupole strength parameters in the TL2 MADX model in order to improve the agreement with the kick measurements.

[REF].

Figures 2.9 and 2.10 compare the measured horizontal and vertical orbit to the new MADX model of TL2 with all the corrections from Table 2.2 in place. The same corrector, CC.DHF0175, and data are used as for the examples from the original model in Figures 2.3 and 2.4. The new version of the model is a clear improvement. In both the horizontal and vertical planes the measured and modelled beam orbit now agree along the full length of the line within a small margin of error. With the original MADX model the mean absolute difference between the measured and simulated positions in the ten BPMs following CC.DHF0175 was 3.9 ± 1.0 mm in the horizontal plane and 3.1 ± 0.8 mm in the vertical plane. The corrected model reduces these differences by an order of magnitude, to 0.2 ± 0.1 mm in the horizontal plane and 0.3 ± 0.1 mm in the vertical plane.

2.6 Matched TL2 Optics

With the corrected MADX model of TL2 in place the new optics for the PFF system were created. The optics were obtained using MADX matching libraries [REF], in which many desired constraints can be defined for the line (each with its own weight) and then one of several optimisation algorithms can be used to change the quadrupole strengths and derive an optics that meets those constraints. Two sets of optics were created for TL2 – a nominal set of optics containing only the constraints from Section 2.4.1 [REF], and a set of PFF optics containing the additional R_{52} and orbit closure related constraints for the PFF system from Section 2.4.2. Without the additional PFF constraints roughly a factor two smaller maximum dispersion and beta values can be achieved in the nominal optics [REF]. However, the beam quality achieved with the PFF optics has been sufficient to use routinely at CTF3 (including for non-PFF running), and so only this optics is documented here.

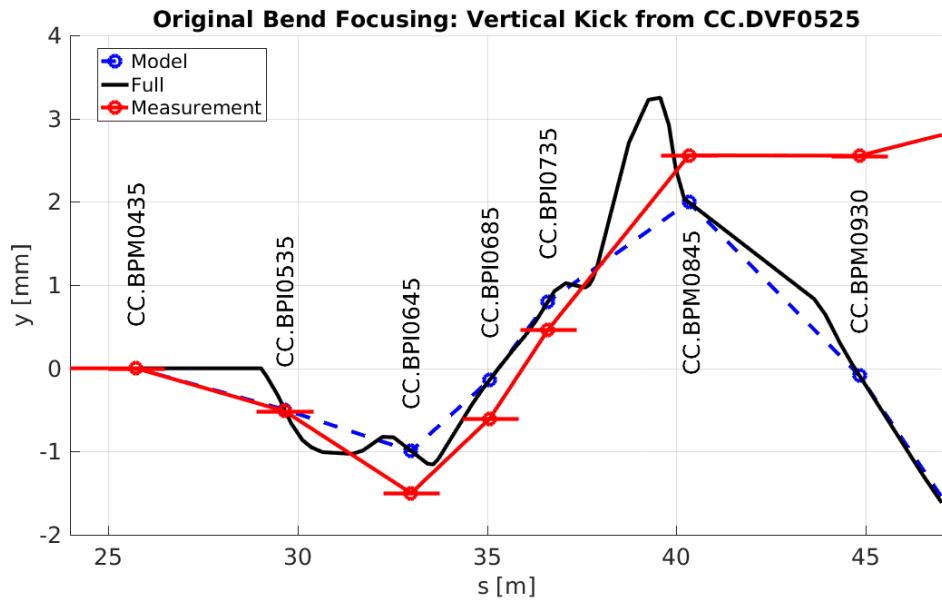


Figure 2.7: Vertical orbit due to a kick from the corrector CC.DVF0525 with the default MADX model of the dipole focusing.

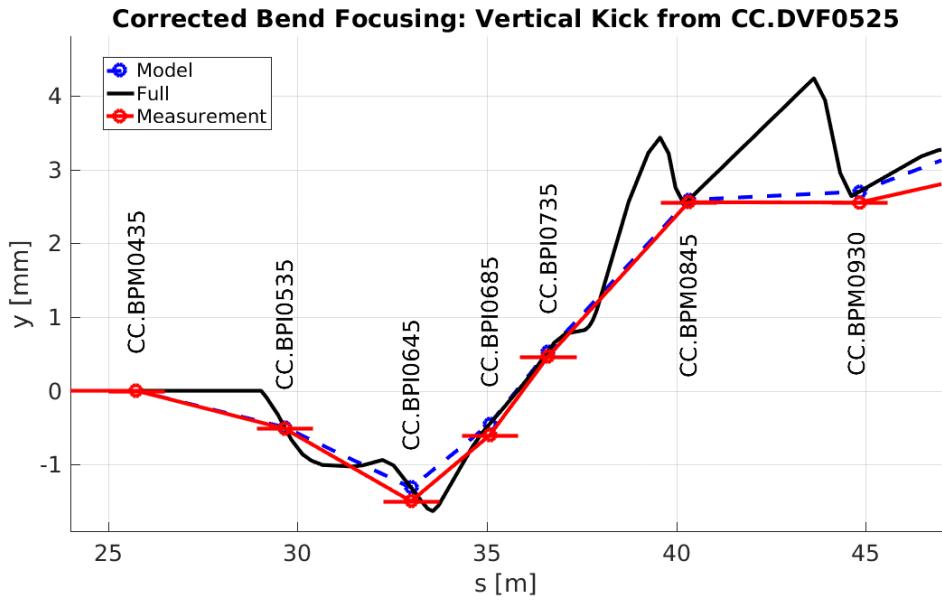


Figure 2.8: Vertical orbit due to a kick from the corrector CC.DVF0525 with the corrected model of the dipole focusing.

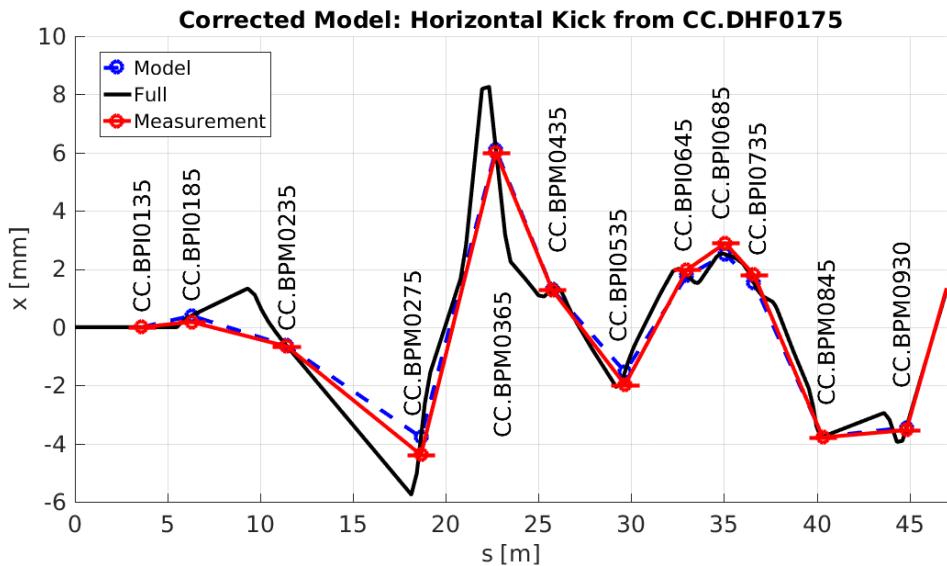


Figure 2.9: Horizontal orbit due to a kick from the corrector CC.DHF0175 compared to the corrected MADX model of TL2.

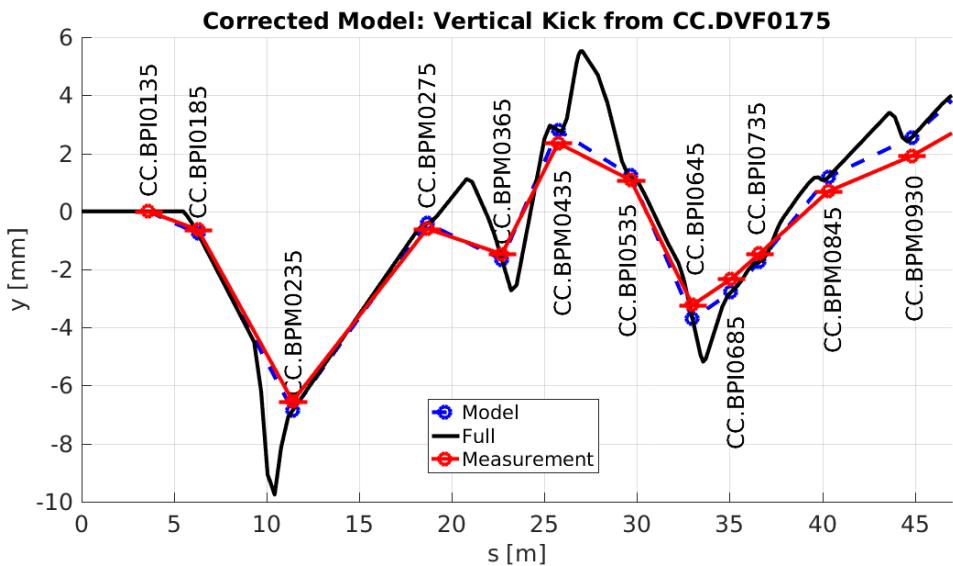


Figure 2.10: Vertical orbit due to a kick from the corrector CC.DVF0175 compared to the corrected MADX model of TL2.

2.6.1 Matching Process

For the purposes of the matching TL2 was split in to three parts [REF] – from the beginning of the line to the exit of the vertical chicane (CC.BHH0210 to CC.BVB0400), from the exit of the vertical chicane to the exit of the horizontal chicane (CC.BVB0400 to CC.BHG0800) and the end of the line (from CC.BHG0800 and in to CLEX). The middle section, containing the PFF kickers and horizontal chicane, is the most critical for the PFF system.

At the beginning of each matching section the initial twiss parameters for that part of the lattice must be defined in MADX. At the start of the first section this is fixed by the properties of the beam leaving the combiner ring (Table 2.1). The initial and final twiss functions of the middle section containing the horizontal chicane were left as free parameters to allow the greatest degree of flexibility for meeting the constraints in the chicane. The optics for the middle section is therefore created first, with the resulting initial and final twiss parameters forming additional matching constraints that must be met at the end of the first section and the start of the final section.

Suitable initial values for the quadrupole strengths (defined by the current sent to each quadrupole from its power supply) must also be chosen to ensure the matching algorithms can converge to a good solution in a reasonable amount of time. For this purpose the currents from the prior optics for TL2 (before the modifications for the PFF system and the model corrections) were used [REF].

To accurately simulate the effect of a kick applied at the PFF kickers they must be modelled together with the quadrupoles within which they are installed but this can not be directly defined in MADX. Instead the quadrupole definition in MADX has been split in to quarters, with zero length kicker elemnets inserted between each quadrupole quarter [REF]. Each kicker element provides a quarter of the angular deflection of the complete kicker. This allows the focusing effect of the quadrupole on the applied kick to be approximated.

All the constraints for each section (taken from Section 2.4) are implemented in the matching scripts. The largest weight is given to achieving a non-zero R_{52} value between the kickers, as without this the optics would be of no use for the PFF system. The next strongest constraints are placed on the dispersion, both ensuring the dispersion is closed and keeping the dispersion well below 2 m around the second kicker. The remaining constraints, for example on the orbit closure, R_{56} and maximum twiss functions, are also included but with lower weights as a useful PFF demonstration could still be achieved if they are not precisely met.

As there are relatively few variables (quadrupoles) compared to the number of constraints it is not straightforward to match optics that meet most of the constraints. The final derivation of the optics in the following section required many matching steps, with the weights altered between each step to drive the optics closer to the desired solution. The matched quadrupole currents and twiss parameters for each matching section are used as the initial conditions for the following matching iteration. Ultimately a compromise had to be accepted that met most, but not all, of the constraints as described below.

2.6.2 PFF Optics

The PFF optics for TL2 that has been used for all the tests presented in this thesis is presented here. Tables 2.3 lists the matched quadrupole currents to apply and Table 2.4 summarises the optics constraints and their matched values. Figures 2.11 to 2.16 plot various optical parameters of interest along the line from CC.QFH0210 to CC.QDL0920. In these plots the position of the two chicanes (vertical and horizontal) are marked by dotted red lines, and the position of the two PFF kickers are marked by dotted green lines. In addition, underneath each figure the position of all quadrupoles, dipoles and the PFF kickers along the line are indicated. Blue convex lenses indicate focusing quadrupoles and blue concave lenses defocusing quadrupoles. Red rectangles mark the position of dipoles, and green squares the position of the two PFF kickers.

Figures 2.11 and 2.12 shows the twiss alpha and beta functions along TL2. The largest beta and alpha functions in both the horizontal and vertical planes are found in the region around the beginning of the vertical chicane, as a result of the long drift space between CC.QFH0230 and CC.QFL0270 (2 m to 9 m on the horizontal axis) preceding it. In this area the maximum beta value is 72.1 m, and the maximum alpha is 75.9. The alpha value is above the desired target of keeping alpha below a magnitude of 30 everywhere along the line. This constraint is primarily placed to minimise the strength of focusing in the line in order to reduce the effect of any remaining optics errors in the model. The vertical chicane therefore becomes a sensitive area for achieving good beam transport through TL2. However, in the rest of the line the beta (up to 32 m) and alpha functions (up to 18) are much smaller than their constrained maximums.

Figure 2.13 shows the horizontal and vertical dispersion along the line. The dispersions and their derivatives are closed at the 10^{-5} level in the vertical chicane, and the 10^{-6} level in the horizontal chicane. This leaves no residual dispersion in the straight sections, as desired. In the vertical chicane the dispersion reaches a peak magnitude of 0.11 m, whilst in the horizontal chicane the peak is 1.16 m. Both are much smaller than the constrained maximums. In addition, the dispersion in the region of the second kicker (K2), which is an aperture restriction as described previously, is 0.98 m and much smaller than the constrained maximum of 2 m.

The key constraint that is not met in the matched PFF optics is R_{56} , where a non-zero value had to be accepted in order to meet the remaining constraints. As seen in Figure 2.14 the R_{56} transfer matrix coefficient across the horizontal chicane is -0.18 m. This leads to additional energy dependent phase jitter following the chicane that is not present prior to the chicane, which has severe consequences for the performance of the PFF system. The extent and mitigation of this effect is the focus of Chapter 4.

The last two figures, 2.15 and 2.16, show the effect of kicking the beam with the PFF kickers. A deflection of +1 mrad is applied from the first kicker, and -1 mrad from the second kicker. This leads to a peak horizontal orbit offset of 3.5 mm inside the horizontal chicane. The orbit is closed at the 10^{-7} level following the chicane, so that the beam's trajectory following the PFF system is independent of the applied kick. The values from the approximated orbit closure expressions in Equations 2.25 and 2.26 are -1.03 and 1.04

respectively (with $R_{11} = 1.2$, $R_{12} = -0.9$, $R_{21} = -0.7$, $R_{22} = 1.3$ and $m = 0.4$). The 3–4% difference from the expected values of -1 and +1 is explained by the 4% stronger strength of CC.IQDH0790 compared to CC.IQDH0490, which was not taken in to account in the derivation of the simplified equations.

Finally, the R_{52} value between the two kickers is 0.74 m. This defines the phase shift resulting from kicking the beam in the chicane, which is the key figure of merit for the PFF system. As shown in Figure 2.16 a kick of 1 mrad provides a phase shift of -10.6° in this optics. This is converted in to the actual range of the PFF system taking in to account the specifications of the kicker amplifiers in Chapter 5. Verifications of the performance of the optics are presented in Chapters 4 and 5.

Quadrupole	Current
CC.IQFH0210	47.00 A
CC.IQDH0220	48.83 A
CC.IQFH0230	31.75 A
CC.IQFL0270	14.44 A
CC.IQDL0280	24.96 A
CC.IQDL0330	24.57 A
CC.IQFH0350	40.07 A
CC.IQDL0370	24.57 A
CC.IQDL0430	0.16 A
CC.IQFL0450	0.17 A
CC.IQDH0490	-6.05 A
CC.IQFL0530	21.94 A
CC.IQDL0550	27.57 A
CC.IQFL0570	30.19 A
CC.IQFL0620	-1.16 A
CC.IQDL0650	-10.46 A
CC.IQFL0680	-5.96 A
CC.IQFL0730	-7.77 A
CC.IQDL0750	-14.81 A
CC.IQDH0790	6.32 A
CC.IQDD0820	63.49 A
CC.IQFD0840	76.53 A
CC.IQFL0910	-11.97 A
CC.IQDL0920	-12.84 A

Table 2.3: Quadrupole power supply currents to set in the new PFF TL2 optics for a beam energy of 135 MeV.

Parameter	Constraint	Value
β_x	Max < 100 m	55.3 m at CC.QFL0270
β_y	Max < 100 m	72.1 m at CC.QDL0330
$ \alpha_x $	Max < 30	43.3 at CC.QFL0270
$ \alpha_y $	Max < 30	75.9 at CC.QFL0330
$ D_x $	Max < 2.5 m	1.16 m at CC.QDL0650
$ D_y $	Max < 1 m	0.11 m at CC.QDL0330
$ D_x $	$\ll 2$ m at K2	0.98 m
D_x	0 m at CC.BHG0800	9×10^{-7} m
D'_x	0 at CC.BHG0800	-6×10^{-6}
D_y	0 m at CC.BVB0400	4×10^{-5} m
D'_y	0 at CC.BVB0400	-6×10^{-5}
$ R_{52} $	$\gg 0$ m from K1 to K2	0.74 m
x	0 m after K2	3×10^{-7} m (1 mrad kick)
x'	0 after K2	-5×10^{-7} (1 mrad kick)
R_{56}	0 m from CC.BHG0500 to CC.BHG0800	-0.18 m

Table 2.4: Summary of constraints and their matched values in the new TL2 PFF optics.

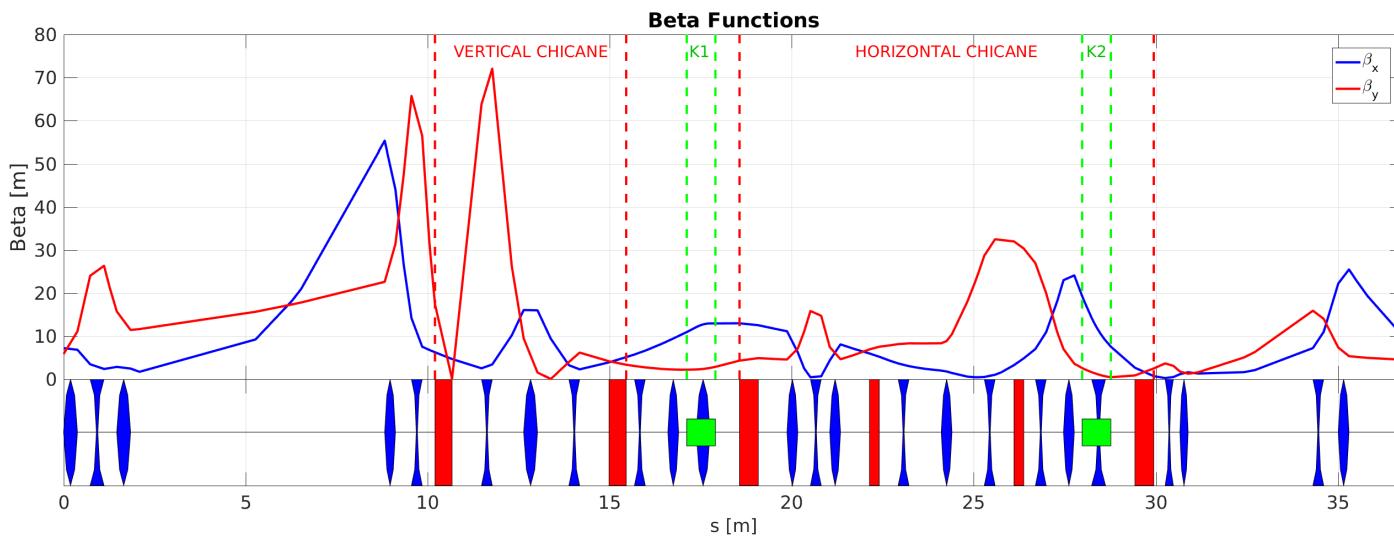


Figure 2.11: Horizontal and vertical beta functions in the new TL2 optics.

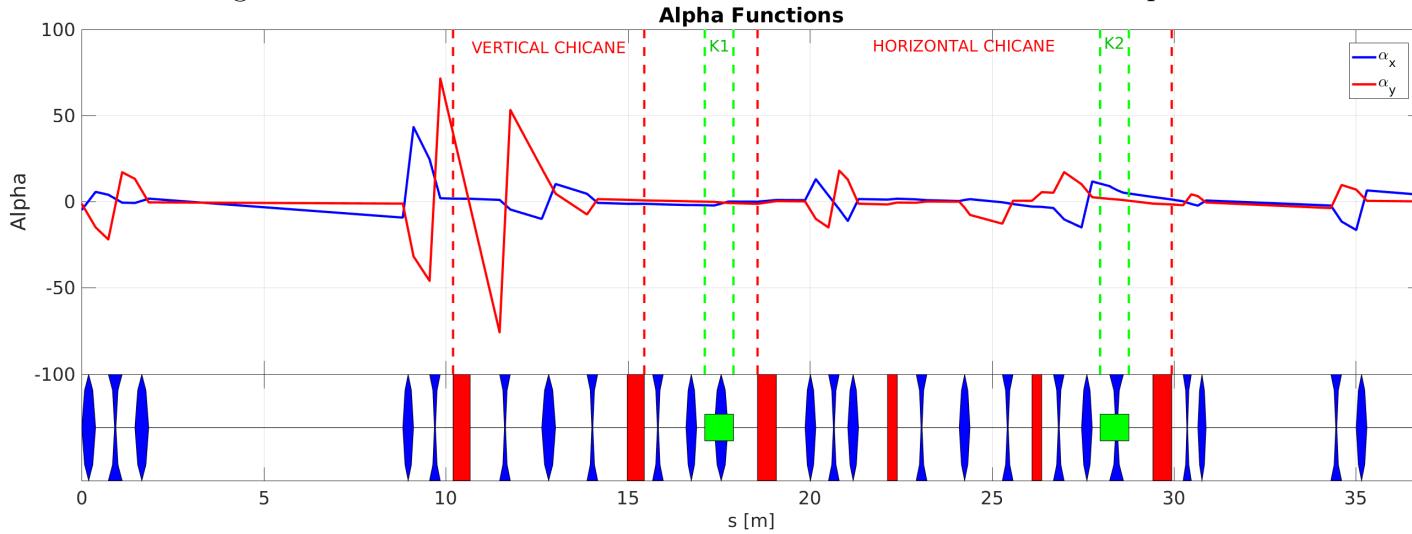


Figure 2.12: Horizontal and vertical alpha functions in the new TL2 optics.

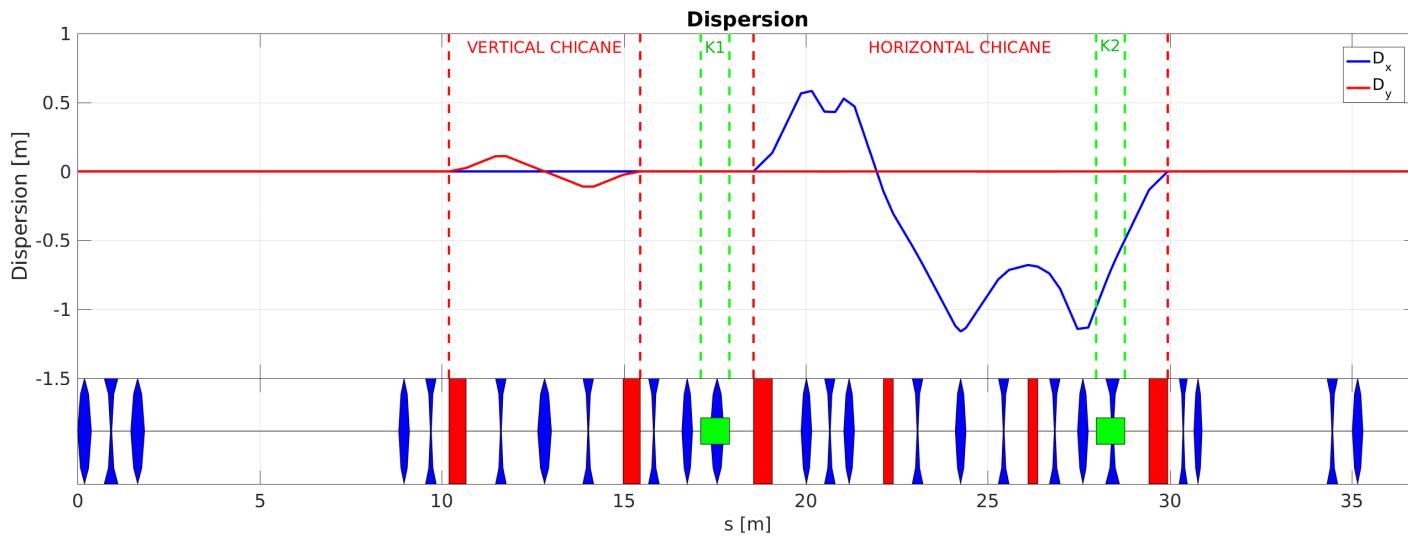


Figure 2.13: Horizontal and vertical dispersion in the new TL2 optics.

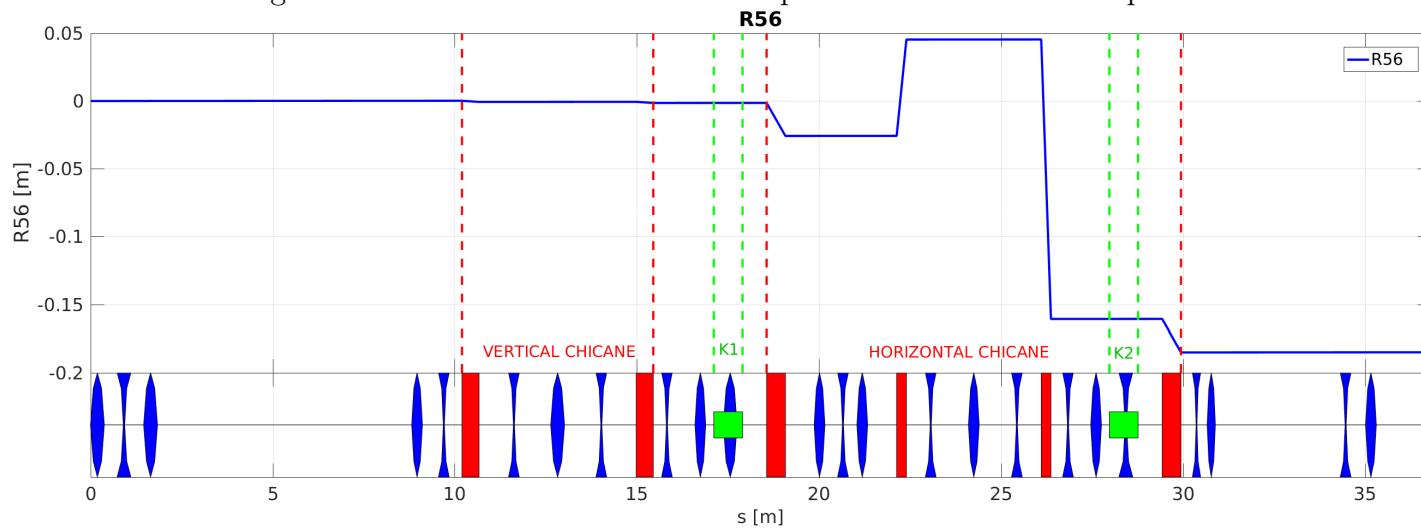


Figure 2.14: R_{56} in the new TL2 optics.

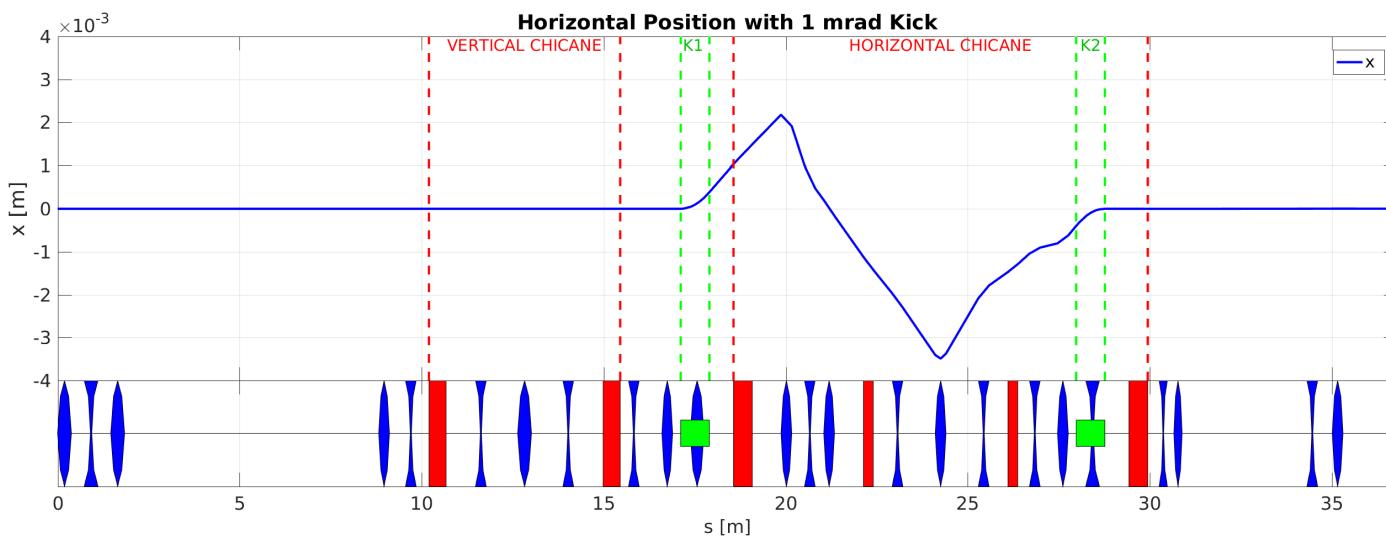


Figure 2.15: Horizontal orbit in the TL2 chicane with a 1 mrad kick applied.




Figure 2.16: Phase in the TL2 chicane with a 1 mrad kick applied.

Chapter 3

Phase Monitor Performance

In order to successfully correct the phase with the PFF system it is clearly necessary to be able to accurately measure the phase, whilst meeting the low latency (380 ns for the complete system) and high bandwidth (30 MHz) requirements of the correction. Purpose-built phase monitors for the PFF system have been designed and constructed by INFN, Italy [REF]. Three of these monitors are currently installed at CTF3 — two in the CT line at the end of the linac and one after the TL2 chicane in the TBL line. The approximate positions of the monitors are shown on the CTF3 layout in Figure [REF]. The two “upstream” phase monitors in the CT line will be referred to as Mon 1 and Mon 2 (Mon 1 being before Mon 2 in the beam line) in this chapter. The “downstream” phase monitor in TBL will be referred to as Mon 3. Mon 1 is normally used as the PFF correction input, with the neighbouring Mon 2 used for performance cross-checks. Mon 3 is then used to measure the effect of the PFF correction. In the rest of the thesis the phase measurements are generally simply referred to as being from either one of the upstream phase monitors, or the downstream monitor.

The chapter begins with an overview of the design and installation of the phase monitors themselves as well as the associated electronics. The remainder of the chapter then focuses on operational procedures, measurements of the performance of the phase monitors and the changes that have been made in order to improve it. The headline result for the chapter is the achieved resolution of below 0.14° in Section 3.9, which is derived to be necessary to be able to achieve the targeted 0.2° corrected phase jitter in Section 3.3. Several effects have also been identified that can still degrade the accuracy of the phase measurement, such as the position dependence seen in Section 3.13. These are the first areas that should be looked at if an improvement in phase monitor performance is necessary for future PFF tests.

3.1 Phase Monitor Design

The phase monitors are cylindrical cavities with a length of approximately 19 cm and an internal diameter of 23 mm, as shown in Figure 3.1. When a charged beam traverses a cavity the interaction of the beam with the cavity walls creates electromagnetic fields inside the

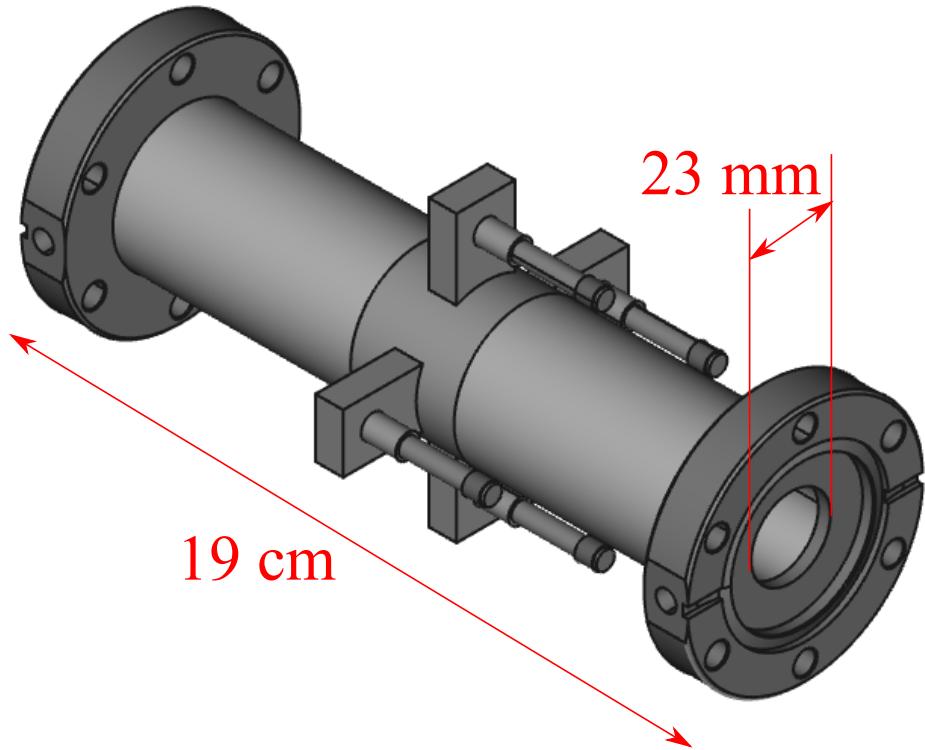


Figure 3.1: Technical drawing of the phase monitor cavity design [REF]. The length and internal diameter of the monitor is shown., Around the centre of the monitor are the four evenly spaced RF feedthroughs.

cavity. The amplitude of the induced fields depends both on the bunch charge and the bunch length [REF]. Small ridges (called notch filters) in the cavity create a volume resonating at 12 GHz (the CLIC combined bunch frequency) that contains the beam induced fields and reflects any stray 12 GHz fields not associated with the beam outside the cavity. This is shown in Figure 3.2. Four rectangular slots, arranged in horizontal and vertical pairs, around the mid-point of the cavity are then used to extract the beam induced resonant fields. The fields leaving the cavity are transported in short rectangular waveguides before a transition to a 50Ω coaxial cable via an RF feedthrough [REF]. The output of the phase monitor cavities is therefore four 12 GHz signals whose time structure depends on the arrival time, or phase, of the drive beam bunches.

The solutions to Maxwell's equations in cavities such as this give a discrete set of transverse electric (TE) and transverse magnetic (TM) modes dependent on the geometry of the cavity [REF]. TE modes are characterised by having only transverse electric field components, and no longitudinal electric field component, whereas TM modes have only transverse magnetic field components and no longitudinal magnetic field component. Each TM and TE mode has an associated cutoff frequency dependent on the number of half-period variations, n and m , in the field horizontally and vertically across the cavity respectively. The amplitude of the 12 GHz signals output from the cavity will contain components of each TM and TE mode with a cutoff frequency below 12 GHz. For cylindrical cavities the cutoff frequency of

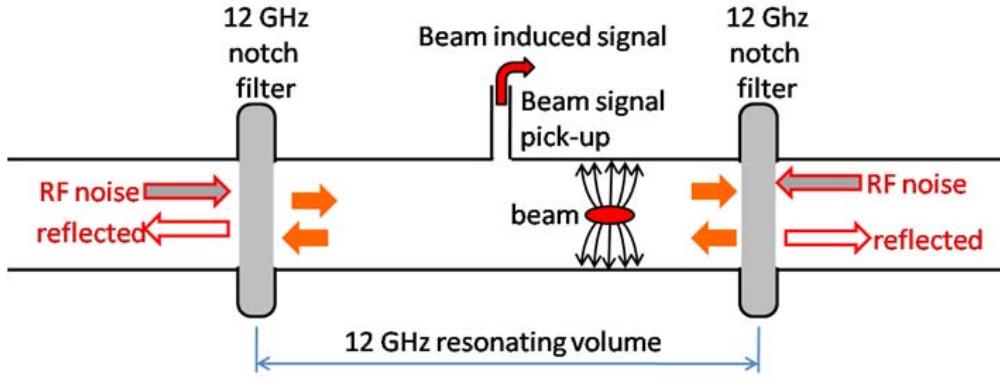


Figure 3.2: Schematic of phase monitor design [REF].

the TM is defined as [REF]:

$$f_{nm} = \frac{c}{2\pi} \frac{p_{nm}}{a} \quad (3.1)$$

And for the TE modes as:

$$f_{nm} = \frac{c}{2\pi} \frac{p'_{nm}}{a} \quad (3.2)$$

Where a is the radius of the cavity, p_{nm} is the m^{th} zero of the Bessel function $J_n(x)$, and p'_{nm} is the m^{th} zero of the derivative of the Bessel function $J'_n(x)$ [REF]. The beam pipe around the location of the phase monitors at CTF3 is usually 4 cm in diameter, and a cavity of this size would support six separate TM or TE modes with a cutoff frequency below 12 GHz [REF]. It is for this reason that the phase monitor diameter was reduced to 23 mm, where only two modes are present: TM01 ($n = 0$, $m = 1$, $p_{01} = 2.4$) and TE11 ($n = 1$, $m = 1$, $p'_{11} = 1.8$).

TM01 is referred to as the monopole mode and TE11 as the dipole mode. The induced field distribution resulting from a bunch entering the cavity for both modes is shown in Figure 3.3. The precise expressions for each field distribution can be found in [REF]. When the beam has been correctly setup it should enter the phase monitor cavity close to its centre. For small (horizontal or vertical) offsets between the beam position and the cavity centre there is no dependence of the monopole mode amplitude on the incoming beam position. However, the amplitude of the dipole mode does depend on the beam position, even for small offsets from the cavity centre. This means the amplitude of any of the four individual RF outputs from the monitor will have a position dependence. This property is used in cavity beam position monitors (BPMs) [REF], but is undesirable for a phase monitor where the measurement should be position independent. For a 1 mm beam position offset the dipole mode is expected to have around 10% the amplitude of the monopole mode [REF].

To remove the unwanted dipole mode the horizontal and vertical pairs of outputs from the cavities can be combined, as the mode is symmetric and has equal magnitude but opposite sign on each side of the cavity. As the CTF3 beam is generally more stable in the vertical plane (as the majority of bends in the beam line are horizontal the energy related orbit jitter resulting from dispersion is larger in the horizontal plane) the vertical pair of RF feedthroughs from the monitors are normally used. The two outputs are summed using 180 degree hybrids [REF] installed next to the phase monitor cavities in the machine hall.

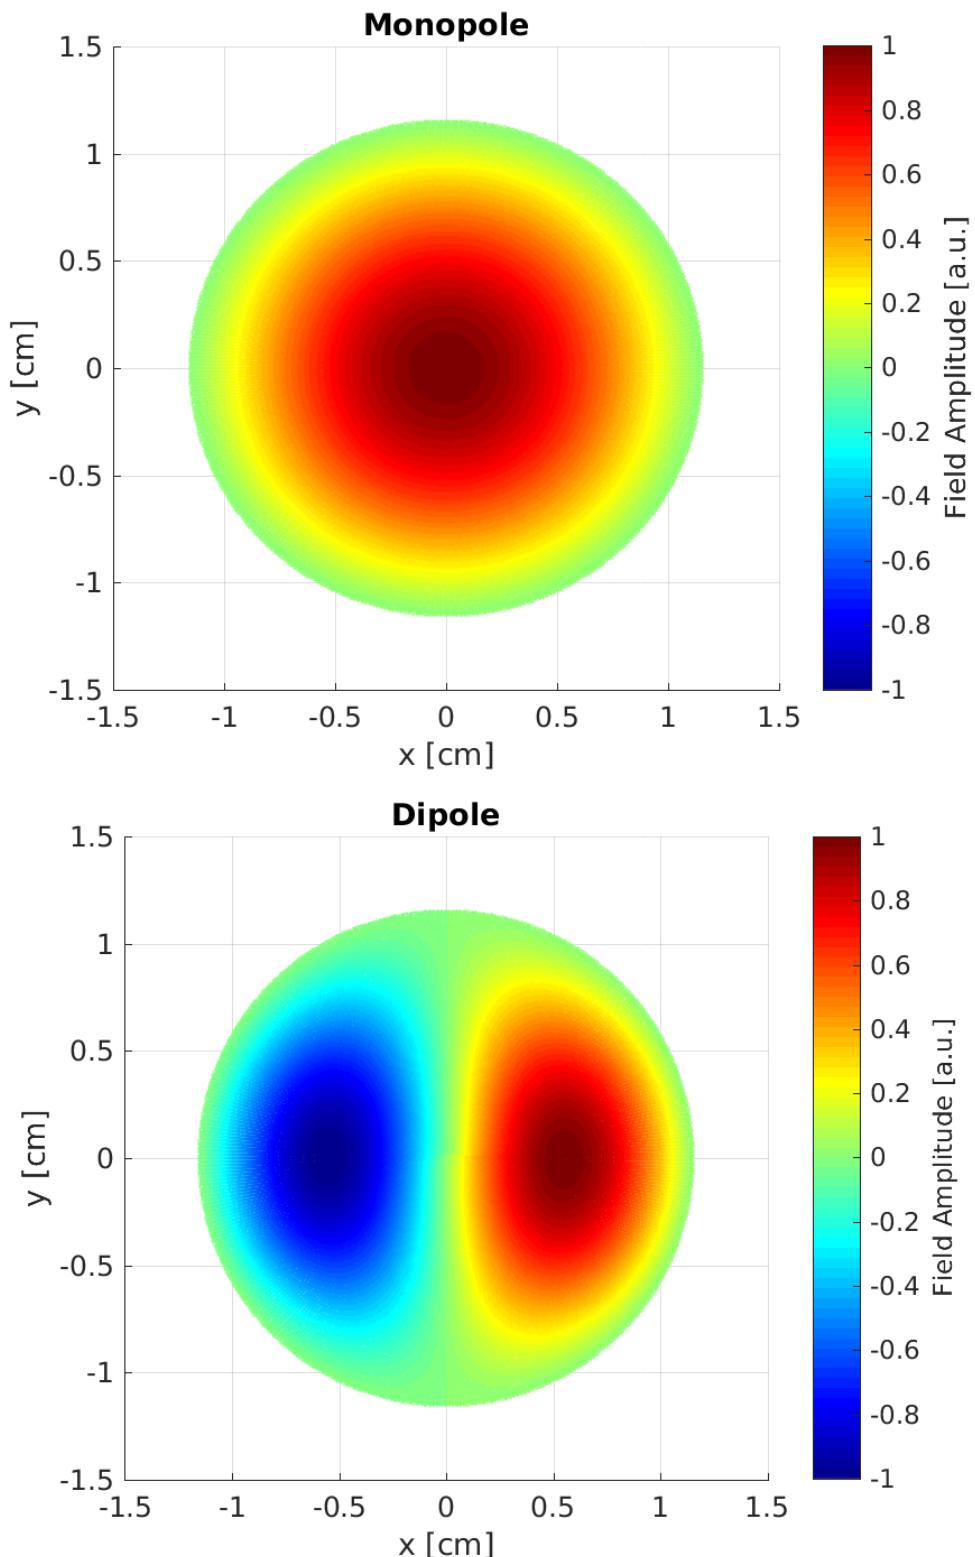


Figure 3.3: Field distribution of Monopole (TM01) and dipole (TE11) modes in the phase monitor cavities. The dipole mode has both horizontal and vertical components - only the horizontal component is shown.

Monitor	Power
Mon 1	27.6 dBm
Mon 2	29.8 dBm
Mon 3	24.5 dBm

Table 3.1: Power of the phase monitor signals (hybrid sum outputs) measured in the klystron gallery prior to being processed by the phase monitor electronics.

The horizontal pair is also instrumented in the same way but the outputs are typically not used. For an ideal (perfectly symmetric) cavity and hybrid this would create an output with only the position independent monopole mode present. In reality small misalignments, for example in the waveguides and RF feedthroughs, cause slight asymmetries in the cavity and signal combination which leaves a small residual dipole component. The signal combination in the hybrids is expected to reduce the dipole amplitude by a further 20 dB at a 1 mm offset, giving a final amplitude around 1% of the monopole mode [REF]. The remaining position dependence of the phase measurement is discussed in Section 3.13. How the output from the hybrids is processed to calculate the phase is discussed in the next section.

Figure 3.4 shows the installation of the upstream phase monitors in the CT line. The installation allows up to three phase monitors to be installed neighbouring each other, with the current two monitors installed in the first and third slots leaving approximately a 20 cm gap between the two. The connections between the four RF feedthroughs on the monitors and the hybrids fixed underneath the monitors can be seen in the figure, with the hybrids combining the horizontal pairs visible and two further hybrids placed on the other side of the beam line for the vertical pairs. The outputs from the hybrids are routed up to the klystron gallery on the floor above the accelerator hall, where they are processed and used for the PFF inputs. The overall phase monitor setup including the cables adds around 50 ns to the overall PFF latency [REF]. For reference the power level of the three phase monitor signals as measured once they reach the klystron gallery are quoted in Table 3.1. These are useful to interpret the results of Section 3.6, for example.

3.2 Phase Monitor Electronics

The output of the phase monitor cavities (or more precisely the combined output from the hybrids) is a sinusoidal signal with a frequency of $\omega_{RF} = 11.994$ GHz (the precise CLIC combined drive beam bunch frequency is 11.994 GHz rather than 12 GHz). This can be expressed as follows:

$$RF(t) = A_{RF}(t) \cos[\omega_{RF}t + \phi(t)] \quad (3.3)$$

Where $A_{RF}(t)$ is the time dependent amplitude (voltage) of the phase monitor signal and $\phi(t)$ is the time dependent beam phase of interest for the PFF system. The beam based signals from the phase monitors are referred to as the RF signals.

This high frequency signal cannot be digitised directly so in order to extract the phase dependence it is combined with a similar reference signal to create a low, or intermediate,

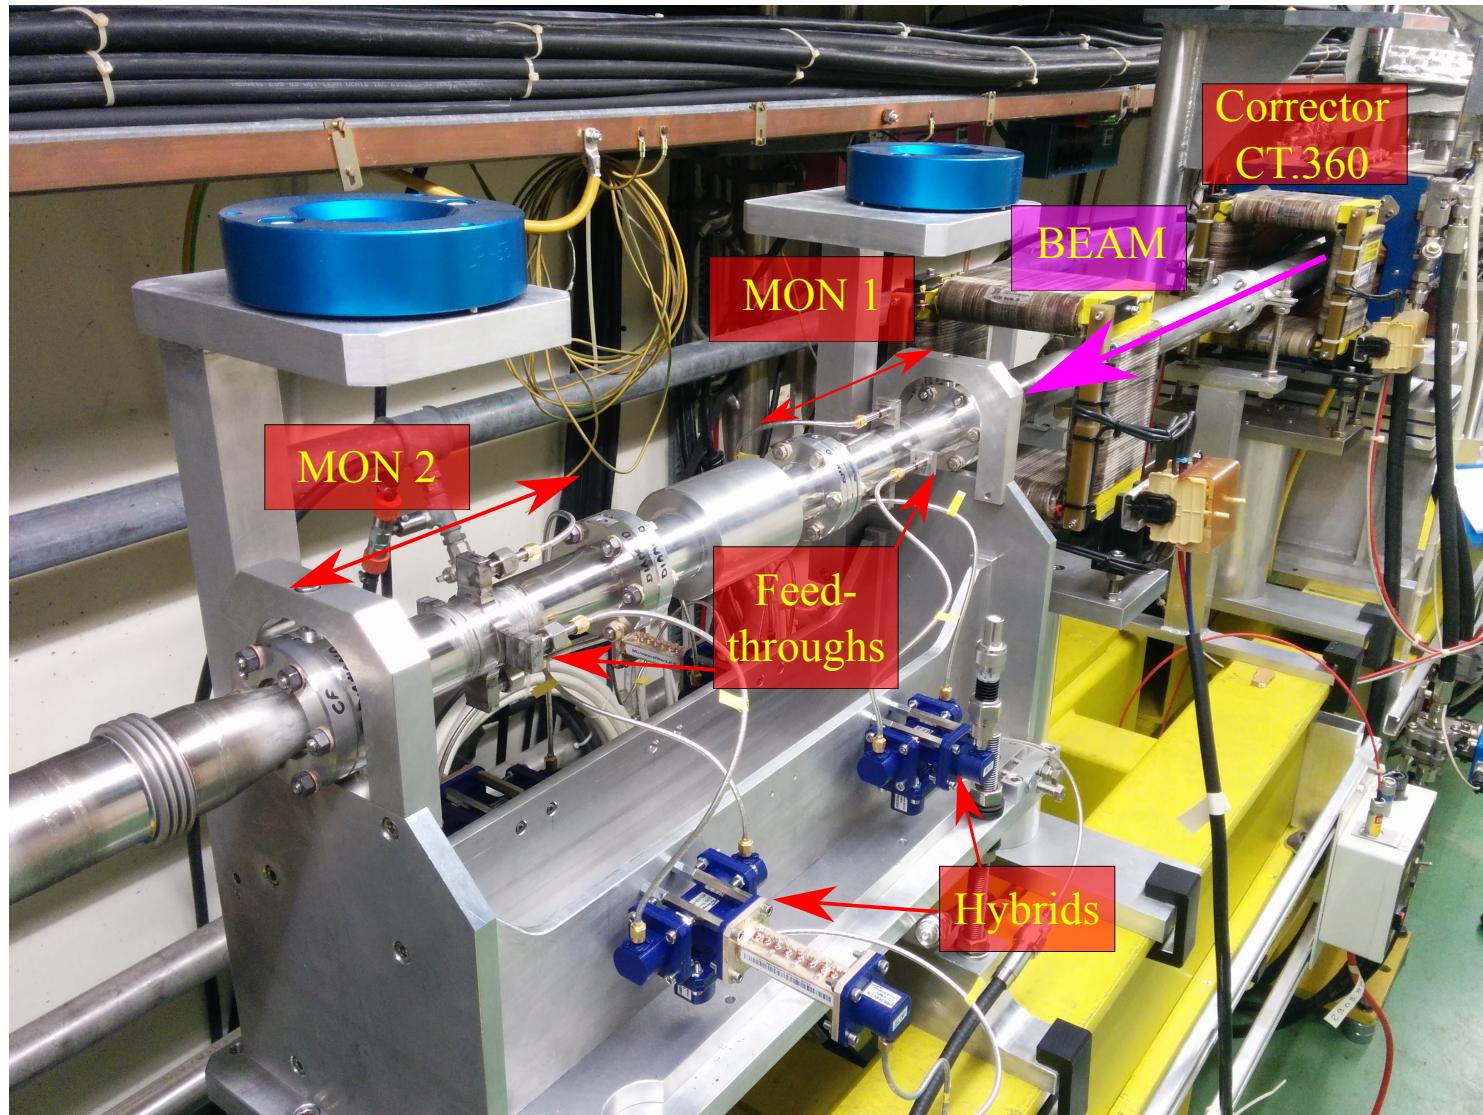


Figure 3.4: Annotated picture of the two upstream phase monitors installed in the machine.

frequency signal that can be digitised and preserves the beam phase information. The devices that “mix” the beam signal with the reference signal are called mixers. The reference signal is referred to as the local oscillator or LO signal. The LO is usually a continuous sinusoidal signal with constant amplitude A_{LO} and frequency ω_{LO} :

$$LO(t) = A_{LO} \cos[\omega_{LO}t] \quad (3.4)$$

The mixer multiplies the RF and LO signals, to create an output with the following dependence on each signal:

$$\text{Mixer}(t) = RF(t) \times LO(t) \quad (3.5)$$

$$\text{Mixer}(t) = A_{RF}(t)A_{LO} \cos[\omega_{RF}t + \phi(t)] \cos[\omega_{LO}t] \quad (3.6)$$

Using trigonometric identities this can be expressed in terms of a high frequency component, with a frequency of $(\omega_{LO} + \omega_{RF})$, and a low frequency component, with a frequency of $(\omega_{LO} - \omega_{RF})$:

$$\text{Mixer}(t) = \frac{A_{RF}(t)A_{LO}}{2} \{ \cos[(\omega_{LO} + \omega_{RF})t + \phi(t)] + \cos[(\omega_{LO} - \omega_{RF})t + \phi(t)] \} \quad (3.7)$$

The high frequency component can then easily be removed using a low pass filter, so that the mixer output becomes:

$$\text{Mixer}(t) = \frac{A_{RF}(t)A_{LO}}{2} \cos[(\omega_{LO} - \omega_{RF})t + \phi(t)] \quad (3.8)$$

Finally, by using a reference LO signal which has the same frequency as the RF signal, which is the case for the phase monitor electronics presented here, this simplifies to:

$$\text{Mixer}(t) = \frac{A_{RF}(t)A_{LO}}{2} \cos[\phi(t)] \quad (3.9)$$

The resulting mixer output is therefore a low frequency signal that depends only on the amplitude of the RF signal and the phase (as well as the constant LO amplitude).

To use the mixer output to calculate the phase the voltage of the RF signal, A_{RF} , must be known. This can be determined by splitting the RF signal between the mixer and a diode detector, whose output is dependent on the power of the signal (the square of the voltage):

$$\text{Diode}(t) = A_{RF}(t)^2 \quad (3.10)$$

The phase can then be reconstructed from the mixer and diode outputs as follows:

$$\frac{\text{Mixer}(t)}{\sqrt{\text{Diode}(t)}} = A \cos[\phi(t)] \quad (3.11)$$

$$\phi(t) = \arccos \left[\frac{\text{Mixer}(t)}{A \sqrt{\text{Diode}(t)}} \right] \quad (3.12)$$

Where A is a calibration constant dependent on the relative amplitudes of the mixer and diode outputs. The electronics for the three PFF phase monitors use mixers and diodes

LO	Power
LO 1	22.6 dBm
LO 2	23.6 dBm
LO 3	25.5 dBm

Table 3.2: Power of the LO inputs to the phase monitor electronics.

as described above to produce signals that are dependent on the phase and amplitude of the output from the phase monitor cavities. Figure 3.5 shows the front panel of one of the sets of the electronics, showing the RF and LO inputs and mixer and diode outputs. The mixer and diode outputs are digitised, with the phase calculation above performed after the digitisation rather than by the phase monitor electronics themselves.

In order to achieve resolutions below 0.2° , as necessary for the PFF system (Section 3.3), the electronics do not work exactly as described above but instead use multiple mixers and diodes. Non-linearities in the response of mixers and diodes are typically worse at higher input powers [REF], but higher input powers are desirable to improve the signal to noise ratio of the output. With a single mixer and diode a compromise must therefore be made between the accuracy (linearity) and resolution (noise) of the device. To be able to operate the mixers at low power whilst maintaining a good signal to noise ratio eight separate mixers and diodes are used in each set of electronics [REF]. The RF and LO inputs are split between the eight mixers and diodes, and then the eight outputs of each type are added together to create the final mixer and diode outputs from the electronics. Figure 3.7 shows a simplified example of this with two mixers and diodes.

Figure 3.7 also shows all the elements present in the generation of the LO reference signal for the electronics. The LO signal for all three sets of electronics is generated from a common 3 GHz source that is locked to the timing of the CTF3 drive beam (which has a 3 GHz acceleration frequency). The 3 GHz signal is split in to three to provide the LO input for each set of electronics. After the split each LO chain includes a phase shifter, bandpass filter, frequency multiplier and an amplifier. The phase shifters are required for the calibration process as discussed in Section 3.7. Initially these were 3 GHz digital phase shifters of type [REF], but these were later replaced with passive mechanical phase shifters for the reasons seen in Section 3.8. The mechanical shifters are of type [REF], and can be changed in units of 0.02 degrees at 4 GHz. After the shifters, a bandpass filter [REF] removes any unwanted out of frequency noise in the LO signal, the frequency multiplier [REF] increases the frequency of the LO to 11.994 GHz and finally the amplifier [REF] is used to boost the power of the LO signals to the level seen in Table 3.2.

Figure 3.6 finally shows the complete phase monitor electronics installation in the racks in the klystron gallery.



Figure 3.5: Front panel of the phase monitor electronics boxes containing the mixers and diodes. The connectors labelled LO and RF take the signals from the LO source and the phase monitors respectively. The connectors labelled A^2 and $A \sin(\phi)$ are the diode and mixer outputs respectively.

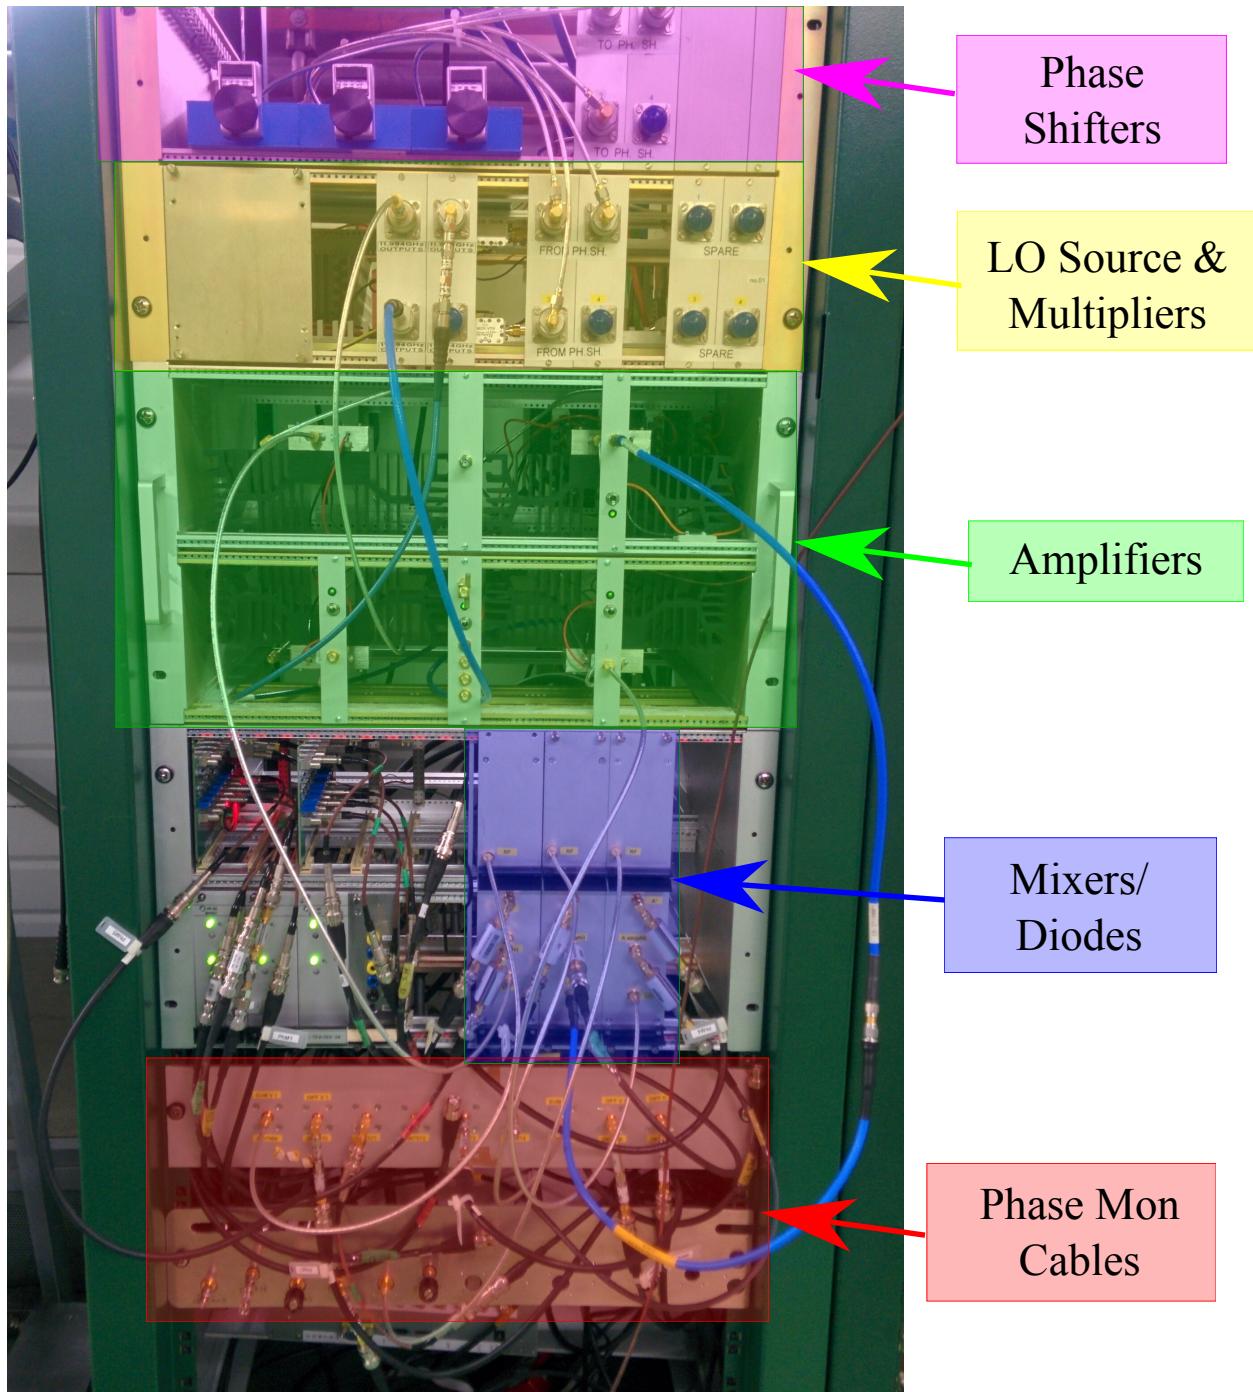


Figure 3.6: Annotated picture of the phase monitor electronics rack in the klystron gallery..

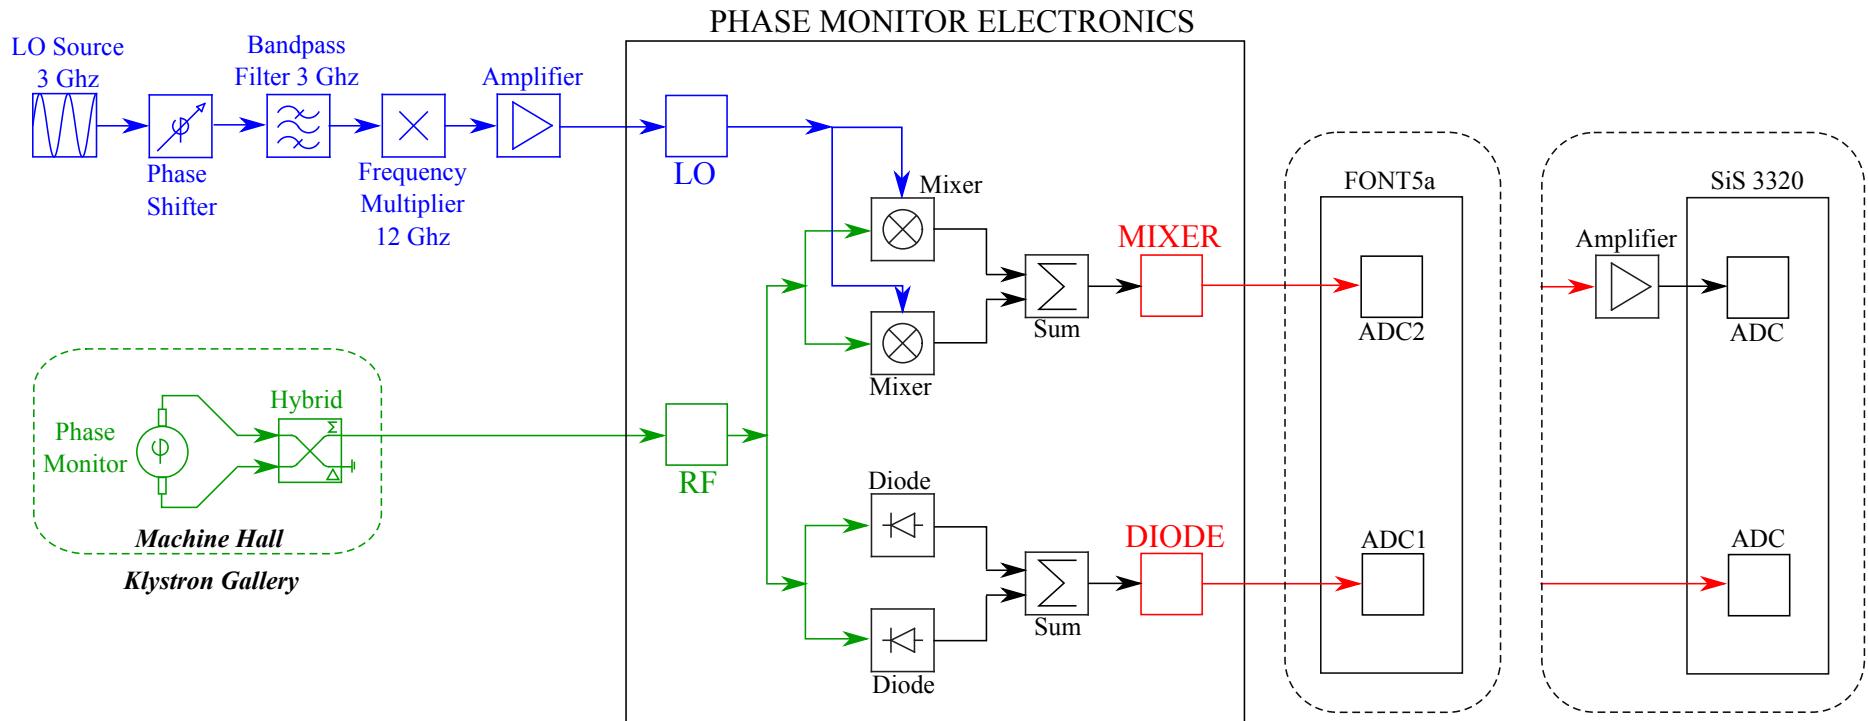


Figure 3.7: Simplified schematic of the phase monitor electronics setup. The components involved in the generation of the LO are shown in blue, the phase monitor beam signal in green and the outputs in red. For the purposes of the diagram only two individual mixers and diodes are shown. In the actual design the input signals are split between eight mixers and diodes, and then combined in the same way to create the two outputs. The outputs are digitised on either the FONT5a board or the SiS digitisers.

3.3 Resolution Definition

The performance of the PFF system clearly depends on the accuracy to which the phase can be measured. Many of the measurements in this chapter are therefore focused on the phase monitor resolution, or more precisely on the resolution of the combined phase monitor and electronics setup. The resolution is defined as the jitter between the measured phase and the true beam phase. This can be calculated by comparing the difference between the measured phase of two monitors. This is why two phase monitors, Mon 1 and Mon 2 are installed neighbouring each other in the upstream system in the CT line. The beam phase should be identical in these two monitors thus their measurements can always be compared to derive the resolution.

The precise derivation of the resolution dependent on the measurement of two monitors is as follows. First, the measured phase, $\phi_x(t)$ and $\phi_y(t)$, in two monitors at time t can be defined as:

$$\phi_x(t) = \phi_b(t) + n_x(t) \quad (3.13)$$

$$\phi_y(t) = \phi_b(t) + n_y(t) \quad (3.14)$$

Where $\phi_b(t)$ is the true beam phase and $n_x(t)$ and $n_y(t)$ is the noise on the measurement at that time. The time dependence will not be written explicitly from this point. These equations assume the beam phase is identical in each monitor, as should be the case for Mon 1 and Mon 2. The variance of each phase monitor measurement can then be derived from the equations above by adding the variance of the beam phase and the noise in quadrature:

$$\sigma_x^2 = \sigma_b^2 + \sigma_{nx}^2 \quad (3.15)$$

$$\sigma_y^2 = \sigma_b^2 + \sigma_{ny}^2 \quad (3.16)$$

Where σ_x and σ_y are the phase jitters measured by each phase monitor, σ_b is the true beam phase jitter and σ_{nx} and σ_{ny} are the phase monitor resolutions. Assuming each phase monitor has the same resolution, σ_n , this can be simplified to $\sigma_x^2 = \sigma_y^2 = \sigma_b^2 + \sigma_n^2$.

The quantity of interest for calculating the phase monitor resolution is the jitter in the difference between the two measured phases, σ_{x-y} . The variance of the difference between two correlated variables is defined as:

$$\sigma_{x-y}^2 = \sigma_x^2 + \sigma_y^2 - 2\sigma_x\sigma_y\rho_{xy} \quad (3.17)$$

$$(3.18)$$

Where ρ_{xy} is the correlation between the phase measurement of x and y . Substituting in the previously derived expressions for σ_x and σ_y this becomes:

$$\sigma_{x-y}^2 = 2(\sigma_b^2 + \sigma_n^2)(1 - \rho_{xy}) \quad (3.19)$$

The correlation coefficient ρ_{xy} depends on the covariance between x and y , $\text{cov}[x, y]$, as follows:

$$\rho_{xy} = \frac{\text{cov}[x, y]}{\sigma_x\sigma_y} = \frac{\text{cov}[x, y]}{\sigma_b^2 + \sigma_n^2} \quad (3.20)$$

$$(3.21)$$

Where the covariance is defined as:

$$\text{cov}[x, y] = \frac{1}{N} \sum_{i=1}^N \phi_{xi} \phi_{yi} \quad (3.22)$$

$$(3.23)$$

Substituting in the expressions for ϕ_x and ϕ_y above and separating the terms in the sum then gives the following expression for the covariance of x and y :

$$\begin{aligned} \text{cov}[x, y] &= \frac{1}{N} \sum_{i=1}^N (\phi_{bi} + n_{xi})(\phi_{bi} + n_{yi}) \\ \text{cov}[x, y] &= \frac{1}{N} \sum_{i=1}^N \phi_{bi}^2 + \frac{1}{N} \sum_{i=1}^N \phi_{bi} n_{xi} + \frac{1}{N} \sum_{i=1}^N \phi_{bi} n_{yi} + \frac{1}{N} \sum_{i=1}^N n_{xi} n_{yi} \end{aligned} \quad (3.24)$$

The first term is the definition of the variance of the beam phase, σ_b^2 . The remaining terms are the covariance between the beam phase and the monitor noises, and the covariance between the two monitor noises. Assuming the noise is uncorrelated all these terms are zero. The remaining equation for the covariance between x and y is therefore simply: $\text{cov}[x, y] = \sigma_b^2$. Finally, the correlation between the phase measurement of x and y becomes:

$$\rho_{xy} = \frac{\sigma_b^2}{\sigma_b^2 + \sigma_n^2} \quad (3.25)$$

Substituting this expression for the correlation in to the derived equation for the variance between the two phase measurements gives the following simple dependence on the phase monitor resolution:

$$\begin{aligned} \sigma_{x-y}^2 &= 2(\sigma_b^2 + \sigma_n^2) \left(1 - \frac{\sigma_b^2}{\sigma_b^2 + \sigma_n^2} \right) \\ \sigma_{x-y}^2 &= 2\sigma_n^2 \end{aligned} \quad (3.26)$$

Finally, the resolution is defined as:

$$\sigma_n = \frac{\sigma_{x-y}}{\sqrt{2}} \quad (3.27)$$

In terms of a resolution calculation these equations only apply to the two upstream phase monitors. All the resolution values quoted in this chapter use this equation and the difference between the measurement of Mon 1 and Mon 2.

However, as the act of the PFF system can also be thought of as subtracting two phases (removing the upstream phase from the downstream phase) the same equations can be directly applied to determine the limitations that the phase monitor resolution places on the PFF performance. Equation 3.26 shows that the lowest possible corrected downstream phase jitter is a factor $\sqrt{2}$ times larger than the phase monitor resolution. In order to reduce the downstream phase jitter to the CLIC target of 0.2° the phase monitor resolution must therefore be better than 0.14° . Equation 3.25 shows that with this 0.14° resolution and a typical beam phase jitter of 0.8° (Section 4.2) the measured correlation between two phase monitor measurements would be 97%.

Digitiser	No. ADCs	Resolution	Input Range	Sampling Rate	Bandwidth
SiS 3320	8	12-bit	± 2.5 V	up to 250 MHz	100 MHz
FONT5a	9	14-bit (13-bit used)	± 0.5 V	up to 400 MHz	1.4 GHz

Table 3.3: Specifications of the ADCs on the FONT5a board and SiS digitisers.

3.4 Digitisation of Phase Monitor Signals

The mixer and diode outputs from the phase monitor electronics must be digitised on analogue to digital converters (ADCs) so the signals can be processed and used for the PFF correction and offline data analysis. Two different types of ADCs have been used to digitise the phase monitor signals — the Texas Instruments ADS5474 ADCs [REF] on the purpose-built FONT5a board used as the PFF controller and a commercially available SiS 3320 digitiser [REF]. The design and use of the FONT5a board is discussed in more detail in Section 5.1. Table 3.3 summarises the specifications of each type of ADC.

The SiS digitisers are used in addition to the FONT5a board as the PFF correction running on the FONT5a board is operated as a standalone system independent from other acquisition systems at CTF3. The PFF algorithm requires only the signals from one of the upstream phase monitors to be connected to the FONT5a board, with the convention being to use Mon 1. Mon 2 and Mon 3 are then normally connected to the SiS digitisers instead. The SiS digitisers are setup with the same trigger and sampling frequency (192 MHz) used for other signals at CTF3, and data from them can be acquired together with other devices using the standard systems in place at CTF3. This allows the Mon 2 and Mon 3 signals to be easily compared to other measurements, such as beam position signals, which has been indispensable for optimising the setup of the PFF system and in particular the phase propagation (Chater 4).

Digitising the phase monitor signals contributes additional noise to the overall phase monitor electronics setup. The purpose of this section is to ensure that the digitiser noise makes only a negligible contribution to the resolution on the phase measurement. The main parameters of interest needed to determine this are the input range and resolution of the ADCs, with the SiS ADCs being 12-bit with a range of ± 2.5 V and the FONT5a ADCs being 13-bit with an input range of ± 0.5 V. The full 5 V peak-to-peak input range of the SiS ADCs is therefore split across $2^{12} = 4096$ values, or ADC ‘counts’, with each count corresponding to roughly 1.2 mV. The equivalent 1 V peak-to-peak range and $2^{13} = 8192$ counts of the FONT5a board corresponds to a factor 10 lower interval of 0.12 mV per ADC count. The voltage jitter added by the digitisation of the phase monitor signals cannot be expected to be better than 1 count. This already indicates that the FONT5a board should give a much lower contribution to the phase resolution than the SiS digitisers.

Figure 3.8 shows the ADC noise, converted from counts in to an equivalent voltage, for both the SiS and FONT5a ADCs. As expected the noise on the FONT5a board is much lower than on the SiS digitisers. The actual ADC jitter values are 1.47 ± 0.04 counts or 0.179 ± 0.005 mV on the FONT5a board, and 1.11 ± 0.03 counts or 1.36 ± 0.03 mV on the

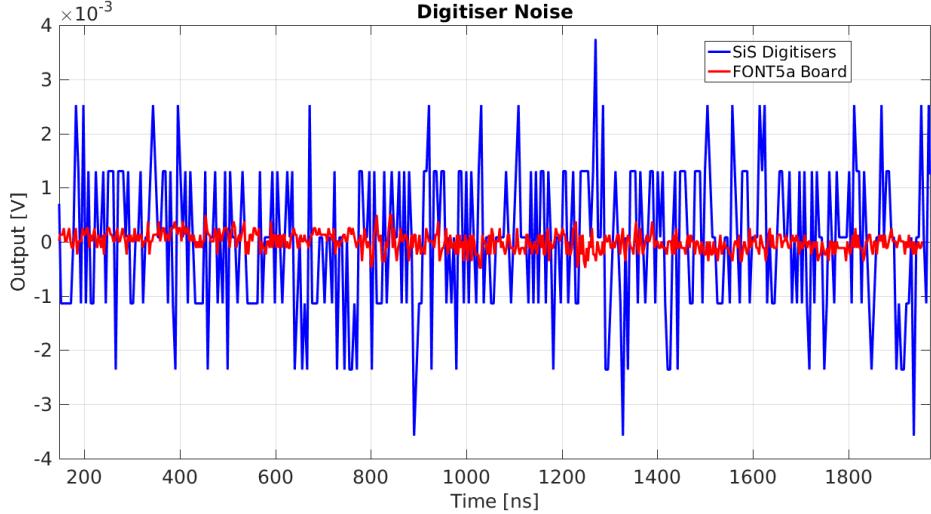


Figure 3.8: Comparison of noise on the output of the SiS and FONT5a ADCs.

SiS digitisers. These values can be converted in to an equivalent phase jitter using the phase reconstruction method described later in Section 3.6.5 and the monitor calibration constants determined in Section 3.7. For reference the peak output of the three phase monitor mixers varies between approximately 400 mV and 500 mV, which is well matched to the input range of the FONT5a ADCs. Taking the worst case scenario of Mon 3, which gives the lowest output voltage, the ADC jitter corresponds to $0.0245 \pm 0.0007^\circ$ on the FONT5a board but $0.198 \pm 0.005^\circ$ on the SiS digitisers. These values are summarised in Table 3.4.

As derived in the previous section the phase resolution must be better than 0.14° in order to achieve a corrected downstream phase jitter of 0.2° with the PFF system. The $\sim 0.02^\circ$ contribution of ADC noise on the FONT5a board is therefore insignificant compared to the resolution requirements. However, although it does not directly impact the PFF performance the $\sim 0.2^\circ$ ADC jitter on the SiS digitisers would greatly degrade the resolution of the measurements of Mon 2 and Mon 3 usually connected to the SiS digitisers and used for offline data analysis of the PFF results.

The high phase jitter contribution from the SiS digitisers originates from the roughly 500 mV maximum mixer output being much lower than the SiS ADC range of ± 2.5 V. In order to rectify this the mixer outputs are boosted by roughly a factor 2.5 in voltage using an amplifier prior to the SiS digitisers. The specifications of the amplifier used are documented in [REF]. With the amplifier in place the peak signal level sent to the SiS digitisers is around 1 V, and the equivalent phase jitter is reduced to $0.078 \pm 0.002^\circ$. This no longer prevents 0.14° resolution from being achieved on measurements using the SiS digitisers, as proven later in Section 3.9. A small further improvement in measured resolution could be achieved using a different amplifier and boosting the peak output voltage closer to 2 V.

Digitiser	Jitter [counts]	Jitter [mV]	Phase Jitter [degrees]
FONT5a	1.47 ± 0.04	0.179 ± 0.005	0.0245 ± 0.0007
SiS 3320	1.11 ± 0.03	1.36 ± 0.03	0.198 ± 0.005
SiS 3320 Amplified	1.11 ± 0.03	1.36 ± 0.03	0.078 ± 0.002

Table 3.4: ADC jitter on the FONT5a board, SiS digitisers and with the mixer outputs amplified prior to the SiS digitisers expressed in terms of ADC counts, volts and equivalent phase jitter.

3.5 Fitting Method

Due to the dependence of the mixer output on $\cos(\phi)$ as seen in Section 3.2, many of the measurements in this chapter require a sinusoidal fit of the form:

$$y = A \sin(bx + c) + d \quad (3.28)$$

The use of sine rather than cosine makes no difference to the fitted amplitude, A , and offset, d which are usually the only parameters of interest in this chapter. It is also convenient to consider a mixer output of zero to correspond to zero phase (rather than 90° as in Equation 3.11). All the fits of this type have been performed using a weighted nonlinear least squares fit implemented in MatLab fitting libraries [?]. Each data point is weighted by the inverse of its standard error squared.

Care must be taken to select suitable initial values for the four parameters in the fit in order to avoid local minima and ensure a reasonable fit. This is particularly important for a sinusoidal fit as there are many solutions with different frequencies and phase offsets that can match the data. The frequency, b , is the most critical parameter but for all the applications in this chapter this is already known, e.g. from the properties of the used phase shifters. Initial values for the three remaining parameters are estimated as follows:

$$A = \frac{\max(y) - \min(y)}{2} \quad (3.29)$$

$$d = \frac{\max(y) + \min(y)}{2} \quad (3.30)$$

$$c = \arcsin\left(\frac{y - d}{A}\right) - bx \quad (3.31)$$

The initial amplitude, A , and offset, d , of the sine curve are simply determined by comparing the minimum and maximum output. These initial estimates are therefore highly biased by any large outliers around the minimum and maximum output, but this is rarely the case for the application here and these simple estimators are sufficient. Rearranging Equation 3.28 gives the expression for c above. Due to its use of \arcsin the equation is only valid in the first and fourth quadrants, between $-\pi/2$ and $+\pi/2$ where the gradient of the sine curve is positive. The y value at each data point is compared to its neighbours to determine whether it is on the rising slope. The initial value of c is the mean value calculated across all the data points that meet this criteria.

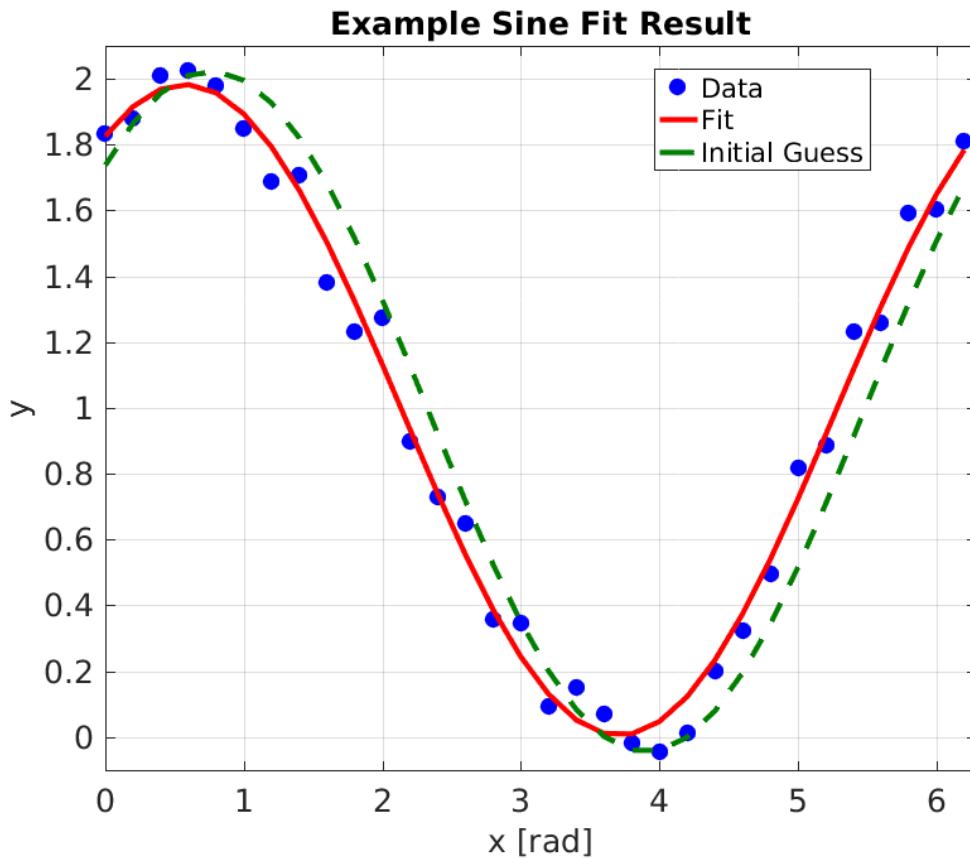


Figure 3.9: Example sine fit to generated data with added random noise.

Figure 3.9 and Table 3.5 show the results of an example fit using this approach. An initial distribution of points with $y = \sin(x + 1) + 1$ is used ($A = b = c = d = 1$), with random noise added. b is assumed to be known, as is normal. The initial estimates for A and d are within a few percent of their true value. The initial estimate for c is within 20% of the correct value. After fitting all four parameters are in agreement with the expected values within error bars.

Parameter	Value	Initial	Fit
A	1	1.03	0.99 ± 0.02
b	1	1	1.00 ± 0.02
c	1	0.81	1.00 ± 0.06
d	1	0.99	0.99 ± 0.02

Table 3.5: Typical upstream phase and energy conditions at CTF3.

3.6 Signal Generator Measurements

Measurements have been taken using a 12 GHz signal generator to determine the performance of the three sets of phase monitor electronics independently from the phase monitors themselves. In particular, these tests were focused on identifying the saturation and cross-talk characteristics of the output mixer and diode signals in order to determine a suitable input power range to use during normal operation.

3.6.1 Experimental Setup

Two changes were made to the setup shown in Figure 3.7 for these tests. Firstly, the beam induced signal from the phase monitors usually connected to the RF port of the mixers is replaced by the output from a 12 GHz signal generator. To be able to reach the same input power levels as the beam signals the signal generator output is amplified. This allowed the input power to the mixers to be varied in a wide range between 0 and 33dBm, or between 0.2 and 10.0 V in terms of voltage. The precise power sent to the mixer is verified between each measurement using a power meter.

Secondly, the diode outputs were amplified during these tests (using the same amplifier introduced in Section ??) by a factor 10 in voltage to reduce digitiser noise in the measurement. The non-amplified peak diode output is therefore 170 mV, rather than the 1.7 V seen in the plots in this section. The ± 500 mV mixer outputs have not been amplified. Usually the mixer output is amplified and the diode not amplified, as in Figure 3.7, for reasons that will become clear later in this chapter.

There are some differences between the properties of the generated signal and the beam signal that would be used in normal operation. Firstly, unlike the pulsed beam signal the used generated signal is continuous. It has been verified that the response of the mixers is equivalent for both the continuous and pulsed signals, at least in terms of output power and saturation levels [REF]. The cross-talk properties are difficult to characterise with beam based measurements alone, but assumed to be similar.

Secondly, the phase of the generated signal does not vary with time, compared to the beam signal which has a large $\sim 40^\circ$ phase sag along the pulse and much larger phase jitter. If the signal generator was used at the same frequency as the beam and LO signals, 11.994 GHz, the mixer output would therefore be constant as it depends only on the static phase as per Equation 3.9. Instead, a generated signal with a slightly lower frequency of 11.991 GHz has been used. From Equation 3.8 it can be seen that in this case the mixer output voltage is sinusoidal, with a frequency equal to the frequency difference between the LO and RF inputs – or $11.994 - 11.991$ GHz = 3 MHz with the setup used here. This has the benefit of being able to see the response of the electronics to all input phases in one measurement, rather than having to take multiple measurements varying the LO phase shifter between each one.

3.6.2 Results

Figures 3.10–3.15 show the mixer and diode outputs for all three sets of electronics at each of the input power levels sent from the signal generator. These will be referred back to and discussed in more detail in the remainder of this section. Some initial observations that are immediately clear from these figures are as follows. All mixer outputs show a sinusoidal oscillation with a frequency of 3 MHz, or 60 samples at the sampling frequency of 192 MHz, as expected. An oscillation with the same frequency is also visible on the diode outputs, with the largest amplitude for the 2nd set of electronics. This is the first hint of the non-ideal characteristics of the diodes. Finally, output of the mixer and diode increases with the input power, as expected. At high input powers the outputs begin to saturate. This is apparent by observing the diode signals, on which the output is much flatter at the highest power levels, without the oscillation seen at low input powers. The characteristics of the mixers are discussed in Section 3.6.3 and the diodes in Section 3.6.4.

3.6.3 Mixer Performance

Sinusoidal Characteristics

Figure 3.16 shows fits to the response of all mixer outputs at an input power of 27 dBm, close to the typical input power from the beam signals when they are connected. Markers show the data points and the lines are sine fits to the data. The phase offset (displacement in peaks) between the output of each mixer holds no significance for the electronics performance. This is set only by the relative phase between the signal generator and the LO at the time the measurement was started. For normal operation with beam the LO phase shifters are changed to match the phasing of each set of electronics (Section 3.7).

The reconstruction of the phase from the mixer output depends on the mixer output being sinusoidal. In particular the maximum mixer output, or equivalently the gradient around zero mixer output (using the small angle approximation) is critical due to the dependence on the amplitude in Equation 3.12. Each set of electronics has a different output amplitude due to slight differences in the LO power for each set of electronics and between the individual components used. At an input power of 27 dBm Mixer 1 has a higher peak output of 510 mV, compared to 410 mV and 380 mV for Mixer 2 and Mixer 3 respectively.

Overall, the agreement between the actual mixer output and the sinusoidal fits at this input power is good. However, there is some distortion away from the ideal sine curve that is most visible around the maximum and minimum mixer output. Figure 3.17 shows the residuals between the mixer outputs and the sine fits across one half wavelength – from maximum output to minimum output. In the figure the plotted residual is the difference between the fit and the data expressed in terms of an equivalent phase offset, $\Delta\phi$, using:

$$\Delta\phi(t) = \arcsin\left(\frac{V_{MIXER}(t) - V_{FIT}(t)}{A}\right) \quad (3.32)$$

Where $V_{MIXER}(t)$ and $V_{FIT}(t)$ are the mixer voltage and fitted voltage at sample t respectively, and A is the fitted mixer amplitude. On the falling slope between the peaks there is

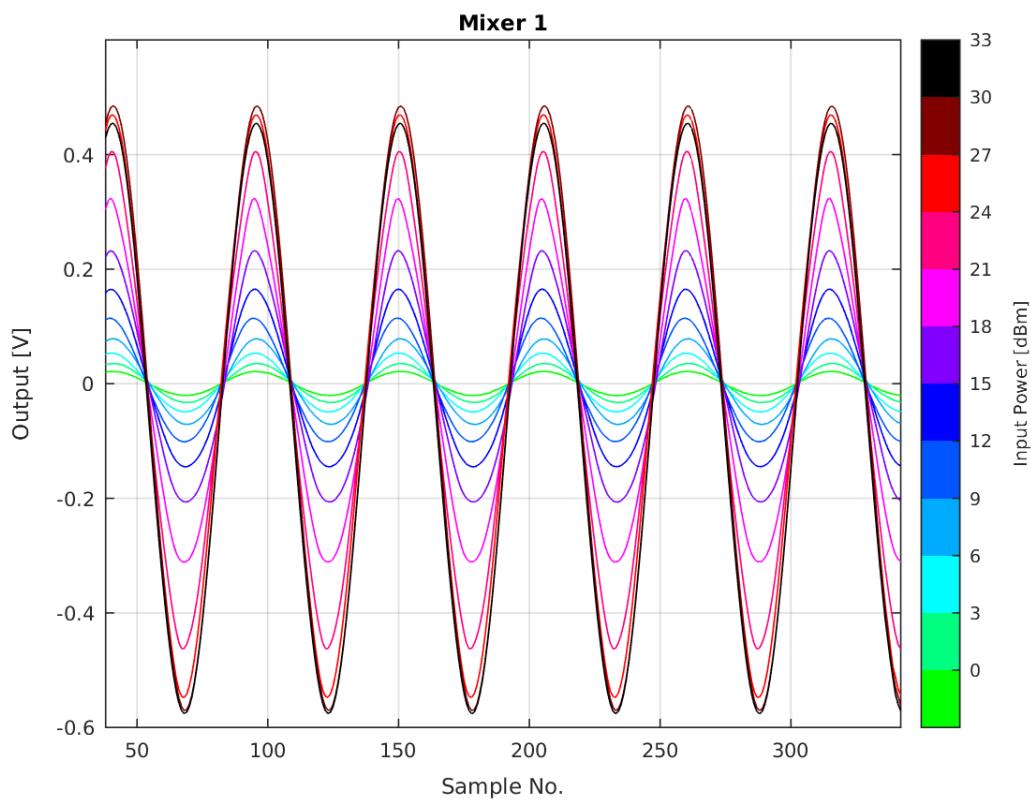


Figure 3.10: Response of Mixer 1 to signal generator input.

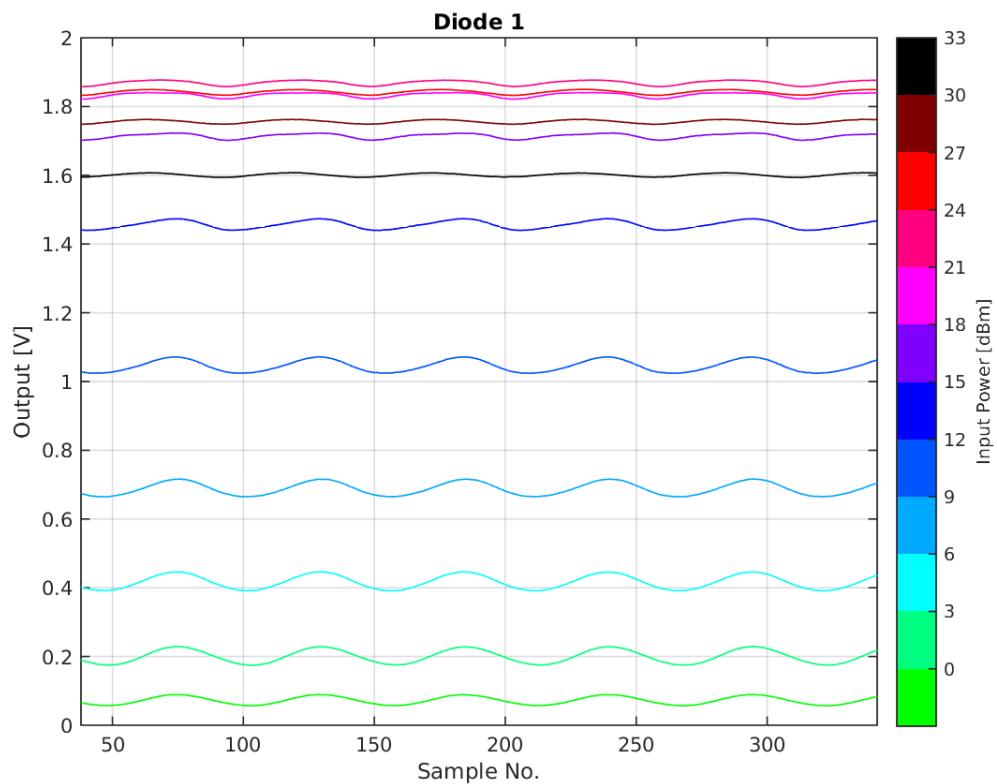


Figure 3.11: Response of Diode 1 to signal generator input.




Figure 3.12: Response of Mixer 2 to signal generator input.




Figure 3.13: Response of Diode 2 to signal generator input.




Figure 3.14: Response of Mixer 3 to signal generator input.




Figure 3.15: Response of Diode 3 to signal generator input.

only a slight oscillation about the ideal sinusoidal behaviour. The deviation from ideal is at the level of $0.25 \pm 0.03^\circ$ and $0.30 \pm 0.04^\circ$ for the first and third mixers, with a slightly larger effect of $0.45 \pm 0.04^\circ$ for the second mixer. This applies within $\pm 80^\circ$ of the zero crossing in the mixer output. Within $\pm 10^\circ$ of the maximum or minimum output the deviation from the sine fit rapidly increases, reaching several degrees for each mixer. For operation with the beam this means that the accuracy of the phase measurement cannot be guaranteed when the LO phase is set so that the mixer is giving close to its maximum or minimum output. This is also true for other reasons, as seen in Section ???. The PFF system can only correct small offsets at the level of around $\pm 5^\circ$ (Section 5.4.1), so the non-ideal response close to peak output is not an issue for the PFF performance.

However, for input powers in the range from 15–21 dBm the non-ideal characteristics of the mixers are larger. One example of this is shown in Figure 3.18, at an input power of 18 dBm. If input powers in this range are used calibrations of the mixer response should normally be restricted to around the zero crossing so that the fitted amplitude gives the best approximation to the true behaviour for small phase offsets.

Dependence on Input Power

The mixer output is expected to linearly depend on the input voltage and the diode on the square of the input voltage. Both these dependencies must hold in order to use Equation 3.12 and obtain a phase measurement that does not depend on the input voltage to the electronics (therefore making the calculated phase insensitive to any possible variations in power along the pulse from the beam signal, for example). For these measurements the input voltage can be calculated using the known input power and $50\ \Omega$ impedance of the electronics.

Figure 3.19 shows the dependence of the mixer output amplitudes on the input voltage. As seen previously the first mixer gives a larger output than the other two mixers. The 2nd and 3rd mixers give a similar response up to an input voltage of 3.5 V (24 dBm). Dashed lines in the figure show a linear fit to the mixer output restricted to the range between 0.45 V and 1.75 V (6 dBm to 18 dBm) in each case, as marked by the vertical black lines. All three mixers give a linear response up to an input voltage of around 3 V (23 dBm), after which the effects of saturation begin to appear. By an input voltage of 5 V (27 dBm) the first and third mixers are almost fully saturated with almost no remaining power dependence in the output. The second mixer begins to enter saturation at the same voltage as the other two mixers but retains a strong power dependence up to a higher input voltage of 7 V (30 dBm).

Asymmetry in Output

One final interesting property of the mixers is that the output is not symmetric about zero, in other words the maximum output voltage is different to the absolute value of the minimum output voltage. This is perhaps most visible looking back to the Mixer 1 output at all power levels in Figure 3.10, where the maximum output is around +0.45 V but the minimum output is around -0.55 V.

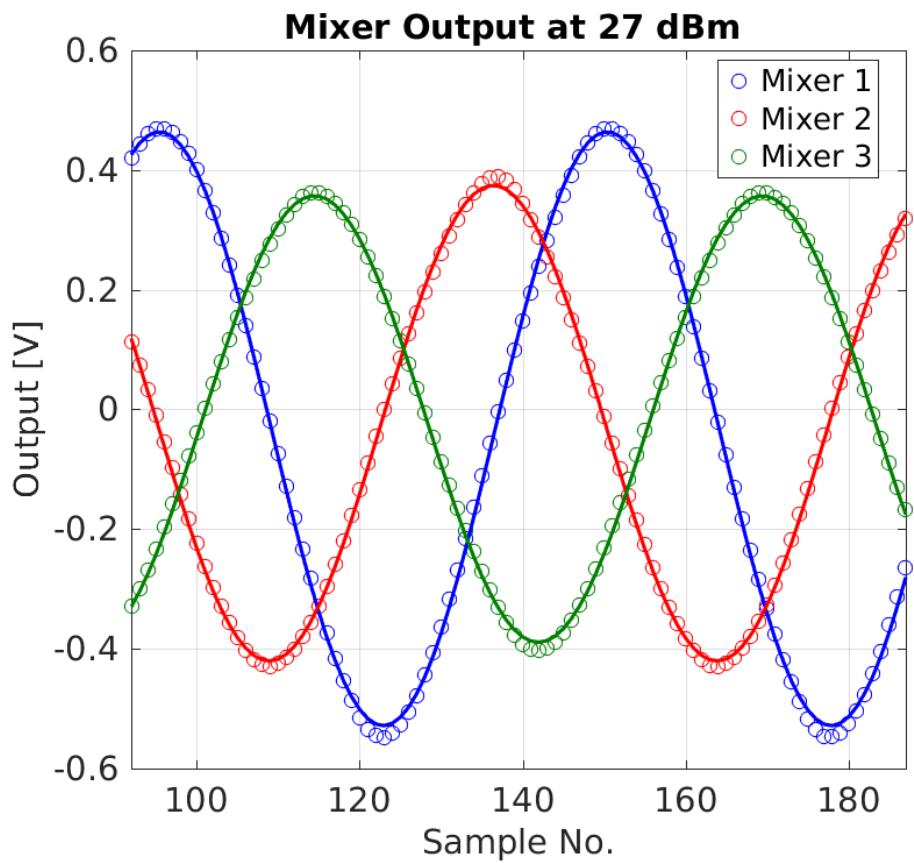


Figure 3.16: Sinusoidal fit to mixer responses at 27 dBm input power.

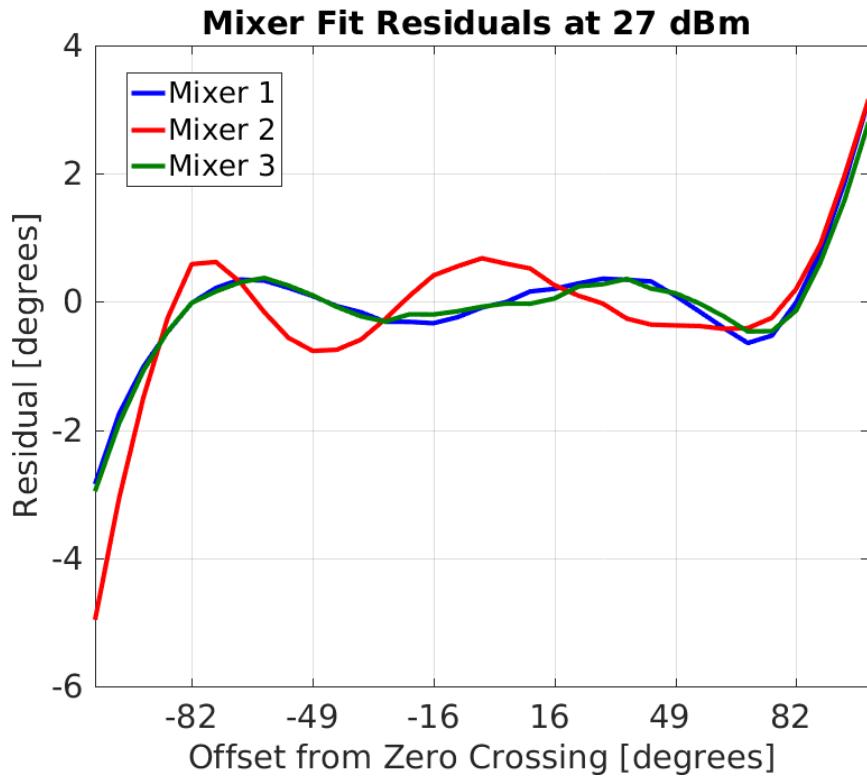


Figure 3.17: Residuals to sinusoidal fit at 27dBm.

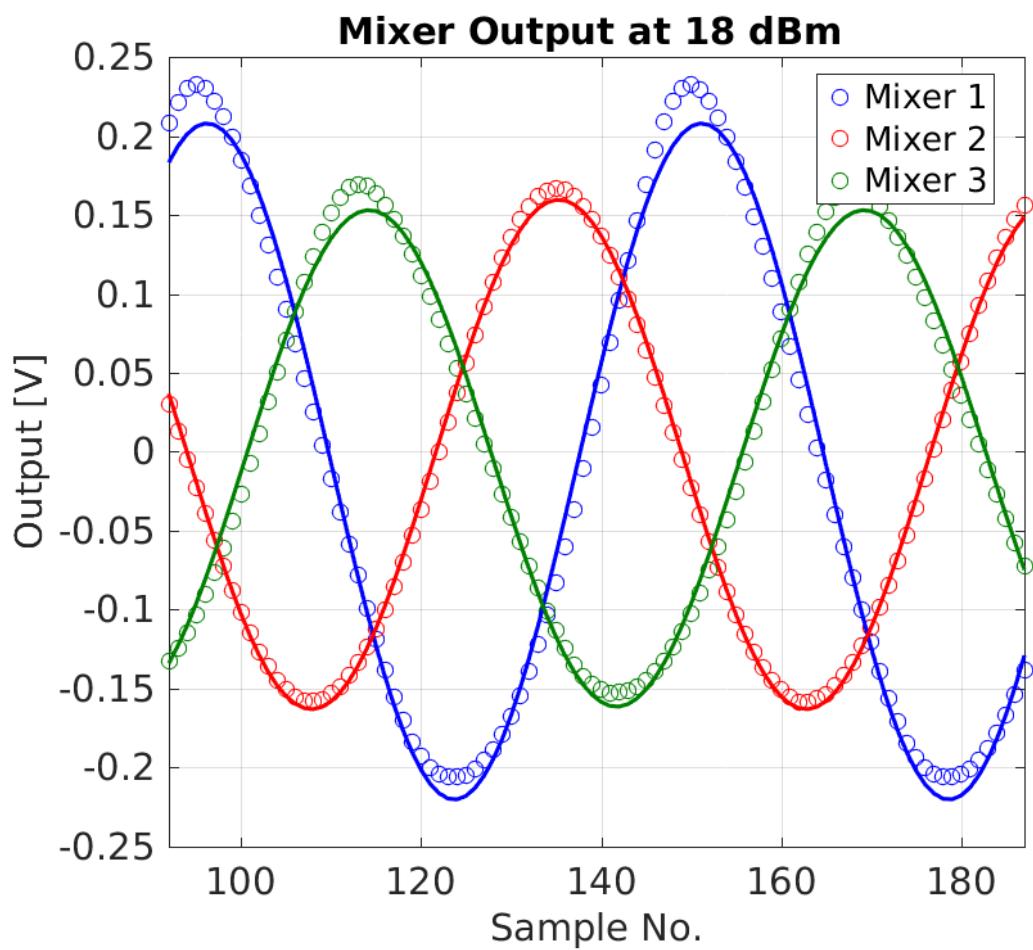


Figure 3.18: Sinusoidal fit to mixer responses at 18 dBm input power.

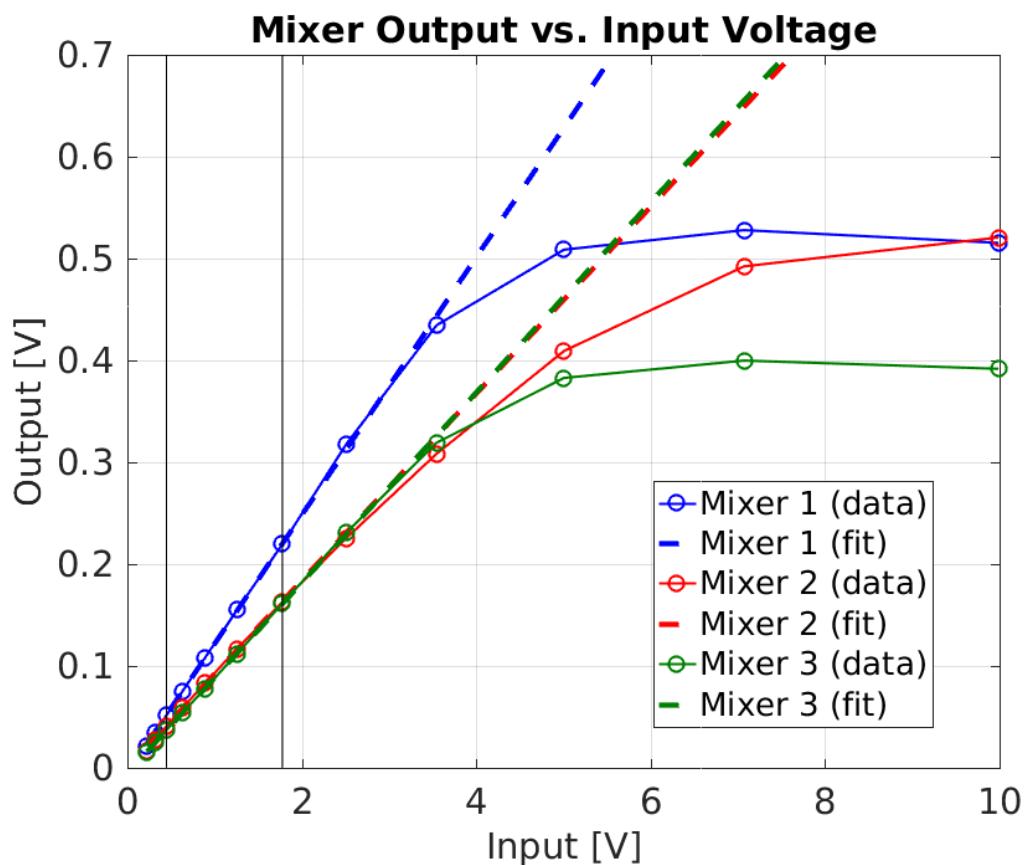


Figure 3.19: Linear fit to mixer output voltage vs. input voltage.

Figure 3.20 shows how the mixer amplitude at maximum and minimum output varies with the input voltage. Mixer 1 asymmetry is largest for Mixer 1 and smallest for Mixer 3. The effect appears to increase in magnitude with the input voltage, with the ~ 100 mV difference mentioned previously for Mixer 1 at an input of 10 V, but differences of only several mV at low input powers. For each mixer the amplitude at maximum output is larger for input voltages up to 2.5 V (21 dBm). Above 2.5 V input voltage this flips, with the minimum mixer amplitude being larger than the maximum amplitude.

For input voltages between 0.45 V and 1.25 V (6 dBm to 15 dBm) the mixer asymmetry has an approximate quadratic dependence on the input voltage, as shown in Figure 3.21. Outside this range there is no simple relationship that can explain the dependence of the asymmetry on the input voltage. One explanation for the asymmetry in the mixer outputs is cross-talk coming from the diode signals. Above 15 dBm the diodes enter saturation, as discussed in the next section, which may explain why the quadratic fit to the mixer asymmetry is only valid at power levels up to this value.

Taking the power dependent asymmetry into account the actual mixer response can be modified from Equation 3.9 to become:

$$\text{Mixer}(t) = m_1 A(t) \sin[\phi(t)] + m_2 A(t)^2 + m_3 A(t) + m_4 \quad (3.33)$$

Where $A(t)$ is the input voltage at time t , and m_1 , m_2 , m_3 , and m_4 are calibration constants.

3.6.4 Diode Performance

Dependence on Input Power

The dependence of the three diode outputs on the input power is shown in Figure 3.22, with the square root of the diode output plotted rather than the diode directly as this is the expected linear relationship. Immediately it is apparent that the diode signals saturate at much lower input voltages than the mixer signals. All three diodes are almost fully saturated at an input of 2 V (20 dBm), with the effects of saturation already beginning to appear above 1.25 V (15 dBm). Figure 3.23 shows a linear fit to the square root of the diode, using the range of input voltages between 0.45 V and 1.25 V (6–15 dBm). Even below saturation the response of $\text{sqrt}(\text{Diode})$ is not well approximated by a linear dependence as desired. However, in the range from 0.30 V to 1.25 V (3 dBm to 15 dBm) a quadratic fit to the diode output directly (not $\text{sqrt}(\text{Diode})$) does give a good approximation to the true dependence of the diodes on the input voltage. This is shown in Figure 3.24.

Cross-Talk

As seen previously in Figures 3.11, 3.13 and 3.15 the diode outputs show a sinusoidal oscillation. Like there is cross-talk from the diode on the mixer outputs, there is also cross-talk from the mixers on the diode outputs. Figure 3.25 shows a sinusoidal fit to the cross-talk on Diode 1 at an input power of 6 dBm. It has the same characteristics as the mixer output, including the slight distortion away from ideal sinusoidal behaviour around the peaks.

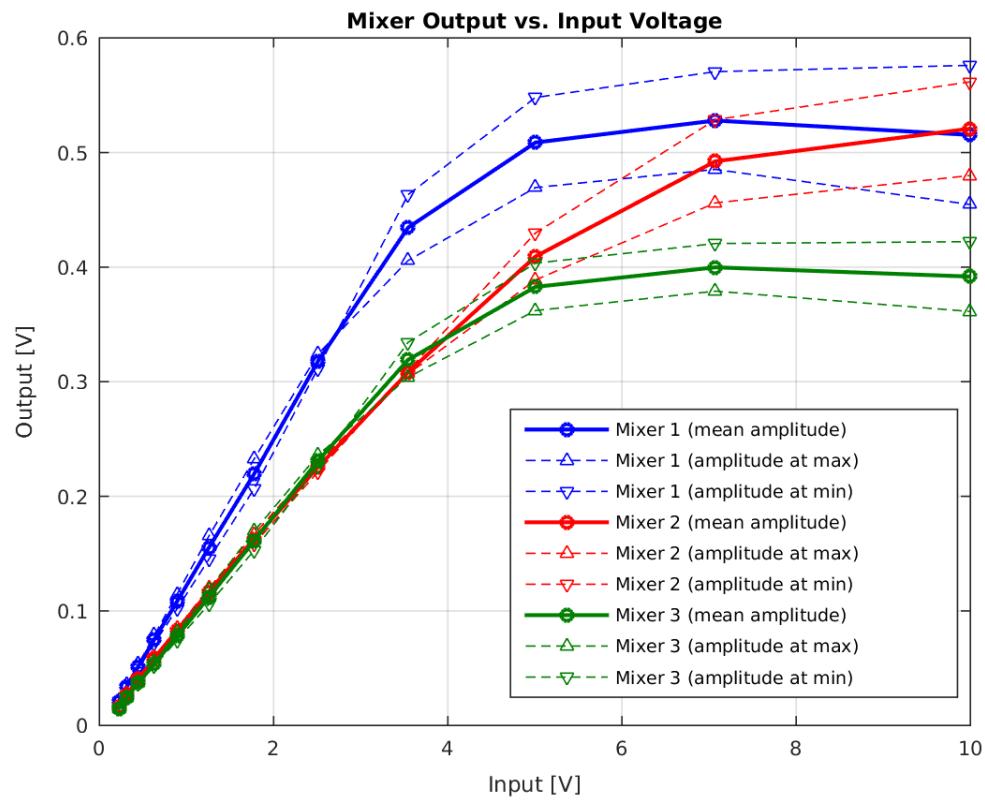


Figure 3.20: Mixer maximum and minimum output voltage vs. input voltage.

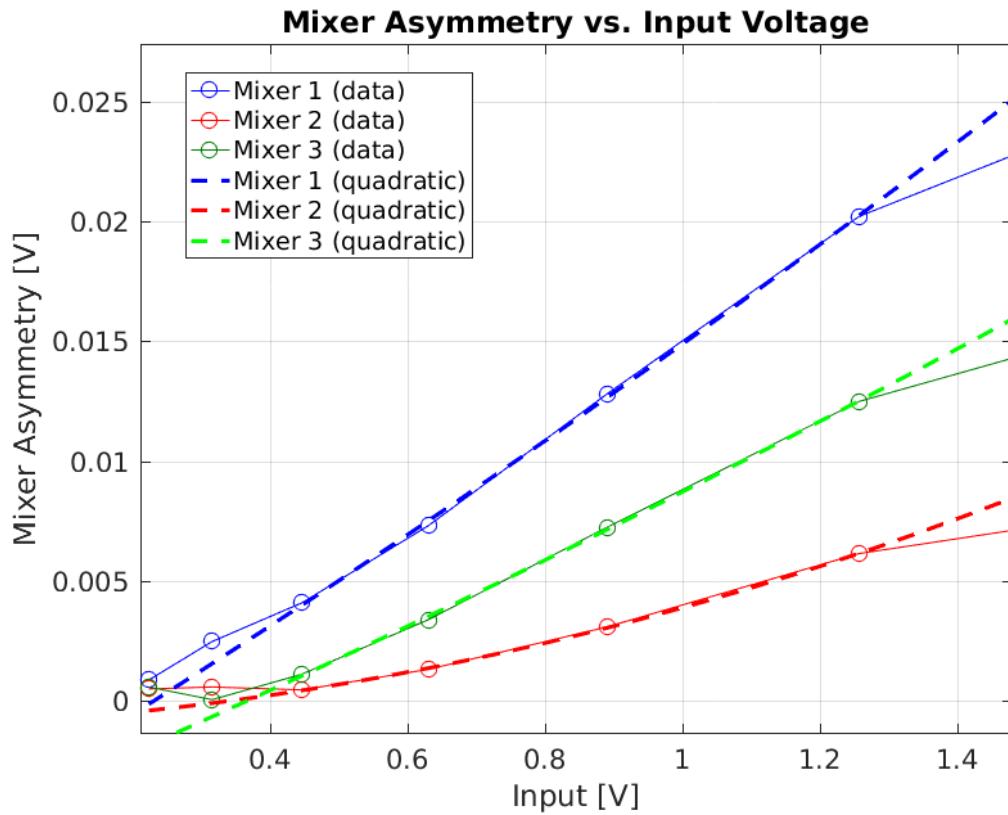


Figure 3.21: Relative amplitude vs. input power of cross-talk on the mixer from the diode.

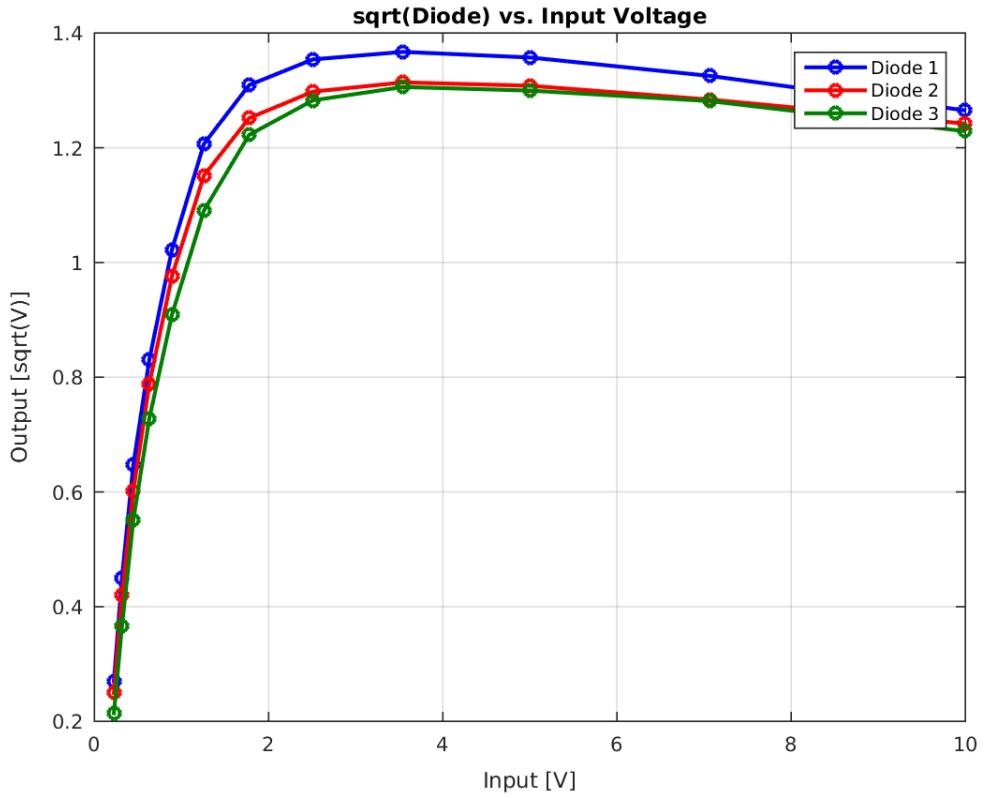


Figure 3.22: $\text{sqrt}(\text{Diode})$ vs. input voltage.

However, as the diodes enter saturation the oscillation is initially distorted, and then has a much smaller amplitude when the diode output is fully saturated. One example of this is shown for the Diode 1 output at 18 dBm in Figure 3.26. The peaks around the maximum output are clearly non-sinusoidal in this case.

Figure 3.27 shows the dependence of the relative amplitude of the cross-talk on the input power. The relative amplitude of the cross-talk means the fitted amplitude of the sinusoidal oscillation on the diode divided by the mean diode output. Up until the diode outputs are fully saturated the relative amplitude of the cross-talk is around a factor two larger for the second diode. For example, at an input power of 12 dBm the relative cross-talk is at around the level of 30% for the second diode, or 15% for the first and third diode outputs. Up to input powers of 15 dBm the relative cross-talk is always above 10%.

Finally, Figure 3.28 compares the oscillation on the diode to the oscillation on the mixer. It can be seen that there is a phase shift between the two, which adds a further complication to the necessary phase reconstruction method. Taking in to account the actual characteristics of the diodes, including the cross-talk and quadratic dependence on the input power, the expected expression for the diode output from Equation 3.10 can be modified to:

$$\text{Diode}(t) = d_1 A(t)^2 + d_2 A + d_3 + d_4 A(t) \sin[\phi(t) + \delta] \quad (3.34)$$

Where d_1, d_2, d_3, d_4 and δ are calibration constants.

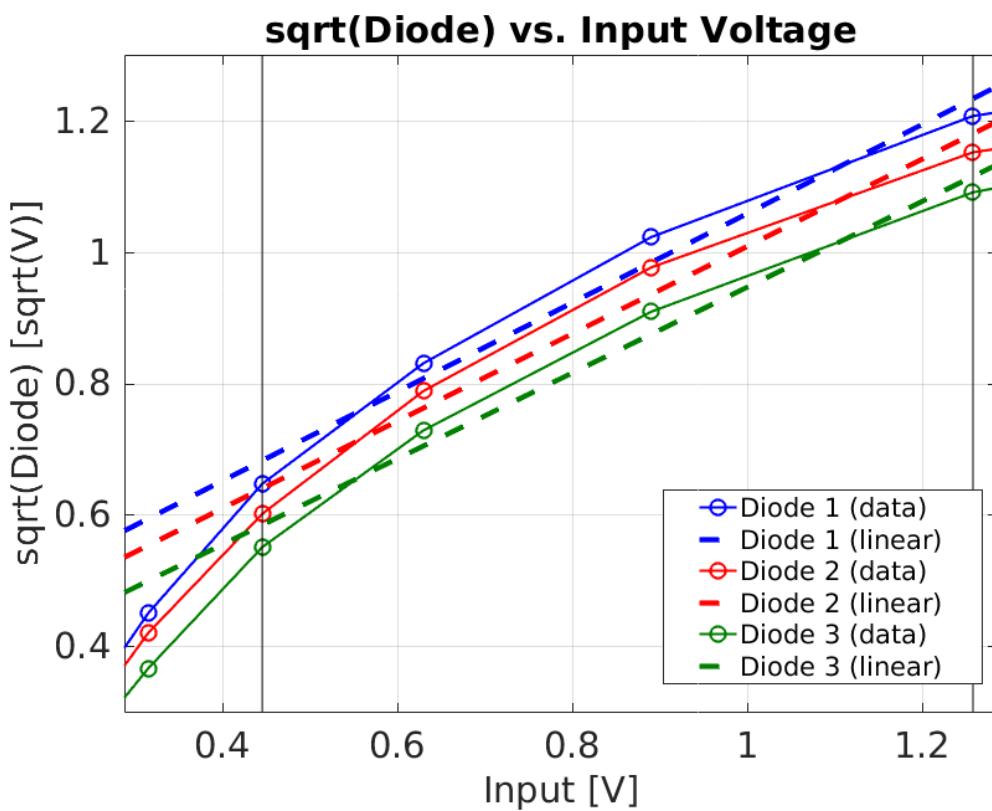
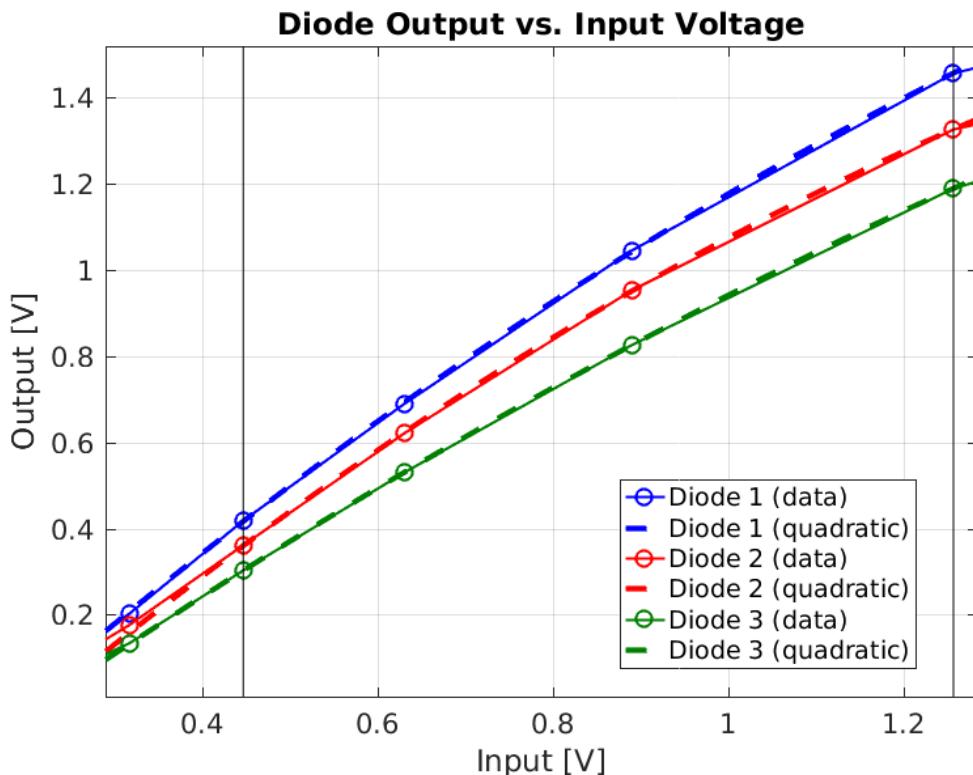
Figure 3.23: Linear fits to $\text{sqrt}(\text{Diode})$ vs. input voltage.

Figure 3.24: Quadratic fits to Diode vs. input voltage.




Figure 3.25: Sinusoidal fit to cross-talk on diode at 6 dBm input power.




Figure 3.26: Sinusoidal fit to cross-talk on diode at 18 dBm input power.

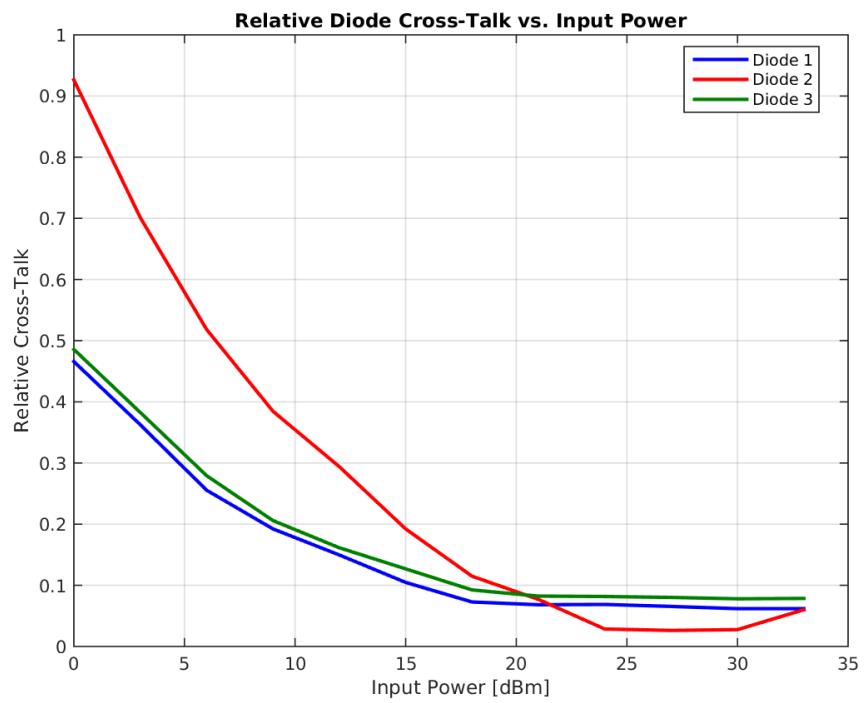


Figure 3.27: Dependence of the relative amplitude of cross-talk on the diode versus the input power.

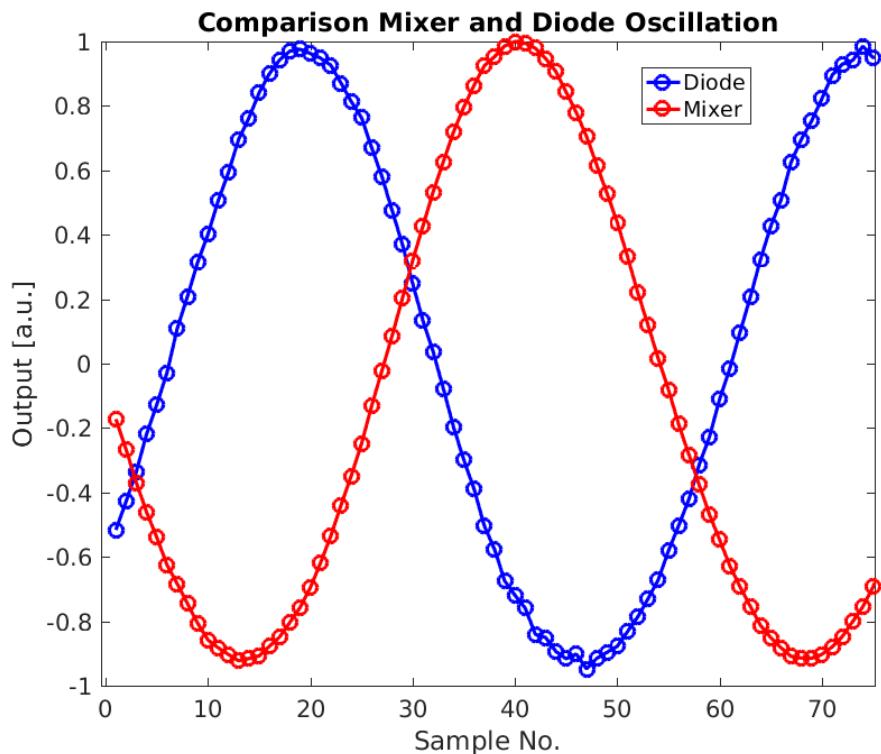


Figure 3.28: Comparison of the oscillation on the mixer and the diode, showing a relative phase offset between the two.

3.6.5 Consequences for Normal Operation

The results of the signal generator tests have several consequences for the setup of the electronics and phase reconstruction during normal operation with the beam induced signals from the phase monitors. Firstly, in order to maximise the signal to noise ratio and yield the best possible resolution on the phase measurement the highest possible input power below saturation should normally be used. The degradation of the phase resolution with the input power is seen for beam based measurements in Section [REF]. However, the diodes begin to enter saturation much earlier than the mixers, at around 15 dBm rather than 23 dBm. This means that in order to be able to use the diode measurement as part of the phase reconstruction the input power would have to be limited to below 15 dBm, 8 dBm lower than would be ideal for the mixer performance. There is no way to use different input powers for the mixers and diodes without a complete redesign of the electronics.

Secondly, the modifications that the cross-talk on the mixer and diode make to Equations 3.33 and 3.34 above makes the needed calculation to reconstruct the phase much more complex than the ideal case using Mixer/sqrt(Diode) in Equation 3.12. In particular, the dependence of the diode output on $\sin(\phi + \delta)$ means there is no simple expression that can be derived to create an input power independent phase measurement. An iterative process would have to be used to estimate the phase instead, converging towards the true diode output without cross-talk after each iteration using the estimated phase value. This may be possible in offline data analysis but would be difficult to implement in the PFF algorithm whilst still meeting latency requirements.

Due to these reasons, and with no possibility to make modifications to the electronics, the decision was eventually taken to not include the diode measurement in the phase reconstruction process. For operation with the beam this means making the assumption that the output power from the phase monitors is constant along the pulse, and that the jitter in the output power is small. This is a good approximation, as seen later in Section [REF]. To reduce the sensitivity to any slow drifts in the output power due to changes in the beam conditions calibrations are taken at regular intervals between measurements.

With this treatment of the electronics outputs the phase is reconstructed as follows:

$$\text{Mixer}(t) = A \sin[\phi(t)] + d \quad (3.35)$$

$$\phi(t) = \arcsin\left(\frac{\text{Mixer}(t) - d}{A}\right) \quad (3.36)$$

Two calibration constants are needed – A and d . A is the fitted amplitude of the sinusoidal mixer output, and d is the asymmetry or offset between the maximum and minimum mixer output. This is a simplified form of Equation 3.33 given the assumption that the power is constant. In reality both A and d have a power dependence.

As any variations in input power are not removed by this method there is a benefit to operating the mixer in a region where the power dependence of the output is reduced. The best phase resolution achieved to date has been achieved with input powers to the electronics in the range between 24.5 dBm and 27 dBm (as stated in [REF]), where the mixers have actually begun to enter saturation, as a result. Operating in this range also has the benefit

of reducing the deviation from ideal sinusoidal behaviour at lower input powers as seen in Section 3.6.3.

All the beam based measurements in the remainder of this chapter and the rest of the thesis use this phase reconstruction approach. Although the diodes are no longer directly used as part of the phase measurement they are still useful for the purposes of the time alignment of signals and for monitoring whether there have been any large changes in input power. The PFF firmware on the FONT5a feedforward controller has also been changed to add the option to not include the diode in the correction calculation (Section ??). The nominal PFF setup, as used for the results presented in Chapter 6, now does not use diode normalisation.

3.7 Calibrations

The remainder of this chapter presents the performance of the complete phase monitor and electronics system for normal operation with the beam, replacing the signal generator with the RF output from the phase monitors.

The first step in using the the phase monitor measurements is to calibrate the mixer outputs. Calibrations of the phase monitor signals are typically taken on a daily basis during data taking periods, as well as additional calibrations when there have been any changes in beam conditions or to the setup of the electronics. These are needed to determine the calibration constants, amplitude and offset, to be able to determine the phase from the mixer output as previously discussed. This section presents typical calibration results for all three phase monitors and discusses aspects such as the stability of the calibration along the pulse and determining the optimal set point for the LO phase shifters.

For beam based measurements calibrations are performed using the LO phase shifters. By varying the LO phase shifter the phase between the LO and the beam signal is changed. During a calibration the phase shifters are moved through 360° at 12 GHz so that the response of the mixer to all phase offsets between the beam and LO can be determined. Normally calibrations are taken at 12 shifter settings across the full 360° range, with 10 pulses acquired at each setting and the whole scan taking approximately 10 minutes. These choices are a compromise between having enough points for a good quality fit whilst reducing the possibility of large drifts in beam phase during the scan which would degrade the fit results. All the calibrations presented use the electronics setup with the mechanical LO phase shifters in place (Section 3.8). The settings on these phase shifters approximately correspond to degrees at 4 GHz, thus a phase shifter change of 120 units corresponds to 360° at 12 GHz.

A calibration from both the SiS digitiser and FONT5a setup will be shown. During operation of the PFF system Mon 1 is usually connected to the FONT5a board as the correction input, whilst Mon 2 and Mon 3 are connected to the SiS digitisers where they can be acquired together with other signals at CTF3. The difference in measured output voltage between the two setups results from the use of an amplifier before the SiS digitisers to reduce

noise in that setup (Section 3.4). Also, when using the FONT5a board the mixer outputs are attenuated by 1 dB to avoid saturating the FONT ADCs. The calibration constants from Mon 1 on the FONT5a board are needed for the PFF gain calculation in Section 5.1.1. All measurements of the upstream and downstream phase after this chapter use Mon 2 and Mon 3 on the SiS digitisers and their respective calibration constants.

3.7.1 Calibration on SiS Digitisers

Figures 3.29–3.31 show the output of the mixer for each phase monitor along the pulse for all the phase shifter settings used during the LO scan. Away from the minimum and maximum output each mixer shows the expected phase sag along the beam pulse resulting from the RF pulse compression used at CTF3 [REF]. The shape of the phase sag as seen on the mixer changes sign depending on whether the LO phase places the mixer on the rising or falling slope of its sinusoidal output. Usually the mixers are operated on the falling slope where the measured phase sag is ‘u’-shaped, rather than ‘n’-shaped, as this is also the convention for other phase dependent signals at CTF3 [REF].

Near the minimum or maximum mixer amplitude the beam phase sag causes a much smaller variation in the mixer output voltage along the pulse. The phase resolution close to the peaks in the mixer output is therefore greatly reduced, as seen in Section ???. LO phase scans are used to not only calculate the calibration constants but also to determine the phase shifter settings that zero the mixer output, where the resolution is maximal. This process is documented in Section 3.7.4.

The noisier appearance of the output on Mon 3 is not an effect of the phase monitors or phase monitor electronics but is rather caused by real differences between the beam phase upstream (Mon 1, Mon 2) and downstream (Mon 3). Reducing the differences between the upstream and downstream phase is the focus of Chapter 4.

Figure 3.32 shows the result of fitting the mixer output versus the phase shifter setting at sample 605 along the pulse. The mixer response is sinusoidal as expected and as seen previously in the signal generator tests. In the signal generator tests there was some visible distortion away from the sine fit around the peaks at some input power levels (Section 3.6.3). There is no visible effect at the powers of the phase monitor signals (Section 3.1). Differences in the peak output of each monitor are expected due to differences in the input power from each phase monitor as well as differences between the sets of electronics. Small offsets between the data and the fit at some shifter settings are caused by drifts in the beam phase during the scan (particularly for Mon 3 where the beam is less stable), as well as human error in setting the shifter values.

The fitted values of the mixer amplitude, A , and offset, d , are found in Table 3.6. These values are used to calculate the beam phase as per Equation 3.36.




Figure 3.29: Mon 1 phase along the pulse for each LO phase shifter setting during the calibration.

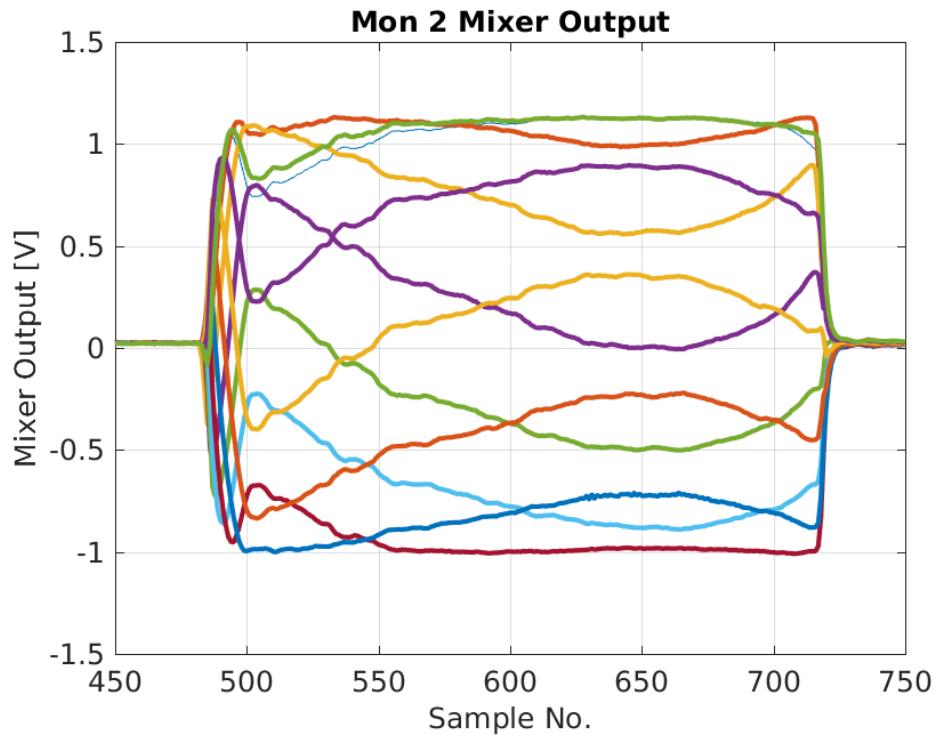


Figure 3.30: Mon 2 phase along the pulse for each LO phase shifter setting during the calibration.

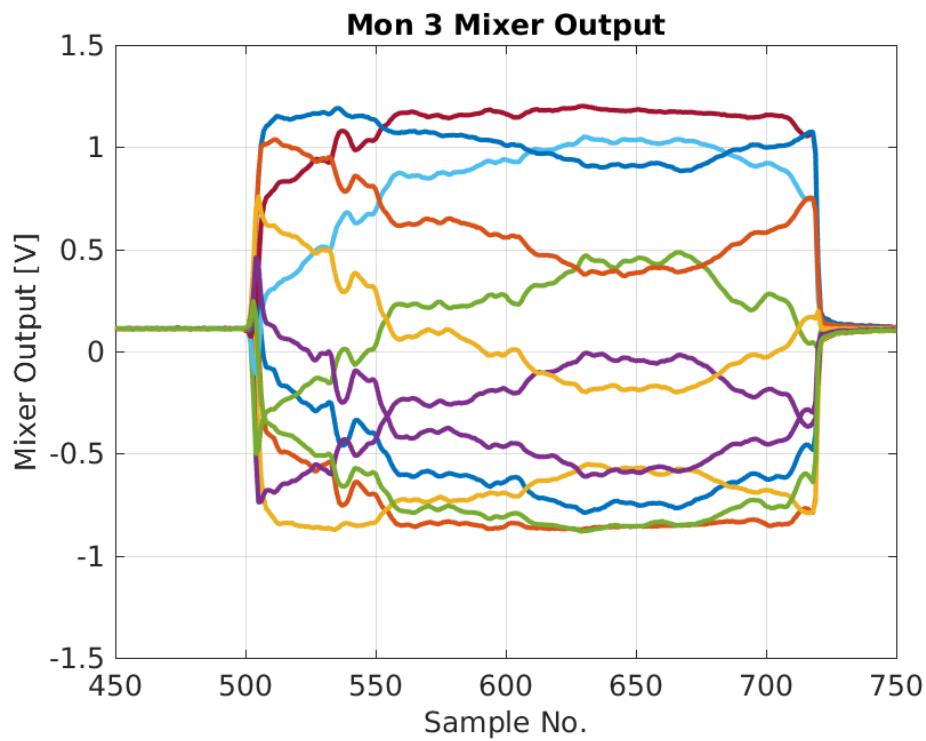


Figure 3.31: Mon 3 phase along the pulse for each LO phase shifter setting during the calibration.

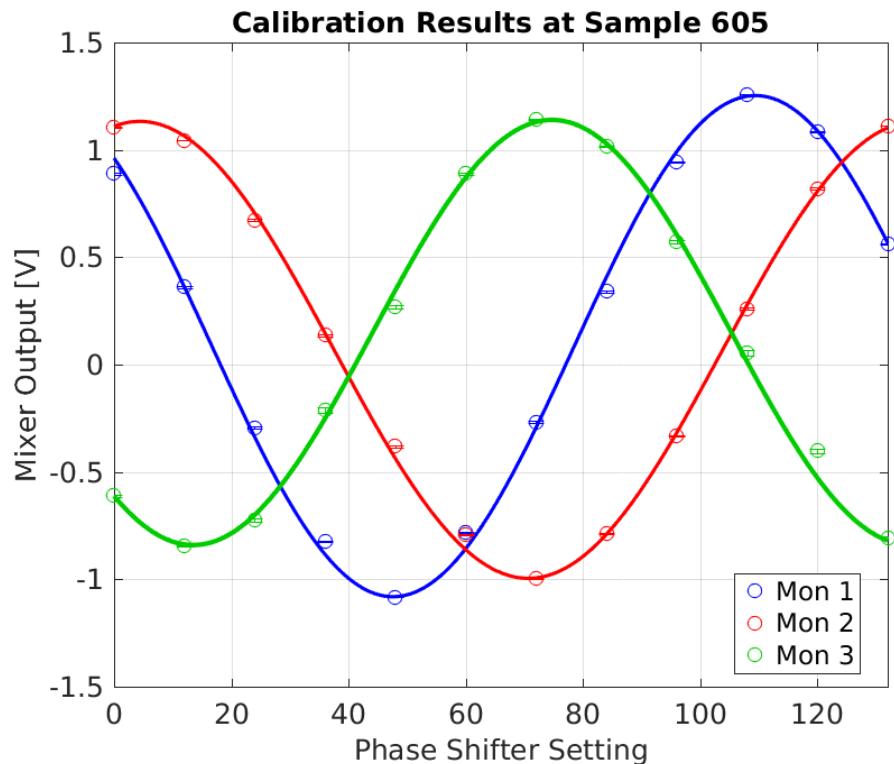


Figure 3.32: Fits to the mixer output vs. LO phase shifter setting at sample 605 on the SiS digitisers.

Monitor	A (amplitude)	d (offset)
Mon 1	1167 ± 10 mV	86 ± 9 mV
Mon 2	1064 ± 6 mV	69 ± 7 mV
Mon 3	990 ± 12 mV	150 ± 10 mV

Table 3.6: Fit parameters from the calibration on the SiS digitisers for each monitor.

Monitor	A (amplitude)	d (offset)
Mon 1	3434 ± 41 counts	-278 ± 34 counts
Mon 2	3144 ± 16 counts	-138 ± 17 counts
Mon 3	2870 ± 43 counts	-44 ± 44 counts

Table 3.7: Fit parameters from the calibration for each monitor on the FONT5a board.

3.7.2 Calibration on FONT5a Board

Figure 3.33 and Table 3.7 show the results of a calibration performed in exactly the same way but on the FONT5a board. The FONT5a board ADCs flip the sign of the input [REF], so a positive ADC output in counts corresponds to a negative input voltage and vice-versa. This explains why the apparent mixer output is on the rising slope at a phase shifter setting of zero on the FONT5a board in Figure 3.33, but on the falling slope on the SiS digitisers in Figure 3.32. For operation of the PFF system this difference must be taken in to account either by operating the mixer on the rising slope as seen on the FONT5a board (which in reality is the falling slope, as desired), or alternatively by using negative gain values for the correction output. The fitted values for d in Table 3.7 are also negative rather than positive as a result of this sign flip. Apart from these differences the overall shape of the mixer output follows the sinusoidal dependence as expected.

The FONT5a ADC outputs are 13-bit, or ± 4096 counts, with an input range of ± 500 mV (Section 5.1). The fitted Mon 1 output, with 1 dB attenuation added after the mixer, of between -3712 counts and +3156 counts therefore corresponds to an input voltage range of between +453 mV and -385 mV. Without the attenuator, which reduces the voltage by roughly 10%, the Mon 1 mixer would saturate the ADC at its peak output. As the contribution of digitiser noise is small on the FONT5a board (Section 3.4) a 1 dB attenuator is also added to the Mon 2 and Mon 3 outputs so that the overall setup for each monitor is the same in this measurement. However, during normal operation of the PFF system Mon 2 and Mon 3 are connected to the SiS digitisers, with their mixer outputs then being amplified rather than attenuated.

3.7.3 Multi-Sample Results

The calibration results on both the SiS digitisers and FONT5a board have been presented at one sample number around the middle of the pulse close to where the phase sag along the pulse is flattest. In this section the variation in the fitted calibration constants along

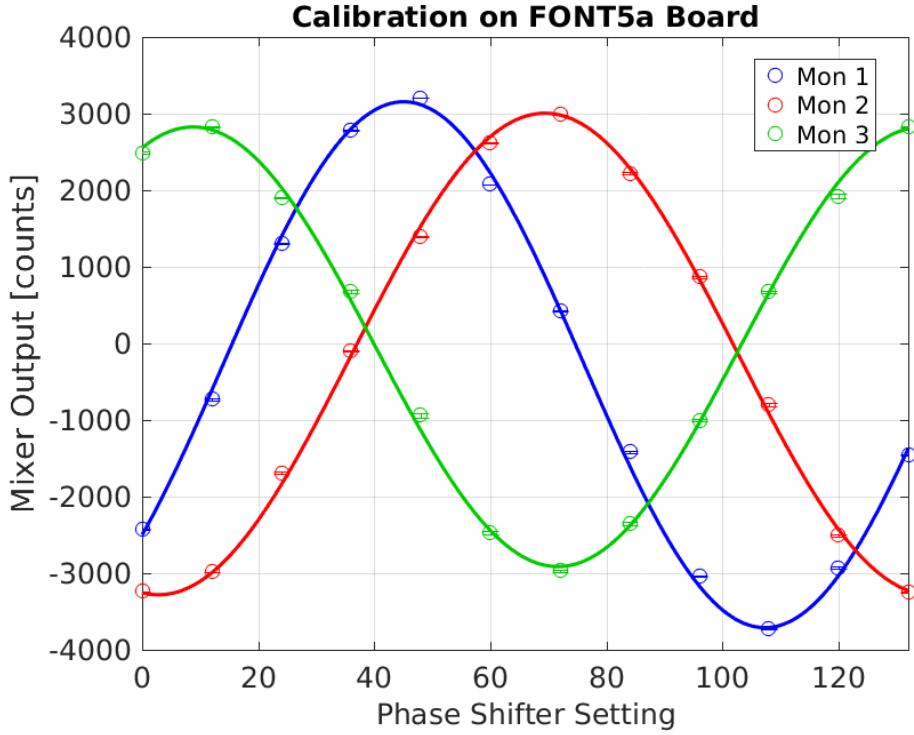


Figure 3.33: Results of a calibration performed on the FONT5a board.

the pulse is discussed. This is particularly important after taking the decision to not use the diodes, as the intended purpose of using the diodes was to normalise the mixer response to give an output independent of the input power. Without using the diodes any variations in input power along the pulse will also create differences in the calibration constants along the pulse.

The current implementation of the PFF algorithm on the FONT5a board uses the mixer multiplied by one gain value that is constant along the full pulse length to create the correction output (Section 5.1.1). It therefore cannot take in to account any variations in calibration constants along the pulse. Offline data analysis is usually performed in the same way so that the quoted resolutions are representative of the values that apply to the implementation of the PFF correction. The effect of taking in to account the variations in calibration parameters along the pulse seen here is shown in Section 3.9.2.

Figures 3.34 and 3.35 show the variation in the fitted calibration amplitude and offset across the full pulse length, using the same calibration on the SiS digitisers presented in Section 3.7.1. Differences in both the amplitude and offset along the pulse are visible. These are summarised in Table 3.8 in terms of the standard deviation of the fitted parameter values along the pulse.

The stability of the fitted amplitude along the pulse is similar for the upstream monitors (Mon 1 and Mon 2), with a jitter of around 8 mV in both cases. As the downstream beam is less stable than the upstream beam the variations in fitted amplitude along the pulse are larger for Mon 3, at the 15 mV level. In terms of a relative difference these values correspond to roughly a 0.7% variation in fitted amplitude for Mon 1 and Mon 2, or 1.5% for Mon 3.

Monitor	A (amplitude)	d (offset)
Mon 1	8.3 ± 0.4 mV	6.7 ± 0.3 mV
Mon 2	7.7 ± 0.4 mV	3.0 ± 0.3 mV
Mon 3	14.7 ± 0.7 mV	6.1 ± 0.3 mV

Table 3.8: Standard deviation in calibration fit parameters along the pulse.

With further optimisation of the downstream beam, as documented in Chapter 4, it should be possible to achieve similar Mon 3 amplitude stability to that seen for Mon 1 and Mon 2.

Absolute stability in the fitted offset along the pulse is similar to that of the amplitude but therefore much larger as a relative difference at the level of several percent. The variation in fitted offset along the pulse is smallest for Mon 2 at around 3 mV. For both Mon 1 and Mon 3 the variation is around a factor two larger, at 6 mV.

3.7.4 Zero Crossing

The full fit to the calibration result that is performed is:

$$\text{Mixer} = A \sin(bx + c) + d \quad (3.37)$$

Where x is the phase shifter setting. A and d are the calibration constants used to reconstruct the phase, with their values already quoted in Table ???. The remaining fit parameters b and c convert the phase shifter setting in to the phase offset between the LO and the beam. As the shifter readings are approximately in 4 GHz degrees, the expected value for b that converts the shifter value in to 12 GHz radians is $(12/4) * (\pi/180) \simeq 50$ mrad/unit.

To obtain the best resolution for the measurement the mixers should be operated where the dependence of the output voltage on the phase is maximal. This means maximising the partial derivative of the mixer output with respect to the phase shifter setting:

$$\frac{\partial \text{Mixer}}{\partial x} = Ab \cos(bx + c) \quad (3.38)$$

This is maximised when $bx + c = n\pi$, where n is any positive or negative integer. The optimal phase shifter settings therefore meet this criteria:

$$x = \frac{n\pi - c}{b} \quad (3.39)$$

Where the mixer output is $\text{Mixer} = A \sin(n\pi) + d = d$. In the case where there is no offset between the minimum and maximum mixer output ($d = 0$) the optimal point to operate the mixers is at zero output. Because of this the optimal shifter setting will be referred to as the zero crossing. In reality the small asymmetry in the mixer output means the optimal shifter setting is where mixer output is d . The effect of operating the mixers away from the zero crossing on the resolution is shown in Section ??.

In addition, as previously mentioned the convention is to operate the mixers on the falling slope where the partial derivative above is negative. The set shifter values are obtained using

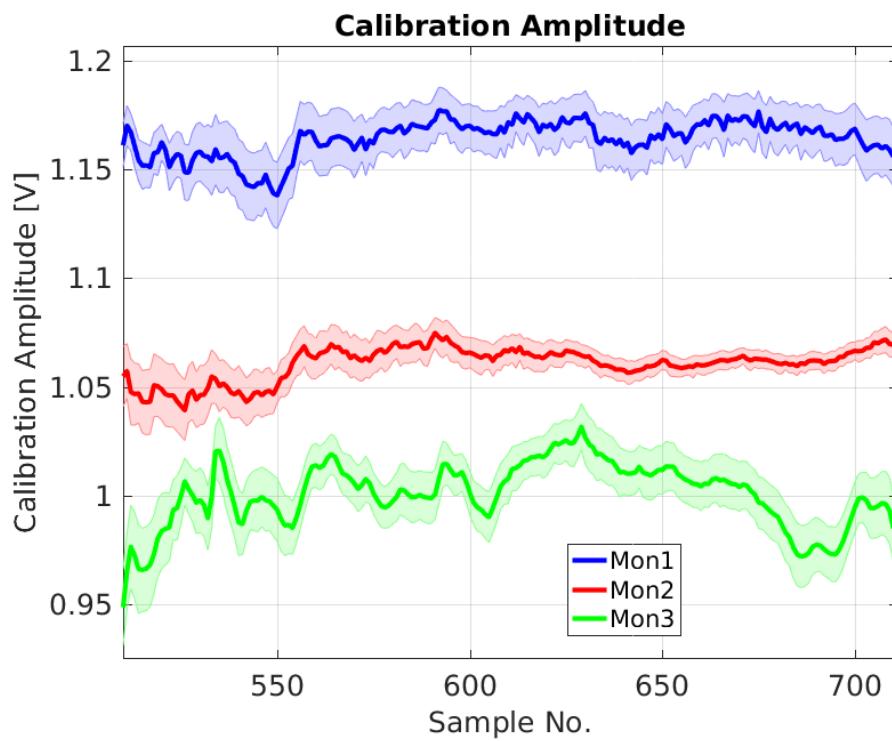


Figure 3.34: Variation in fitted amplitude along the pulse for each phase monitor.

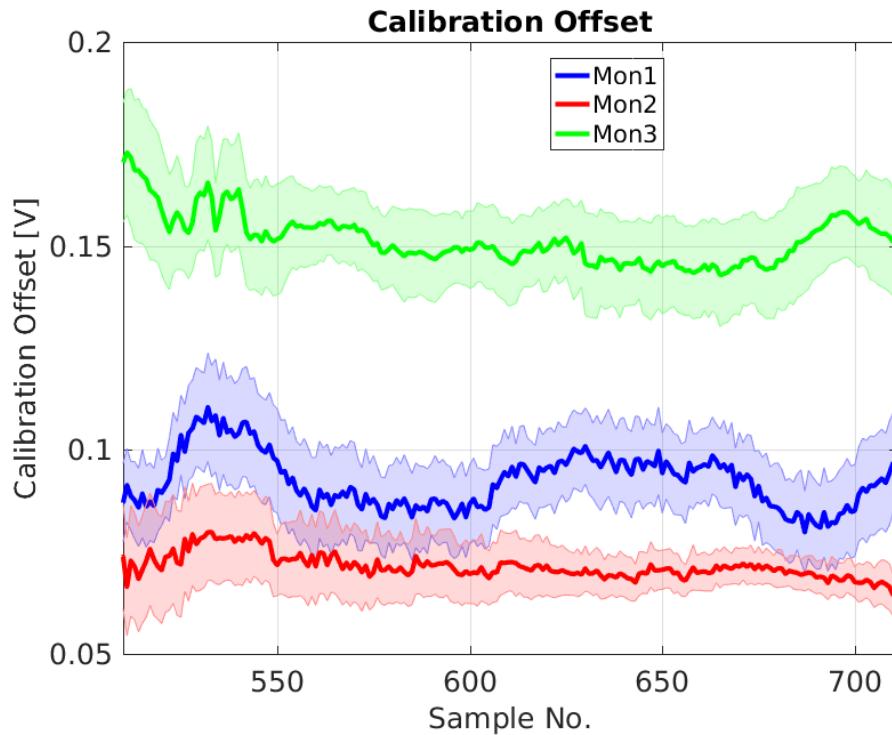


Figure 3.35: Variation in fitted offset along the pulse for each phase monitor

Monitor	b	c	Zero Crossing
Mon 1	50.8 ± 0.5 mrad/unit	2.29 ± 0.04 rad	16.7 ± 0.9 units
Mon 2	47.4 ± 0.5 mrad/unit	1.36 ± 0.04 rad	37.5 ± 0.8 units
Mon 3	51.3 ± 0.6 mrad/unit	4.02 ± 0.04 rad	105.5 ± 1.5 units

Table 3.9: Phase shifter setting to obtain the zero crossing for each mixer output and the fit parameters needed to calculate them.

the smallest positive integer n that leads to this criteria being met. Table 3.9 shows an example of values for the fit parameters b and c and the calculated phase shifter settings to be on the zero crossing for each monitor. These values are taken from the same calibration and sample number presented in Section 3.7.1 (on the SiS digitisers). The reader may be interested to compare where the shifter settings fall along the mixer output traces during the calibration in Figure ??.

Due to the large phase sag along the beam pulse it is clearly not possible to be at the zero crossing of the mixer for the full pulse length. For the PFF system the region of interest is the central portion of the pulse where the phase sag is flattest. The shifters are therefore set to zero the mixer output in this region, giving best resolution in the central part of the pulse but degraded resolution near the start and end of the pulse. In addition to this, slow drifts in the beam phase, particularly downstream, mean that the shifters must routinely be changed to stay at the zero crossing. With the current setup using mechanical phase shifters this must be done by hand, with no possibility to implement an automatic feedback on the shifter settings, for example.

3.8 Phase Shifter Noise

In the first tests with all three phase monitors in their final positions at CTF3 the measured phase resolution was far in excess of the 0.14° derived to be necessary to achieve a 0.2° corrected downstream phase jitter in Section 3.3. Figure 3.36 shows a typical example of the measured resolution along the pulse at that time with an achieved resolution of around 0.4° , three times larger than required. The lowest downstream phase jitter that could be achieved with the PFF system in these conditions is above 0.55° , only about 30% smaller than the initial upstream phase jitter at CTF3 (Chapter 4). To be able to achieve a large reduction in downstream phase jitter with the PFF prototype the source of the poor phase monitor resolution had to be identified and removed.

The first hint towards identifying the cause of the degraded resolution came by comparing the measured phase jitter from Mon 1 and Mon 2, with one example shown in Figure 3.37. Note that for all the results in this section it is not the absolute value of the phase jitter (which depends mostly on the beam conditions in that dataset) that is important but rather the difference between the measured phase jitter in each monitor. The measured phase jitter along the pulse in Mon 2 with a mean of $1.38 \pm 0.01^\circ$ is 1.7 times larger than the $0.83 \pm 0.01^\circ$ phase jitter in Mon 1. Jitter values from this and all the other datasets presented in this

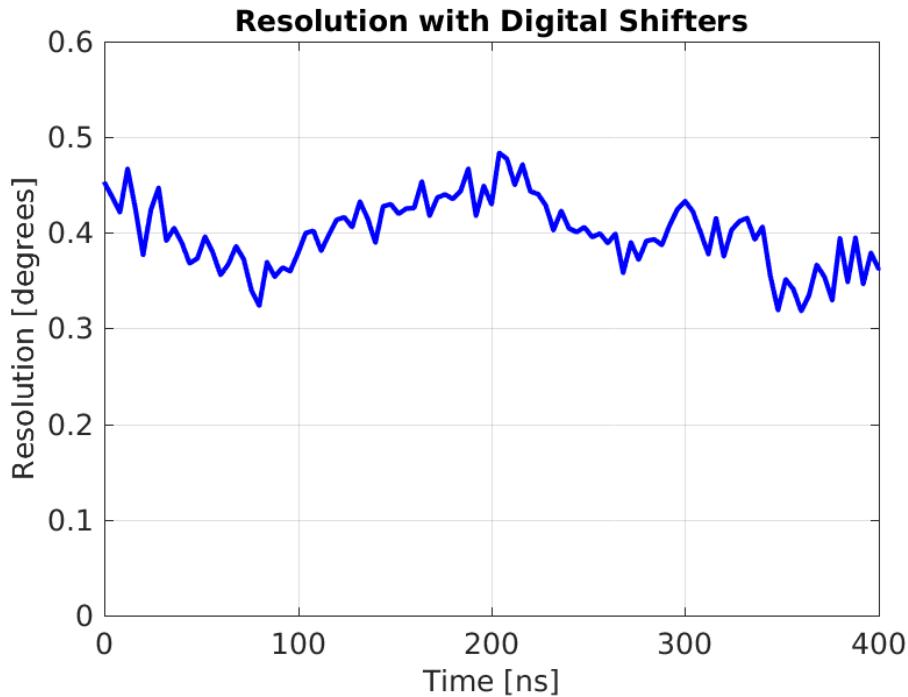


Figure 3.36: Phase monitor resolution using initial setup with digital phase shifters in place.

section are shown in Table 3.10. As Mon 1 and Mon 2 are neighbouring each other in the beam line the actual phase jitter in the two should be close to identical.

The overall phase monitor and electronics setup can be roughly split in to three systems – the RF signal from the beam (dependent on the phase monitors themselves and the hybrids), the LO reference signal and the mixers. An issue with any one of these systems could explain the larger phase jitter measured from Mon 2 (usually connected to Mixer 2 and LO 2) compared to the phase jitter from Mon 1 (usually connected to Mixer 1 and LO 1). To determine whether the issue was with the RF signal of one of the monitors measurements were taken with Mon 1 and Mon 2 swapped between the three mixers.

Figure 3.38 shows the measured phase jitter along the pulse when Mon 2 is moved on to Mixer 1 (and LO 1) and with Mon 1 moved on to Mixer 2 (and LO 2). The colours in all the plots in this section correspond to the mixer that measurement is from – blue for Mixer 1, red for Mixer 2 and green for Mixer 3. With the phase monitors swapped the higher measured phase jitter stays with Mixer 2 with the same ratio of 1.7, in this case $1.48 \pm 0.01^\circ$ on Mixer 2 (connected to Mon 1) and $0.89 \pm 0.01^\circ$ on Mixer 1 (connected to Mon 2). This rules out that the difference in phase jitter is coming from the phase monitors themselves, and suggests the problem is with either Mixer 2 or the reference signal LO 2 for that mixer. The same exercise can then be repeated with Mon 2 moved on to Mixer 3 (and LO 3) and the nominal setup used for Mon 1. Mixer 3 gives similar results to Mixer 2, with 1.7 times larger phase jitter than Mixer 1. This is seen in Figure 3.39. The remaining task is therefore to identify whether the additional source of noise is from Mixer 2 or LO 2, and from Mixer 3 or LO 3.

This was determined by swapping the LO reference signals between the mixers. Fig-

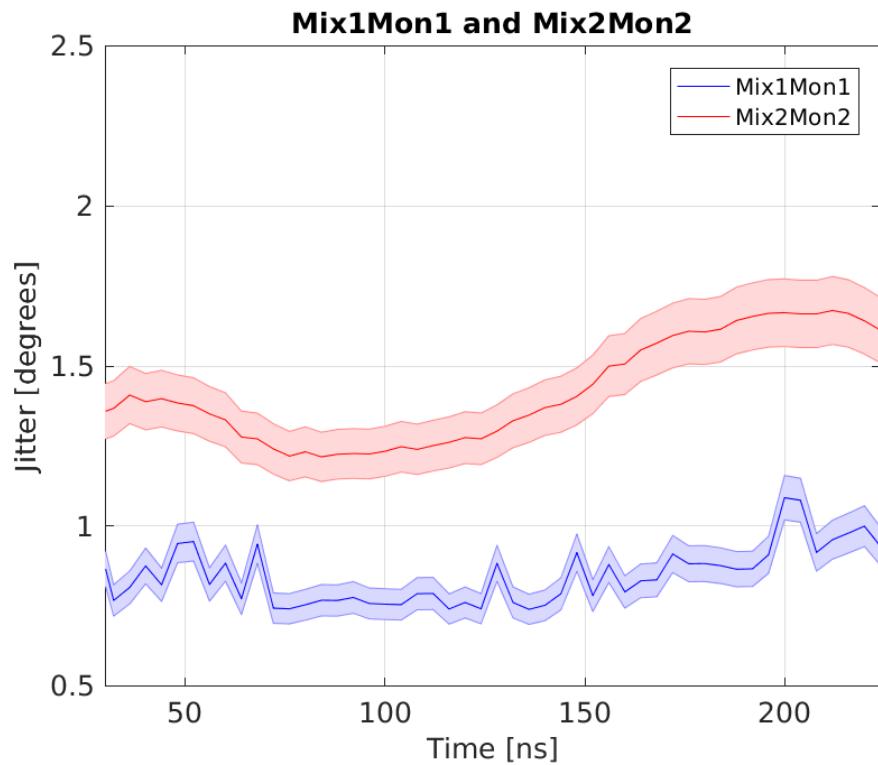


Figure 3.37: Phase jitter along the pulse with the nominal electronics setup – Mon 1 connected to the first mixer, and Mon 2 connected to the second mixer.

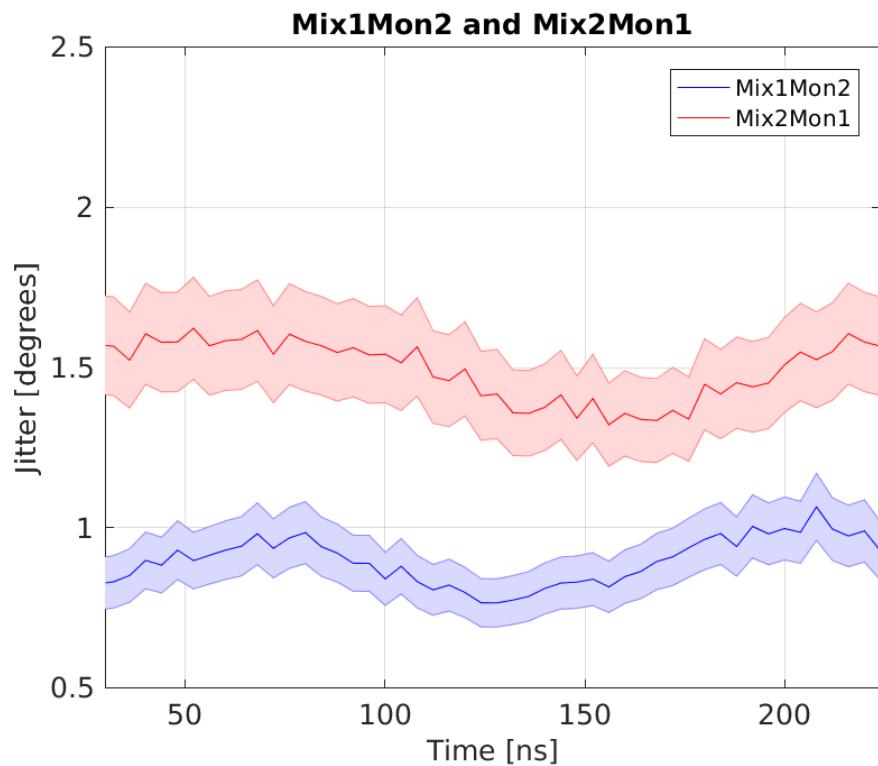


Figure 3.38: Phase jitter along pulse with Mon 2 connected to the first mixer and Mon 1 connected to the second mixer.

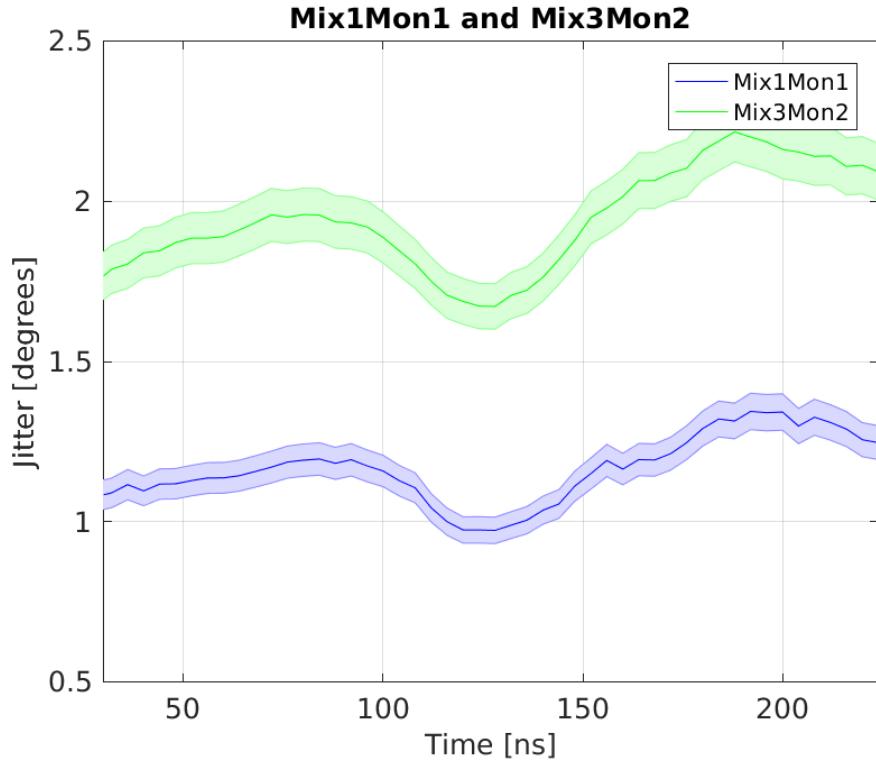


Figure 3.39: Phase jitter along pulse with Mon 1 connected to the first mixer and Mon 2 connected to the third mixer.

Figure 3.40 shows the measured phase jitter with LO 2 connected to Mixer 1, and LO 1 connected to Mixer 2. In this case the lower measured phase jitter stays with LO 1, with $1.10 \pm 0.01^\circ$ jitter on Mixer 2 (with LO 1) and $1.68 \pm 0.01^\circ$ on Mixer 1 (with LO 2). In Figure 3.41 the same is seen with LO 3 connected to Mixer 1, and LO 1 connected to Mixer 3, with the lower phase jitter coming from Mixer 3 (LO 1) in this case. In Table 3.10 it can also be seen that the lowest phase jitter in every dataset always comes from the phase monitor for which LO 1 was used.

In this way an issue with LO 2 and LO 3 was identified. The generation of the LO reference signals (Section 3.2) consists of a 3 GHz source that is common to all three signals, plus separate phase shifters, frequency multipliers (to create the 12 GHz reference) and amplifiers for each LO signal. The digital phase shifters were thought to be the most likely culprit to be adding noise in to the system, and preliminary tests replacing one of the digital phase shifters with a mechanical alternative provided an immediate gain in resolution. Finally, all three digital shifters were replaced with mechanical phase shifters (the specifications of the shifters were introduced in Section 3.2). With this setup the measured phase jitter for Mon 1 and Mon 2 is the same irrespective of which mixer or LO the signal is connected to, as shown in Figure 3.42. The final achieved resolution is presented in the next section.

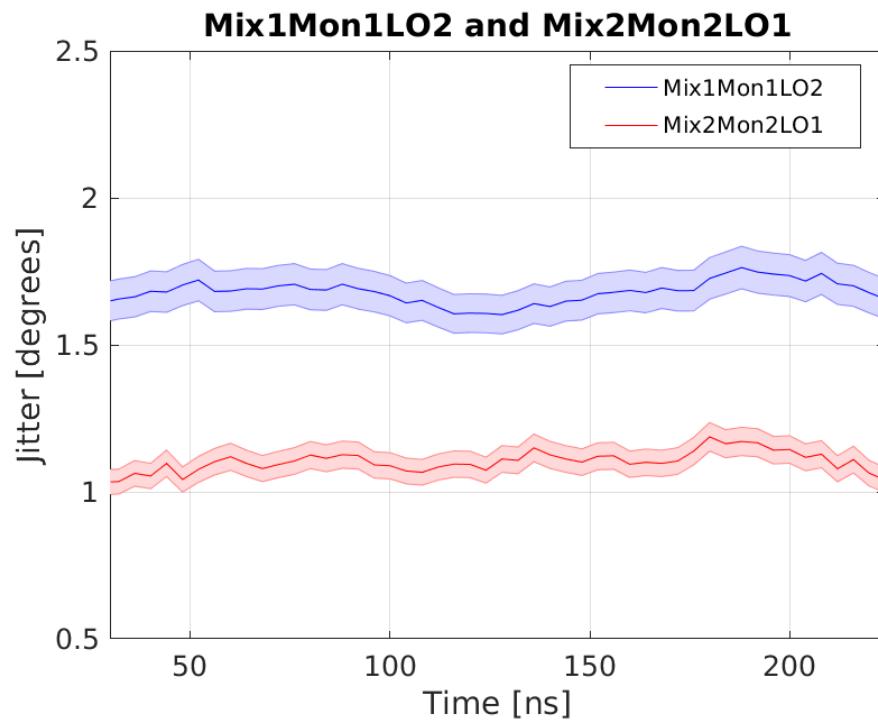


Figure 3.40: Phase jitter along pulse with the LO swapped between the first mixer and the second mixer.

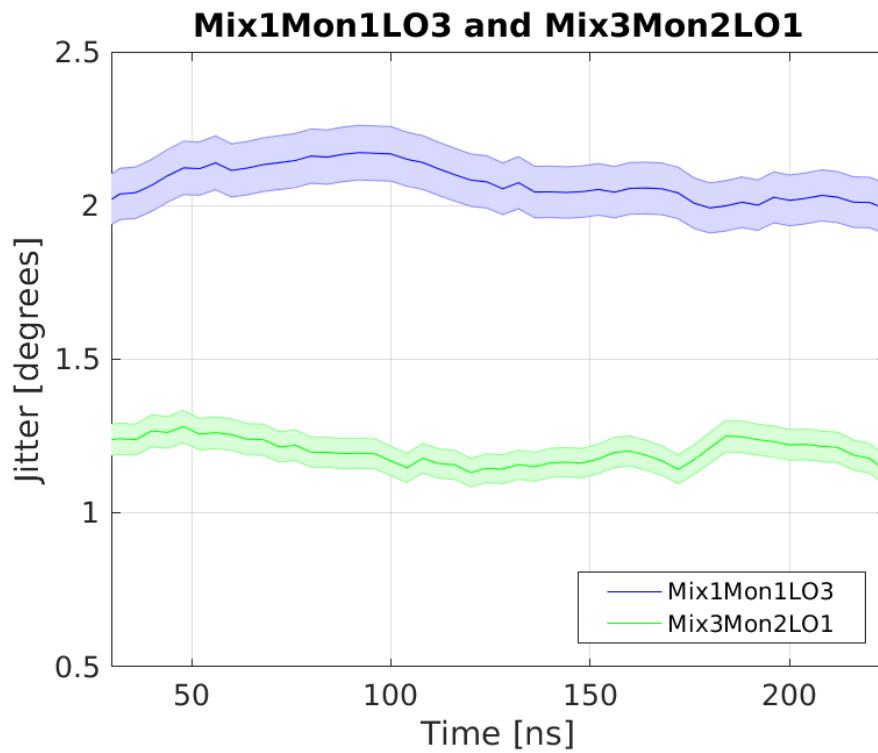


Figure 3.41: Phase jitter along pulse with the LO swapped between the third mixer and the first mixer.

Mon 1			Mon 2		
Mixer	LO	Jitter	Mixer	LO	Jitter
1	1	$0.83 \pm 0.01^\circ$	2	2	$1.38 \pm 0.01^\circ$
2	2	$1.48 \pm 0.02^\circ$	1	1	$0.89 \pm 0.01^\circ$
1	1	$1.15 \pm 0.01^\circ$	3	3	$1.91 \pm 0.01^\circ$
1	2	$1.68 \pm 0.01^\circ$	2	1	$1.10 \pm 0.01^\circ$
1	3	$2.07 \pm 0.01^\circ$	3	1	$1.20 \pm 0.01^\circ$

Table 3.10: Comparison of phase jitter along the pulse for each measurement with different setups of the electronics. Each row corresponds to the results of one dataset. The left hand side of the table shows the results from Mon 1 in that dataset, and the right hand side of the table the results from Mon 2. Bold text indicates the lower jitter value in that dataset, all of which use LO 1.




Figure 3.42: Phase jitter along pulse after installation of the mechanical phase shifters.

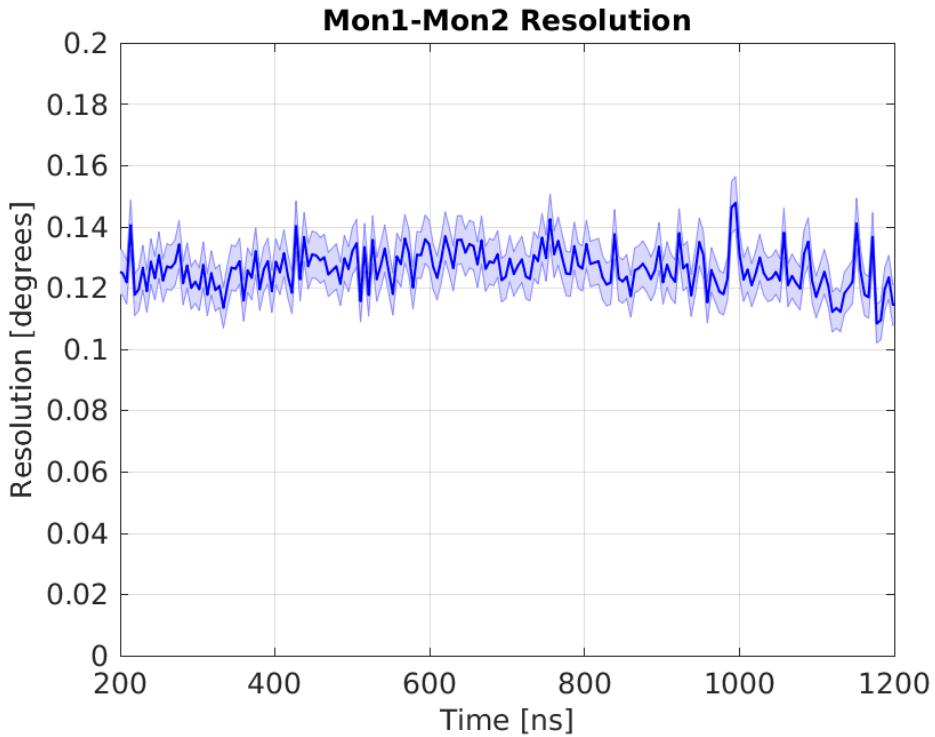


Figure 3.43: bestResolution

3.9 Resolution Measurements

This section presents resolution measurements with the mechanical phase shifters in place and the same overall electronics setup that was used to achieve the best PFF results in Chapter 6. All measurements use the upstream phase monitors Mon 1 and Mon 2. There is no way to directly verify the resolution of the downstream phase monitor Mon 3, but it should match that of Mon 1 and Mon 2 and the achieved phase stabilisation in Chapter 6 suggests it does have below 0.2° resolution.

3.9.1 Best Resolution

Figure 3.43 shows the best phase resolution that has been achieved to date sample by sample along the pulse. Each point corresponds to the phase jitter in the difference between the measured Mon 1 and Mon 2 phase divided by $\sqrt{2}$ as per Equation 3.27. The resolution is quite stable along the pulse, with a mean value of $0.1257 \pm 0.0005^\circ$. This is below the 0.14° needed to be able to theoretically achieve a 0.2° correction with the PFF prototype, as derived in Section 3.3. The achieved resolution corresponds to a theoretical limit of around 0.18° in the corrected downstream phase jitter.

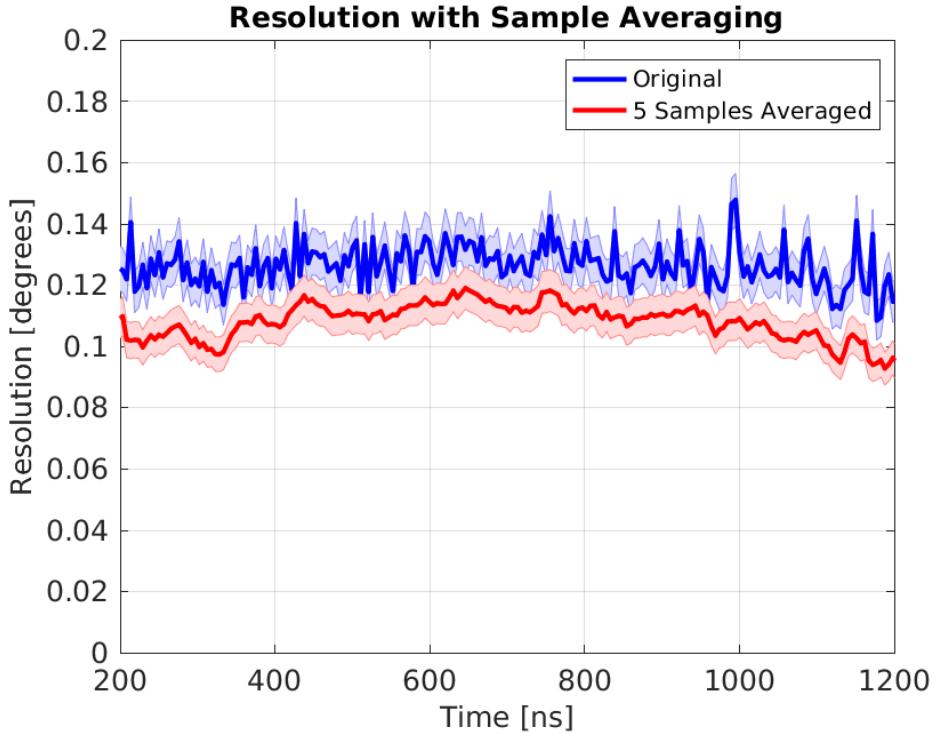


Figure 3.44: Effect on resolution by averaging samples.

3.9.2 With Sample Averaging

The phase monitor signals are digitised with a much higher sampling rate, 357 MHz on the FONT5a board or 192 MHz on the SiS digitisers, than the PFF correction bandwidth of around 30 MHz. This means the results from several samples could be averaged to reduce noise on the measurement whilst maintaining the same correction bandwidth. For example, for a measurement on the SiS digitisers, such as those shown here, 5 samples can be averaged to effectively create a signal with 38.4 MHz sampling rate, reduced from the initial 192 MHz. The effect of doing this is shown in Figure 3.44. The resolution is decreased from $0.1257 \pm 0.005^\circ$ as seen previously to $0.1077 \pm 0.0005^\circ$ with 5 samples averaged. Currently the phase propagation (Chapter 4) is the main limiting factor for the PFF performance rather than the phase monitor resolution. However, it would be possible to implement sample averaging in the PFF firmware on the FONT5a board if deemed necessary for future tests [REF].

3.9.3 Dependence of Resolution on LO Phase

The process of setting the mixers on their zero crossing after calibrations was documented in Section 3.7.4. This is necessary to operate the mixers where there output voltage is most sensitive to variations in the phase (close to the peaks of the mixer output there is negligible change in output voltage for small changes in phase). The dependence of the phase resolution on the beam phase across the full $\pm 180^\circ$ range is shown in Figure 3.45. The plotted phase is the offset between the LO phase shifter setting and the calculated optimal setting. More

than 50° away from the zero crossing there is a large degradation in resolution, reaching above 1 degree at the $+90^\circ$ peak.

However, for the PFF system the phase resolution only needs to be guaranteed within its correction range, which is close to $\pm 6^\circ$ as seen in Section 5.4.1. In Figure 3.46 it is seen that there is no noticeable degradation in resolution in the range between $\pm 15^\circ$.

3.10 Bandwidth

The overall targeted bandwidth for the prototype PFF correction is 30 MHz, which means each individual hardware component must have a bandwidth in excess of 30 MHz. To estimate the bandwidth of the phase measurement the delay loop has been used to create a sharp jump in the phase mid-way along the pulse. Normally the delay loop (whose length is 140 ns) is used with 1.5 GHz beam at CTF3, with seven 180 degree phase switches in the bunching at 140 ns intervals. Alternating 140 ns sub-trains are deflected in to the delay loop with a 1.5 GHz RF deflector, and then recombine with the 140 ns sub-trains bypassing the delay loop to double the drive beam current [REF]. If the delay loop is instead used with 3 GHz beam alternating bunches, rather than sub-trains, enter the delay loop. The first 140 ns of the pulse after the delay loop contains only bunches that bypassed the loop, so is a 1.5 GHz beam with half the initial beam current. After the first 140 ns there is a fast transition back to 3 GHz beam with the same initial current as delayed bunches merge with the bypassing bunches. By varying the length of the delay loop (using a wiggler [REF]) the phase of the bunches leaving the delay loop can be changed, also creating a fast phase jump at the transition to 3 GHz beam.

Figure 3.47 shows the response of the Mon 3 mixer to a beam pulse with a phase step setup in this way. The first 140 ns of the pulse contains only bunches bypassing the delay loop, as described, followed by the fast transition to 3 GHz beam and the phase step. The time taken for the mixer output to respond to the fast phase change can be related to the bandwidth by this approximate relationship [REF]:

$$\text{BW} = \frac{350}{t} \quad (3.40)$$

Where BW is the bandwidth in MHz and t is the rise time of the signal in ns. The transition between the two phase states as seen on the mixer output occurs between 156.6 ns and 167.0 ns, a rise time of 10.4 ns corresponding to a bandwidth of approximately 34 MHz. In addition, the falling edge of the mixer output at the end of the pulse also takes around a 10 ns, also leading to a bandwidth estimate in the region of 35 MHz.

3.11 Comparison of Measured Phase Along Pulse

It has been shown that the phase jitter is at the same level in each monitor after the installation of mechanical phase shifters, and that the achieved resolution is better than the level needed to be able to theoretically achieve a 0.2° PFF correction. Nevertheless there are




Figure 3.45: Resolution.




Figure 3.46: Resolution.

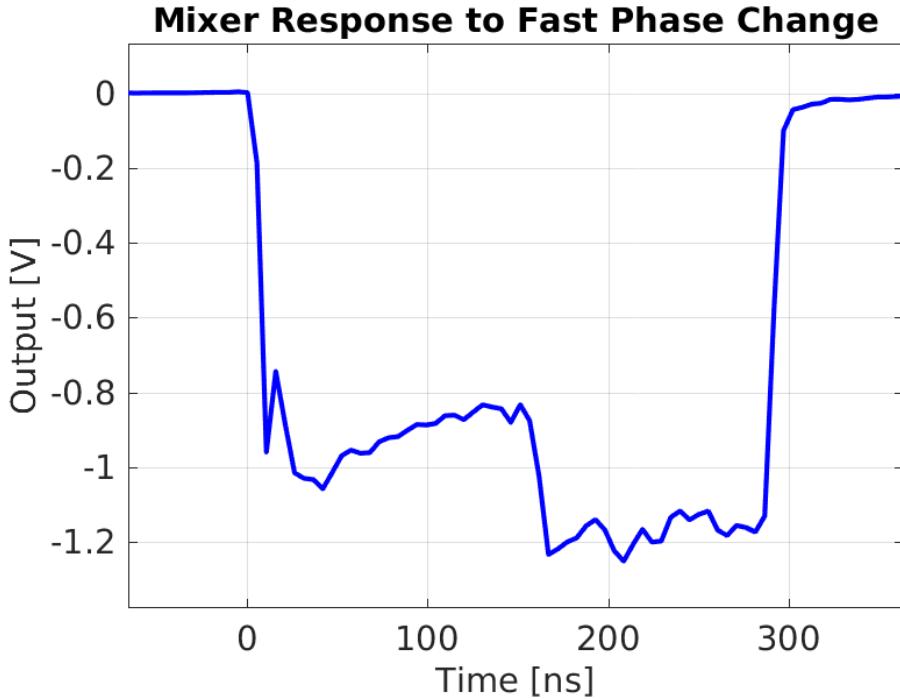


Figure 3.47: Response of Mixer 3 to a jump in phase in the middle of the pulse.

several remaining effects that have been identified which may limit the PFF performance in a way not described by the resolution. These are the focus of the remainder of this chapter, starting with differences between the measured shape of the phase along the pulse here. Neither the phase jitter nor the resolution are sensitive to any static differences in the measured phase in each monitor. For the PFF system to remove not only phase jitter but also to flatten the phase sag along the pulse the shape of the upstream and downstream phase must be the same.

Figure 3.48 shows an example of the measured phase along the pulse in all three monitors, as well as a fourth phase measurement taken from one of the power extraction and transfer structures (PETS) in the TBL line after Mon 3 [REF]. The PETS measurement provides an independent cross-check of the response and calibration of the three PFF monitors. All four phase measurements show approximately the same overall phase sag of around 45° along the full pulse length, as desired.

Nevertheless, there are differences between the measurements. As described later in Chapter 4 there are many mechanisms by which the beam phase can change between the upstream and downstream monitors. Small discrepancies between the Mon 1 and Mon 2 phase compared to the Mon 3 and PETS phase are therefore not unexpected. However, as Mon 1 and Mon 2 are neighbouring each other in the beam line their measurements are expected to agree almost perfectly but this is not the case. At the start of the pulse the Mon 2 phase is around 6° lower than the Mon 1 phase, whereas at the end of the pulse the Mon 2 phase is around 2° higher than the Mon 1 phase, for example.

Two possible explanations could have been an error in the fitted calibration amplitude or a time offset between the two signals. The fact that the Mon 2 phase is lower at one end

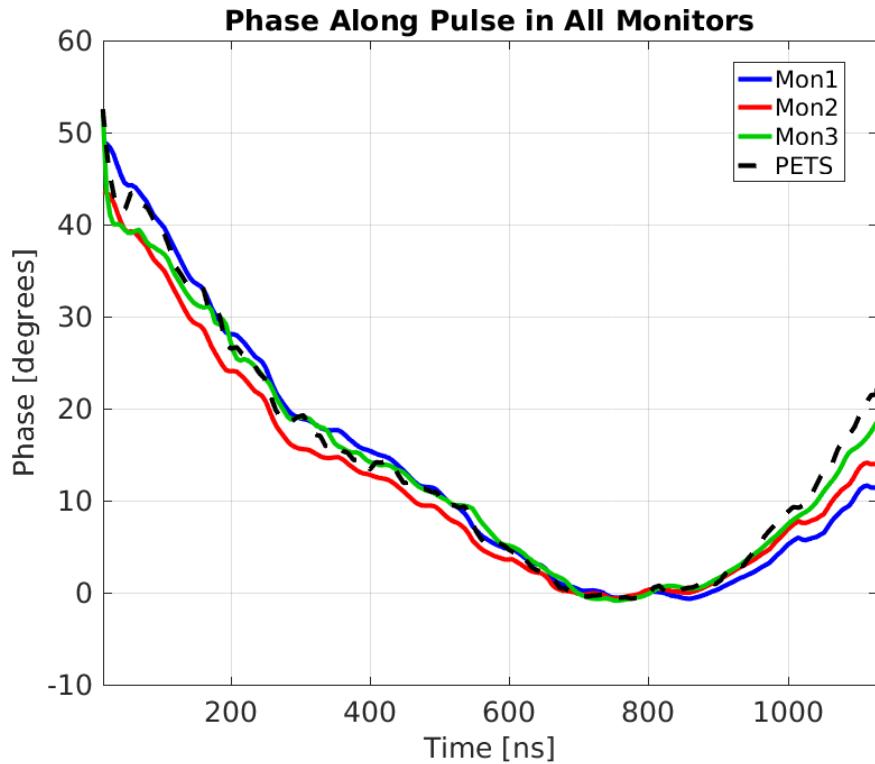


Figure 3.48: Comparison of phase along the pulse in the three PFF phase monitors and an alternative phase measurement from a PETS.

of the pulse but higher at the other means that the difference is not simply a scale factor between the two, and therefore cannot be described by an error in calibration amplitude. Neither can it be explained by differences in the calibration constants along the pulse, as seen in the next section. Similarly, by observing small variations along the pulse, for example at the end of the pulse after a time of 1000 ns in the figure, it can be seen that the two measurements are well aligned in time.

The blue line in Figure 3.50 shows the difference between the Mon 1 and Mon 2 phase along the pulse. The offset between the two has a clear linear dependence on the time along the pulse. As a result, one way in which the agreement between the two measurements can be improved is by rotating Mon 2 with respect to Mon 1 about a time of 850 ns where the offset between the two is zero. With an optimal rotation of 2.1° the mean absolute offset between Mon 1 and Mon 2 between 370 ns and 1080 ns can be reduced from $1.44 \pm 0.08^\circ$ to $0.12 \pm 0.01^\circ$. The red line in the figure shows the difference along the pulse after performing this rotation. Figure 3.48 then shows a comparison between the mean phase along the pulse in Mon 1 and Mon 2 after the rotation, with the original measurement from Mon 2 also shown for reference. The agreement in shape between the two is now excellent, with only a slight discrepancy at the end of the pulse.

The source of this apparent rotation between Mon 1 and Mon 2 is not clear and could be an effect of the monitors themselves, the calibration process or a real difference between the beam phase in each monitor. If there was a real rotation in the phase sag between monitors or if there was an issue with Mon 1, used as the PFF correction input, it would

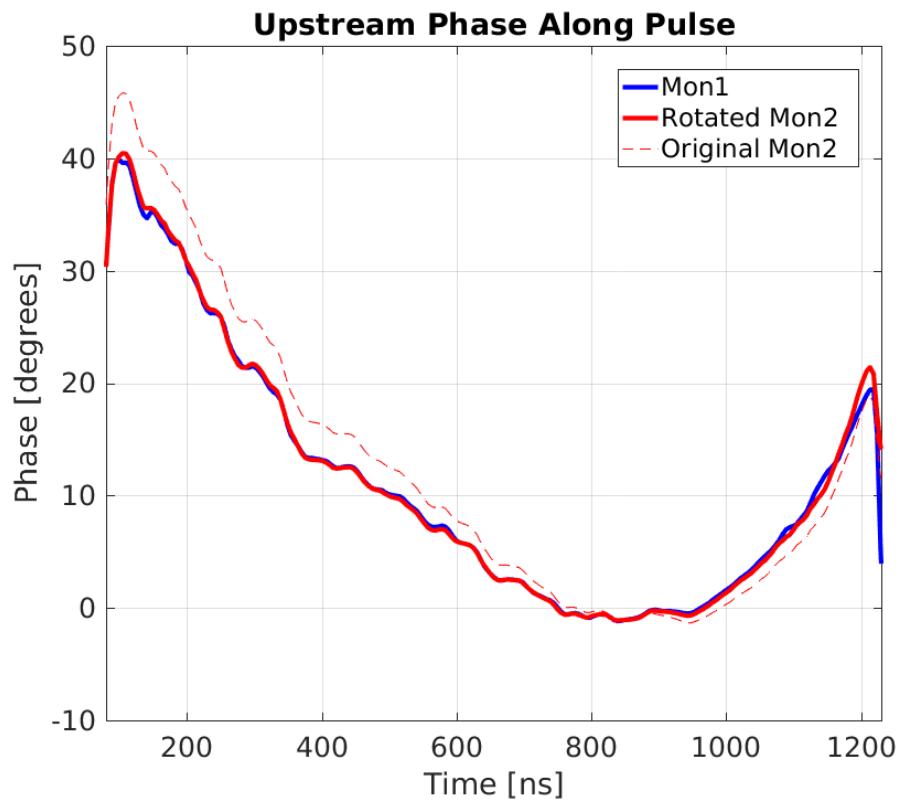


Figure 3.49: Comparison of Mon 1 and Mon 2 phase along the pulse, with Mon 2 rotated by 2.1° about a time of 850 ns on the horizontal axis.

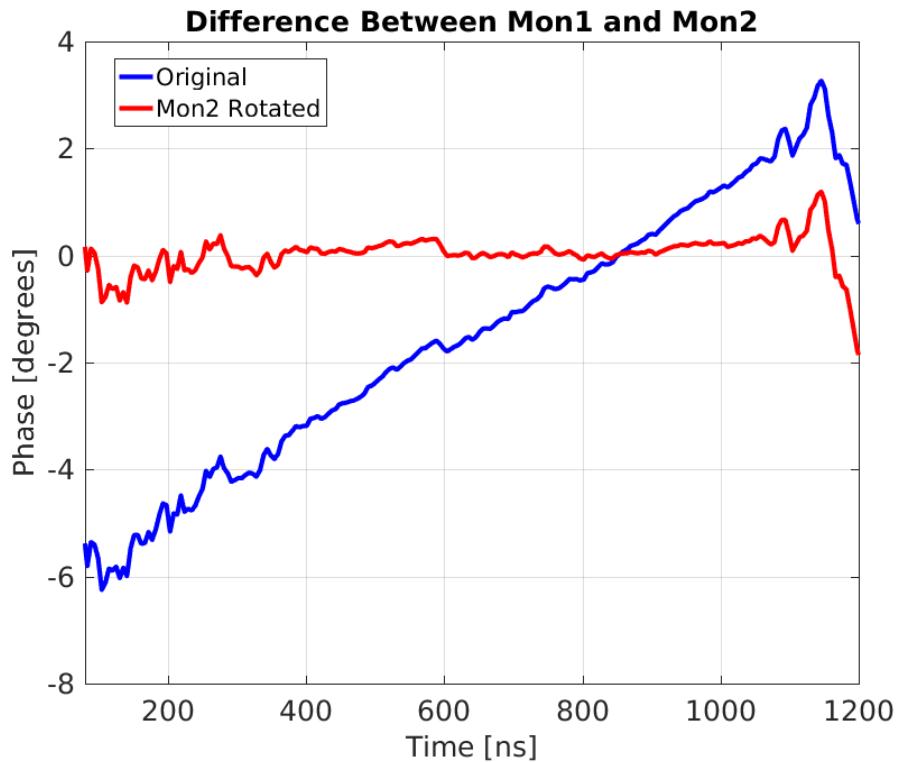


Figure 3.50: Difference between Mon 1 and Mon 2 phase along the pulse with and without Mon 2 rotated.

not be possible to completely flatten the phase sag along the pulse with the PFF system. The excellent flattening of the pulse that has in reality been achieved with the PFF system (Chapter 6) suggests this may be an issue affecting only Mon 2, as the correction input is based solely on Mon 1. An issue with the measured pulse shape from Mon 2 may also help to explain differences between the shape of the applied kick and the upstream phase (measured by Mon 2) seen later in Section 5.4.3.

3.12 Effect of Variations in Calibration Constant

In Section 3.7.3 it was shown how the calibration parameters vary along the pulse. In the PFF algorithm, and usually also in data analysis, single values for each calibration constant are used across the full pulse length. This will lead to small discrepancies between the measured phase and the true beam phase. Examples of the size of this effect are shown in Figures 3.51, for Mon 2, and in Figure 3.52 for Mon 3, comparing the measured phase along the pulse when using single value calibration constants or varying calibration constants along the pulse. In the range between 600 ns and 1000 ns around the middle of the pulse where the phase sag is flattest (the part of interest for the PFF system) using a varying calibration constant along the pulse makes almost no difference to the measured upstream phase. This is unsurprising as the calibration constants normally used are calculated from one sample in this range. Outside this range the difference between the two reaches a peak of 1.5° near the start of the pulse.

For the downstream phase the differences are larger, including up to 1.5° offsets between the two methods in the range of interest for the PFF system. However, the example shown is without fully optimised downstream beam conditions achieved in Chapter 4 and so represents a worst case scenario for normal PFF operation. The PFF correction quality also does not directly depend on the downstream phase measurement (apart from for the gain calculation, as derived in Section 4.1), so these differences can be removed in offline analysis if deemed necessary.

3.13 Dependence on Position

As discussed in Section 3.1 the phase monitor output from the two vertical RF feedthroughs are summed in hybrids, and it is this sum signal that is connected to the mixers. This reduces, but does not remove completely, the dependence of the phase monitor signal power on the beam position in the monitor. In this section the remaining position dependence of the phase measurement is determined. Understanding the magnitude of the effect is particularly important considering the diodes are no longer used to power normalise the mixer output as originally intended.

The magnetic corrector placed roughly 1 m prior to Mon 1 (labelled CT.DHD0360) has been used to scan the beam position in the upstream phase monitors both horizontally and vertically. Around 1 m following Mon 2 there is a quadrupole (CT.QFF0420) followed by a

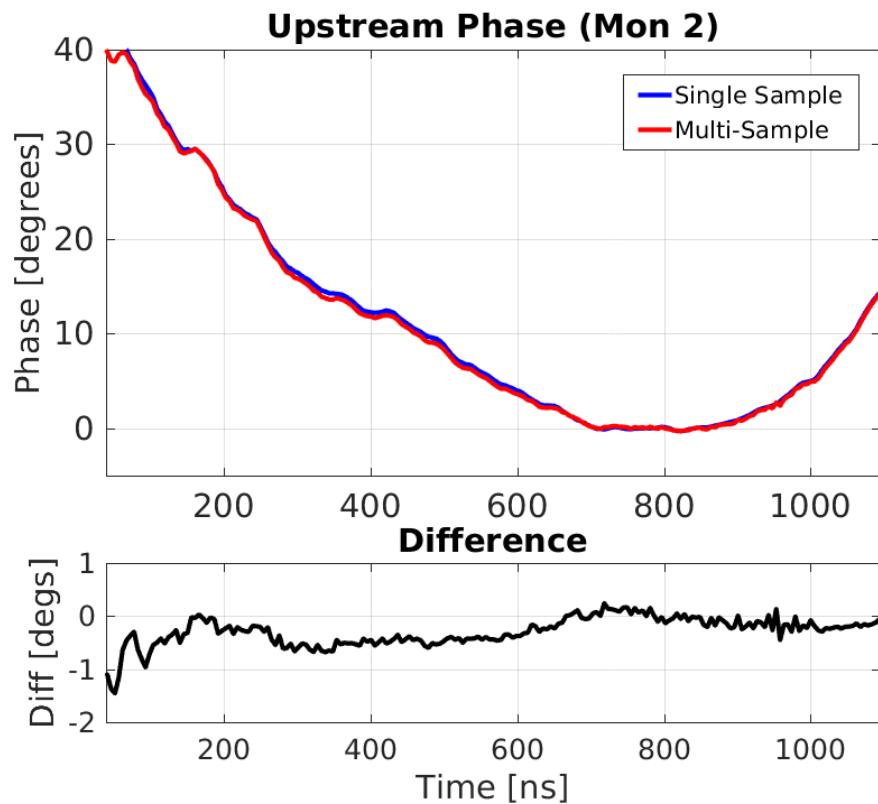


Figure 3.51: Effect of using a varying calibration constant on the upstream phase.

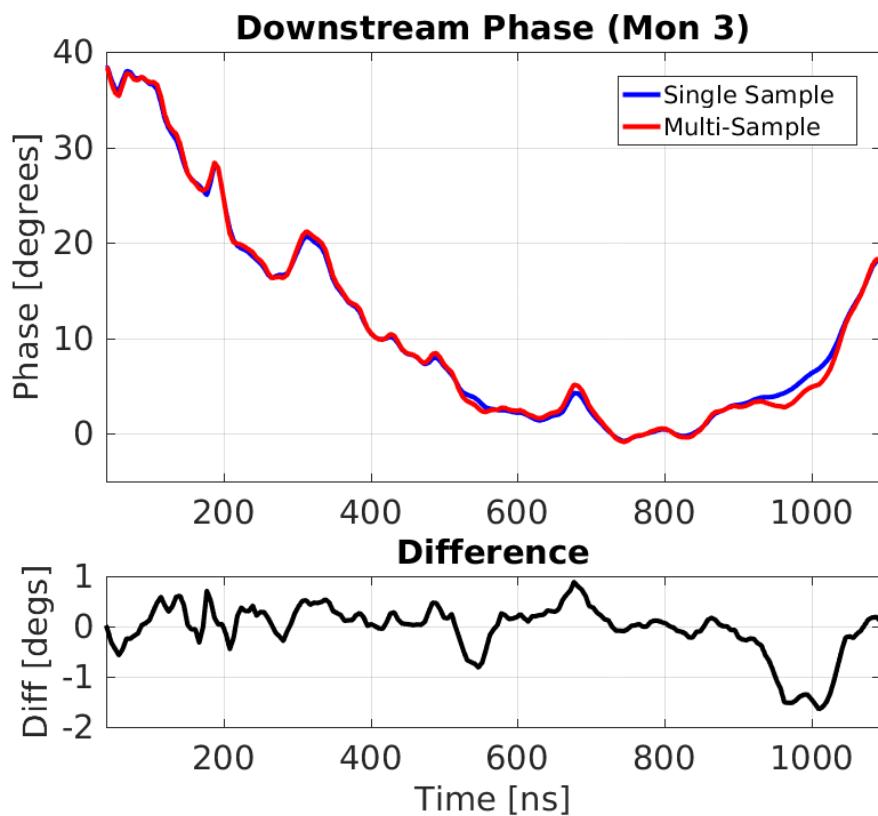


Figure 3.52: Effect of using a varying calibration constant on the downstream phase.

Device	Distance From CT.0360	Label
Mon 1	104.5 cm	s_{M1}
Mon 2	144.5 cm	s_{M2}
CT.BPM0430	357.0 cm	s_{430}

Table 3.11: Distance of the upstream phase monitors and following BPM CT.0430 to the corrector CT.0360 before the phase monitors.

beam position monitor or BPM (CT.BPM0430). For the purposes of this measurement the quadrupole is turned off so that the whole 3.5 m length of beam line between the corrector and the BPM is a drift space. It is then straightforward to reconstruct the position offset in the two phase monitors based on the BPM measurement without having to rely on the MADX model to take in to account the effect of the quadrupole on the beam orbit. The position of the corrector with respect to the phase monitors can be seen in Figure 3.4. Distances between the elements are shown in Table 3.11.

In these conditions a position offset of x_{430} measured in the BPM (in horizontal or vertical) corresponds to the following position offsets, x_{M1} and x_{M2} , in each phase monitor:

$$x_{M1} = r_{M1}x_{430} \quad (3.41)$$

$$x_{M2} = r_{M2}x_{430} \quad (3.42)$$

Where r_{M1} and r_{M2} are the position offsets in the phase monitors per unit position offset in the BPM, given by:

$$r_{M1} = \frac{s_{M1}}{s_{430}} \quad (3.43)$$

$$r_{M2} = \frac{s_{M2}}{s_{430}} \quad (3.44)$$

Where s_{M1} , s_{M2} and s_{430} are the distances between the corrector and the phase monitors (s_{M1} , s_{M2}) or BPM (s_{430}). Substituting in the values from Table 3.11 gives values of $r_{M1} = 0.29$ and $r_{M2} = 0.41$.

Figure 3.53 shows the results of a horizontal position scan in the upstream phase monitors, with the phase plotted against the horizontal position in the BPM. Position scans of this type require a long time to complete in order to be able to acquire enough statistics at multiple corrector settings. The scan presented represents a data taking period of around one hour, for example. This means real drifts in beam phase (aside from any position dependent effects on the monitor signals) are unavoidable during the scan. The first impression of Figure 3.53 is that the only visible change in measured phase is due to actual drifts in beam phase rather than being as a result of the changes in position. However, Figure 3.54 shows the response of the phase monitor diode signals to the horizontal position in the BPM during the scan. Even though the diodes are heavily saturated a dependence of their output, and so a dependence of power, on the position is visible. This suggests there should also be a visible effect on the phase.

To remove the effects of drifts in the beam phase the difference between the measured Mon 1 and Mon 2 phase can be considered instead. The true beam phase, ϕ_b , should be

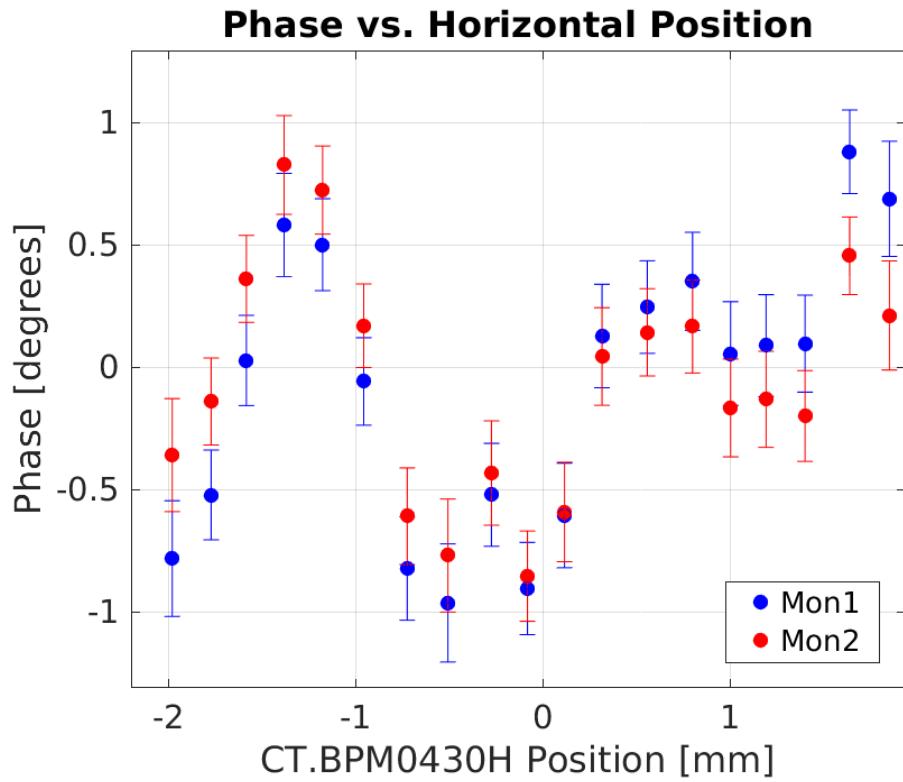


Figure 3.53: Mon 1 and Mon 2 phase dependence on horizontal position during scan.

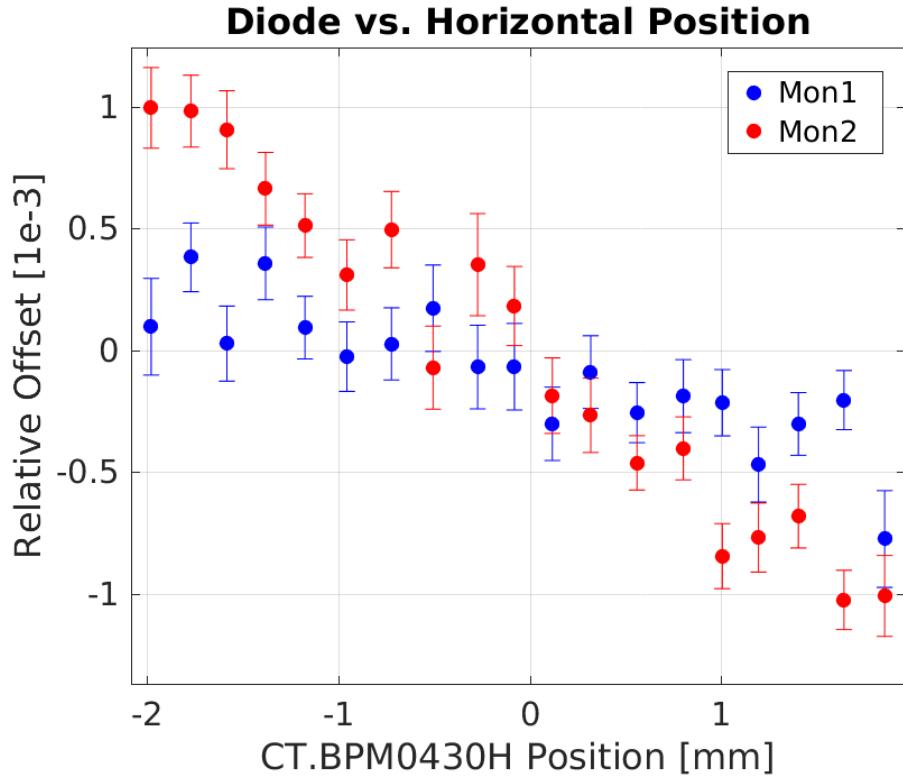


Figure 3.54: Mon 1 and Mon 2 diode dependence on horizontal position during scan.

identical in Mon 1 and Mon 2, so any differences between their measurements ϕ_{M1} and ϕ_{M2} during the scan can be modelled as coming from the beam position in the monitor, as follows:

$$\phi_{M1} = \phi_b + c_{M1}x_{M1} \quad (3.45)$$

$$\phi_{M2} = \phi_b + c_{M2}x_{M2} \quad (3.46)$$

Where c_{M1} and c_{M2} are constants expressing the phase shift per unit position offset in each monitor. The difference between the measured phase in Mon 1 and Mon 2 is therefore given by:

$$\phi_{M2} - \phi_{M1} = c_{M2}x_{M2} - c_{M1}x_{M1} \quad (3.47)$$

c_{M1} and c_{M2} cannot be determined individually from the scan results alone using this approach, so instead it will be approximated that the strength of the position dependence is the same in Mon 1 and Mon 2. Letting $c_{M1} = c_{M2} = c$ finally gives:

$$\phi_{M2} - \phi_{M1} = c(x_{M2} - x_{M1}) \quad (3.48)$$

$$c = \frac{\phi_{M2} - \phi_{M1}}{x_{M2} - x_{M1}} \quad (3.49)$$

The expressions for x_{M1} and x_{M2} in terms of the BPM position x_{430} derived earlier can then be substituted in to give:

$$c = \frac{\phi_{M2} - \phi_{M1}}{x_{430}(r_{M2} - r_{M1})} \quad (3.50)$$

The gradient of the phase difference versus the position offset in the BPM, $(\phi_{M2} - \phi_{M1}) / x_{430}$, is the only remaining parameter left to calculate in order to estimate c .

Figure 3.55 shows how the difference between the Mon 1 and Mon 2 phase measurement depends on the horizontal position in the BPM during the scan. By removing the actual beam phase drifts the dependence of the phase measurements on the beam position becomes clear. The fitted gradient is $-0.22 \pm 0.01^\circ$ per mm offset in the BPM. Using Equation 3.50 this corresponds to a phase shift of $-1.84 \pm 0.07^\circ$ per mm offset in the phase monitors themselves.

Exactly the same process can be repeated in the vertical plane, and the results of doing this are shown in Figure 3.56. Although the effect is smaller there is still a visible dependence of the phase on the vertical position, in this case with a gradient of $0.06 \pm 0.01^\circ$ per mm offset in the BPM corresponding to $0.53 \pm 0.07^\circ$ per mm vertical offset in the phase monitors.

The position dependence of the measurement appears large when quoted like this but it must be remembered that millimetre scale changes in beam orbit are rare during normal operation. Orbit jitter at CTF3 around the location of the upstream phase monitors is typically at the 0.02 mm level [REF]. Taking the calculated position dependence of Mon 1 and Mon 2 this corresponds to only an additional measured phase jitter of roughly 0.04° , using the larger dependence in the horizontal plane. This is only a small contribution to the overall phase monitor resolution of around 0.13° as seen in Section 3.9. However, it is important to consider that any change in beam setup that alters the trajectory through the monitors will lead to a change in output power and therefore a change in calibration constants. Calibrations are always repeated when the CTF3 setup is changed to take this in

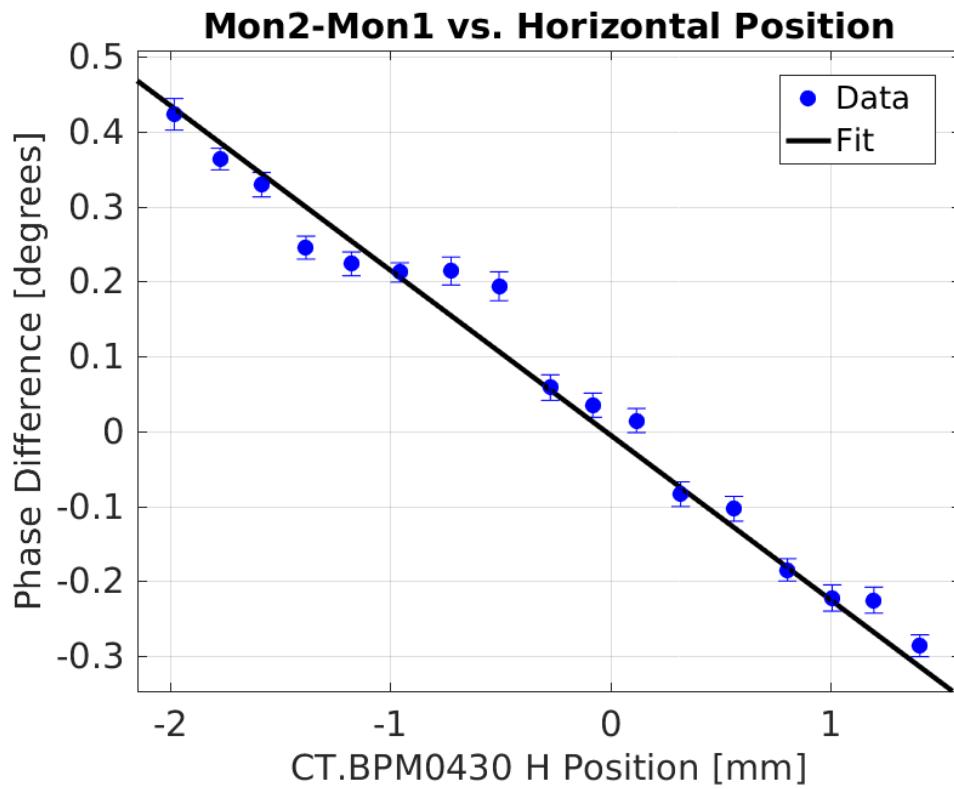


Figure 3.55: Fit to difference between Mon 1 and Mon 2 phase versus horizontal position.

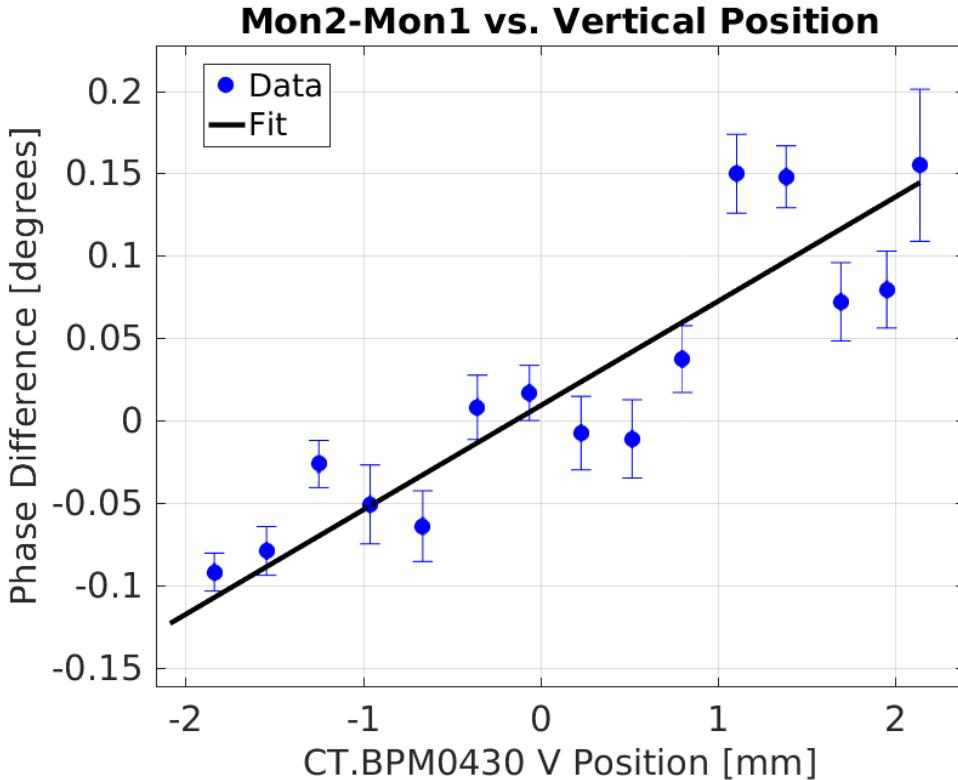


Figure 3.56: Mon 1 and Mon 2 dependence on horizontal position during scan.

to account. Position scans around the region of the downstream phase monitor Mon 3 would be much more difficult to perform and have not yet been attempted. The effect should be of a similar magnitude, although the downstream beam is less stable than upstream so the contribution to measured phase jitter is likely to be slightly larger.

Chapter 4

Phase Propagation

The PFF system uses the upstream phase measurement to correct the downstream phase. In this scheme any differences between the initial upstream phase and the downstream phase cannot be removed by the PFF system. The PFF system is therefore very sensitive to the “phase propagation”, or the extent to which the upstream phase predicts the actual downstream phase. As at CTF3 the upstream and downstream phase monitors are separated by roughly 150 m of beam line there are many potential sources that may change the downstream phase with respect to the upstream phase. This chapter derives the requirements that the PFF system places on the phase propagation before describing its original status during the first PFF tests and the extensive work that has been needed to improve it.

4.1 Feedforward Algorithm

In the PFF system the voltage sent to the kickers in the TL2 chicane is varied depending on the upstream phase. The corrected downstream phase, ϕ_{PFF} , can therefore be simply modelled as subtracting the upstream phase, ϕ_u , from the initial downstream phase ϕ_d :

$$\phi_{PFF} = \phi_d - g\phi_u \quad (4.1)$$

Where g is the applied correction gain. The corrected downstream phase jitter, σ_{PFF} , can then be defined using the standard result of subtracting correlated variances [REF]:

$$\sigma_{PFF}^2 = \sigma_d^2 + g^2\sigma_u^2 - 2g\rho_{ud}\sigma_u\sigma_d \quad (4.2)$$

Where ρ_{ud} is the correlation coefficient between the upstream phase and the downstream phase.

Setting the partial differential of Equation 4.2 with respect to the gain equal to zero yields an expression for the theoretical optimal gain value to apply:

$$\frac{\partial\sigma_{PFF}^2}{\partial g} = 2g\sigma_u^2 - 2\rho_{ud}\sigma_u\sigma_d = 0 \quad (4.3)$$

$$g = \rho_{ud} \left(\frac{\sigma_d}{\sigma_u} \right) \quad (4.4)$$

In the case where the phase propagation is perfect and the downstream phase is identical to the upstream phase the optimal gain factor is 1, as expected. Alternatively, if there is no correlation between the upstream and downstream phase the PFF system could only act to increase the downstream phase jitter, thus the optimal gain would be zero. If the downstream phase jitter is amplified with respect to the upstream phase jitter ($\phi_d = \text{const} \times \phi_u$), this can be removed by the PFF system via the dependence of the optimal gain on the upstream–downstream phase jitter ratio.

Substituting the optimal gain value back in to Equation 4.2 gives an expression for the theoretical limit on the corrected downstream phase jitter using the PFF system:

$$\sigma_{PFF} = \sigma_d \sqrt{1 - \rho_{ud}^2} \quad (4.5)$$

With the optimal PFF setup the achievable corrected downstream phase jitter has no dependence on the upstream phase jitter. It depends only on the initial downstream phase jitter and most strongly on the upstream–downstream phase correlation. If non-optimal gains are used there is also a dependence on the upstream phase jitter and the set gain value as per Equation 4.2. Figure 4.1 shows how the achievable corrected downstream phase jitter depends on the upstream–downstream phase correlation. A factor 2 reduction in the initial downstream phase jitter requires a correlation of 86.6%, for example.

Equations 4.2, 4.4 and 4.5 are used extensively in the remainder of this thesis to determine the beam conditions needed to achieve a 0.2° correction at CTF3, as well as to calculate the expected effect of the PFF system given the beam conditions and PFF setup.

4.2 Characteristics of Uncorrected Phase Jitter

This section summarises the status of the phase propagation during the first PFF tests to demonstrate why the work in this chapter was necessary and to provide a point of comparison to the improved conditions later achieved in Section 4.6.

Mean Phase

Figure 4.2 compares the mean phase upstream and downstream across one dataset of 213 pulses, or roughly 5 minutes. The downstream phase jitter is more than double the upstream phase jitter — at $1.86 \pm 0.09^\circ$ compared to $0.81 \pm 0.04^\circ$ upstream. Theoretically the PFF system could still remove the amplified downstream phase jitter providing the jitter is still well correlated with the upstream phase, as derived in the previous section. Unfortunately, as shown in Figure 4.3, there is only very weak correlation between the upstream and downstream phase, 0.34 ± 0.06 in this case.

Using Equation 4.5, with the optimal PFF gain applied an initial downstream phase jitter of $1.86 \pm 0.09^\circ$ could be reduced to $1.51 \pm 0.09^\circ$ with the PFF system in these conditions with 0.34 ± 0.06 upstream–downstream phase correlation. This is only a modest improvement of around 20%, and far from the CLIC target of 0.2° phase stability. Also, the limitations of

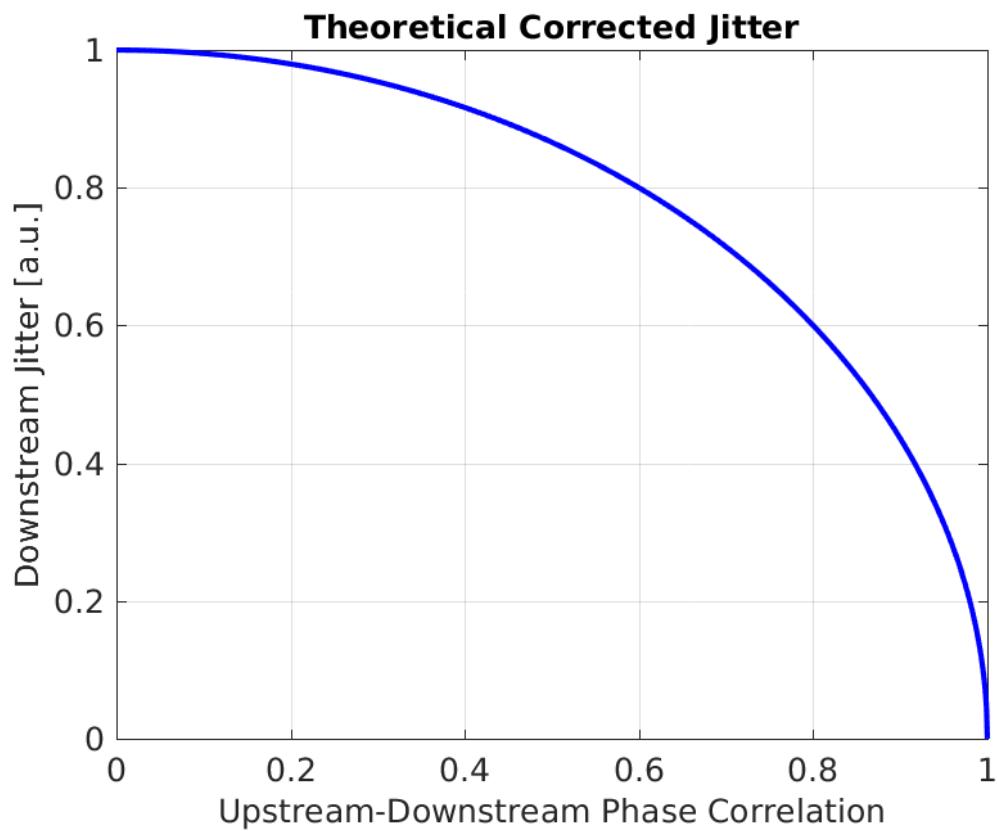


Figure 4.1: Theoretical corrected downstream phase jitter in optimal conditions with an initial downstream jitter of 1 in arbitrary units.

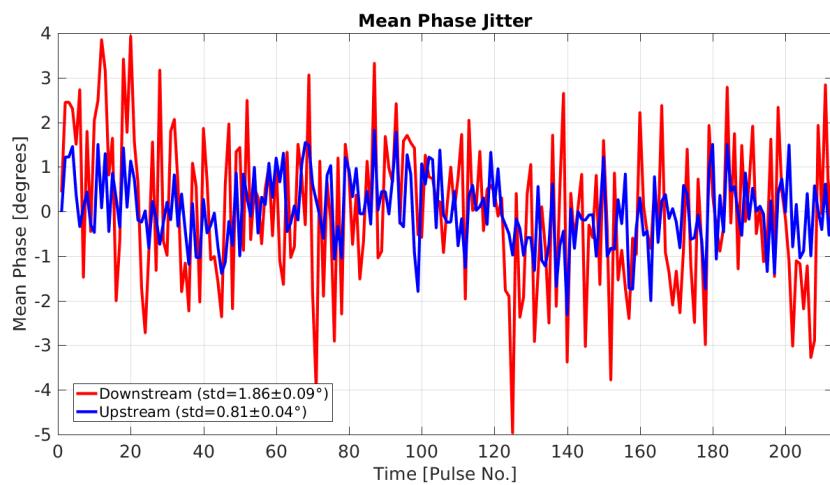


Figure 4.2: Mean phase jitter upstream and downstream during early PFF tests.




Figure 4.3: Correlation between the downstream phase and the upstream phase during early PFF tests.

the actual PFF system, such as in the correction range (Section 5.4.1), means that in reality even this 20% reduction would be difficult to achieve in reality.

Assuming the downstream phase jitter could be reduced to the same level as the upstream phase jitter, at 0.8° , the upstream–downstream phase correlation required to be able to make 0.2° corrected phase jitter theoretically possible is 97%. The target of this chapter is to achieve both this correlation and downstream phase jitter.

Phase Along Pulse

In addition to the differences in the mean phase there are also large discrepancies between the upstream and downstream phase along the pulse. Figure 4.4 shows one example of this. The overall phase sag along the pulse is approximately equivalent for both the upstream and downstream phase. However, there are many oscillations along the downstream phase, up to 10° peak-to-peak, that are not present upstream. Clearly these differences also need to be removed to enable phase jitters close to the CLIC target to be achieved with the PFF prototype at CTF3.

As well as the static variations in the downstream phase along the pulse, there are also large differences in the phase jitter along the pulse. This is shown in Figure 4.5. The point by point upstream phase jitter along the pulse is quite flat at around 1 degree, slightly larger than the jitter on the mean. Alternatively, the downstream phase jitter has large variations in stability along the pulse. In some regions the downstream jitter is a factor 4 higher than

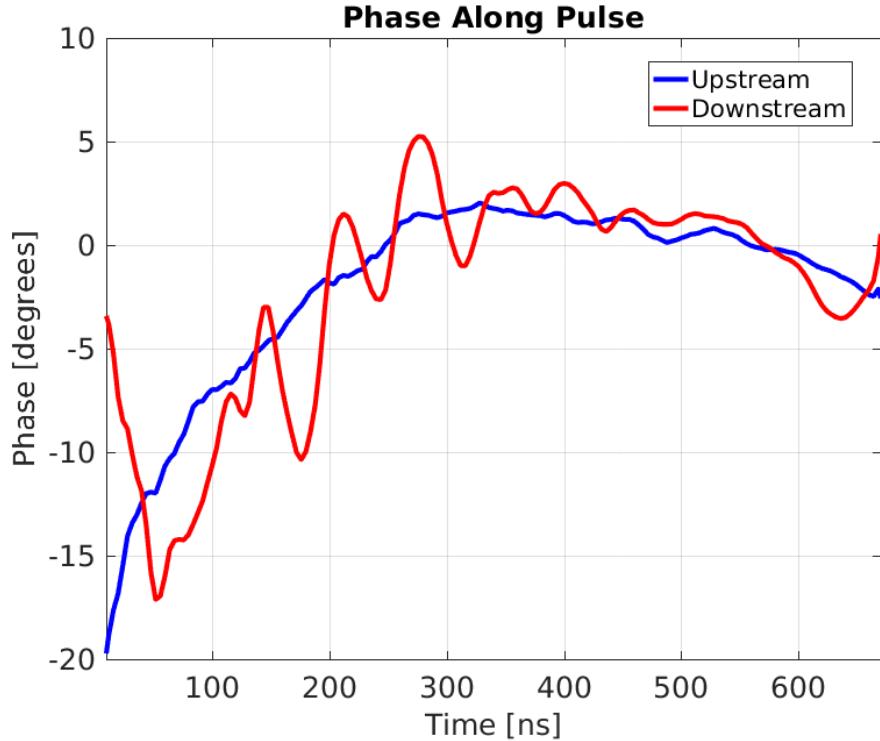


Figure 4.4: Upstream and downstream phase along the pulse during early PFF tests.

the upstream phase jitter, whilst in others it is only 50% higher.

4.3 First Order Energy Dependencies

The transfer matrix coefficient R_{56} defines the phase shift between two points in the lattice resulting from a beam energy offset (Section ??). Due to issues with meeting all the optics constraints for the PFF system a non-zero R_{56} value of close to -0.2 m had to be tolerated in the TL2 chicane (Section ??). As a result a dependence of the downstream phase on the beam energy is expected. Naturally, this was the first place to look to try to understand and improve the poor upstream–downstream phase correlation and high downstream phase jitter.

4.3.1 Beam Energy Variations

The best way to measure variations in the beam energy at CTF3 is via the beam position in one of the chicanes or rings, where the position after a bending magnet depends on the beam energy. The beam position offset, Δx , in a BPM (beam position monitor) in these regions is linked to the relative energy offset, $\Delta p/p$, as follows:

$$\frac{\Delta p}{p} = \frac{\Delta x}{D_x} \quad (4.6)$$

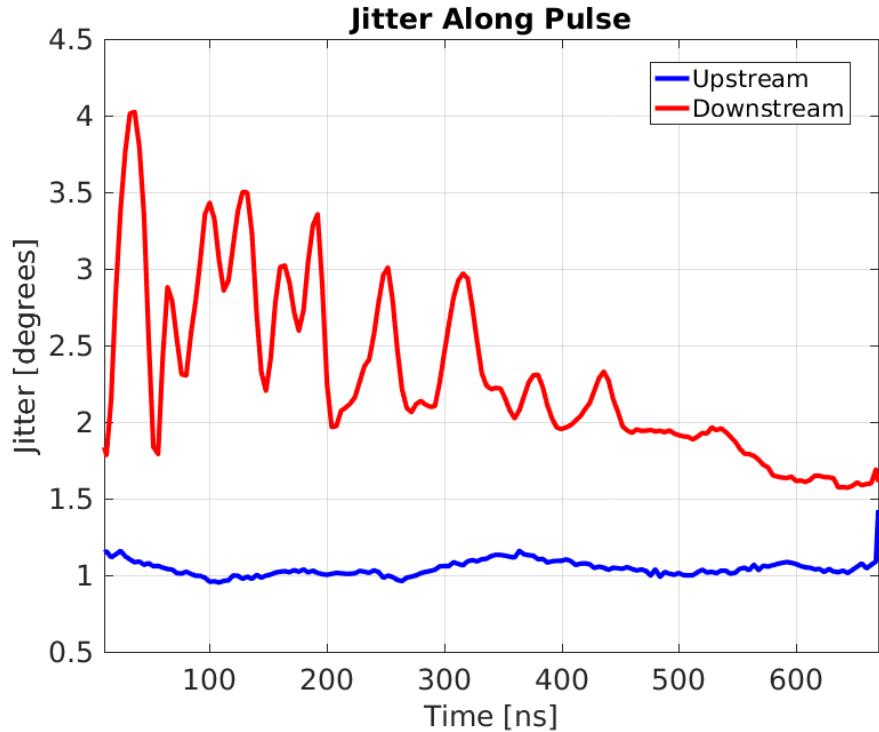


Figure 4.5: Upstream and downstream phase jitter along the pulse during early PFF tests.

Where D_x is the dispersion at that location given by the machine optics. Regions where D_x is non-zero are referred to as dispersive regions.

For the current beam setup used for the PFF system (bypassing the stretching chicane and delay loop, see Figure 1.5) the first dispersive BPM after the CTF3 linac is the BPM CT.0608 in the transfer line TL1. TL1 is introduced in more detail in Section 4.4.1. The BPM CT.0608 is placed roughly 1 m after the first dipole magnet in TL1, with a horizontal dispersion of $D_x = -0.61$ m at this point according to the CTF3 MADX model. There are no quadrupoles or other elements between the dipole and CT.0608, which reduces the sensitivity of this dispersion value to any inaccuracies in the model. All quoted energy measurements in this chapter are obtained using the horizontal position in BPM CT.0608, converted in to the energy using the equation above.

It should be noted that whilst this provides an accurate measurement of the relative energy jitter and variations along the pulse, it cannot be used to determine the absolute energy value. The mean beam position in CT.0608 is non-zero due to device misalignments, for example. This mean position offset is subtracted in all cases so that the resulting energy values are centred around zero, although the actual mean energy offset may be non-zero. In certain simulations in this chapter non-zero mean energy offsets are used to improve the agreement with the data, as indicated.

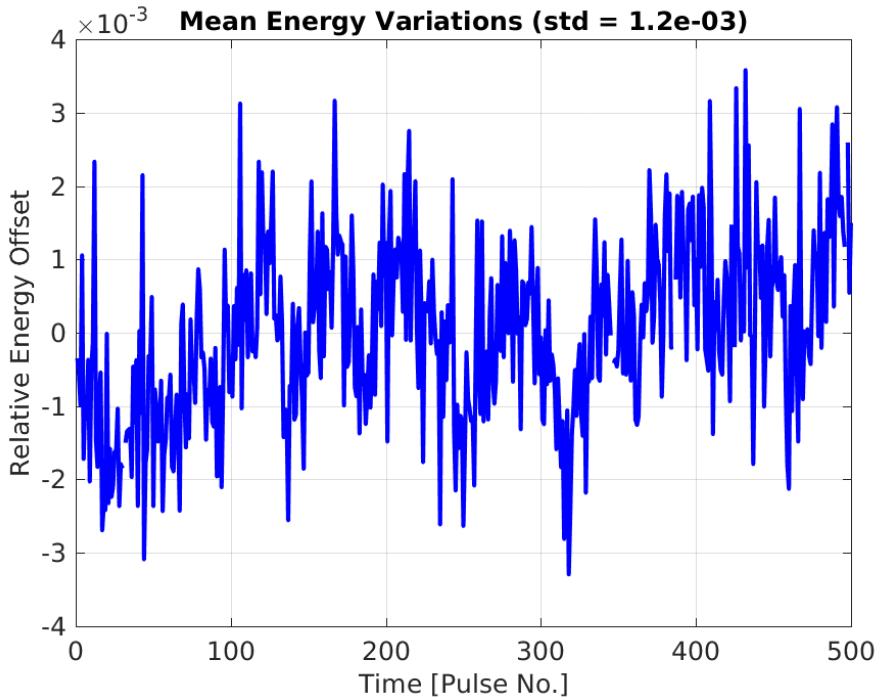


Figure 4.6: Example of variations in mean energy across 500 pulses, with a relative jitter of 1.2×10^{-3} .

Mean Beam Energy

Figure 4.6 shows an example of variations in the mean beam energy on medium to long timescales at CTF3, in this case 500 pulses or roughly 10 minutes. Relative energy jitter on the mean is typically at the level of 1×10^{-3} , and varies by roughly $\pm 50\%$ about this value. For shorter timescales, in datasets up to a couple of hundred pulses or a few minutes in length, the relative energy jitter is routinely closer to the 0.5×10^{-3} level (seen later in Figure 4.23, for example). The main sources of drifts in the mean energy at CTF3 are variations in beam current and temperature dependent effects on the klystron phase and RF power feedback loops [REF]. Recent improvements have demonstrated below 1×10^{-3} energy stability also on longer timescales (see Section 4.5.3), which will aid future PFF tests.

Energy Variations Along the Pulse

Energy stability point by point along the pulse is at the same level as the mean energy stability, with relative jitter at, or slightly above, the 1×10^{-3} level, as seen in Figure 4.7. Typically the energy stability is slightly worse at the start and end of the pulse.

In addition to the jitter there are also static variations in the energy along the pulse. Figure 4.8 shows one example of this. Peak-to-peak variations in the energy along the pulse are normally several times larger than the energy jitter, at the level of $2\text{--}5 \times 10^{-3}$. New feedbacks have also been recently developed at CTF3 to reduce variations in energy along the pulse (Section 4.5.3).

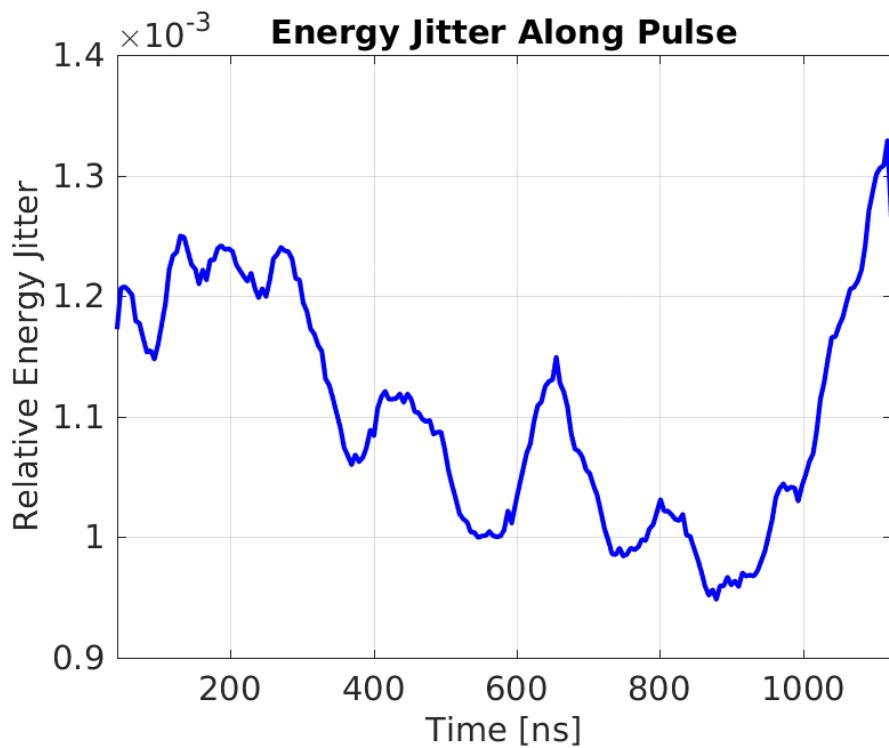


Figure 4.7: Example of relative energy jitter along the pulse.

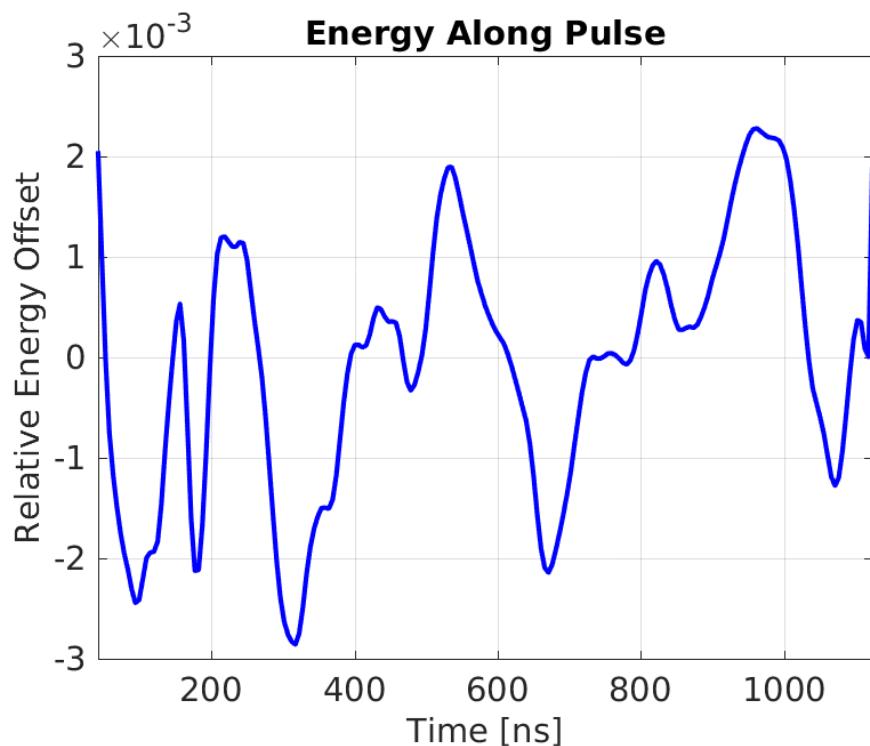


Figure 4.8: Example of variations in energy along the pulse.




Figure 4.9: Dependence of the mean downstream phase on the beam energy.

4.3.2 Correlation between Phase and Energy

Figure 4.10 shows one example of the typical dependence of the mean downstream phase on the beam energy during the first PFF attempts. It's immediately clear that there is a strong relationship between the two, with a correlation in this case of 0.82 ± 0.03 .

Alternatively, in Figure 4.9 there is almost no dependence of the upstream phase on the energy. However, the calculated correlation coefficient of 0.19 ± 0.06 is statistically significant and this has consequences for the discussions in the remainder of this chapter. Normally the small upstream phase–energy correlation typically varies between 0 and 0.4 at CTF3 depending on the conditions, in particular the energy jitter, at that time. In certain setups there can also be high correlations between the upstream phase and energy (Section ??).

By itself having a high correlation between the downstream phase and the energy is not an issue for the PFF system. The problem is the difference between the upstream phase–energy correlation and the downstream phase–energy correlation, which leads to low correlation between the upstream and downstream phase. The high downstream phase jitter is also a problem for the PFF system due to its limited correction range, as previously mentioned.

The goal of this section is to determine whether the non-zero optics R_{56} value between the upstream and downstream phase monitors is sufficient to explain both the large amplification in downstream phase jitter and the low upstream–downstream phase correlation seen here. To start, the effect of R_{56} on the downstream phase jitter and upstream–downstream phase correlation will be more formally defined.

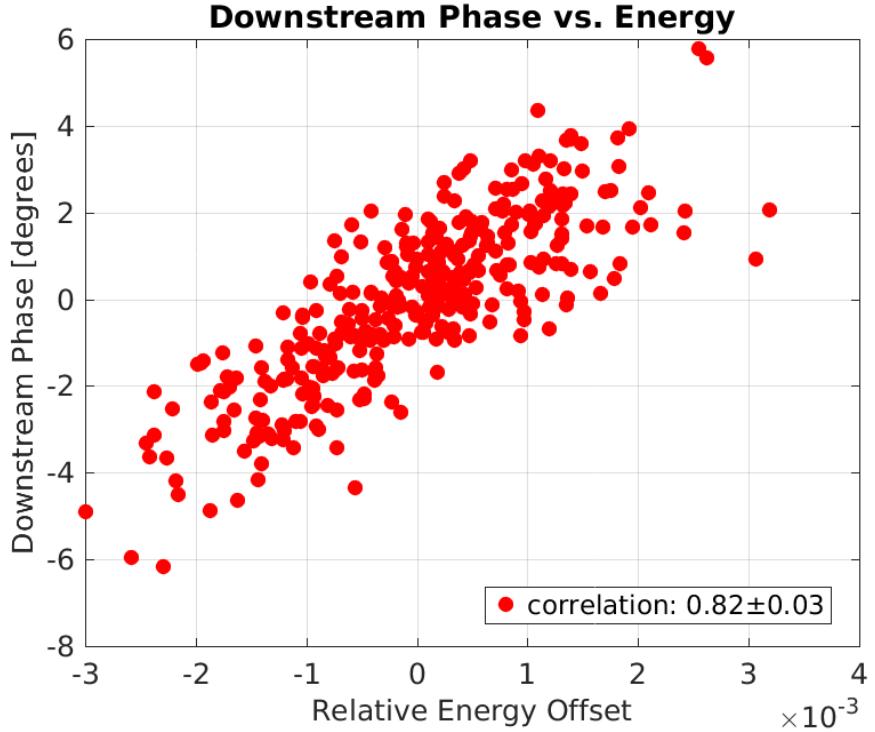


Figure 4.10: Dependence of the mean downstream phase on the beam energy.

4.3.3 R56

Assuming energy is the only source of differences between the upstream and downstream phase, the downstream phase, ϕ_d , can be expressed in terms of the optics transfer matrix coefficient R_{56} (Section 2.1) as follows:

$$\phi_d = \phi_u + R_{56} \frac{\Delta p}{p} \quad (4.7)$$

Where ϕ_u is the upstream phase, $\Delta p/p$ is the relative energy offset and R_{56} is the R56 value between the upstream and downstream phase monitors, defined by the machine optics. The units of R56 in the equation above are 12 GHz radians per unit relative energy offset ($\Delta p/p = 1$). MADX uses units of metres and this value is what will be referred to in this chapter. To obtain the R56 value to use in the equation above the MADX value must be multiplied by the conversion factor $2\pi/0.025$, where 0.025 m is the 12 GHz wavelength.

In terms of jitters Equation 4.7 becomes:

$$\sigma_d = \sqrt{\sigma_u^2 + R_{56}^2 \sigma_p^2 + 2R_{56}\rho_{up}\sigma_u\sigma_p} \quad (4.8)$$

Where σ_d is the downstream phase jitter, σ_u is the upstream phase jitter, σ_p is the relative energy jitter and ρ_{up} is the correlation between the upstream phase and the energy. This follows from the result of adding correlated variances. Clearly, any non-zero R56 between the upstream and downstream phase monitors introduces an additional energy component to the downstream phase that increases the downstream phase jitter.

The effect of R56 on the upstream-downstream phase correlation, ρ_{ud} , can also be defined starting from the definition of the correlation coefficient:

$$\rho_{ud} = \frac{\text{cov} [\phi_u, \phi_d]}{\sigma_u \sigma_d} \quad (4.9)$$

Where $\text{cov} [\phi_u, \phi_d]$ is the covariance between the upstream and downstream phase, given by:

$$\text{cov} [\phi_u, \phi_d] = \frac{1}{N} \sum_{i=1}^N \phi_{ui} \phi_{di} \quad (4.10)$$

By inserting the definition of the downstream phase from Equation 4.7 and separating the terms in the sum this becomes:

$$\text{cov} [\phi_u, \phi_d] = \frac{1}{N} \sum_{i=1}^N \phi_{ui}^2 + R_{56} \frac{1}{N} \sum_{i=1}^N \phi_{ui} \frac{\Delta p}{p} \quad (4.11)$$

The first term is now the variance of the upstream phase, σ_u^2 , and the second term is R_{56} multiplied by the covariance between the upstream phase and the energy, $\text{cov} [\phi_u, \frac{\Delta p}{p}]$, which can be expressed in terms of the correlation between the upstream phase and the energy, ρ_{up} :

$$\text{cov} \left[\phi_u, \frac{\Delta p}{p} \right] = \rho_{up} \sigma_u \sigma_p \quad (4.12)$$

Therefore, Equation 4.11 becomes:

$$\text{cov} [\phi_u, \phi_d] = \sigma_u^2 + R_{56} \rho_{up} \sigma_u \sigma_p \quad (4.13)$$

Finally, substituting Equations 4.8 and 4.13 into Equation 4.9 gives:

$$\rho_{ud} = \frac{\sigma_u + R_{56} \rho_{up} \sigma_p}{\sqrt{\sigma_u^2 + R_{56}^2 \sigma_p^2 + 2 R_{56} \rho_{up} \sigma_u \sigma_p}} \quad (4.14)$$

Considering that in this model the only difference between the upstream and downstream phase results from the R56, it is perhaps obvious that the best conditions for the PFF correction are obtained when the R56 coefficient between the upstream and downstream phase monitors is zero. In these conditions $\sigma_d = \sigma_u$ and $\rho_{ud} = 1$. This can be more formally defined by using the expression for the theoretical corrected downstream phase jitter when using the optimal gain factor as derived in Section ??:

$$\sigma_{PFF} = \sigma_d \sqrt{1 - \rho_{ud}^2} \quad (4.15)$$

All these quantities have been derived above and inserting them in to this equation gives the following expression for the corrected downstream phase jitter in terms of the R56:

$$\sigma_{PFF} = |R_{56}| \sigma_p \sqrt{1 - \rho_{up}^2} \quad (4.16)$$

As expected the achievable corrected downstream phase jitter is minimised when $R_{56} = 0$. Note that this equation does not take in to account the effects of the phase monitor resolution,

which in reality limits the minimum achievable downstream phase jitter to $\sigma_{PFF} = 0.2^\circ$ [REF].

In principle, the beam conditions for the PFF correction can also be improved by reducing the relative energy jitter (σ_p) or by increasing the upstream phase-energy correlation (ρ_{up}). Reducing the relative energy jitter decreases the additional phase jitter created by non-zero R56. Increasing the upstream phase-energy correlation (ρ_{up}) reduces the effect that non-zero R56 has on the upstream-downstream phase correlation (ρ_{ud}). For example, if $\rho_{up} = 1$ the source of all upstream phase jitter is energy jitter. In that case although non-zero R56 would further increase the downstream phase jitter, the additional jitter would be well correlated with the upstream phase and the upstream-downstream phase correlation would not be affected. In practice σ_p and ρ_{up} are defined by the CTF3 injector and can not be varied with a great degree of flexibility, so having zero R56 is the only way to obtain ideal conditions for the PFF correction. However, recent improvements at CTF3 have reduced the relative energy jitter, as discussed later in Section 4.5.3. High upstream phase-energy correlations can also be created at CTF3 but not without greatly amplifying the upstream phase jitter, which causes issues for the PFF system due to its limited correction range (Section ??).

An interesting side note of Equations 4.8 and 4.16 is that the best beam conditions for the PFF correction are not given by minimising the initial downstream phase jitter in the case where the upstream phase-energy correlation, ρ_{up} , is non-zero. As seen in the previous section in normal conditions there is a small correlation between the upstream phase and the energy at CTF3, typically around $\rho_{up} = 0.2$. In these conditions the downstream phase jitter can in theory be reduced to below the level of the upstream phase jitter by using a negative R56 to remove the energy component in the downstream phase. Differentiating Equation 4.8 gives the minimum downstream phase jitter to be obtained when $R_{56} = -\rho_{up}\sigma_u/\sigma_p$. However, using this R_{56} value would degrade the upstream-downstream phase correlation and increase the achievable corrected downstream jitter, which is always minimised when $R_{56} = 0$ as in Equation 4.16. This is significant for the R_{56} optimisation attempts presented later in this chapter, as it must always be kept in mind that the goal is to maximise the upstream-downstream phase correlation rather than to create the most stable downstream phase (with the PFF system off).

4.3.4 Effect of R56 in TL2

Unfortunately, it was not possible to find optics for the TL2 chicane that met all the PFF requirements and thus an R56 in the chicane of close to -0.2 m had to be accepted (Section ??). All other lines at CTF3 nominally have zero R56 [REF], and therefore don't introduce additional phase jitter via energy, at least to first order and to within the accuracy of the CTF3 MADX model. The overall R56 between the upstream phase monitors (in the CT line) and the downstream phase monitors (in the TBL line, labelled CB, after the TL2 chicane) is therefore -0.2 m also. Whether this can explain the low upstream-downstream phase correlation and high downstream phase jitter seen in Section 4.2, as well as what residual R56 between the two monitors can be tolerated to be able to achieve CLIC-level phase stability at CTF3, is discussed in this section.

Parameter	Value
R_{56}	-0.2 m
σ_u	0.8°
ρ_{up}	0.2
σ_p	0.001

Table 4.1: Typical upstream phase and energy conditions at CTF3.

Equations 4.8 and 4.14 can be used to estimate the downstream phase jitter and upstream-downstream phase correlation in the conditions at CTF3. Typical values for the various parameters in the equations have already been presented in previous sections, and these values are summarised in Table 4.1. The value $R_{56} \simeq -0.2$ was obtained in Section ?? as previously mentioned, the value $\sigma_u \simeq 0.8^\circ$ in Section 4.2 and the values $\rho_{up} \simeq 0.2$ and $\sigma_p \simeq 0.001$ in Section 4.3.2.

With these parameter values the residual R56 of -0.2 m reduces the upstream-downstream phase correlation to below 10%, and amplifies the downstream phase jitter to above 3 degrees. Therefore, R56 transforming energy jitter in to downstream phase jitter can by itself explain the low upstream-downstream phase correlation and high downstream phase jitter seen in Section 4.2. In order to increase the upstream-downstream phase correlation to 97% and reduce the downstream jitter to 0.8 degrees (the conditions needed to achieve 0.2 degrees corrected downstream phase jitter at CTF3) the R56 between the upstream and downstream phase monitors must be removed.

Figure 4.11 and 4.12 show the expected downstream phase jitter and upstream-downstream phase correlation for residual R56 values from -0.2 to +0.2 m between the upstream and downstream phase monitors again using the equations derived in the previous section. The horizontal black line in each figure marks the requirements needed to reduce an initial upstream phase jitter of 0.8 degrees to the CLIC target of 0.2 degrees. The red line, "corrected", in Figure 4.11 shows the theoretical corrected downstream phase jitter using the PFF correction across the range of R56 values. Note the slight asymmetry in the phase jitter and correlation curves, which is caused by the non-zero correlation between the upstream phase and the beam energy. In order to obtain CLIC level phase stability at CTF3 the residual R56 between the upstream and downstream phase monitors must be reduced from the initial -0.2m to $0 \pm 1.3\text{ cm}$.

To interpret the results of the R56 optimisation attempts presented in the remainder of this chapter it is useful to understand how varying the correlation between the upstream phase and the energy (ρ_{up}) changes the dependence of the upstream-downstream phase correlation (ρ_{ud}) on the residual R56. In particular, in Section 4.5.2 a machine setup that increases ρ_{up} to 90% was used. Figure 4.13 shows how ρ_{ud} varies with ρ_{up} values between 0% and 40% (typical of normal operation) and with the higher correlation of 90%. With high correlations between the upstream phase and energy there is no longer a well defined peak in ρ_{ud} versus the residual R56 value. Instead there is an almost constant high upstream-downstream phase correlation with positive R56 values, and a large anti-correlation for negative R56 values (as in this case the residual R56 acts to flip the sign of the phase jitter).

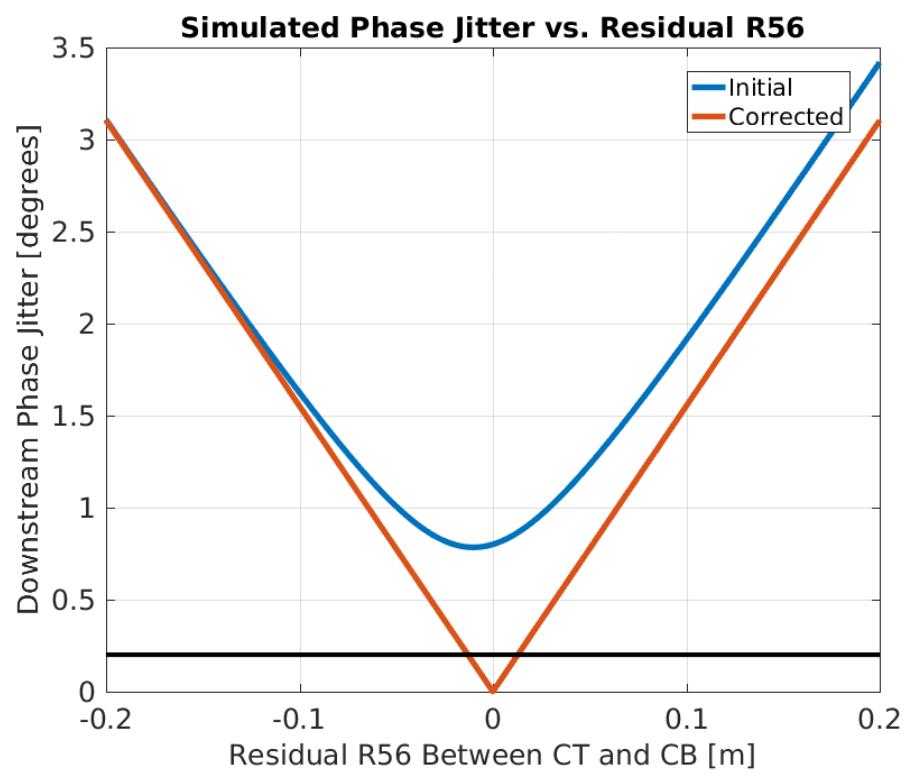


Figure 4.11: Downstream phase jitter vs. residual R56 between monitors.

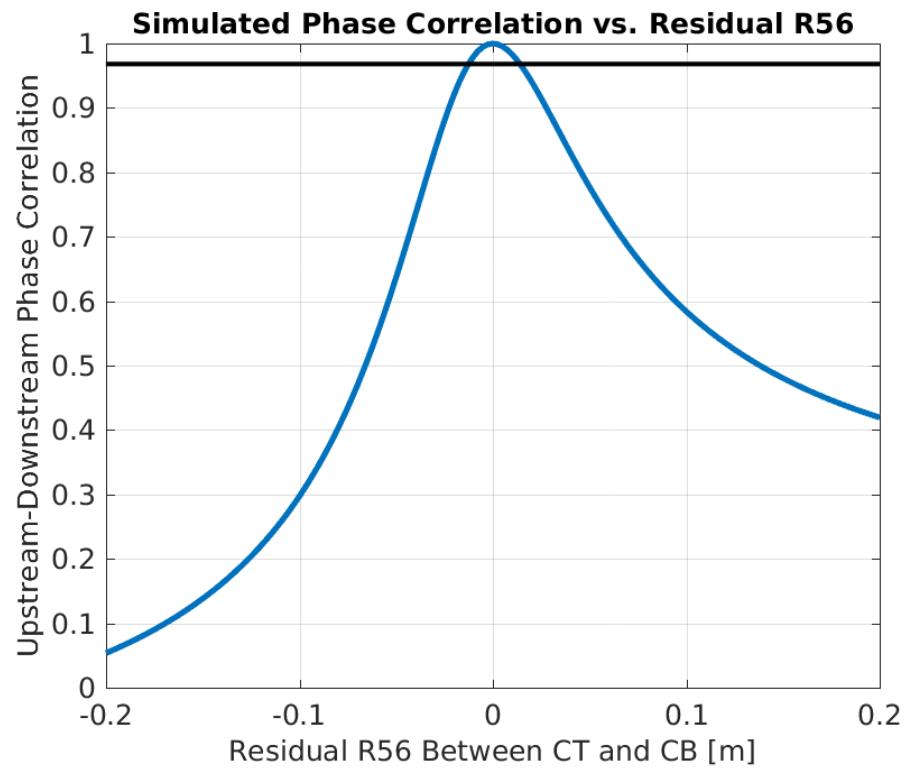


Figure 4.12: Phase correlation vs. residual R56 between monitors.

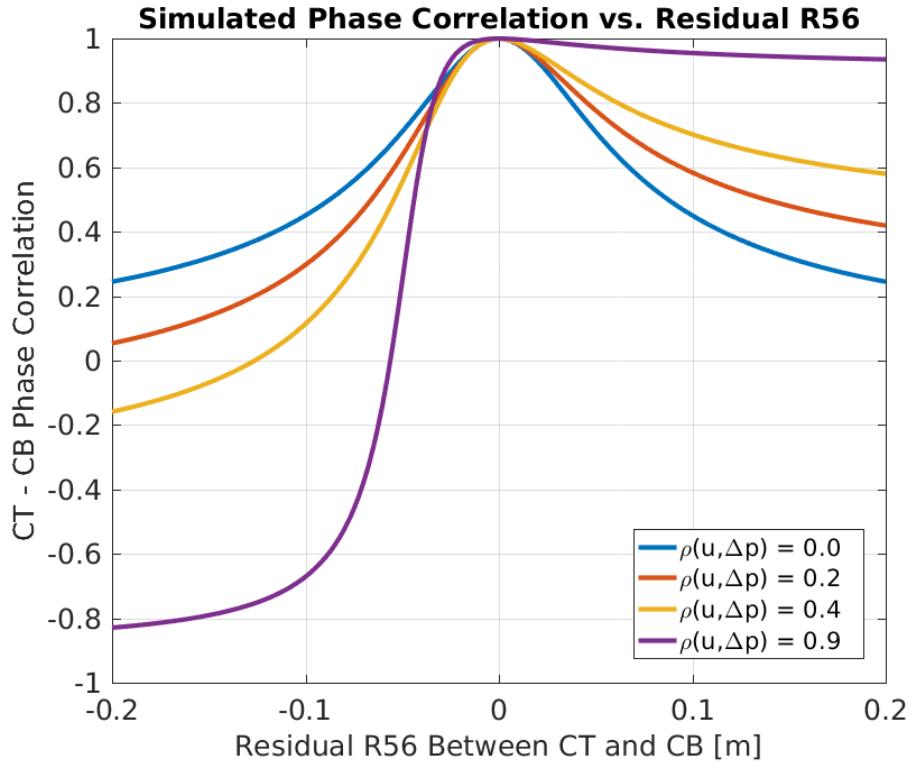


Figure 4.13: Phase correlation vs. residual R56 between monitors for different upstream phase-energy correlations.

In Figure 4.14 plotting the theoretical downstream jitter with $\rho_{up} = 0.9$ gives a clear demonstration that the best conditions for the PFF correction are always with $R_{56} = 0$ rather than with the lowest possible initial downstream jitter, as mentioned in the previous section. In fact, as these conditions relax the requirements on the residual R56 needed to achieve high upstream-downstream phase correlations (as seen in the previous figure) it may be easier to achieve a large factor reduction in the downstream phase jitter with the PFF system with a high ρ_{up} machine setup. This has been attempted and is presented in Section ??.

4.4 Mitigation of First Order Energy Dependence

The discussion in the previous section proves that with a residual R56 value of -0.2 m between the upstream and downstream phase monitors it is impossible to achieve the goals of the PFF prototype at CTF3. However, due to the highly constrained optics in TL2 it has already been seen in Chapter 2 that it was not possible to find optics for the PFF chicane that yield zero R56 whilst also meeting requirements for both the PFF system and transverse matching (dispersion and beta functions). The only way to create a total R56 of zero between the upstream and downstream phase monitors is therefore to add positive R56 to one of the other beam lines at CTF3 in order to compensate for the negative R56 in the TL2 chicane.

The previous transfer line TL1, which transports the beam from the CT line (where the

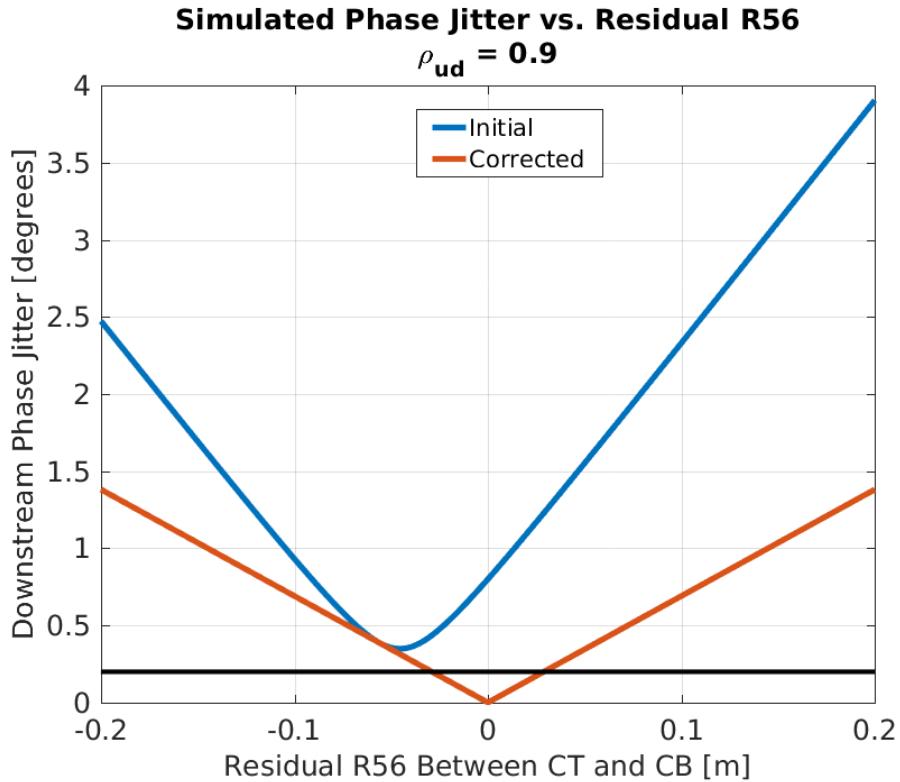


Figure 4.14: Initial and corrected downstream phase jitter vs. residual R56 between monitors with high correlation between the upstream phase and the beam energy.

upstream phase monitors are installed) to the combiner ring (see Figure 1.5), has been used to achieve this. The layout of the TL1 transfer line is shown in Figure 4.15. It consists of: 4 dipoles (bending the beam horizontally) of 2 different types, 13 quadrupoles of 5 different types, 7 magnetic correctors, 1 sextupole (usually not used) and 8 BPMs (the dispersive BPM after the first dipole in TL1, labelled CT.BPI0608, is the device that has been used to determine correlations between the phase and energy in this chapter). The total length of TL1 is approximately 30 m.

Preliminary attempts to reduce the residual R56 between the upstream and downstream phase monitors using TL1 yielded correlations up to 60% and reduced the downstream phase jitter to 2°. Conditions similar to these were used for the first PFF tests (Chapter ??) before the energy related effects discussed in this chapter were fully characterised, but in these tests only a modest reduction of 30% in downstream phase jitter was possible due to the limitations of the phase propagation shown here. At this time only a few different optics for TL1 were available in R56 steps of 10 cm. As the total residual R56 must be reduced to the centimetre level to make a correction down to 0.2 degrees jitter theoretically possible with the PFF system, new sets of optics for TL1 were required.

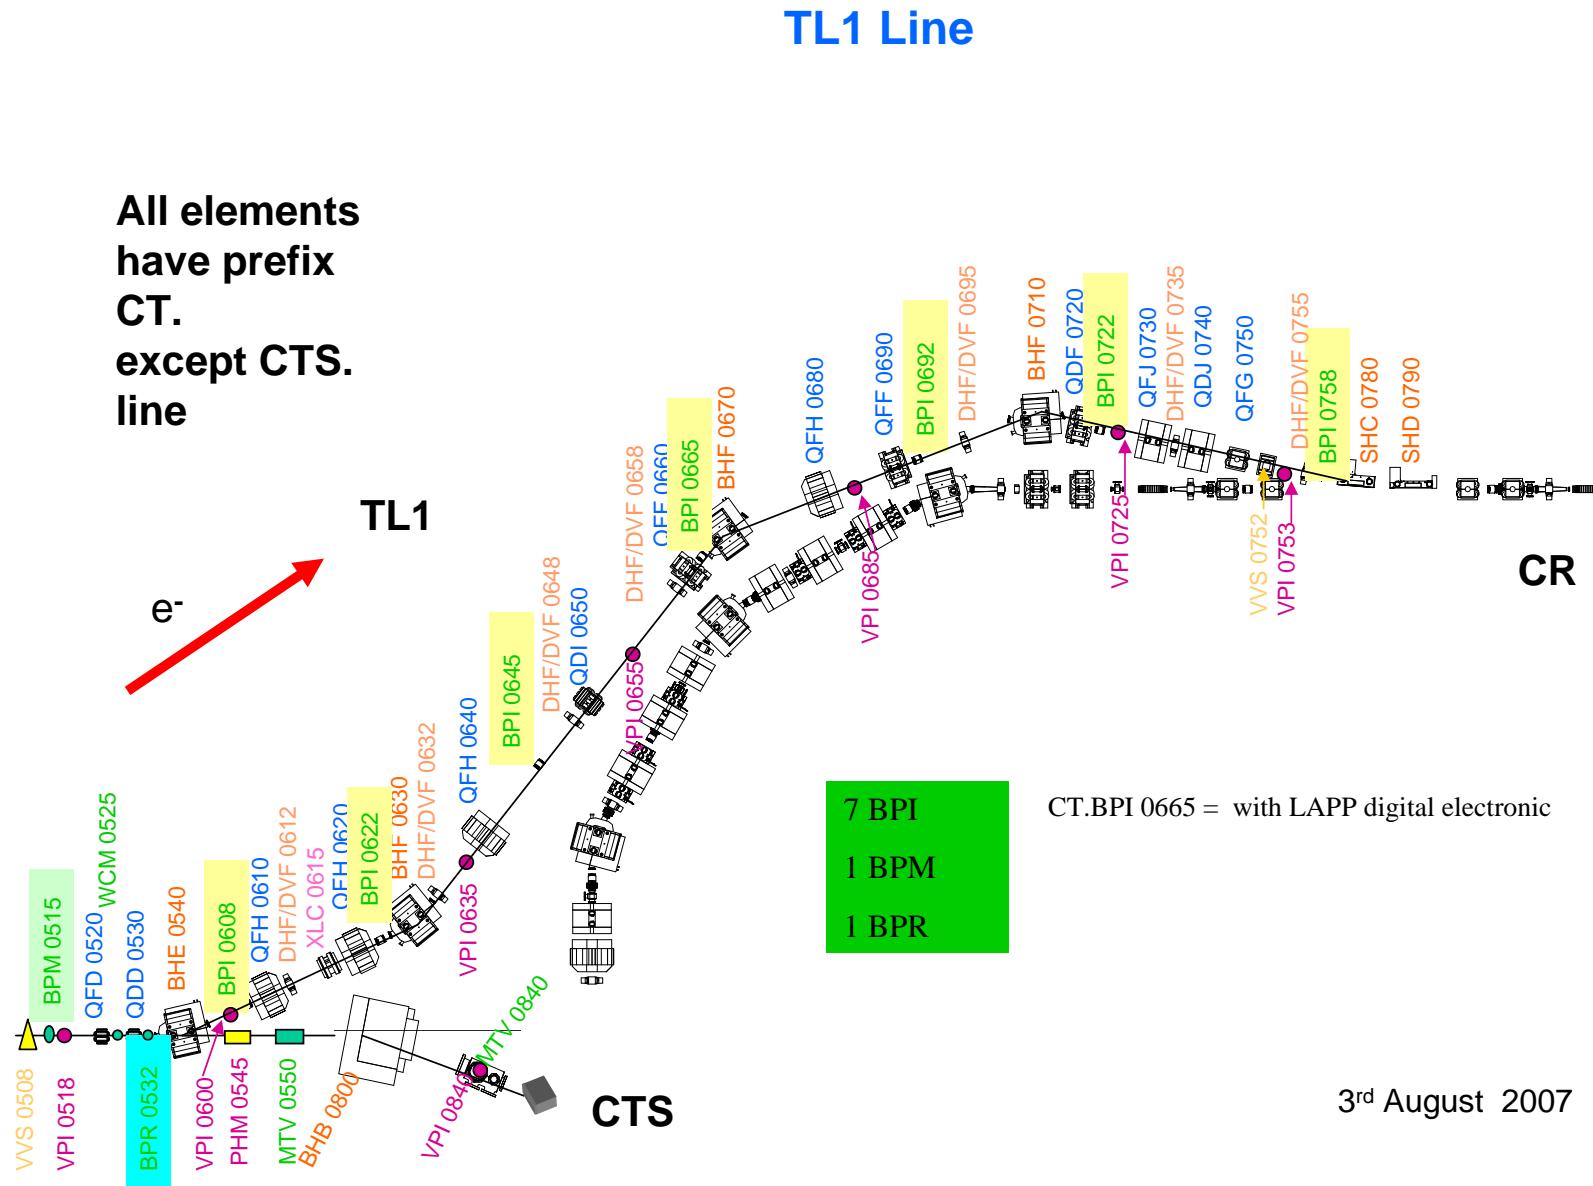


Figure 4.15: TL1 [REF].

Parameter	TL1 Injection	CR Injection
β_x	8.81 m	4.08 m
β_y	13.94 m	5.41 m
α_x	-0.74	-0.31
α_y	-0.45	-0.21
D_x	0 m	-0.03 m
D_{px}	0	0.02

Table 4.2: Initial and final conditions for optics matching in TL1.

4.4.1 Matched Optics for TL1

Although in theory only one set of optics with $R_{56} = +0.2$ m in TL1 is required to compensate for the $R_{56} = -0.2$ m in TL2, in practice errors in the MADX model of CTF3 plus the effect of higher order energy dependencies (see Section 4.5) means it is not possible to know precisely what the optimal R56 to set in TL1 will be, and it is also possible that this value will vary with time. To determine the optimal value of R56 to set it is also useful to scan the R56 value in TL1 across a wide range of values and then fit the maximum resulting upstream-downstream phase correlation.

To allow this, MADX has been used to match optics for TL1 with R56 values ranging from -0.3 m to +0.6 m in steps of 0.5 cm (a total of 181 sets of optics). The optimal R56 value should always be guaranteed to be in this range, and the step size of 0.5 cm allows the residual R56 to be zeroed to within one centimetre as derived to be necessary in Section 4.3.4. As well as the different R56 values, each set of optics must maintain the same initial and final conditions, so that the injection of the beam in to the combiner ring is not affected. Values for the beta functions, alphas and dispersion at the start of TL1 and at the combiner ring injection are summarised in Table 4.2. As well as the initial and final conditions, the maximum beta functions and dispersion in TL1 are constrained to ensure a reasonable beam size throughout the line — the horizontal and vertical beta function is limited to a maximum of 35 m, and the horizontal dispersion to a maximum absolute value of 1.25 m. Around the septum used for injection in to the combiner ring the horizontal beta function is further limited to a maximum of 10 m. The strengths of the 13 quadrupoles in TL1 are varied to meet all these constraints.

Figures 4.16 shows the matched R56 value in TL1 across the range of targeted values. Each matched R56 value is within 10 microns of the desired result. Figure 4.17 then shows an example of how the strength of one of the quadrupoles must be varied in order to achieve each R56 value. If the dependence of each quadrupole strength on the R56 value was continuous the relationship could be fitted to create a set of tuning knobs to allow R56 to be set to any arbitrary value in TL1. However, as seen in the figure there are many discontinuities. The quadrupole strengths for each set of optics are therefore saved to a lookup table, with a MatLab function created to read the table and set the quadrupole currents in the machine appropriate for the specified R56 value. As already mentioned 0.5 cm precision in R56 should be adequate for the PFF requirements, but the discontinuities mean new optics would have to be matched if optics with an R56 value not included in the discrete set used here were

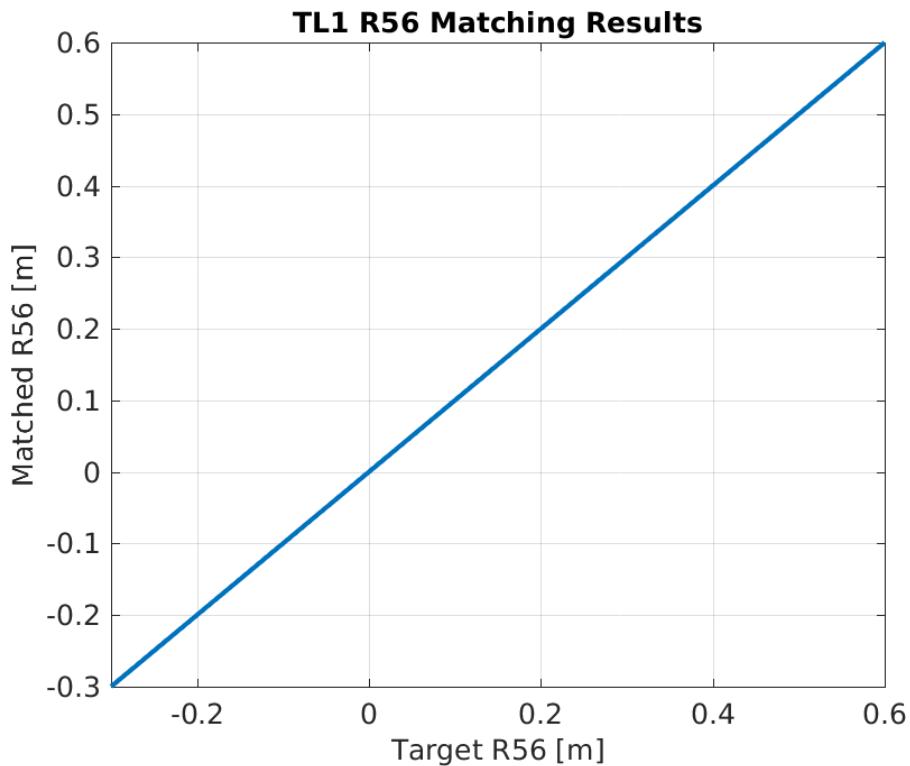


Figure 4.16: Matched R56 values for TL1.

required.

For reference Figures 4.18 4.19 and 4.20 show how the horizontal and vertical beta functions and horizontal dispersion changes in TL1 for each set of optics. For all R56 values each parameter converges to the same value at the start and end of TL1, as needed to ensure that changing the R56 does not impact injection in to the combiner ring. The maximum horizontal and vertical beta functions in TL1 roughly increase with the set R56 value, but in all cases are kept below the set limit of 35 m in the matching procedure. The dispersion pattern in TL1 also changes with the set R56 value, though in most cases the maximum absolute dispersion is around 1 m and only the location of the peak dispersion along the line changes. Again, for each set of optics the maximum absolute dispersion is limited within the set constraint of 1.25 m.

Commissioning of the new TL1 optics in CTF3 was straightforward and in general they can be set with the quadrupole strengths at their nominal matched values without causing issues for the beam quality. At the extremities of the range of optics (close to $R56 = -0.1$ m and $R56 = +0.6$ m) some slight beam losses do begin to occur, but this is not a problem for PFF operation where the required R56 is only 0.2 m. However, for each set of optics the magnetic correctors in TL1 may need to be changed to recover the nominal beam orbit, thus taking in to account slight misalignments in elements along the line.

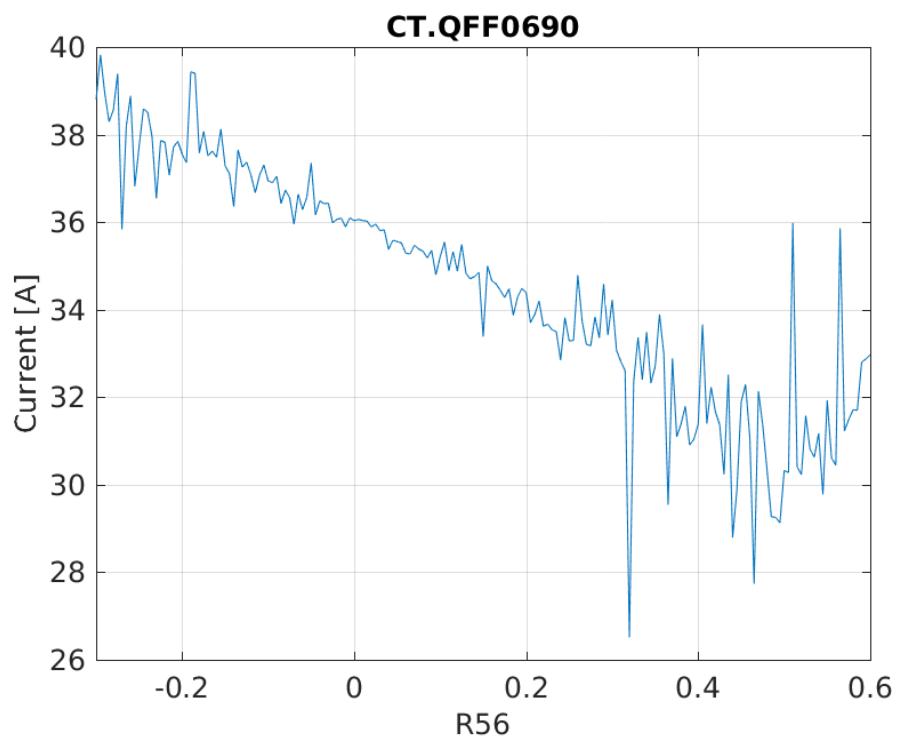


Figure 4.17: Current vs. R56 for the CT.QFG0750 quadrupole in TL1.




Figure 4.18: Horizontal beta in TL1 for all R56 optics.

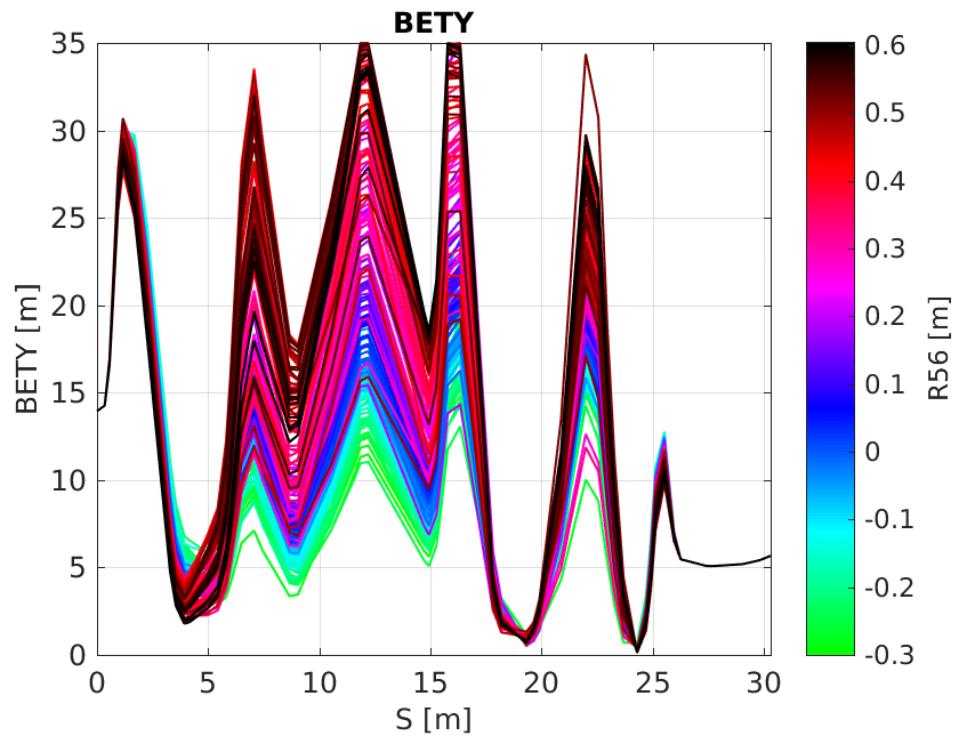


Figure 4.19: Vertical beta in TL1 for all R56 optics.

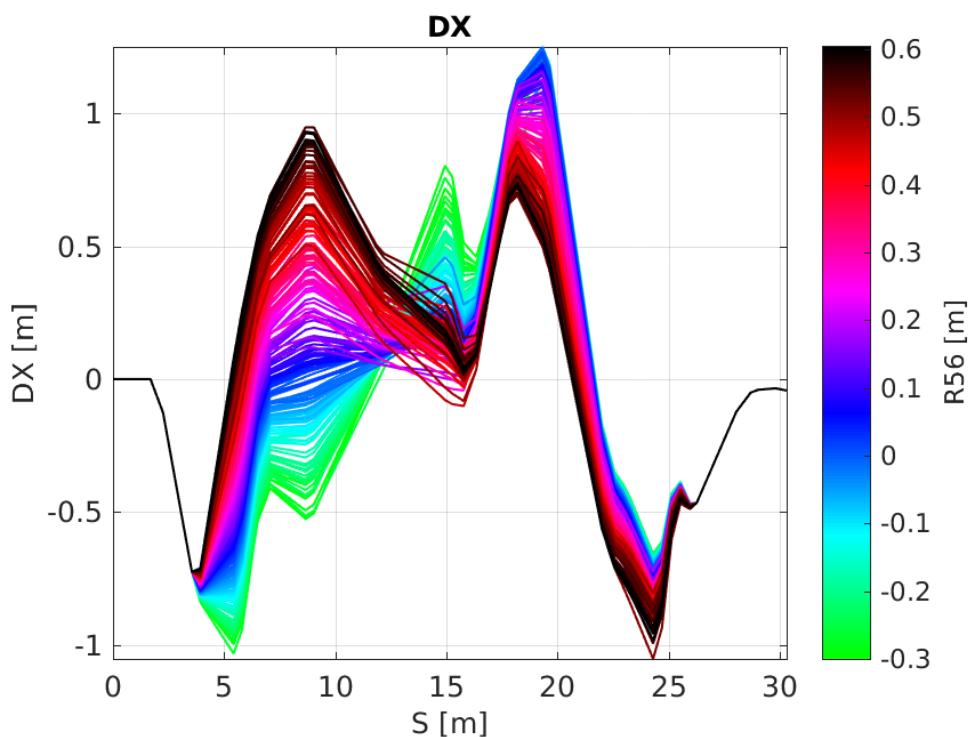


Figure 4.20: Dispersion in TL1 for all R56 optics.

4.4.2 Scans of R56 in TL1

The sets of matched optics from the previous section can be used to perform scans of the R56 value in TL1 to observe how the downstream phase is affected. Scans of this type must be performed prior to all PFF data taking periods in order to optimise the beam conditions (maximise the upstream-downstream phase correlation) for the correction. More recently scans of R56 in TL1 have been performed whilst varying the beam energy, which produces cleaner results and highlights additional factors that must be taken in to account during the optimisation process, as will be shown in Section 4.5.

As a starting point the simplest case, where only the TL1 optics is changed during the scan and all other parameters in the machine are left unchanged, is presented in this section. This also highlights some of the difficulties in maintaining beam conditions at CTF3, which is discussed further in Section 4.5 and extensively in the context of the PFF correction in Section 6.2. Figures 4.21 and 4.22 show one example of an R56 scan performed across the full range of available optics – from -0.1 m R56 in TL1 to +0.6 m. The R56 is incremented by 2.5 cm between datasets, to give a total of 29 R56 points in the scan, with the whole scan taking approximately one and a half hours to complete. With the knowledge gained from measurements of this type it is no longer necessary to scan the R56 across the full range to determine the ideal value, thus the optimisation of the phase propagation for PFF attempts can now be achieved on much shorter time scales.

Mean Phase

Only the mean phase jitters and correlation will be considered here, features along the pulse are discussed in later sections for other scans. Figure 4.21 shows the mean phase jitter during the scan both upstream and downstream. Although the noise in the measurement is quite large, the downstream phase jitter is reduced from above 2.5 degrees with zero R56 in TL1, to below 1 degree and close to the level of the upstream phase jitter by adding positive R56 in TL1. The optimal R56 value is approximately 0.175 m, in close agreement with expectations considering the -0.2 m R56 in TL2. The upstream-downstream phase correlation, in Figure 4.22, is also maximised at this point, from an initial correlation of 20% with zero R56 to up to 80%. In terms of the PFF system, increasing the upstream-downstream phase correlation from 20% to 80% improves the theoretical correction from a 2% reduction in downstream phase jitter to a 40% decrease (Equation 4.5).

As the upstream phase monitors are prior to TL1, changing the TL1 optics has no effect on the upstream phase jitter. All differences in the upstream phase jitter between datasets are caused by drifts in the CTF3 injector, typically changes in either klystron phases or beam current. Although the overall stability of the upstream phase jitter during this scan is good, it does vary between 0.5 degrees and 1.2 degrees. In addition to the upstream phase there are also differences in the relative energy jitter and upstream phase-energy correlation during the scan, as seen in Figure 4.23. The relative beam energy jitter varies between 0.4×10^{-3} and 1.0×10^{-3} and the upstream phase-energy correlation between -0.5 and +0.5. All these parameters influence the downstream phase, as per the equations in Section 4.3.3.

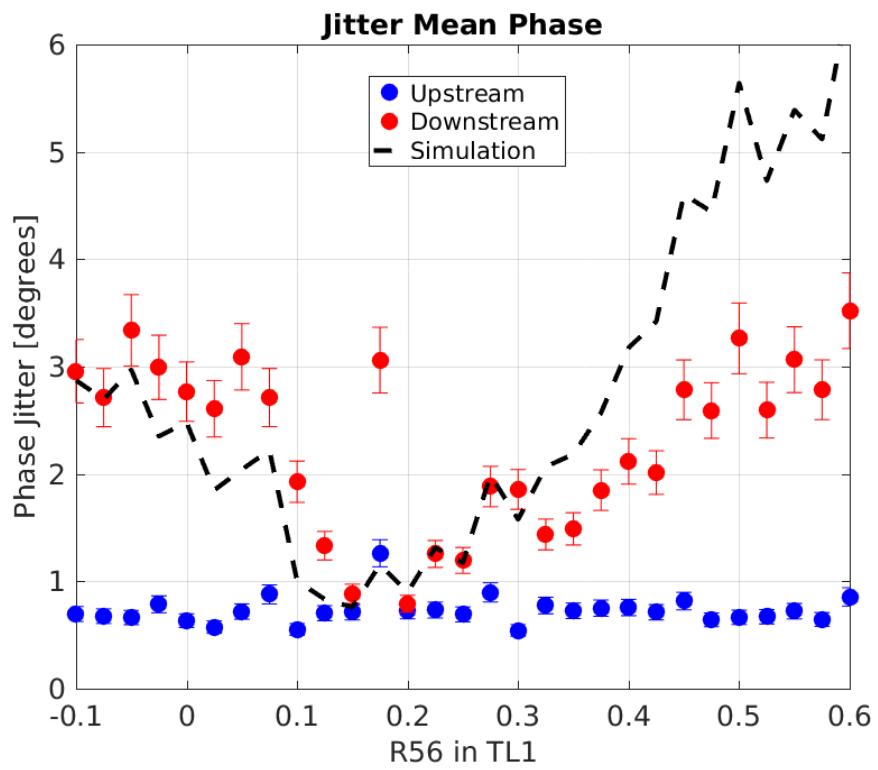


Figure 4.21: Phase jitter during scan of R56 in TL1.

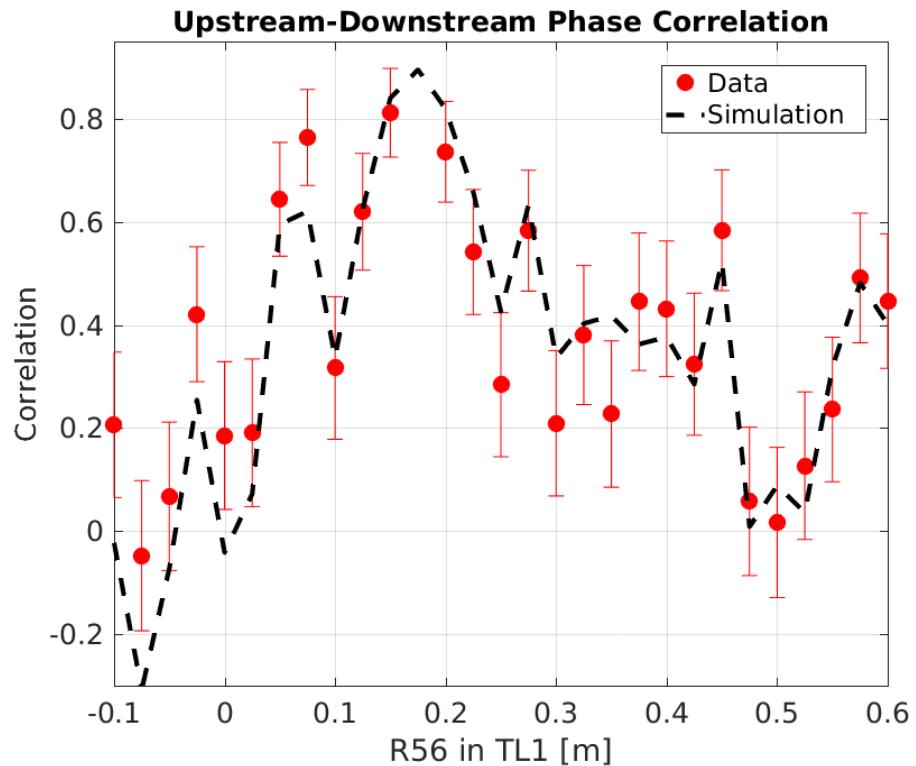


Figure 4.22: Correlation during scan of R56 in TL1.

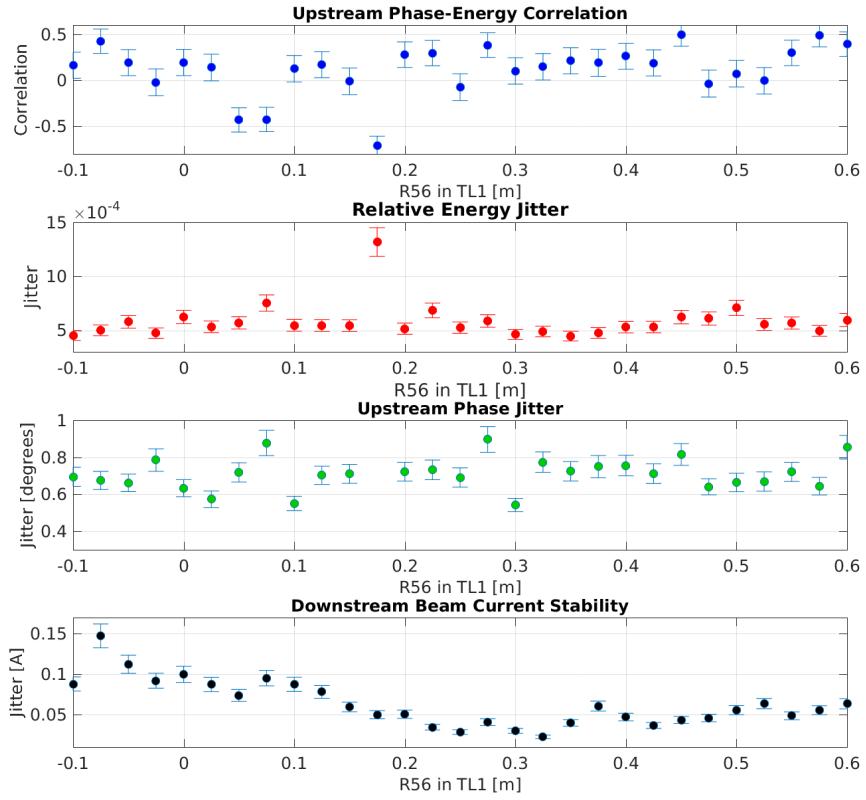


Figure 4.23: Upstream and downstream beam conditions during the R56 scan.

The differences in the upstream phase and energy conditions between datasets partially explains the apparent spread of the data points away from the expected clean distribution. The black “simulation” lines in Figures 4.21 and 4.22 represent the expected downstream phase jitter and upstream-downstream phase correlation at each point in the scan given the upstream phase jitter, relative energy jitter and upstream phase-energy correlation at that time (using Equations 4.8 and 4.14). The correlation simulation in Figure 4.22 has been scaled so that the peak value is in agreement with the data, at 0.8. The majority of the data points follow the scaled simulated distribution, with several remaining outliers. For the downstream phase jitter (which uses the simulated result directly with no scaling) the agreement with the simulation is generally good for R56 values below 0.3 m. However, above 0.3 m the actual phase jitter seen in the scan is smaller than the simulation. One possible explanation for this are the changes in downstream beam current between datasets, which varies by a factor 3 during the scan (bottom plot in Figure 4.23). Small beam losses between measurements may change the phase jitter in a way that is not characterised by the R56. Possible other sources are discussed in Sections 4.5 and 4.7.

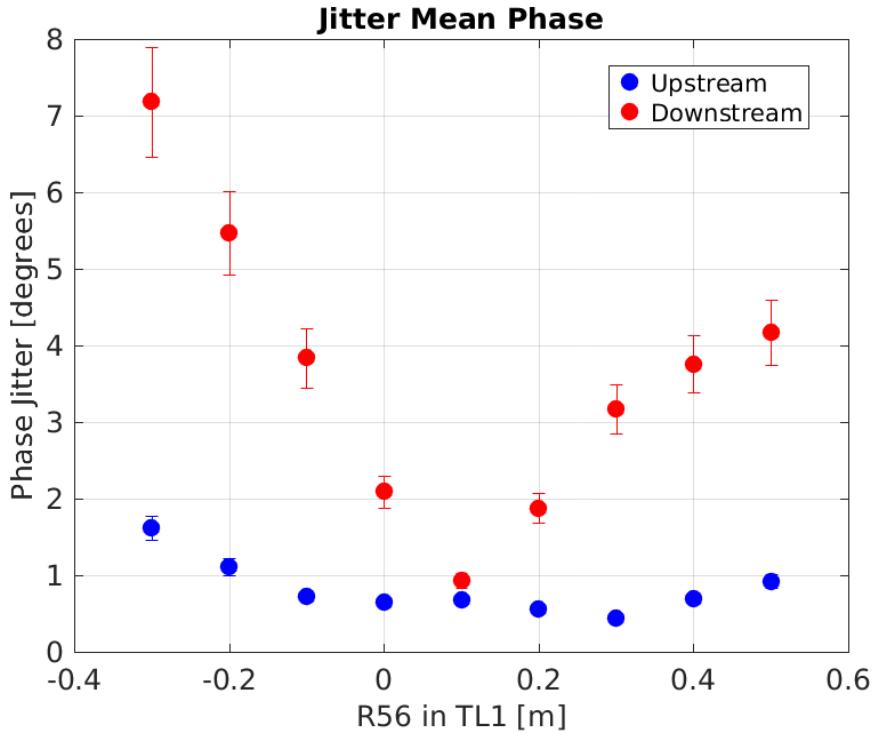


Figure 4.24: Mean phase jitter during R56 scan 2.

Results from Other Scans

Figures 4.24 and 4.25 show the results of two further scans of R56 in TL1, both taken a few days after the scan previously shown. For both scans the mean downstream phase jitter can again be decreased to the level of the upstream phase jitter by varying the R56 in TL1, and the upstream-downstream correlation increased to 80%. However, the optimal optics to use is different for each scan — the scan in Figure 4.24 has an optimal R56 value of around 0.1 m whereas for the scan in Figure 4.25 the optimal value quite close to zero, around 0.04 m. Both values are also different to the scan previously shown, which had an optimal R56 setting of 0.175 m.

With R56 alone and the model of the phase propagation used to derive the equations in Section 4.3.3 there is no mechanism for the optimal R56 value to vary with time. The best conditions for the phase propagation should always be provided with zero residual R56 between the upstream and downstream phase monitors. As the optics in all beam lines between the upstream and downstream phase monitors (apart from TL1) were unchanged between each scan, the optimal R56 value in TL1 should also be the same for each scan in this model. The most likely explanation is a sensitivity to higher order energy dependencies.

4.5 Higher Order Energy Dependencies

In the same way the first order optics dependencies are described by the 6×6 R-matrix, the second order effects are described using a three dimensional $6 \times 6 \times 6$ T-matrix. R_{56}

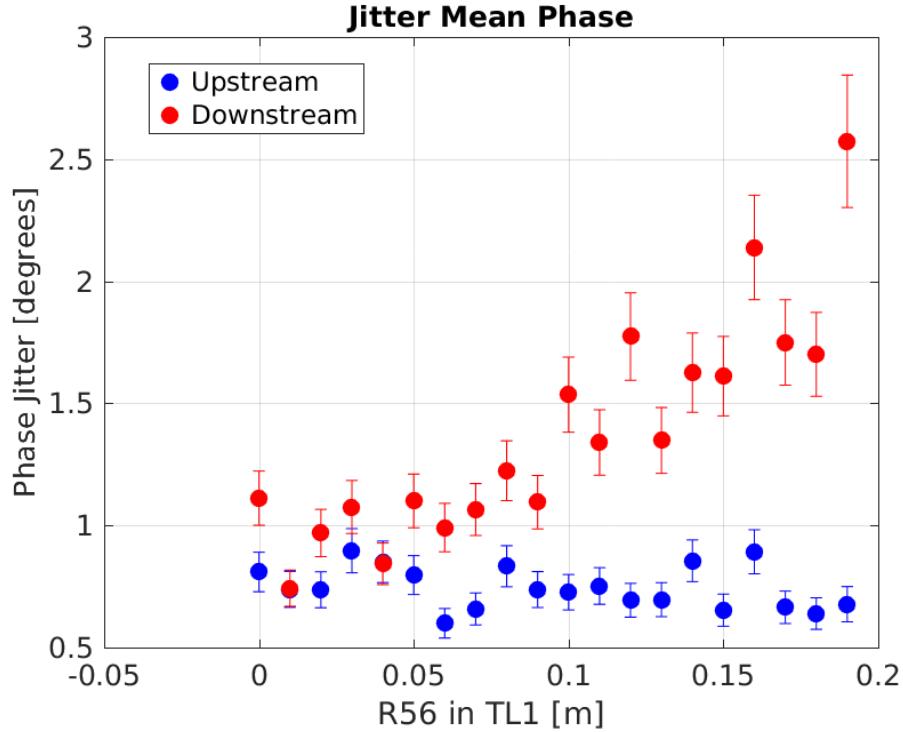


Figure 4.25: Mean phase jitter during R56 scan 3.

is the relevant first order transfer matrix coefficient for the energy related effects on the phase propagation, as already discussed, and it then follows that the relevant T-matrix coefficient for second order energy dependencies is T_{566} . By including the second order term the dependence of the downstream phase on the energy from Equation 4.7 becomes:

$$\phi_d = \phi_u + R_{56} \left(\frac{\Delta p}{p} \right) + T_{566} \left(\frac{\Delta p}{p} \right)^2 \quad (4.17)$$

T_{566} introduces another source of energy dependent phase jitter which is independent from the first order R_{56} value. The ideal case for the phase propagation would be to have both zero R_{56} and zero T_{566} between the upstream and downstream phase monitors. However, constraints are not placed on the T_{566} in the optics at CTF3 and it is therefore typically non-zero. It may be possible to create optics with zero, or at least reduced, T_{566} for the TL1 and TL2 lines at CTF3 in the future but this has not yet been pursued, thus it will be treated as a fixed property of the optics here. In this case an expression for the R_{56} value that minimises the downstream phase-energy dependence can be derived:

$$R_{56} = -2T_{566} \left(\frac{\Delta p}{p} \right) \quad (4.18)$$

This is obtained by zeroing the partial differential of Equation 4.17 with respect to $\Delta p/p$.

The above dependence of the R_{56} value on the beam energy offset has many consequences. Firstly, it provides a mechanism by which the apparent optimal R_{56} value in TL1 can vary with time (and be non-zero), as was seen comparing the results of different R_{56} scans in the previous section. CTF3 does experience drifts in beam energy (Section 4.3.1), creating

small offsets between the actual beam energy and the energy that the optics has been set for (i.e. the strength of bending and focusing elements in the accelerator). In other words, it is possible for the mean of $\Delta p/p$ to be non-zero. The optimal R_{56} value to use in TL1 is therefore expected to drift with the beam energy when higher order phase-energy dependencies are included.

Secondly, energy variations along the beam pulse and jitter in the beam energy mean that the phase propagation cannot be perfectly optimised by varying the R_{56} alone. Due to the energy dependence in Equation 4.18, any energy variations along the pulse cause the optimal R_{56} value to set in TL1 to also vary along the beam pulse. There are static variations in the mean energy along the pulse (e.g. as seen in Section 4.3.1) at CTF3, so this means the phase propagation can never be completely optimised along the full pulse length when T_{566} is non-zero.

4.5.1 Simulated Effect of T_{566} on the Downstream Phase

When only the first order effect of R_{56} is considered the dependence of the downstream phase on the energy is linear, with the gradient depending on the residual R_{56} value between the upstream and downstream phase monitors. The downstream phase versus beam energy offset with only first order R_{56} term included is shown in Figure 4.26 for each set of TL1 optics, demonstrating this effect. The optimal R_{56} value of +0.2 m in TL1 minimises the phase-energy dependence for all energy offsets.

By running MADX with varying energy offsets set in the initial conditions the expected effect of the higher order energy dependencies on the downstream phase can be seen. This is shown in Figure 4.27, and the impact of the higher orders is immediately clear. The optimal R_{56} to use in TL1, where the phase has the minimal dependence on the energy (at the peaks of the phase-energy curves in the figure), now depends on the energy offset. Figure 4.28 shows that the dependence of the optimal R_{56} value on the relative energy offset is linear, as expected from Equation 4.18. The plotted R_{56} value does not exceed +0.6 m or go below -0.3 m as only the available sets of optics for TL1 are considered, which are restricted to this range. This also creates small non-linearities in the central region of the plot, although there is also a contribution from effects above second order.

MADX does not output the optics T_{566} coefficient directly but it can be approximated using a quadratic fit to the downstream phase vs. energy curves seen in Figure 4.27. The fit coefficients then give the T_{566} and R_{56} values, as per Equation 4.17. An example of this is shown in Figure 4.29 for the nominal $R_{56} = 0$ optics in TL1. Again, as the results from MADX also includes effects above second order there is a slight discrepancy between the quadratic fit and the MADX output. However, including up to the second order energy dependence is enough to characterise the true behaviour and the slight modifications induced by higher orders are beyond the scope of the discussion here.

Figure 4.30 then shows the fitted T_{566} coefficient for all the sets of matched optics in TL1. The changes in T_{566} across the range of TL1 optics are much smaller than the (intentional) differences in R_{56} , varying between -13.1 m and -15.4 m. Optics around the usually optimal

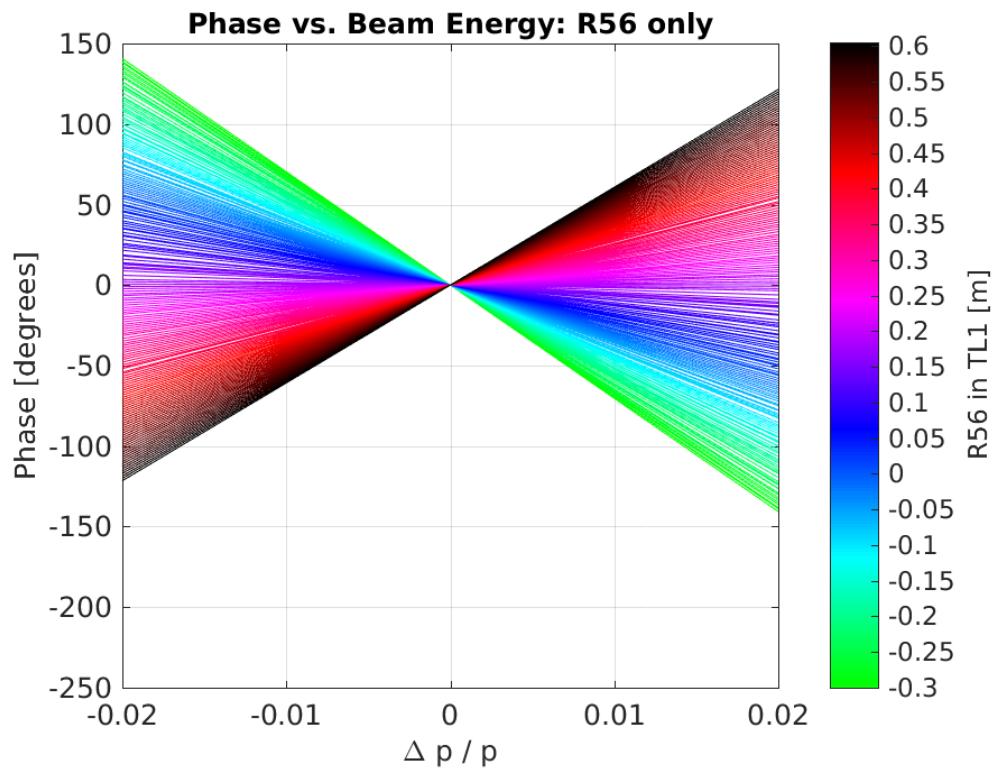


Figure 4.26: Phase shift between the upstream and downstream phase monitors for all sets of TL1 optics when only R_{56} is considered.




Figure 4.27: Phase shift between the upstream and downstream phase monitors for all sets of TL1 optics including higher orders.




Figure 4.28: Dependence of the optimal optics to use in TL1 on the beam energy offset.




Figure 4.29: Quadratic fit to the MADX phase shift output for different energy offsets, giving values for R_{56} and T_{566} .




Figure 4.30: T_{566} coefficient for all sets of TL1 optics.

R_{56} value of 0.2 m in TL1 are close to where the second order effects are minimal, with T_{566} values around -13.4 m. The T_{566} coefficients are approximately two orders of magnitude larger than the R_{56} but as $(\Delta p/p)^2 \ll (\Delta p/p) \ll 1$ the effect on the phase is smaller than for non-optimised R_{56} . For example, for a typical relative energy offset of 1×10^{-3} a residual R_{56} of 0.2 m between the upstream and downstream phase monitors leads to a phase shift of more than one degree. For the same energy offset the phase shift resulting from the second order T_{566} term is approximately 0.1 degrees. However, the key point for the phase propagation is that the first order dependence can be removed by zeroing R_{56} between the upstream and downstream phase monitors, whereas for all the available sets of optics the second order contribution will remain at roughly the same magnitude.

To determine the consequences of the T_{566} for the PFF system the effect it has on the upstream-downstream phase correlation and downstream phase jitter must be calculated. This was done analytically using the equations in Section 4.3.3 for the first order R_{56} , but for the second order terms a simple Monte Carlo simulation approach has been used. Correlated random distributions are created in MatLab to match the typical CTF3 upstream phase and energy conditions — namely $\sigma_p = 0.001$, $\sigma_u = 0.8^\circ$ and $\rho_{up} = 0.2$. The simulated downstream phase for each set of TL1 optics is then calculated using Equation 4.17 and the known R_{56} and T_{566} values. The jitter of this simulated downstream phase and its correlation with the initial upstream phase distribution give the values shown in the following figures.

The solid lines in Figure 4.31 show the downstream phase jitter versus the residual R_{56} between the upstream and downstream phase monitors in the case where only the first order R_{56} term is included (red) and when both the R_{56} and the second order T_{566} are included (blue). The effect of including T_{566} is very small, with the downstream phase jitter only

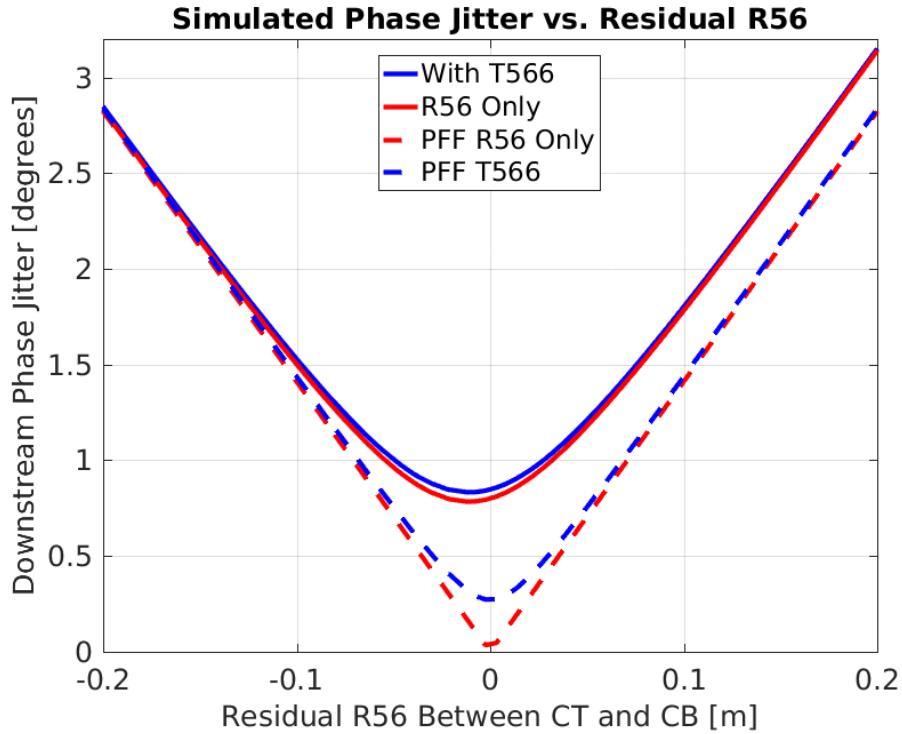


Figure 4.31: Downstream phase jitter vs. residual R_{56} in TL1 including the effects of T_{566} with $\sigma_u = 0.8^\circ$, $\rho_{up} = 0.2$ and $\sigma_p = 1 \times 10^{-3}$.

increasing from 0.80° to 0.85° degrees at the optimal residual R_{56} of zero. The effect of including T_{566} on the upstream-downstream phase correlation is much more significant for the PFF correction. The maximum achievable correlation (excluding the effects of the phase monitor resolution) is reduced from $\rho_{ud} = 1$ with only the first order term to $\rho_{ud} = 0.95$ when T_{566} is included. This is shown in Figure 4.32. This increases the achievable corrected downstream phase jitter with the PFF system from zero to 0.27° (again excluding the phase monitor resolution).

With the initial conditions and optics used here it would therefore be impossible to achieve 0.2° phase stability at CTF3. However, the relative energy jitter of 1×10^{-3} is in fact somewhat pessimistic for the conditions that can be achieved at CTF3, at least on short time scales. Figure 4.33 shows how the maximum achievable upstream-downstream phase correlation varies with the relative energy jitter. An upstream-downstream phase correlation of 97% is required to make achieving 0.2° downstream phase jitter at CTF3 possible. This can be achieved with a relative energy jitter of 0.85×10^{-3} if the R_{56} is perfectly optimised. In good conditions, especially with recent developments (Section 4.5.3), the CTF3 energy jitter can be reduced to around 0.5×10^{-3} , in which case correlations up to 99.6% are theoretically achievable. It should therefore still be possible to achieve the necessary conditions for the PFF system at CTF3 even after taking in to account T_{566} .

One final consequence of the non-zero T_{566} between the upstream and downstream phase monitors is the effect of energy variations along the pulse, or equivalently cases where the mean value of $\Delta p/p$ is non-zero. All previous calculations have assumed the energy jitter to be about a mean $\Delta p/p$ of zero, but this can not be true for all points along the CTF3




Figure 4.32: Upstream-downstream phase correlation vs. residual R_{56} in TL1 including the effects of T_{566} with $\sigma_u = 0.8^\circ$, $\rho_{up} = 0.2$ and $\sigma_p = 1 \times 10^{-3}$.

pulse due to the variations in mean beam energy along the pulse seen in Section 4.3.1. The effect of mean energy offsets on the upstream-downstream phase correlation is shown in Figure 4.34, in this case assuming a relative energy jitter of 0.5×10^{-3} about the offset mean value and zero R_{56} . Typically the energy variation along the CTF3 pulse is at the $\pm 2 \times 10^{-3}$ level, and this by itself can cause the correlation to drop below 90%. As a result the achievable corrected downstream phase jitter with the PFF system will also vary along the pulse. Without reducing either the energy variations along the pulse, or changing the optics to decrease the magnitude of T_{566} , it is unlikely that 0.2° sample-by-sample jitter along the pulse can be achieved for more than very short portions of the pulse where the energy is optimal. Reducing energy variations along the pulse to below $\pm 1 \times 10^{-3}$ would allow correlations above 96% to be achieved across the full pulse length, increasing the feasibility of hitting the 0.2° target. New feedbacks have been developed at CTF3 to try to achieve this (Section 4.5.3).

4.5.2 R56 Scans whilst Varying Beam Energy

By intentionally varying the CTF3 beam energy to artificially increase the energy jitter during an R_{56} scan the energy dependent effects in both the upstream and the downstream phase are amplified. This has the benefit of increasing the visibility of the higher order effects, but it also improves the results of the scan in general by reducing the sensitivity to other small drifts in beam conditions. In this section the results of an R_{56} scan in which




Figure 4.33: Best possible upstream-downstream phase correlation vs. beam energy jitter both with and without including the effects of T_{566} .




Figure 4.34: Upstream-downstream phase correlation vs. relative beam energy offset.

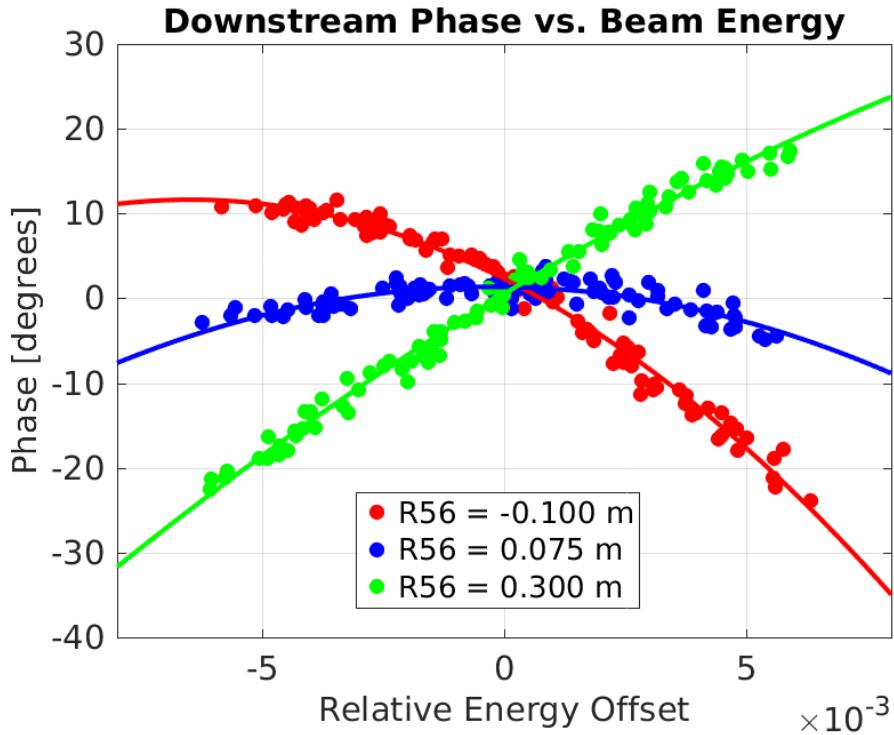


Figure 4.35: Downstream phase vs. energy for three different R_{56} settings in TL1.

the R_{56} value in TL1 was varied between -0.1 m and $+0.3$ m whilst the beam energy was varied by approximately 1% peak-to-peak are discussed. The resulting relative energy jitter of 3×10^{-3} in these conditions is 3–5 times larger than the relative energy jitter in nominal conditions. Direct observations of the effect of higher order energy dependencies during the scan will be presented first, before discussing the overall results of the scan to expand upon the conclusions from the R_{56} scans shown in previous sections.

Energy Dependence

Figure 4.35 shows the dependence of the mean downstream phase on the beam energy for three of the R_{56} values set in TL1 during the scan – the lowest value of -0.1 m, the maximum value of 0.3 m and a mid-range value of 0.075 m, which gives the lowest downstream phase jitter during the scan as seen later in this section. This plot mimics the simulated result from MADX in Figure 4.27. With the increased energy jitter the non-linear dependence of the downstream phase on the energy is clear. The first order effect means changing the R_{56} value changes the gradient of the phase-energy dependence about the central energy. The effect of T_{566} means there is an energy dependent R_{56} value that minimises the gradient. For example, with an R_{56} of -0.1 m set in TL1 energy offsets around -5×10^{-3} give the minimum energy dependence.

The coefficients of a quadratic fit to the curves in Figure 4.35 (plus their equivalents for other points in the scan) give estimates for the residual R_{56} and T_{566} values between the upstream and downstream phase monitors. The results of doing this are shown in Figure 4.36 for the fitted R_{56} values, and Figure 4.37 for the fitted T_{566} values.

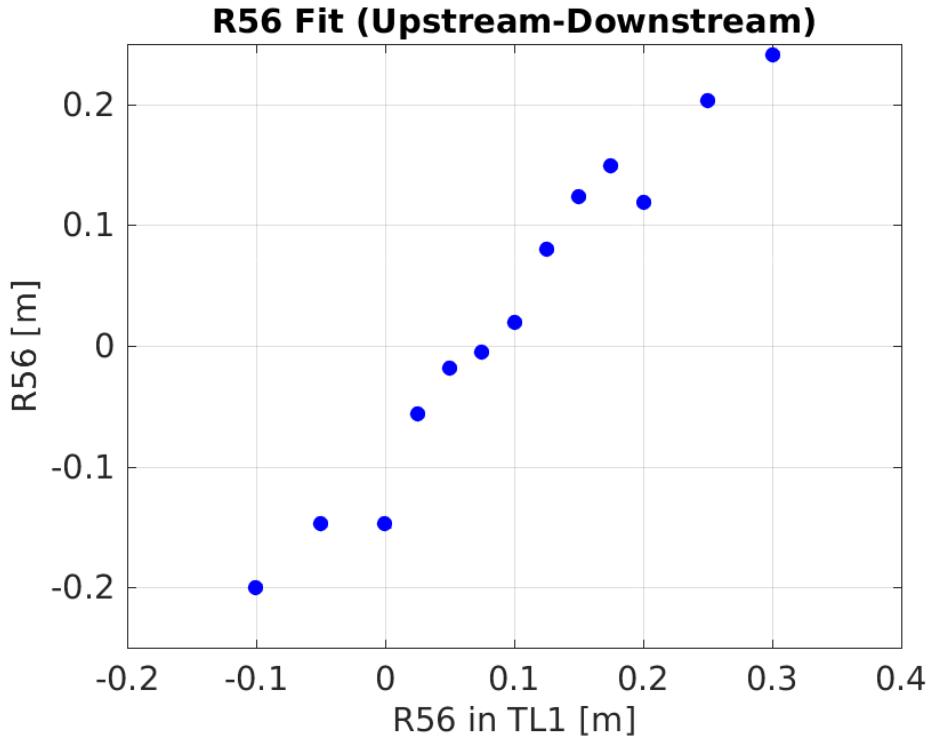


Figure 4.36: Fitted R_{56} values for all points in the scan.

The fit R_{56} values roughly follow a linear dependence on the set R_{56} optics in TL1, as expected. With an R_{56} value of -0.2 m in TL2, varying the R_{56} between -0.1 and +0.3 m in TL1 would be expected to give a total residual R_{56} of between -0.3 m and +0.1 m. The fitted range is between -0.20 m and +0.25 m, so is +0.1 m offset compared to expectations. The fitted R_{56} values depend on the absolute mean energy offset, which is not easily verifiable from the BPM measurement measurement alone. Offsetting the mean relative energy offset from zero to 2×10^{-3} would bring the fitted R_{56} values close to the expected range.

For R_{56} values in TL1 up to 0.2 m the fitted T_{566} values in TL1 in Figure 4.37 are close to the simulated values typically around -13.5 m in Figure 4.29. For the points above 0.2 m R_{56} in TL1 in the scan, the fitted T_{566} value is smaller than expected. However, the measured first and second order energy dependent effects are overall in very good agreement with expectations given the constraints of measurement and accuracy of the MADX model.

Mean Phase

Figure 4.38 shows the mean phase jitter during this R_{56} scan both upstream and downstream. Simulations of the expected phase jitter given the beam conditions and optics are also shown, both for the case where only R_{56} is considered and when both R_{56} and T_{566} are taken in to account. Varying the beam energy during the scan has the effect of increasing the upstream phase jitter from its typical level of 0.8 degrees to 2.0 degrees. The correlation between the upstream phase and the beam energy is also increased from 20% in normal conditions to above 90% whilst varying the beam energy. An example of this is shown in Figure 4.39 for the dataset at $R_{56} = -0.1$ m in TL1. The likely source of the upstream phase-energy

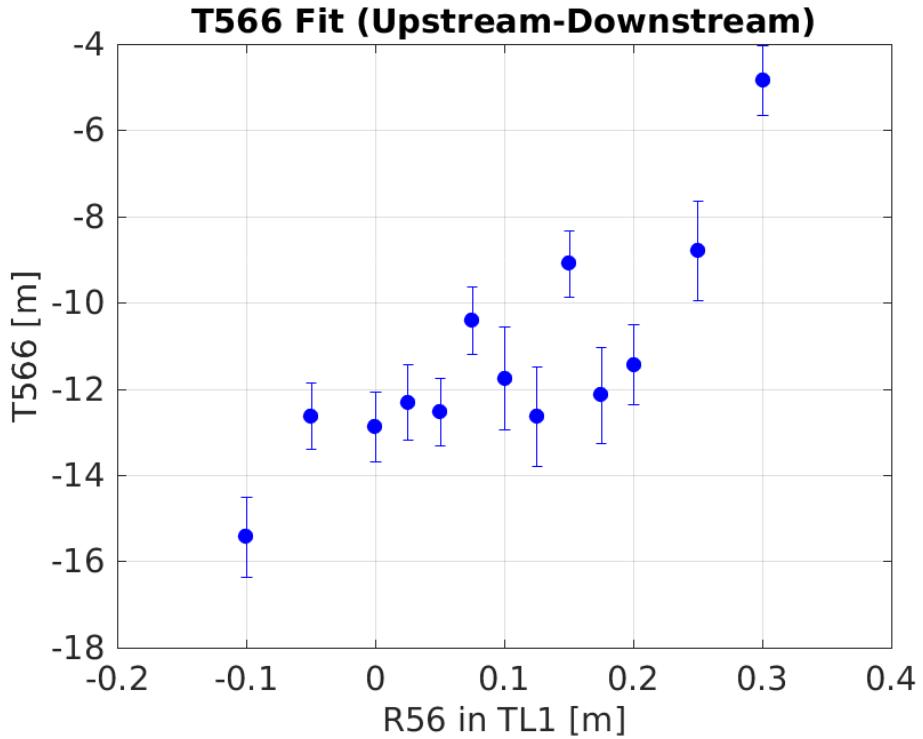


Figure 4.37: Fitted T_{566} values for all points in the scan.

dependence is the energy variation leading to differences in beam orbit through the stretching chicane in the CTF3 linac (see Figure 1.5).

The downstream phase jitter is reduced to close to the level of the upstream phase jitter for R_{56} values between 0.5 m and 0.1 m in TL1. Now knowing that the T_{566} can cause a dependence between the optimal R_{56} value in TL1 and the beam energy it is not completely unexpected that the lowest downstream phase jitter is not at the 0.2 m expected due to the optics in TL2. An mean relative energy offset of -2×10^{-3} can lead to the minimum jitter being shifted to 0.075 m as seen in the scan. This offset has been used to create the simulated “ T_{566} Sim” lines in Figures 4.38 and 4.40. The simulation including this energy offset and the effects of T_{566} follows the actual downstream phase jitter during the scan much more closely than the simulation including only R_{56} . There are still differences between the data and the T_{566} simulation, in particular in the range between $R_{56} = 0.125$ m and 0.175 m in TL1. Some potential sources of additional downstream phase jitter are investigated in Section 4.7, but these can not explain the differences seen in this scan so there are remaining effects that have not yet been identified.

The upstream-downstream phase correlation during the scan is shown in Figure 4.40. As the correlation between the upstream phase and the beam energy is greatly increased as a result of varying the beam energy as previously discussed, there is no longer a clear singular peak in the upstream-downstream phase correlation during the scan. Instead, the correlation quickly flips between a highly correlated state and a highly anti-correlated state. This effect was previously seen in Figure 4.13 as a consequence of the R_{56} equations derived in Section 4.3.3. The actual upstream-downstream phase correlation during the scan is in good agreement with the simulation including the effects of T_{566} (again with an assumed

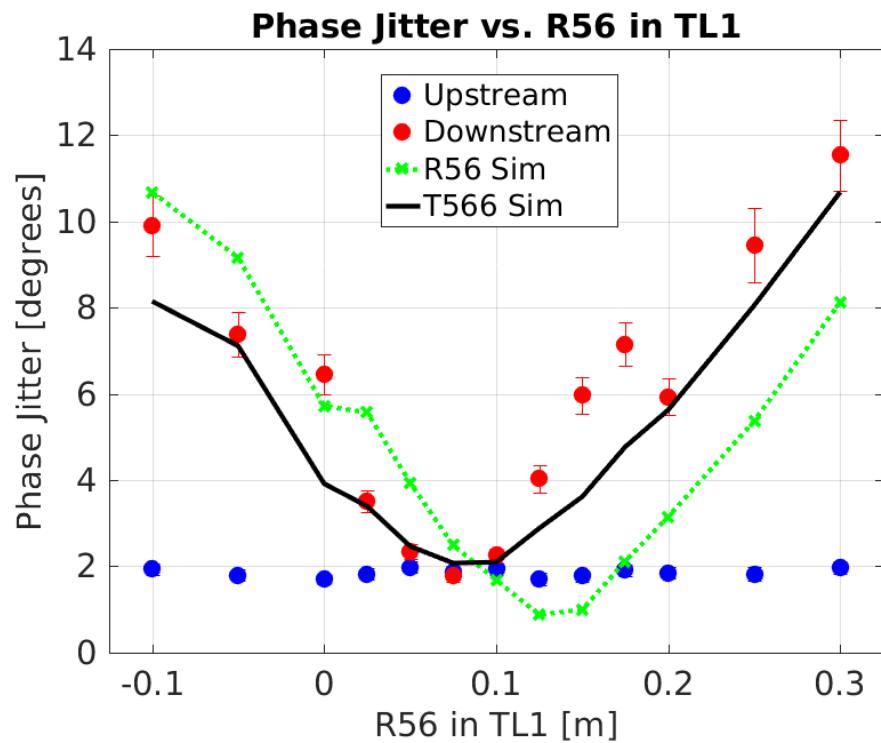


Figure 4.38: Phase jitter for different R56 whilst varying beam energy.

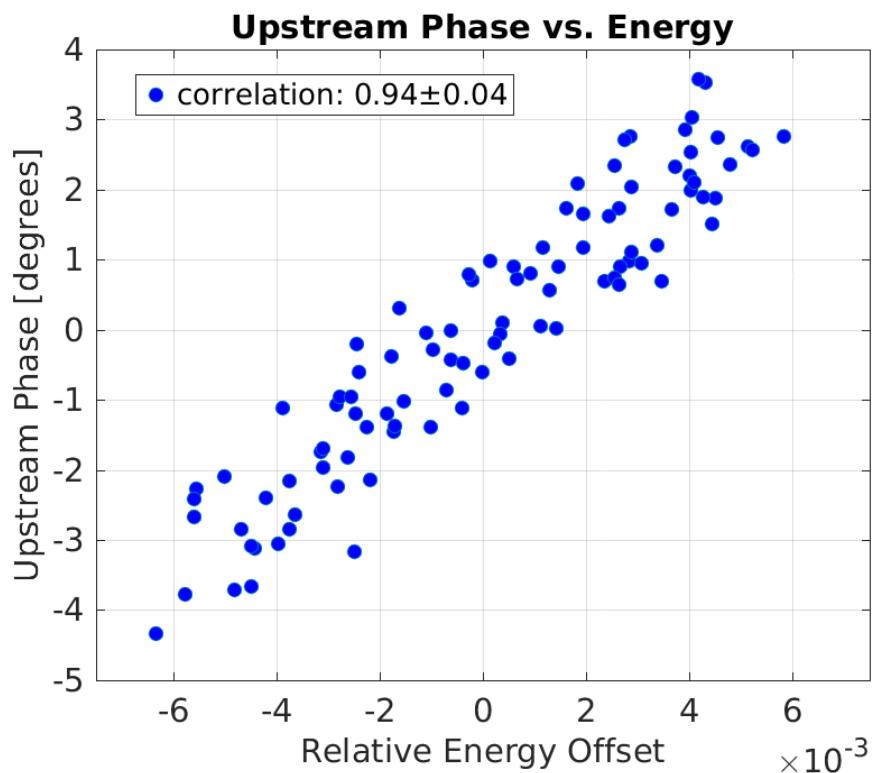


Figure 4.39: Upstream phase-energy correlation whilst varying the beam energy.

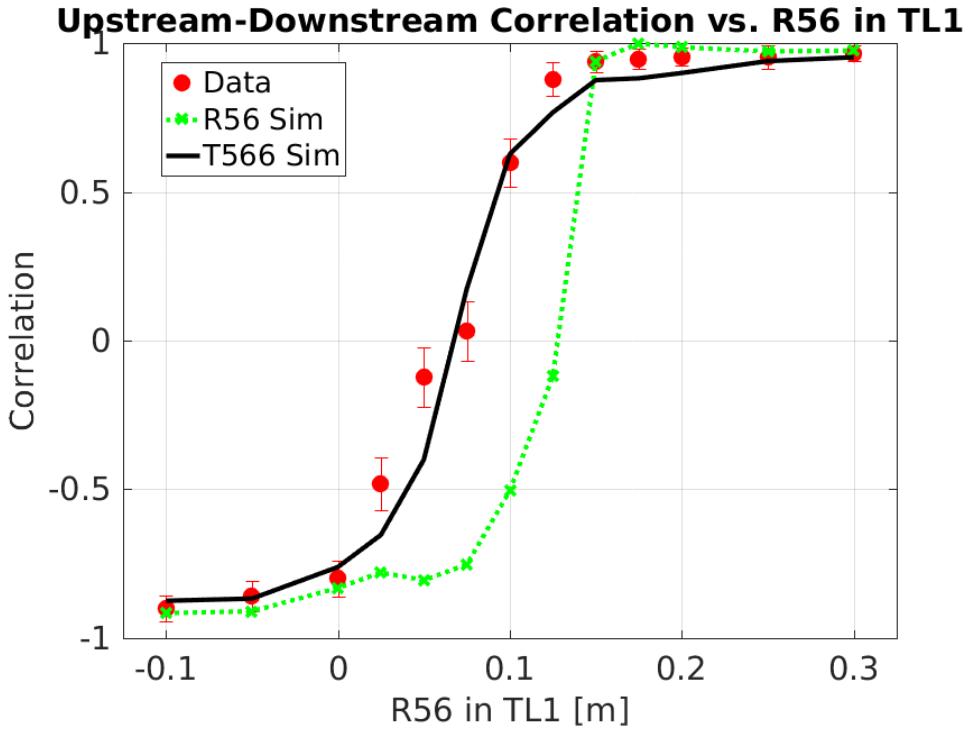


Figure 4.40: Upstream-downstream phase correlation for different R_{56} in TL1 whilst varying beam energy.

mean energy offset of -2×10^{-3}). The optimal R_{56} for the phase propagation in terms of the upstream-downstream phase correlation is in the region around $R_{56} = 0.175$ m in TL1. However, for these R_{56} values the downstream phase jitter is much larger than the upstream phase jitter due to the effects of the upstream phase-energy correlation and T_{566} in the high energy jitter conditions during the scan. This makes it more difficult to precisely define the best R_{56} optics to use based on the results of a scan of this type alone.

Phase Along the Pulse

As well as the mean phase it is interesting to look at the effect of varying the R_{56} on the phase along the pulse. To understand differences in the downstream phase along the pulse it is important to know the properties of the beam energy along the pulse during the R_{56} scan. Figures 4.41 and 4.42 show typical examples of the mean energy and the energy jitter along the pulse. During this scan the mean beam energy was constantly varied about its initial value, as noted previously. This has no effect on the mean beam energy along the pulse, but does increase the energy jitter along the pulse. The relative beam energy offset along the pulse (Figure 4.41) varies by 3.5×10^{-3} peak-to-peak. The energy jitter along the pulse varies between 4×10^{-3} and 3×10^{-3} , with better stability towards the end of the pulse.

Figure 4.43 shows the mean phase along the pulse for each R_{56} setting in TL1 during the scan. Any difference in the mean (rather than the jitter) along the pulse with the R_{56} value should originate from variations in the mean energy along the pulse. If the energy along the pulse was constant changing the R_{56} would only affect the phase jitter and would not

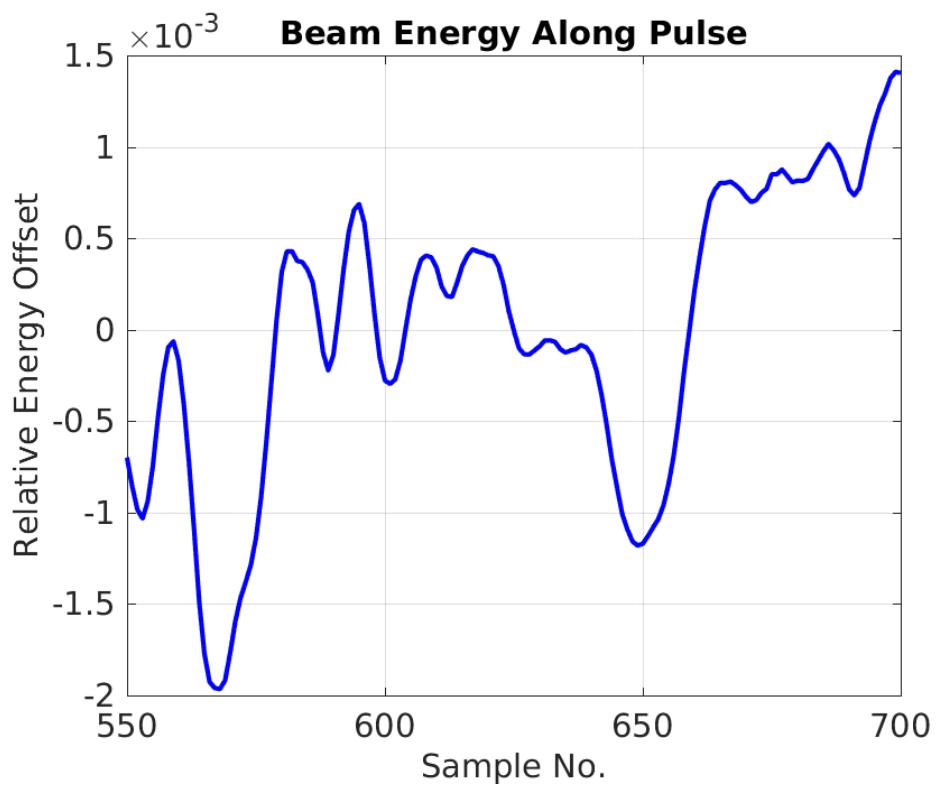


Figure 4.41: Typical variations in mean energy along the pulse during the R_{56} scan.

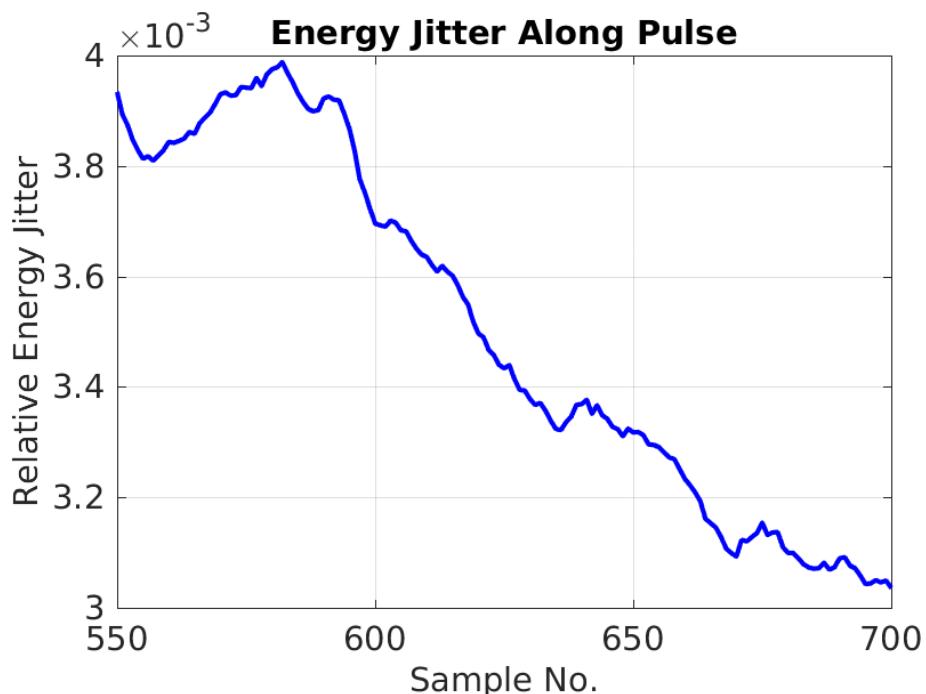


Figure 4.42: Typical energy jitter along the pulse during the R_{56} scan.

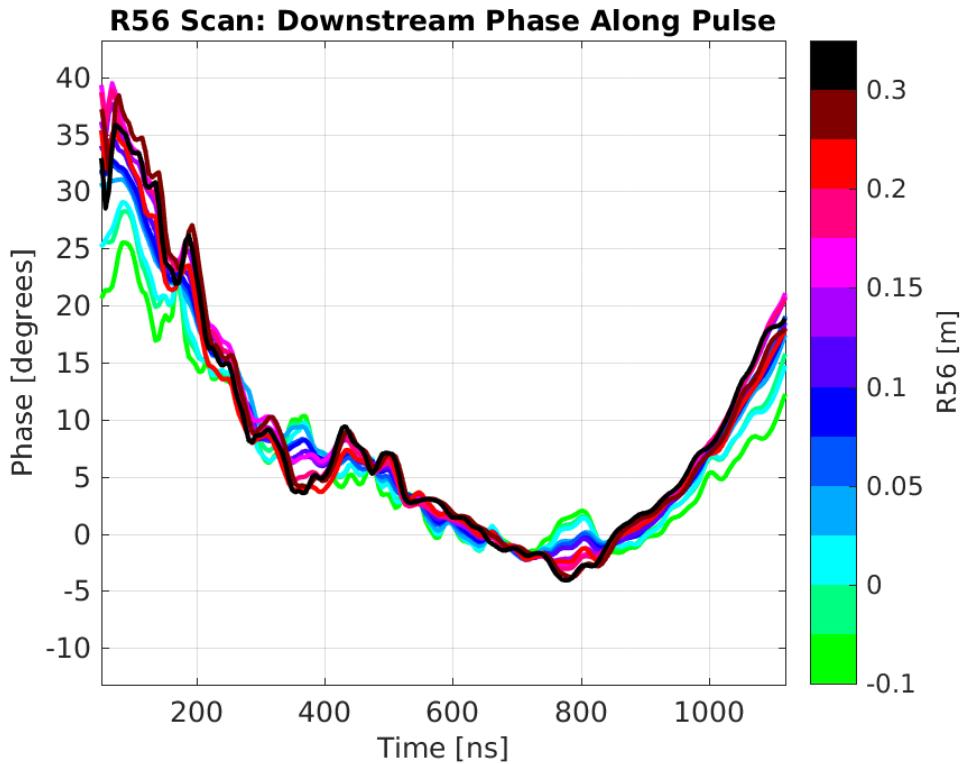


Figure 4.43: Mean downstream phase along the pulse during the R_{56} scan.

change the mean pulse shape. The clear change in certain features along the pulse in the downstream phase is therefore an indication of energy variations in these regions. Perhaps the best example of this is the oscillation around a time of 800 ns, where the phase is flat close to the optimal R_{56} value of 0.1 m but swings upwards when a negative R_{56} in TL1 is used or downwards for R_{56} values above 0.15 m.

The difference between the phase along the pulse for two different settings of R_{56} in TL1 should be proportional to the beam energy along the pulse. Figure 4.44 plots the difference between the $R_{56} = +0.3$ m optics and the roughly optimal $R_{56} = +0.175$ m optics, and compares this to the beam energy along the pulse. Both lines are mean subtracted and normalised to give equivalent amplitudes in arbitrary units. Overall, the differences in phase along the pulse resulting from using non-optimal R_{56} in TL1 are very well matched with the energy variation along the pulse, as expected.

Figure 4.45 shows the downstream phase jitter along the pulse for each R_{56} optics in TL1 used during the scan. Like the mean phase jitter, the jitter along the pulse is lowest for R_{56} values between 0.05 m and 0.1 m in TL1 (blue). Lower R_{56} values (green) and higher R_{56} values (purple, red and black) give higher jitter. Close to the optimal R_{56} value many of the variations in jitter along the pulse are reduced, although in all cases some features remain.

Differences in the phase jitter along the pulse between two different R_{56} optics in TL1 should also be related to the beam energy. One might expect the features in the downstream phase jitter along the pulse to match the shape of the variations in energy jitter along the pulse shown previously. Figure 4.46 compares the difference in phase jitter along the pulse




Figure 4.44: Difference between the mean phase along the pulse with $R_{56} = 0.3$ m and 0.175 m compared to the beam energy along the pulse.




Figure 4.45: Downstream phase jitter along the pulse during the R_{56} scan.

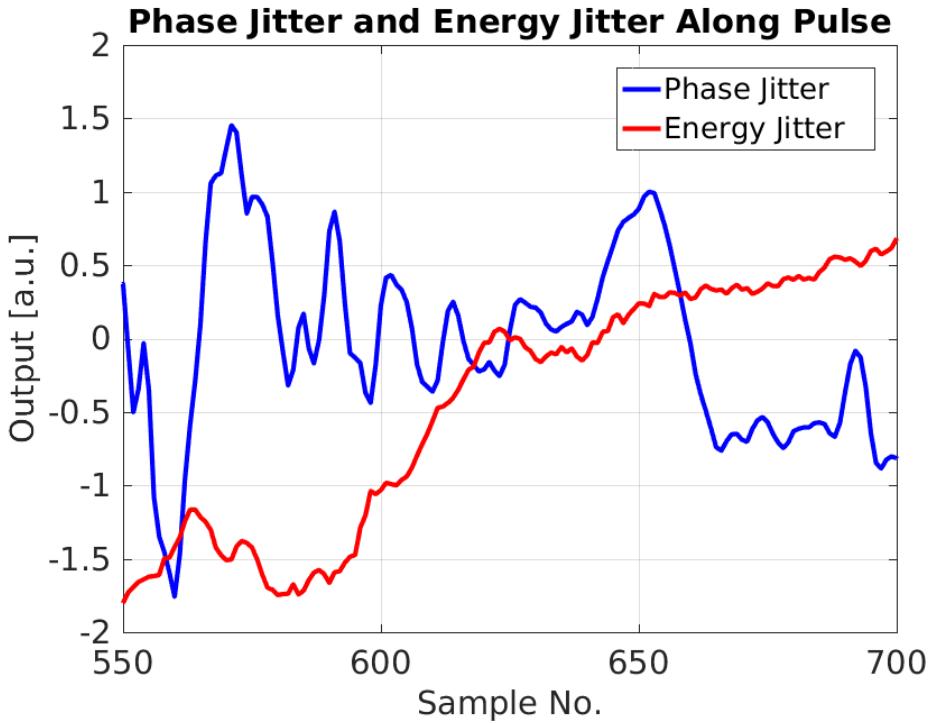


Figure 4.46: Difference between the phase jitter along the pulse with $R_{56} = 0.3$ m and 0.175 m compared to the energy jitter along the pulse.

for the $R_{56} = 0.075$ m and $R_{56} = -0.1$ m optics to the energy jitter along the pulse. There is no clear similarity between the two. In Figure 4.47 the phase jitter along the pulse is compared to the mean energy along the pulse instead. In this case most of the features in the downstream phase jitter are also present in the energy along the pulse. Static variations in the mean energy along the pulse are therefore more critical for the phase propagation than differences in energy jitter along the pulse. This is true even in the conditions used in this scan, in which the energy jitter was artificially increased.

4.5.3 Mitigation of Higher Order Dependencies

The first order energy dependent effects on the downstream phase due to the R_{56} in TL2 have largely been successfully removed by adding positive R_{56} to the TL1 line. If necessary, further improvements could be made by creating further sets of TL1 optics in smaller R_{56} steps. Although the higher order energy dependent effects due to T_{566} have been identified they are more difficult to remove or reduce. To remove the effect completely new optics would have to be matched with constraints included to zero the T_{566} value between the upstream and downstream phase monitors. This would likely require the use of sextupoles, which are typically left unpowered at CTF3 due to difficulties in commissioning optics that include them [REF]. It also depends on the accuracy of the MADX model of CTF3 to second order, although the result in Figure 4.37 is promising in this regard, at least in the R_{56} range of interest.

Alternatively, the second order effects can be reduced by decreasing the energy jitter

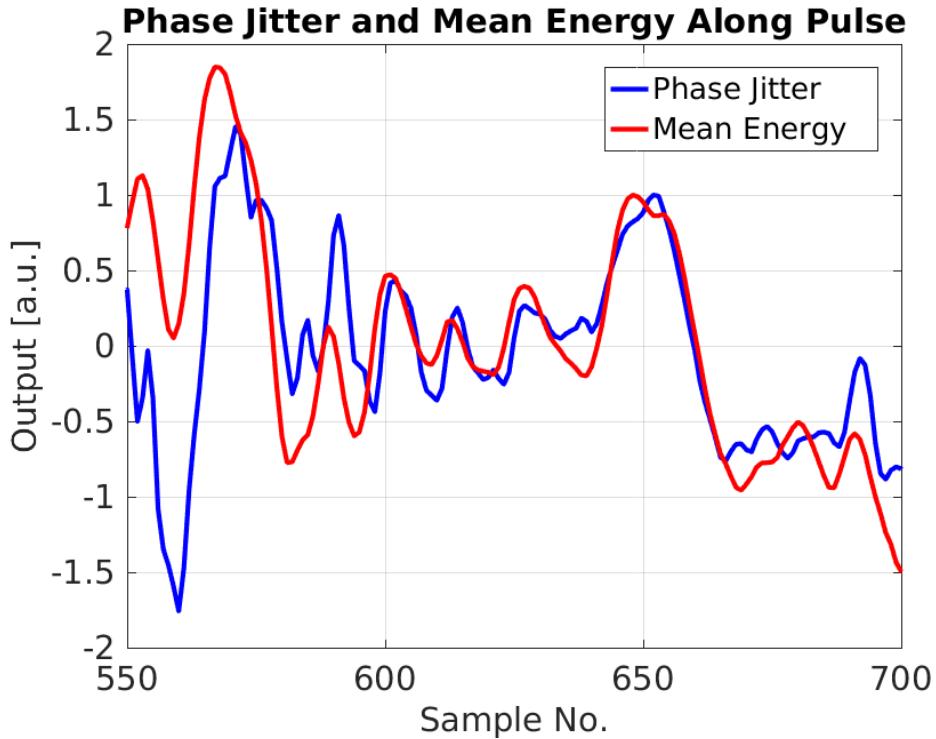


Figure 4.47: Difference between the phase jitter along the pulse with $R_{56} = 0.3$ m and 0.175 m compared to the mean beam energy along the pulse.

and variations along the pulse at CTF3. Many improvements have recently been made at CTF3 to achieve this via the implementation of several new feedbacks [REF LukasIPAC]. New feedbacks on the phases of each klystron in the CTF3 injector as well as on the beam gun current are now routinely in operation, and these have already improved the mean energy jitter to 0.5×10^{-3} on short time scales of several minutes, or 0.8×10^{-3} on longer time scales. Hardware changes and further improvements to the implementation of these feedbacks should allow the mean energy stability to be improved further in 2016.

In addition, although they are not yet run online in normal operation, new feedbacks have also been implemented and commissioned that can smooth energy variations along the pulse, or directly smooth the upstream phase along the pulse itself. The energy flattening feedback [REF Tobias] varies the waveform of the last klystron in the CTF3 linac based on the measurement of the same dispersive BPM used to measure the beam energy in this chapter — CT.0608 in TL1. This flattens the energy along the pulse by varying the energy gain along the pulse in the last accelerating structure. In addition, prior to the best phase propagation conditions currently achieved (presented in the next section) a similar feedback was used to smooth features in the upstream phase along the pulse by varying the waveform of the first klystron in the CTF3 linac [REF Davide]. Feedbacks of this type will be critical to be able to improve the PFF performance in 2016.

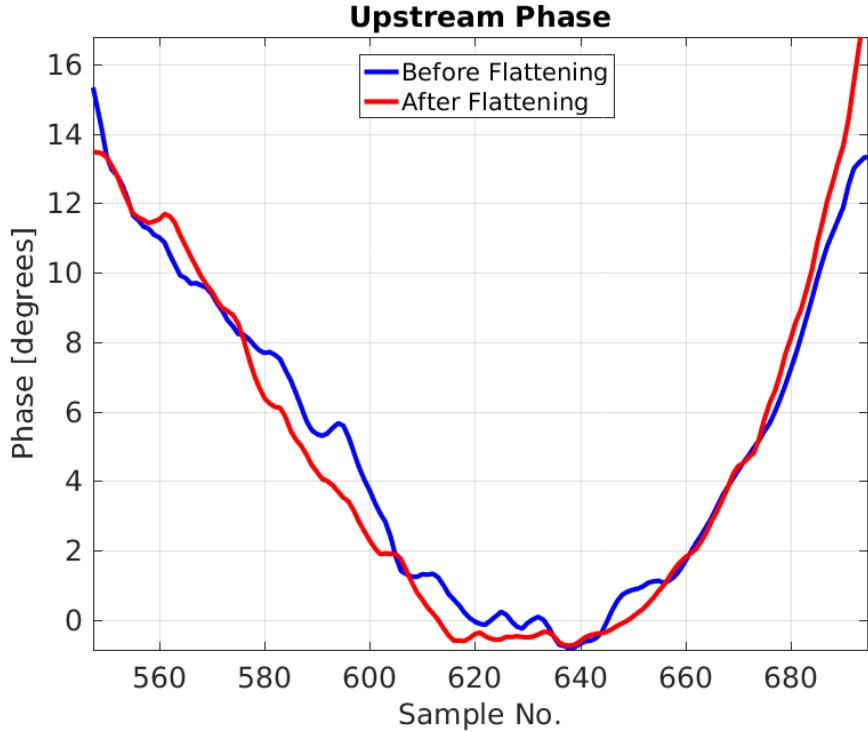


Figure 4.48: Effect of feedback used to smooth variations along the upstream phase.

4.6 Best Phase Propagation

This section summarises the best phase propagation conditions that have been achieved to date after the extensive work to identify and reduce the energy dependence of the downstream phase. The results shown here are taken from the same dataset with which the best PFF results were achieved, presented later in Section 6.1. At this time the $R_{56} = +0.1$ m optics was used in TL1. The achieved conditions are a remarkable improvement compared to the original status presented in Section 4.2. The reproducibility of these conditions is discussed in the context of the PFF results in Section 6.2.

Mean Phase

The best phase propagation conditions are achieved when the correlation between the phase and the beam energy is the same upstream and downstream. Figure 4.49 shows the correlation of the upstream and downstream phase with the beam energy. With the optimised conditions almost all correlation between the downstream phase and the beam energy is removed, with a correlation of 0.2 ± 0.1 . This agrees with the upstream-downstream phase correlation of 0.1 ± 0.1 within error bars.

Figure 4.50 shows that all drifts in the upstream phase are also present downstream, and that the upstream and downstream mean phase jitters are also the same within errors bars — $0.69 \pm 0.06^\circ$ upstream and $0.74 \pm 0.06^\circ$ downstream. The standard deviation of the residuals between the upstream and downstream phase is $0.27 \pm 0.02^\circ$, which corresponds to below

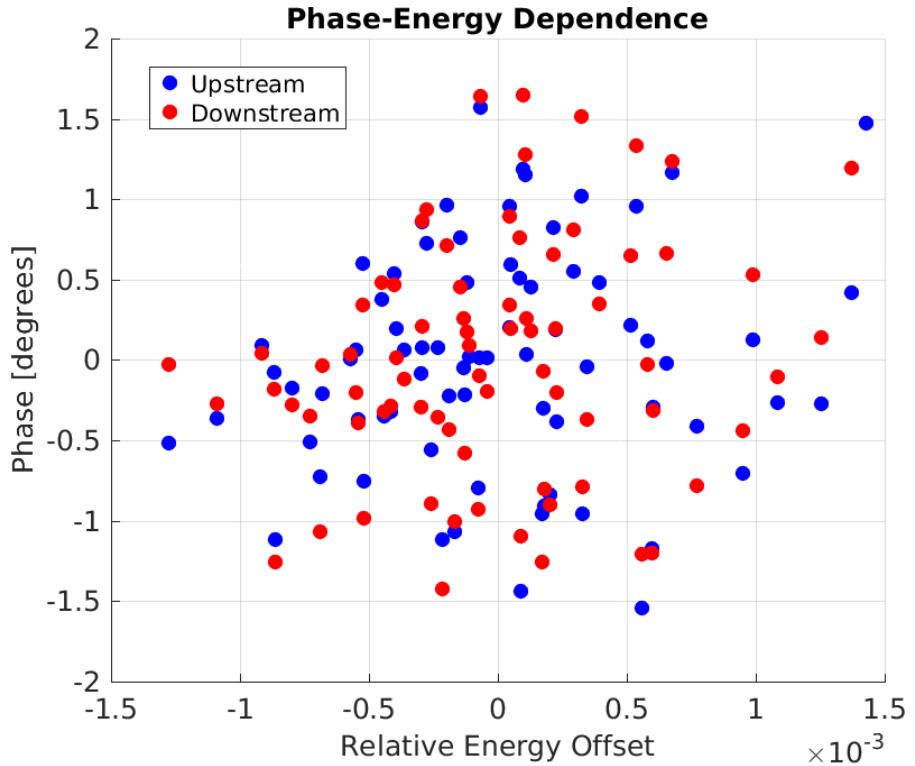


Figure 4.49: Upstream and downstream phase-energy correlation.

0.2 degrees actual beam jitter between the two phases when the phase monitor resolution is taken in to account. The correlation between the mean upstream and downstream phase is $93 \pm 4\%$, as shown in Figure 4.51. This is very close to the targeted 97% that would make a reduction in downstream jitter to 0.2° possible with the PFF system.

Phase Along Pulse

Figure 4.52 compares the phase along the pulse upstream and downstream in the optimal conditions. For reference, the mean phases presented above were calculated in the region between 530 ns and 950 ns. This is the flat central part of the pulse where the PFF system can provide the maximum reduction in phase jitter, as seen in Section 6.1. The overall agreement in shape between the upstream and downstream phase is excellent, with almost all energy dependent features also removed from the mean downstream phase along the pulse. There are still some remaining differences in shape, particularly at the start and end of the pulse before and after the 530–950 ns range.

By viewing the phase jitters and correlations along the pulse the remaining areas for improvement start to become apparent. In Figure 4.53, showing the phase jitter along the pulse, it is clear that the start of the pulse is less stable. For times between 100 ns and 600 ns the downstream phase jitter is as much as 60% larger than the upstream phase jitter. However, there is a 400 ns portion of the pulse where the downstream phase jitter has been successfully reduced to the level of the upstream phase jitter. The best PFF results can be expected in this region.

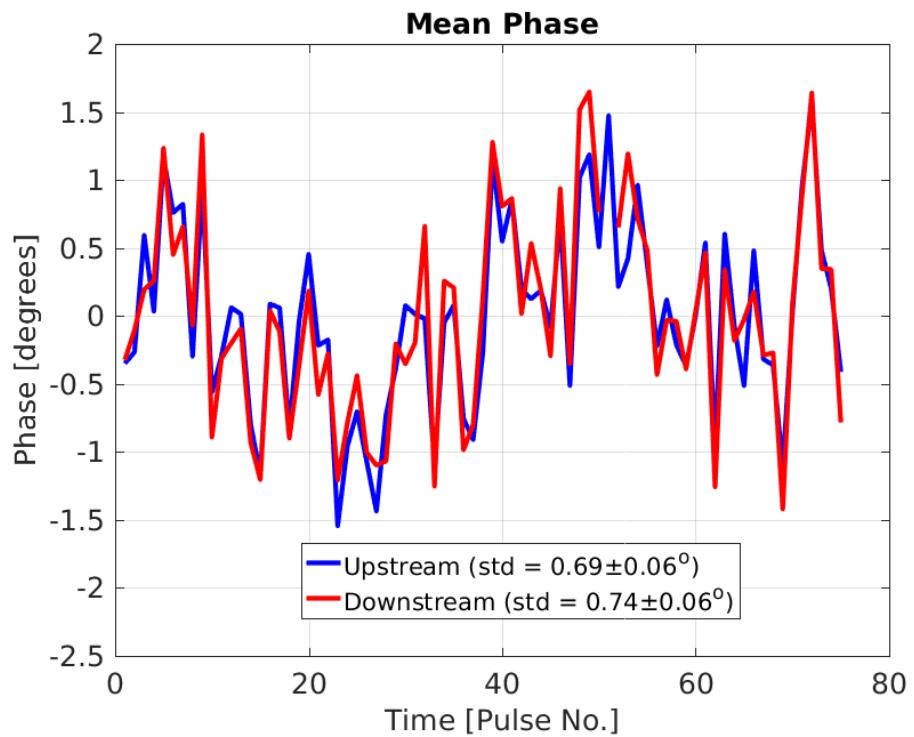


Figure 4.50: Mean phase vs. time upstream and downstream.

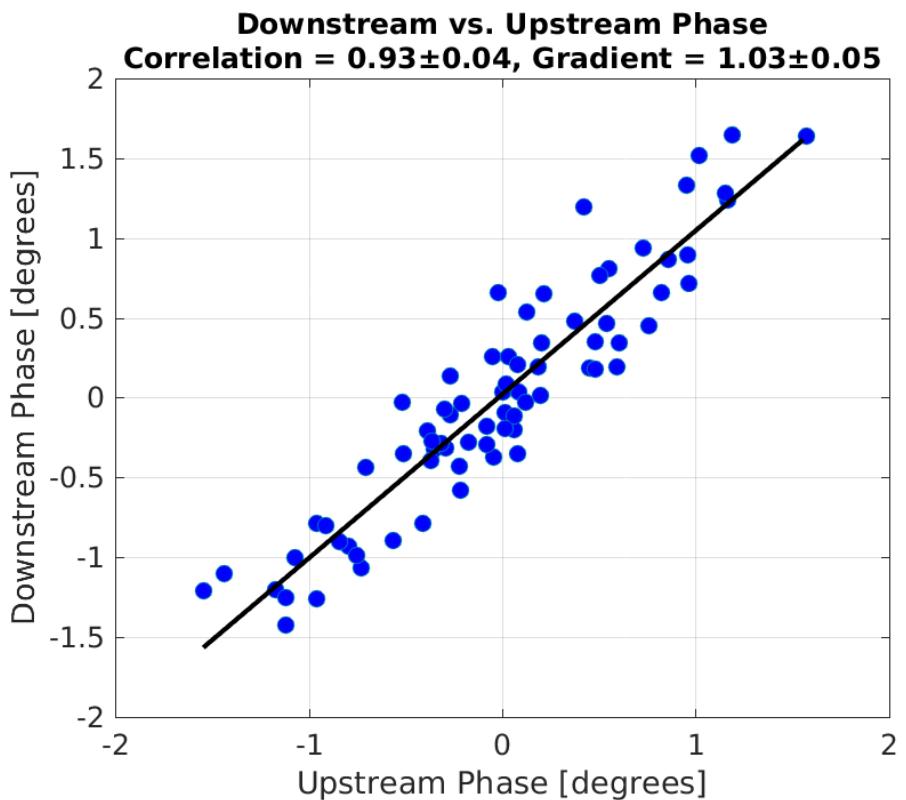


Figure 4.51: Upstream-downstream phase correlation.

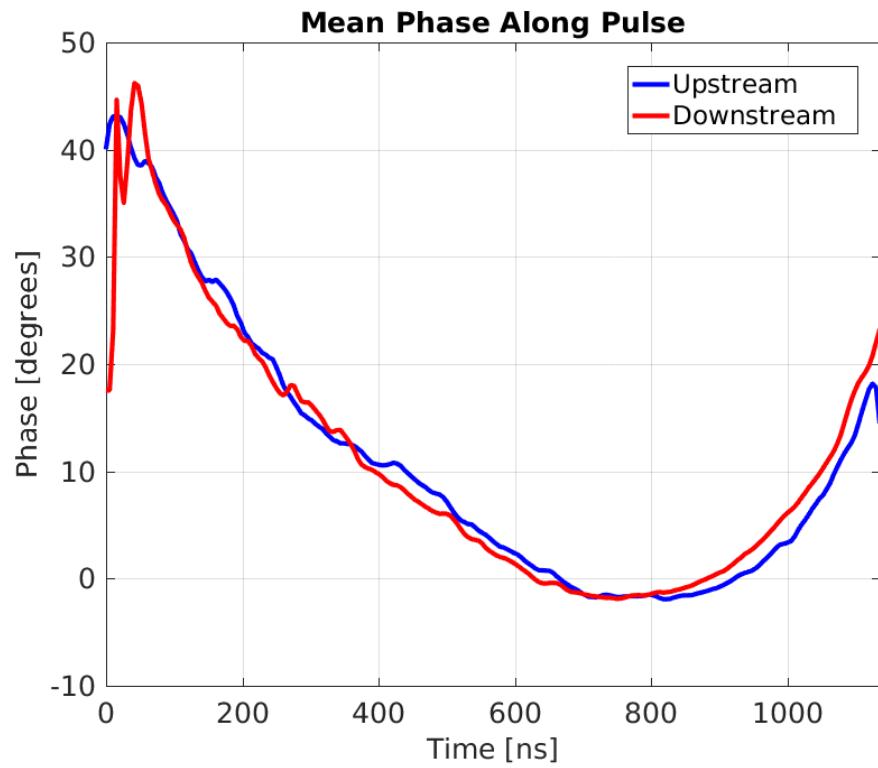


Figure 4.52: Phase along the pulse upstream and downstream.

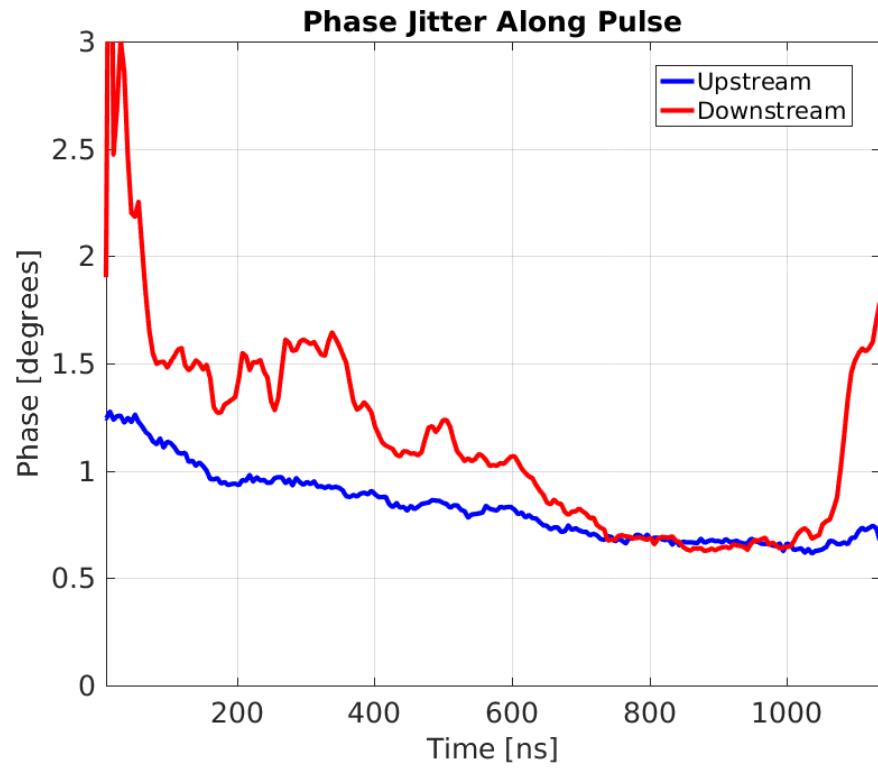


Figure 4.53: Phase jitter along the pulse upstream and downstream.

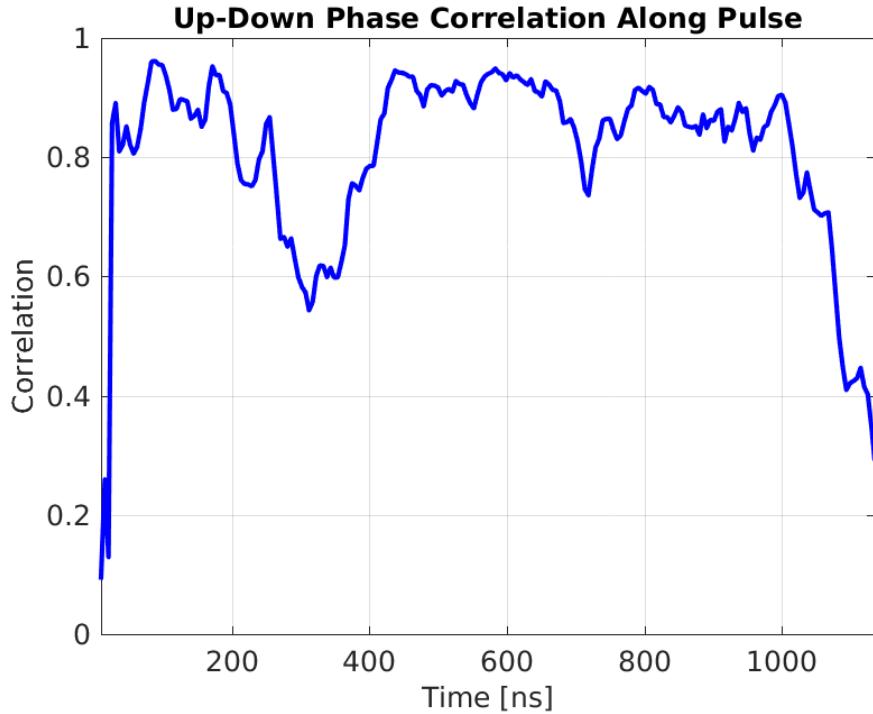


Figure 4.54: Upstream-downstream phase correlation along the pulse.

The upstream-downstream phase correlation along the pulse, in Figure 4.54, is at the same level as the mean phase correlation in the region between 450 ns and 650 ns. However, there are also parts of the pulse where the upstream-downstream phase correlation is greatly reduced – to 55% at 300 ns and 70% at 700 ns, for example. Differences earlier in the pulse are not critical for the PFF system, but the feature at 700 ns is in the area where the correction is attempted.

By comparing the upstream-downstream correlation along the pulse to the correlations with energy along the pulse, in Figure 4.55, the problem becomes clear. Due to energy variations along the pulse and the effect of higher order energy dependencies, as previously discussed, the phase propagation can not be optimised across the full pulse length. As a result there are parts of the pulse where the downstream phase-energy correlation is still high, up to 75%. Areas where there is a large difference between the upstream phase-energy correlation and the downstream phase-energy correlation correspond to regions of degraded upstream-downstream phase correlation in Figure 4.53. This is the key area where improvements are needed in 2016 in order to further improve the PFF performance.

4.7 Possible Other Sources of Phase Jitter

The energy related effects on the downstream phase have been the main focus of attempts to improve the phase propagation for the PFF system, and as shown in the previous section this by itself has lead to upstream-downstream correlations in excess of 90% to be achieved at CTF3. Further optimisations will be needed to achieve the 97% correlation required to

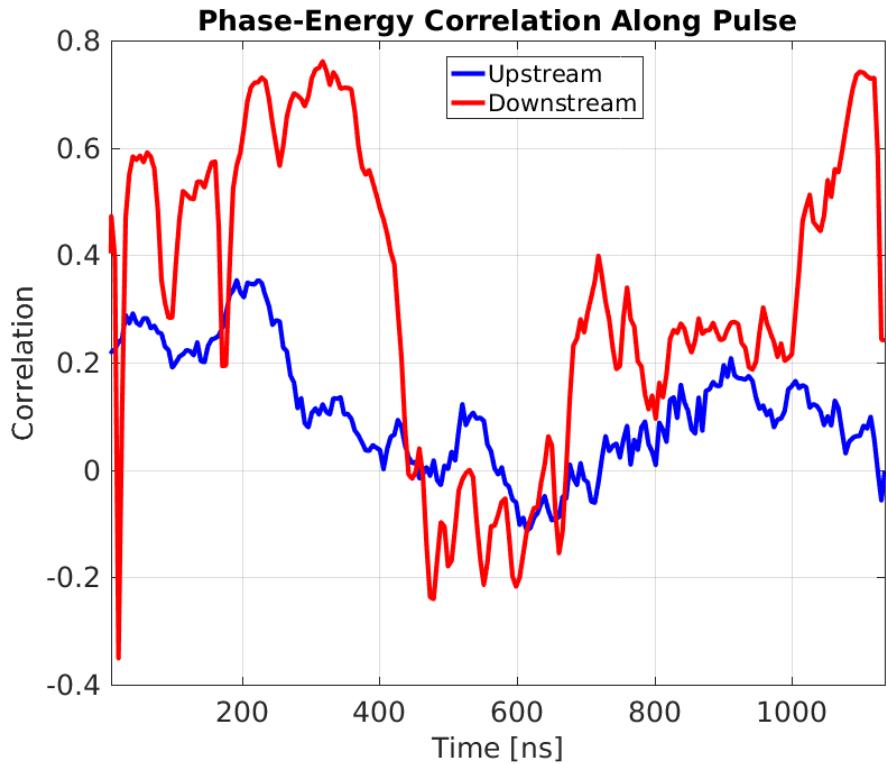


Figure 4.55: Upstream and downstream phase-energy correlation along the pulse.

achieve the CLIC target of 0.2 degrees phase jitter with the PFF prototype, however. This can partly come from further improvements of the CTF3 injector setup and stability in 2016, which will help to reduce any remaining effects from T_{566} by reducing beam energy jitter, drifts and variations along the pulse. At correlations above 90% any remaining small differences in the performance of the upstream and downstream phase monitors may also become significant, but this has been addressed in Chapter 3 so will not be discussed again here.

Preliminary measurements have been taken to investigate whether there may be any other instabilities at CTF3, apart from energy jitter, that can change the downstream phase and reduce the upstream-downstream phase correlation. The most likely culprits are devices between the upstream and downstream phase monitors that have a strong effect on the beam orbit. Any change in beam orbit can change the path length between the upstream and downstream phase monitors, and thus shift the downstream phase with respect to the upstream phase. The main elements for which this could be significant include the dipoles in TL1 and the combiner ring, as well as the two septum magnets used at the combiner ring injection and extraction. Instabilities on the power supplies for one of these devices could be an additional source of uncorrelated downstream phase jitter.

The current applied from the power supplies for each of these devices has been varied to determine their effect on the downstream phase. In some cases one power supply drives multiple devices, meaning these devices can not be changed independently from each other and any jitter in their strengths should be correlated (assuming the jitter source is the power supply, rather than a separate issue with the device itself). There are four power supplies

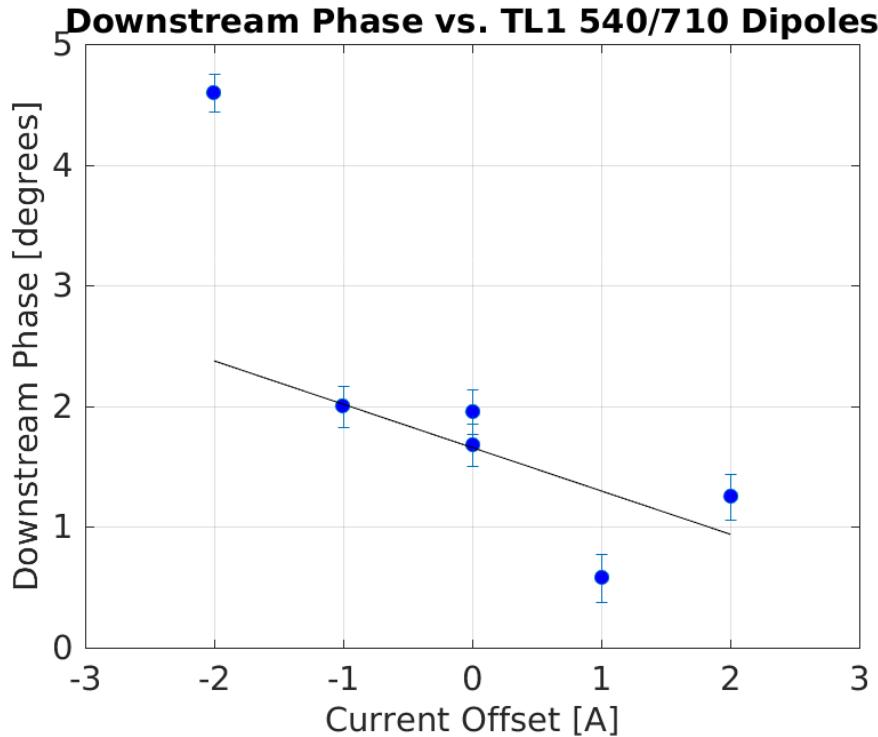


Figure 4.56: Phase vs. strength of first and last dipole in TL1 (CT.0540 and CT.0710).

that control the strength of the devices of interest in the following groups:

- **Power supply 1:** The first (CT.0540) and last (CT.0710) dipole in TL1.
- **Power supply 2:** The second (CT.0630) and third (CT.0670) dipole in TL1.
- **Power supply 3:** The combiner ring injection and extraction setpa.
- **Power supply 4:** All combiner ring dipoles.

Figures 4.56–4.59 show the effect of changing each of these power supplies on the downstream phase. A linear fit to the response is also shown. Where the response is non-linear (for the combiner ring dipoles, for example) the fit is only applied in the central region around the nominal device setting, to give an approximate gradient that is relevant for small offsets. In some cases, particularly for the 540/710 dipoles in TL1, the phase dependence is not clear. In all cases the goal is to determine whether there is an effect that may be significant for the PFF system, which can then be investigated further later, rather than quoting precise numbers.

The power supplies at CTF3 give a relative stability in the supplied current of approximately 10^{-4} . Assuming this stability the effect of each device on the downstream phase jitter can be determined by using the fitted gradients from the plots previously shown and the known current setting for each device. These values are summarised in Table 4.3. By far the strongest potential source of phase jitter appears to be the middle two dipoles in TL1 (CT.0630 and CT.0670), which by themselves could contribute 0.17 ± 0.03 degrees phase jitter. The explanation for these two dipoles having a much larger effect on the phase than the

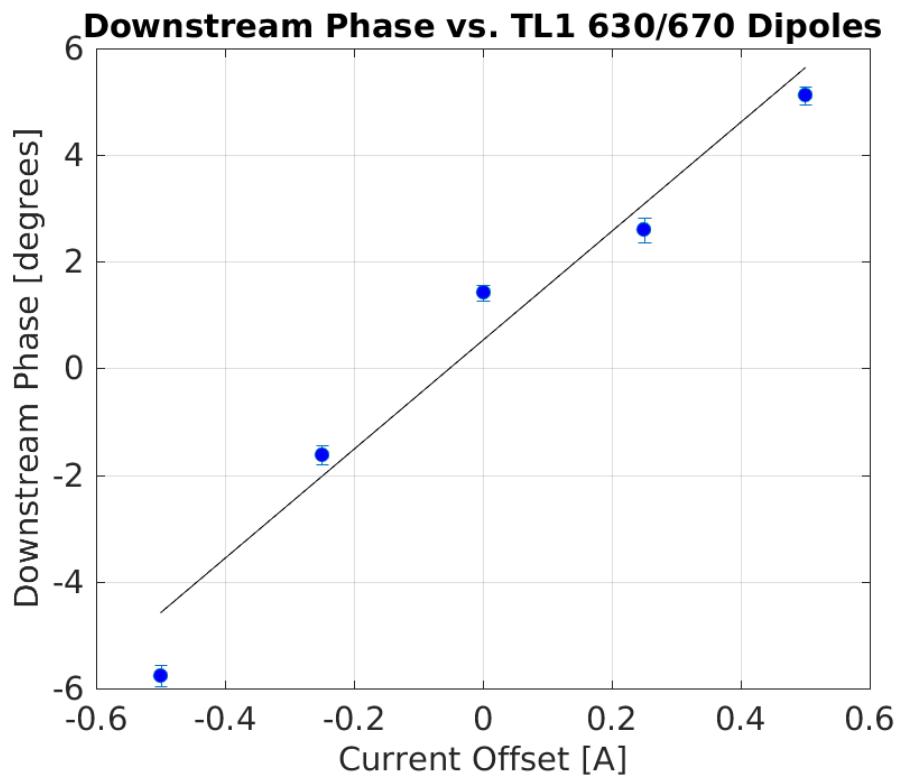


Figure 4.57: Phase vs. strength of first and last dipole in TL1 (CT.0630 and CT.0670).

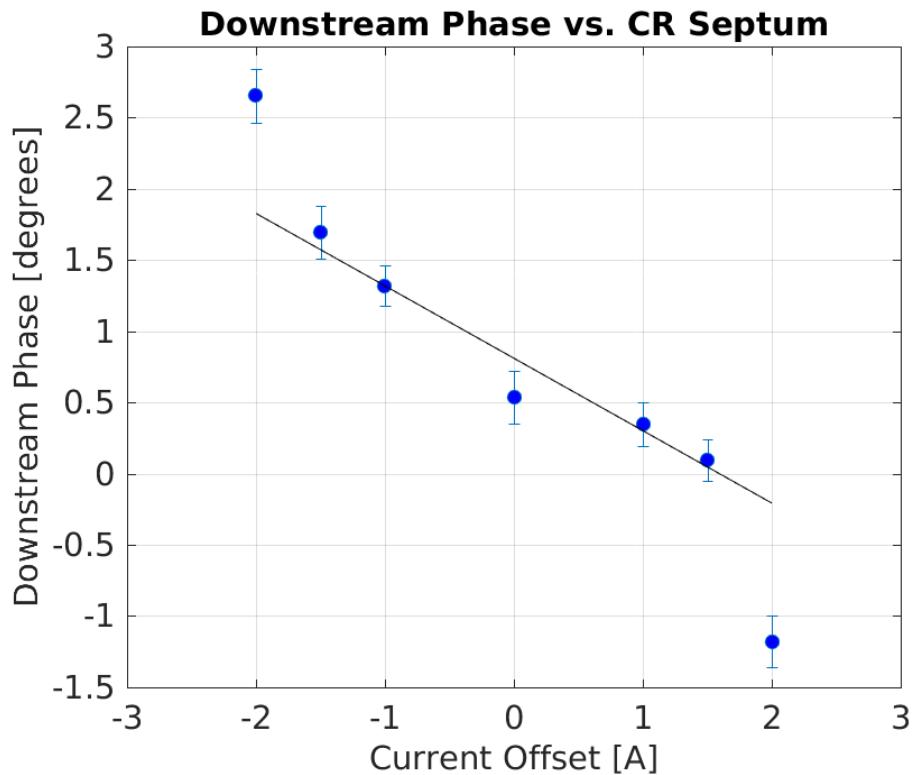


Figure 4.58: Phase vs. strength of the combiner ring injection and extraction septum magnets.

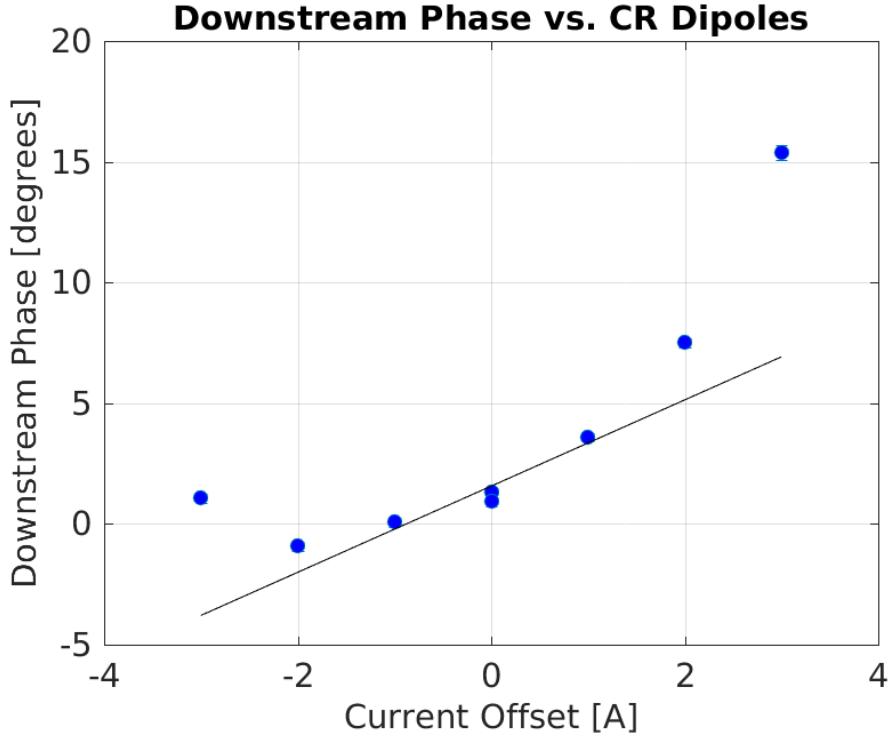


Figure 4.59: Phase vs. strength of the combiner ring dipoles).

Device	Current	Fit Gradient	Estimated Phase Jitter
TL1 540/710 Dipoles	133 A	$-0.4 \pm 0.3^\circ/\text{A}$	$0.005 \pm 0.004^\circ$
TL1 630/670 Dipoles	164 A	$10 \pm 2^\circ/\text{A}$	$0.17 \pm 0.03^\circ$
CR Septa	890 A	$-0.5 \pm 0.1^\circ/\text{A}$	$0.05 \pm 0.01^\circ$
CR Dipoles	156 A	$1.8 \pm 0.7^\circ/\text{A}$	$0.03 \pm 0.01^\circ$

Table 4.3: Current setting, dependence of the downstream phase on the current and estimated contribution to downstream phase jitter for the dipoles and septa in TL1 and the combiner ring.

first and last dipole in TL1 is not known and will need to be verified by repeated measurements and checks of the expected response in the CTF3 MADX model. The combiner ring devices contribute roughly 5 times less phase jitter than the CT.0630 and CT.0670 dipoles. Combining the estimated phase jitters resulting from these devices in quadrature gives an overall contribution of $0.18 \pm 0.03^\circ$ additional downstream phase jitter.

Modelling the downstream phase as $\phi_d = \phi_u + x$, where x is a generic additional source of jitter, the downstream jitter and upstream-downstream phase correlation are given by:

$$\sigma_d = \sqrt{\sigma_u^2 + \sigma_x^2} \quad (4.19)$$

$$\rho_{ud} = \frac{\sigma_u}{\sigma_d} \quad (4.20)$$

These are simplified forms of the equations in Section 4.3.3 in the case where the additional jitter source is uncorrelated with the upstream phase (which is not the case for the energy dependent phase jitter). Assuming an initial upstream phase jitter of 0.8° plus a $\sigma_x =$

0.18° source of jitter resulting from the power supply stabilities previously discussed, the downstream phase jitter is increased slightly to 0.82° and the upstream-downstream phase correlation reduced to 98%. This by itself does not prevent the PFF system from theoretically achieving 0.2° phase jitter, for which $\rho_{ud} = 0.97$ is required. However, in combination with the effects of R_{56} , T_{566} and the phase monitor resolution, each of which can individually reduce the correlation to 0.97 or below depending on the conditions, this highlights that the 0.2° target will be very difficult to achieve with the PFF prototype at CTF3.

Chapter 5

Setup and Commissioning of the PFF System

Figure 5.1 shows the complete PFF electronics setup in the klystron gallery, consisting of three racks. The cables from the phase monitors and kickers in the accelerator hall on the floor below enter the rack from behind. They are connected to the patch panels at the lower left side of the figure, for the phase monitors, and below the amplifier in the middle rack, for the kickers. Signals are then directed from the patch panels to the relevant pieces of hardware in the PFF chain. The function of the phase monitor electronics has already been introduced in Chapter 3. This chapter discusses the design, setup and performance of the remaining PFF hardware — the FONT5a board and the kicker amplifiers. The second part of the chapter then presents remaining measurements that had to be taken to commission the PFF system, such as setting up the correction timing and verifying the ability to shift the phase in the TL2 chicane using the kickers and PFF optics.

5.1 Feedforward Controller (FONT5a Board)

This section describes the design, firmware and operational aspects of the FONT5a board, which acts as the PFF controller. Its role is to digitise the signals from the first upstream phase monitor (Mon 1), process them and calculate the appropriate voltage with which to drive the kicker amplifiers in order to correct the phase downstream. In addition it must ensure the timing of the correction output is such that the drive output from the amplifiers reaches the kickers precisely in time with the beam.

The FONT5a board is a custom built digital board designed and constructed by the FONT group at Oxford University [REF]. More details of the design of the board and the internal components can be found in [REF]. It was initially built as the controller for prototypes of an IP (interaction/beam collision point) position feedback for future linear colliders, either ILC or CLIC. Tests of this feedback are ongoing at ATF2, KEK in Japan [REF]. At the heart of the FONT5a board is a Xilinx Virtex-5 field programmable gate array (FPGA) [REF], an integrated circuit with customisable logic that the user can program as desired for his or her application. Exactly the same hardware can therefore be used for the

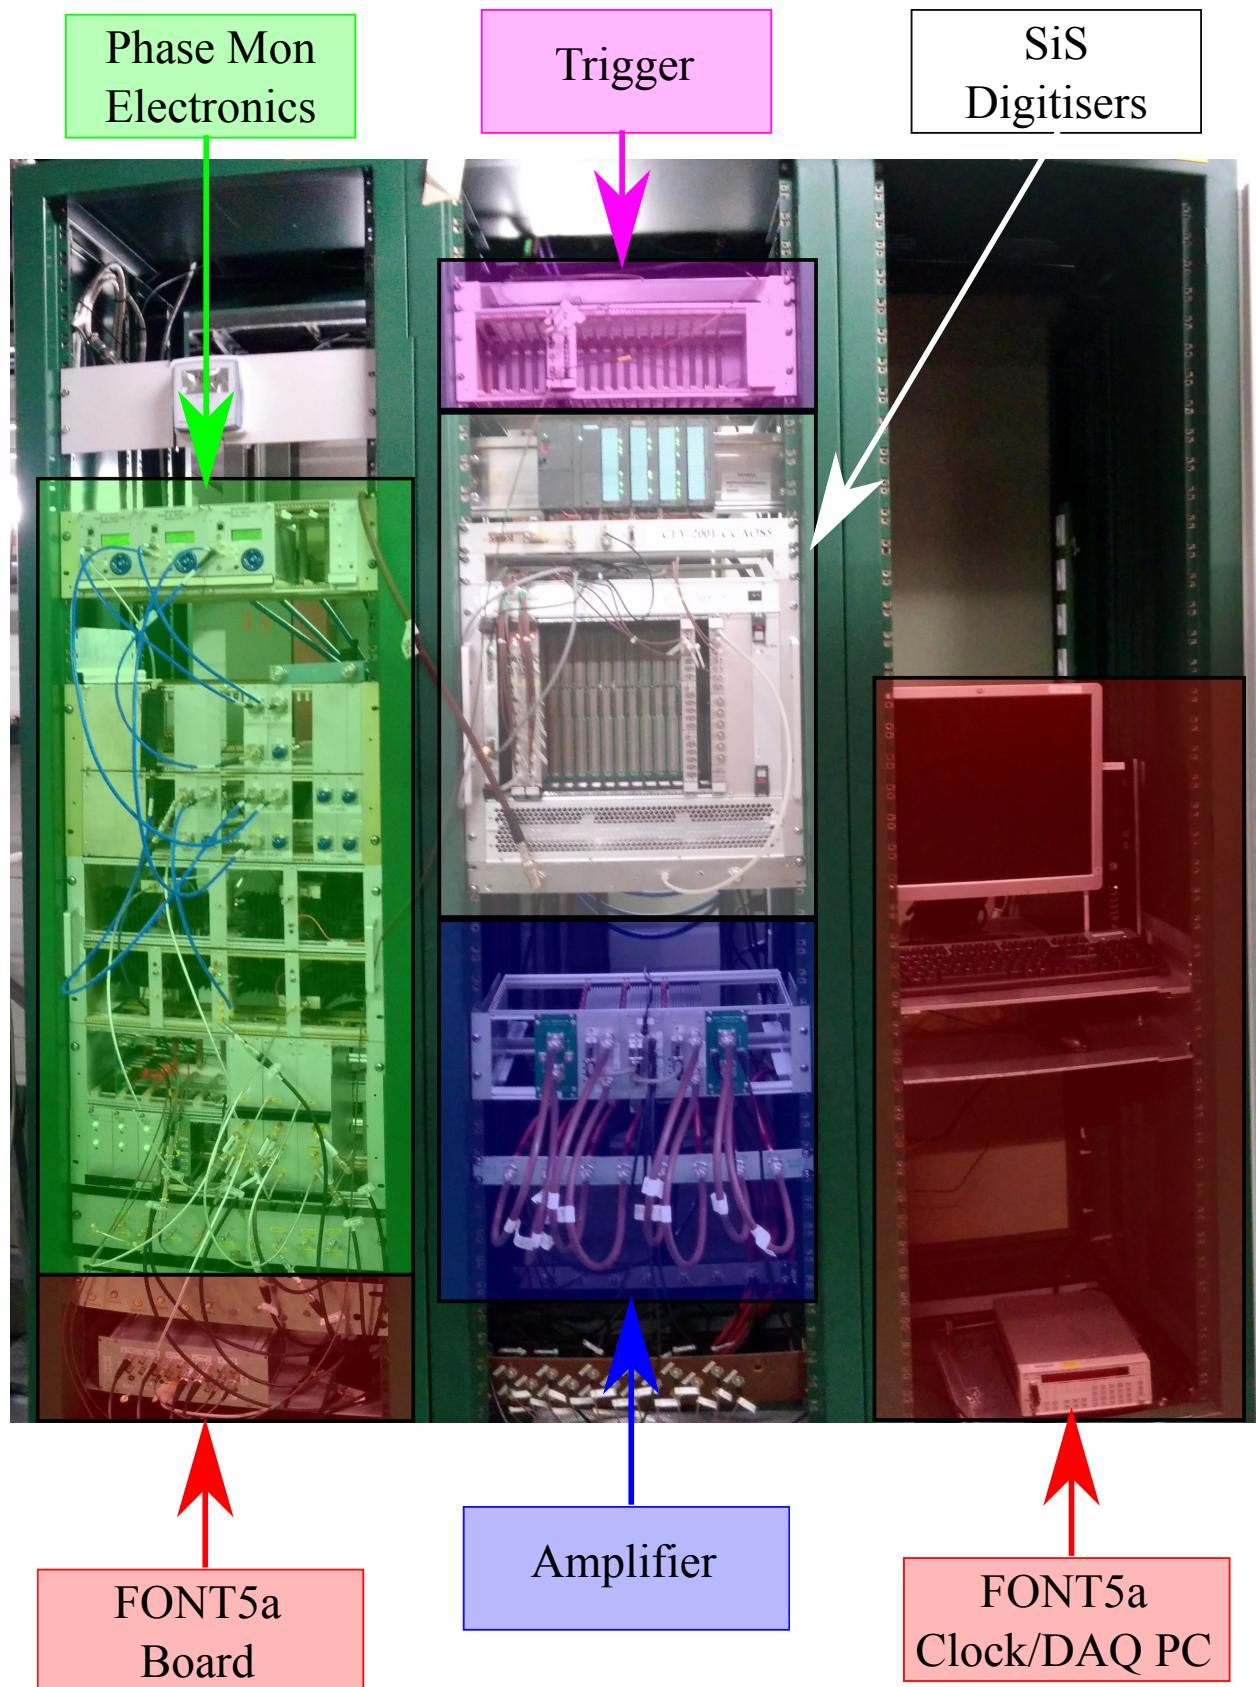


Figure 5.1: The PFF system electronics racks. Each major component of the system is labelled, including the phase monitor electronics, FONT5a board and kicker amplifier.

PFF and IP feedback controllers, with different firmware loaded on to the FPGA to take in to account the different signal processing required for the input and output signals in each application.

An annotated picture of the FONT5a board front panel is shown in Figure 5.12. Apart from the FPGA (inside the board and therefore not seen in the figure) the main components of the FONT5a board are nine analogue to digital converters (ADCs) and four digital to analogue converters (DACs). In addition there is a serial RS232 port used to communicate with a LabVIEW data acquisition and control system (DAQ), a JTAG connector used to program the firmware on the FPGA should it need to be changed, and several other inputs and outputs mostly used for the timing and triggering of the board.

The outputs from the phase monitor electronics, mixer and diode signals, are connected to the ADC inputs on the FONT5a board. Usually only the Mon 1 outputs are connected to the FONT5a board, with the diode on ADC 1 and the mixer on ADC 2, although the signals from additional monitors can be connected if needed. The ADCs are 14-bit with an input range of ± 0.5 V. Due to the expected noise on the ADCs the least significant bit is discarded and only the most significant 13 bits are processed by the FPGA [REF]. The processed ADC outputs therefore have a range of ± 4096 values or “counts”, corresponding to 0.12 mV per ADC count. More details on the digitisation of the phase monitor signals were given in Section 3.4.

Each ADC output can have a non-zero mean voltage in its baseline noise. To be able to remove these intrinsic voltage offsets each analogue input is combined with the output from a DAC, labelled “trim DACs” (separate from the DACs used as the PFF correction outputs). By varying the trim DAC voltages it can be ensured that the baseline output of each ADC is 0 V, so that its full ± 0.5 V range can be used for the signal of interest.

For the PFF system the ADCs are usually clocked at a sampling rate of 357 MHz, generated by an external clock generator [REF] and connected to the “FST CLK” input on the FONT5a board. 357 MHz is used for historical reasons, as this is the sampling rate for which experience has been gained for the IP feedback applications. The start of the ADC sampling window is determined by an external 0.8 Hz trigger derived from the CTF3 timing and therefore locked to the beam. The trigger is connected to the “DIG IN B” input on the board, and arrives [TODO: how much] before the beam to take in to account the ADC warmup time.

After being processed on the FPGA the calculated correction outputs are sent to the DAC outputs, which are 14-bit with an output range of ± 2 V to match the required range for the input of the PFF kicker amplifiers. Only DAC 1 and DAC 2 are used for the PFF system, the other two DACs are unused. In addition the FONT5a board generates the trigger for the amplifier, which is sent from the “AUX OUT A” output and must arrive at the amplifier [TODO: how much] before the DAC output signals.

Finally, the FONT5a board is controlled using a LabVIEW data acquisition system (DAQ) documented in [REF], with communication between the board and the DAQ via the serial RS232 port. An example screenshot from the DAQ is shown in Figure 5.3. It provides functionality to change all the setup parameters in the FONT5a firmware for the PFF system

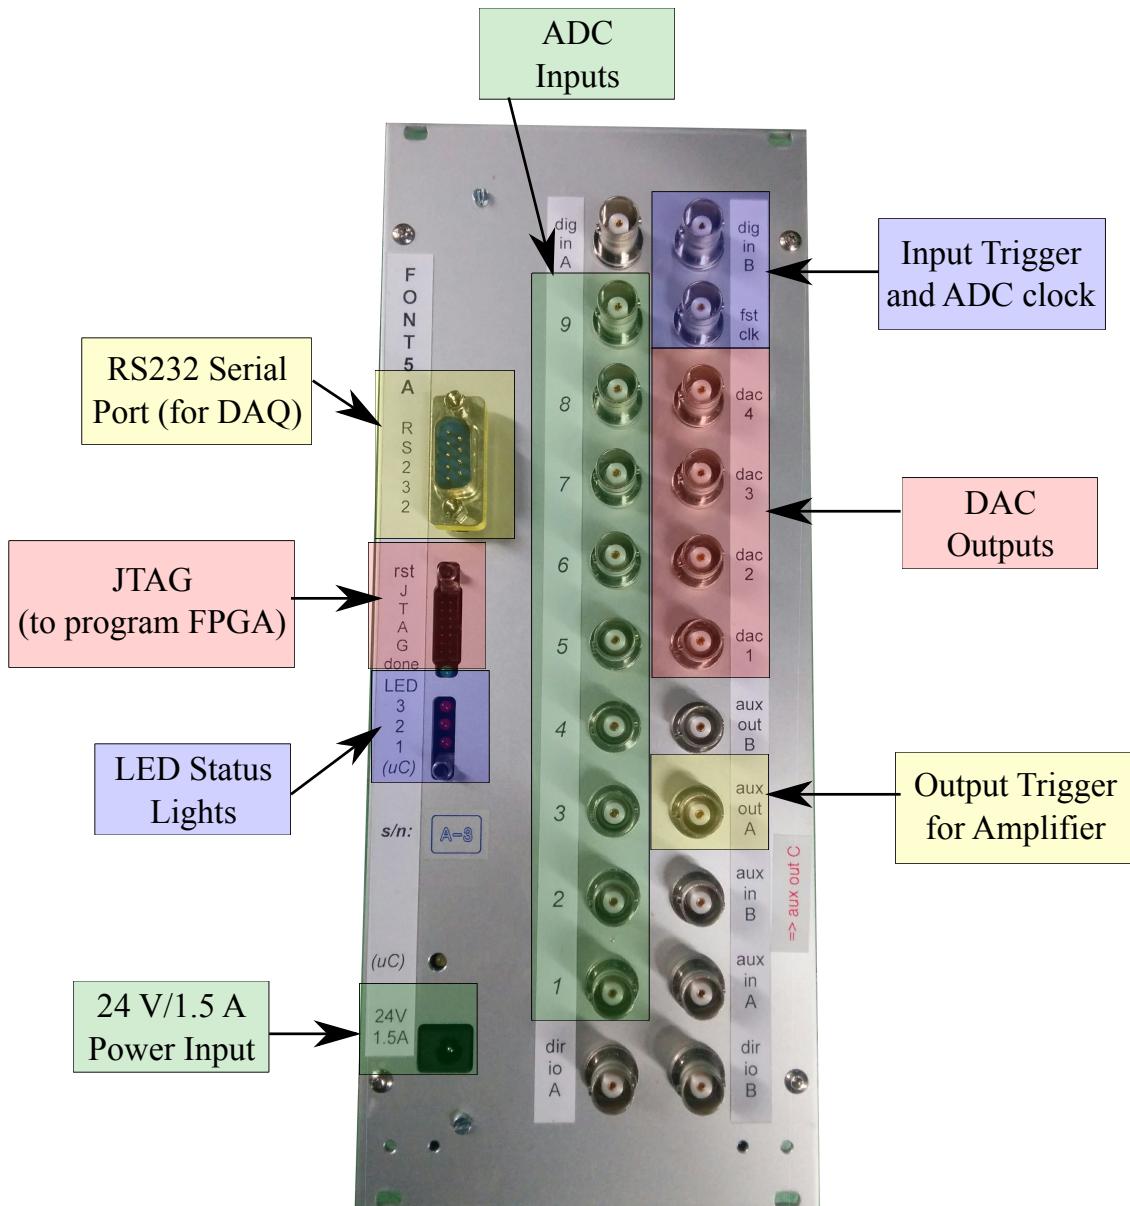


Figure 5.2: Front panel of the FONT5a board. All the connectors relevant to PFF operation are highlighted, the remaining connectors are not used at CTF3. All connectors are BNC apart from the power connector, RS232 serial port used to communicate with the LabVIEW data acquisition system (DAQ) and the JTAG connection used to program the firmware on the FPGA. The use of the ADCs, DACs and timing (trigger and clock) connections is summarised in Figure 5.4.



Figure 5.3: Diode output along the pulse with the IIR filter off and on.

setup, view the current ADC inputs and DAC outputs in real time and to save data directly from the FONT5a board. However, as the FONT5a board and DAQ currently run as a standalone system at CTF3, PFF data is usually saved via the CERN control system and SiS digitisers where data from other devices, such as BPMs, can be saved in sync with the phase monitor signals as discussed in Section 3.4. The DAQ runs on a Windows PC next to the racks used for the phase monitor electronics, FONT5a board and amplifier. The PC can be connected to via remote desktop to allow the FONT5a board to be controlled in the CTF3 control room.

5.1.1 Implementation of PFF Correction

The diagram in Figure 5.4 shows all the connections to and from the FONT5a board that were introduced in the previous section, from the phase monitor and timing inputs on the left side of the diagram to the outputs for the amplifier and DAQ on the right side of the figure. The central portion of the figure shows a simplified version of how the FONT5a firmware calculates and applies the PFF correction. All the parameters shown in blue, and several others not shown in the diagram but described later in this section, must be correctly set (via the DAQ) for operation of the PFF system.

The basic logic of the firmware is as follows. The ADC timing, sampling rate and start time, are determined by the input trigger and external 357 MHz clock. During the set sampling window the ADC outputs are processed, including a number of channel offsets and filters. The processed ADC 2 (Mixer) output is then split into two, to create the two strands that become the DAC 1 and DAC 2 outputs. Note that as the Mixer is used directly (rather than $\arcsin(\text{Mixer})$) this uses the small angle approximation. The effect of this is discussed in Section 5.1.4. If the diode (ADC 1) is being used in the calculation the ADC 2 output is also multiplied by $1/\sqrt{\text{ADC 1}}$, with values taken from a lookup table rather than being calculated directly to save latency [REF]. Prior to being sent to the DACs the two output

strands are multiplied by set gain factors and then can be delayed by a set amount of time. The two calculated DAC outputs are then connected to the signal inputs on the amplifiers, where they are amplified and eventually sent to the kickers to deflect the beam and correct the phase. Meanwhile the input trigger for the FONT5a board is used to derive the trigger for the amplifier, with a variable delay that can also be set in the DAQ. The overall latency of the FONT5a board, from the arrival of the signal at the ADCs to the output of the calculated correction at the DACs, is around 22 clock cycles (at 357 MHz) or 60 ns [REF].

All the parameters and controls that must be adjusted during the PFF system setup, and their respective values where relevant, are listed below for reference. These are only introduced in brief here, but parameters that are either non-trivial to derive or are critical for the PFF performance are described in more detail in later sections and chapters as indicated. The values given are in FONT units as they are set in the DAQ with each parameter expressed by up to a 14 bit number, and the size of each control chosen to give a reasonable degree of flexibility around the expected set point.

Input Timing

The following parameters must be adjusted in order to set the properties of the ADC sampling as desired:

Trig in delay: The Trig in delay allows the start of the ADC sampling window to be delayed with respect to the arrival of the external trigger. Timing of the trigger and correction outputs to the amplifier (Trig out delay and K1/K2 delay) are relative to this delay, therefore changing the Trig in delay value does not effect the synchronisation of the correction output with the beam. The only requirement is to ensure that the full acquired upstream phase monitor signals arrive within the sampling window, and that the trigger arrives early enough to account for the 10 μ s warmup time of the ADCs [REF]. The trigger is set to arrive 17 μ s before the beam [REF], so a trig in delay of around 2500 clock cycles of 357 MHz, or 7 μ s, is used to define the time at which the ADCs are turned on (with them then ready 10 μ s later when the beam arrives).

Enabled channels: The FONT5a board has 9 ADCs but only two are usually needed for the PFF system (for the mixer and diode of Mon 1 connected to ADC1 and ADC2, with the other two monitors normally connected to the SiS digitisers, see Section ??). To avoid hitting the limits of the baud rate of the RS232 port on the PC the remaining ADC channels can be disabled so their data is not transmitted.

No. samples: The length of the ADC sampling window, in number of samples, can also be varied. Typically 900 samples are used, covering a time window of 2.5 μ s with the 357 MHz clock. If the signal from more than two ADCs is needed the number of samples can be reduced to avoid hitting the baud rate limit of the RS232 port on the PC, the only requirement is that the time window is long enough to encompass the full 1.1 μ s beam pulse length.

5.1 Feedforward Controller (FONT5a Board)

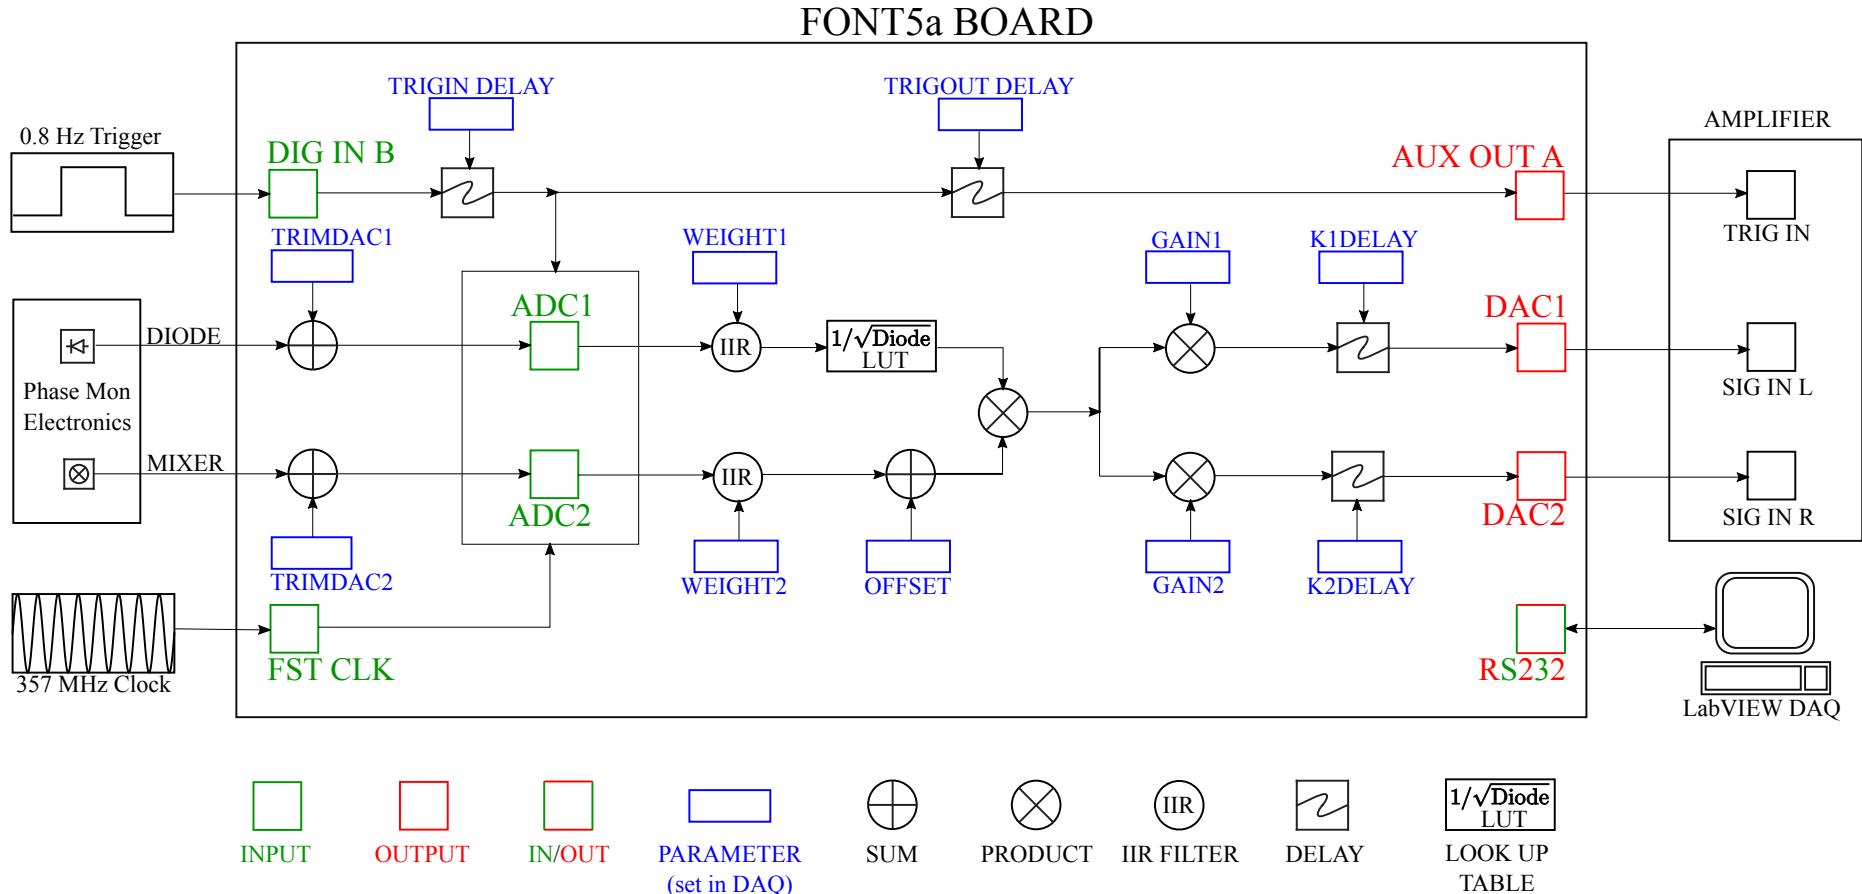


Figure 5.4: Schematic of connections to and outputs from the FONT5a board, as well as the PFF calculation in firmware in the case where the diode is used. If the diode is not used the ADC1 input is not required.

ADC Signal Processing

The following parameters must be correctly set in order to remove offsets and droop in the ADC outputs:

Trim DACs: As mentioned in Section 5.1 the FONT5a board includes trim DACs in order to remove intrinsic voltage offsets in the output of each ADC. The trim DAC output is added to the ADC inputs, and each trim DAC value must be set correctly in the DAQ. For ADC1 (Diode) and ADC2 (Mixer) the Trim DAC values are 1650 and 1400 counts, corresponding to 8 mV and 7 mV [REF], respectively.

Filter Weights: IIR filters are implemented in the FONT5a firmware in order to remove droop in the ADC response, see Section 5.1.2. The filter weights for each ADC can be adjusted in the DAQ. The correct values are 50 for ADC1 (Diode) and 56 for ADC2 (Mixer).

Channel Offset: To maximise the effect of the PFF system it is necessary to zero the mean upstream mixer (ADC2) output in the central region of the pulse where the correction is being attempted (the start and end of the pulse can not be fully corrected due to the large phase sag, so the central portion is usually used). Ideally this should be adjusted by varying the Mon1 phase shifter, but this can not be done remotely so the channel offset on the FONT5a board can be used instead to make small adjustments from the control room. The channel offset adds a static offset in counts to the ADC2 output, allowing the mixer output to be zeroed at any point along the pulse. The consequences of using a non-optimal Channel Offset are discussed in Section 6.2.1. For large phase offsets giving close to maximum mixer output the phase monitor resolution is degraded (Section ??) and the small angle approximation used in the FONT5a phase reconstruction (Section 5.1.4) becomes invalid, in which case the phase shifter should still be used to zero the phase.

Output Mode

The following parameters control the nature of the DAC output signals:

Feedforward enable: The DAC output can be enabled or disabled, as required.

Interleaved mode: With interleaved mode enabled the DAC outputs from the FONT5a board are only sent for half of the triggers, for example being applied to all the odd pulses but not sent for all the even pulses. This is incredibly useful for interpreting the PFF results, as well as being used for many of the other tests presented in this thesis, as it allows a comparison between beam conditions with and without an applied kick at the same time. In this case the effects of any slow drifts should be equivalent in both the kicked and non-kicked data, thus any differences between the two should be a real effect of the PFF system. Data with the DAC output disabled can also be used to simulate the expected effect of the PFF system in those conditions. All the PFF results in Chapters ?? and 6 use interleaved data.

DAC Output Mode: The DAC output to the amplifier can be sent in two modes — Sample-by-sample or Constant DAC. In Sample-by-sample mode the DAC output is as needed for the PFF system, being shaped along the pulse by the reconstructed phase

multiplied by the set gain as shown in Figure 5.4. In Constant DAC output mode a constant output voltage can be sent to the amplifier across the full length of the ADC sampling window (or for a shorter time with an applied gate, see below). The majority of the results presented in the remainder of this chapter use constant DAC outputs for verification of the amplifier, optics, correction range and correction timing.

K1, K2 const DAC: In constant DAC mode the DAC output can be varied between ± 4096 counts, or ± 2 V sent to the amplifier, and can be changed independently for each output.

Gain: In sample-by-sample mode the PFF gain can also be set independently for each correction output, with each being a 14-bit value (± 8192 units). The conversion between the gain in FONT units and the real applied gain is derived in Section 5.1.1. An applied gain of 624 units corresponds to a real gain of approximately 1 (with a gain of 624 an upstream phase offset of 1 degree corresponds to a downstream phase correction of approximately -1 degrees).

Diode Mode: The FONT5a firmware provides three modes for the treatment of the diode signal on ADC1 — normalisation, gating and unused. With diode normalisation enabled the PFF system reconstructs the phase as originally envisaged using Mixer/sqrt(Diode). Due to the issues with the phase monitor diodes as discussed in Chapter 3 the option to not include the diode in the PFF calculation and only include the mixer was later added, which is now the nominal setup. Rather than leaving the diode completely unused, it is usually used to gate the correction output (see below).

Overflow Mode: The PFF correction output can behave in three ways in the case where the calculated output is outside the maximum range of ± 4096 DAC counts. In the first iteration of the PFF firmware the calculated correction output would overflow, causing sign flips in the output in the regions where the correction range was exceeded [REF]. This behaviour can still be applied in the current firmware if desired. However, in normal operation the output is set to Saturate, so that any calculated values outside ± 4096 DAC counts are sent as the maximum ± 4096 DAC counts or ± 2 V to the amplifier. A final option to provide no output at samples where the calculated output is outside ± 4096 DAC counts, is also provided.

Output Timing

The following parameters control the timing of the DAC output signals:

Trig out enable: The trigger sent to the amplifier can be enabled or disabled as required. Clearly the trigger must be enabled for any correction output to have an effect on the beam (the PFF correction output can be turned on with the amplifier trigger disabled for testing purposes).

Trig out delay: The timing of the trigger sent to the amplifier can be delayed with respect to the start of the ADC sampling window. This must be adjusted so that the arrival of the $1.1\ \mu s$ beam pulse at the kickers is aligned with the $1.4\ \mu s$ time during which the

amplifier is powered and the correction output can be applied [REF]. The trig out delay can be varied between 0 and 127, where each unit corresponds to 26 clock cycles or around 70 ns [REF]. With a trig out delay of zero the trigger is output when the ADCs turn on, or around 10 μ s before the beam (see Input Timing above). The trigger to the amplifier must arrive around 1.5 μ s before the beam, so a trig out delay of 110 is typically used. The precise correction timing is set by the K1 and K2 delays below.

K1, K2 delay: The K1 and K2 delay are used to fine tune the timing of the two correction outputs (DAC1/K1 for the first kicker and DAC2/K2 for the second kicker), and can be varied by up to 32 ADC clock cycles (2.8 ns per clock cycle at 357 MHz). The optimal delays are 7 clock cycles for K1 and 7–8 clock cycles for K2. The importance of the correction timing and derivation of these values is presented in Section 5.5.

Gate enable: The correction output can be restricted to a certain sample range by applying a “gate”. The gate can be defined either as a custom sample range picked by the user or the diode signal (ADC1) can be used. Diode gating is typically used so that no output is sent to the amplifier outside the time of the beam pulse. Using a custom sample range has been useful for early PFF tests and to apply a constant kick along only part of the beam pulse (this is used in Section 5.5.3, for example).

5.1.2 ADC Droop Correction

Although the FONT5a (and previous FONT5) boards have been used extensively for IP feedback tests by the FONT group prior to its application for the PFF system, it had not been routinely used to process long, microsecond like, pulse lengths such as the CTF3 beam signals. During the first tests of the FONT5 board at CTF3 (prior to the FONT5a board being available) it was immediately apparent that there was a large droop in the measured phase monitor diode outputs using the FONT ADCs. An example of this, taken from the FONT5a board, is shown in Figure 5.5, remembering that as the diodes are highly saturated the response should be close to a perfect square wave. The measured diode output droops by 200 counts (approximately 15%) across the pulse length, with this difference also visible as an offset in the baseline after the pulse.

The droop emerges as a result of the use of AC coupling on the ADC input transformers for electrical isolation. This involves using a capacitor, the current across which is dependent on dV/dt (V being voltage and t time), to remove the DC component (or mean offset) from a signal [REF]. In particular for the diode channel the output is increasingly well described by a DC signal on the flat top as you move away from the leading edge of the pulse, with the capacitor causing droop in the response as a result. There is also a small effect on the mixer output, as will be seen later. [TODO: Is the effect on the diode and mixer really different? Is this correctly taken in to account by the IIR filter?]

In the simplest case the droop should be well described by an exponential decay of the form $A \exp(-t/\tau)$, where t is the time or sample number along the pulse and τ is the decay time constant. For the FONT5 board this only gave a rough approximation of the true droop characteristics due to non-linear properties of the input transformers. On the




Figure 5.5: Diode output along the pulse with the IIR filter off and on.

updated FONT5a board the transformers were changed to both reduce the magnitude of the droop and give closer to the expected exponential decay [REF]. Figure 5.6 shows an exponential fit to the Mon1 diode output as seen on the FONT5a board (ADC1). Apart from a small deviation at the beginning of the pulse the agreement is excellent, with residuals to the fit of only up to 2 counts compared to a signal magnitude of 1250 counts. For ADC1 the fitted decay time constant is $\tau = 838 \pm 16$ samples, or $2.35 \pm 0.05 \mu\text{s}$ (with one ADC sample every 2.8 ns at 357 MHz). Each ADC on the FONT5a board has slightly different droop characteristics, with the decay time constant for ADC2 being 938 ± 18 samples, for example (calculated with the diode moved on to ADC2).

In the case where the diode is used in the phase reconstruction as originally envisaged the ADC droop would propagate in to the applied correction and create an output to the amplifier that increases with time with respect to the ideal value (as the Mixer is divided by $\text{sqrt}(\text{Diode})$). The effect on the correction is much smaller without diode normalisation but it still slightly distorts the measured Mixer input to ADC2. Therefore, a digital IIR filter has been implemented in the FONT5a board firmware to remove the exponential droop in the ADC outputs [REF]. This works by recursively adding the expected droop to the ADC output based on the known decay constants, so that:

$$y(t) = x(t) + \frac{1}{\tau} \sum_{i=1}^t x(i-1) \quad (5.1)$$

Where $y(t)$ is the filtered ADC output at sample t , $x(t)$ is the original unfiltered output at sample t , and τ is the decay time constant. Rather than being hard-coded in the firmware the applied decay constant in the filter for each ADC is calculated using an 8-bit ± 64 filter

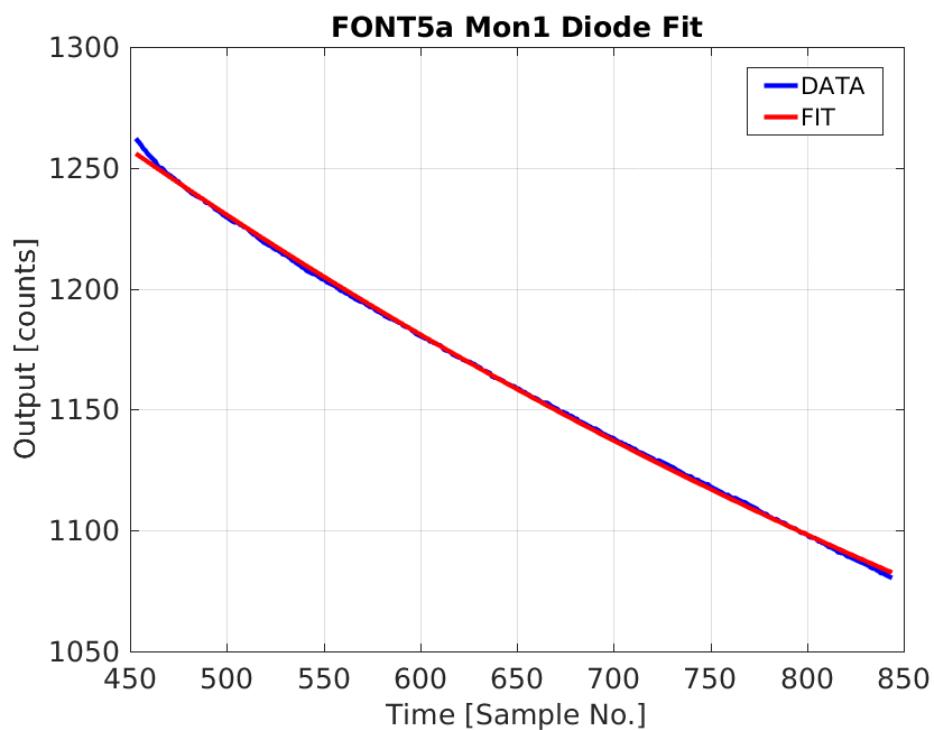


Figure 5.6: Exponential fit to diode droop.

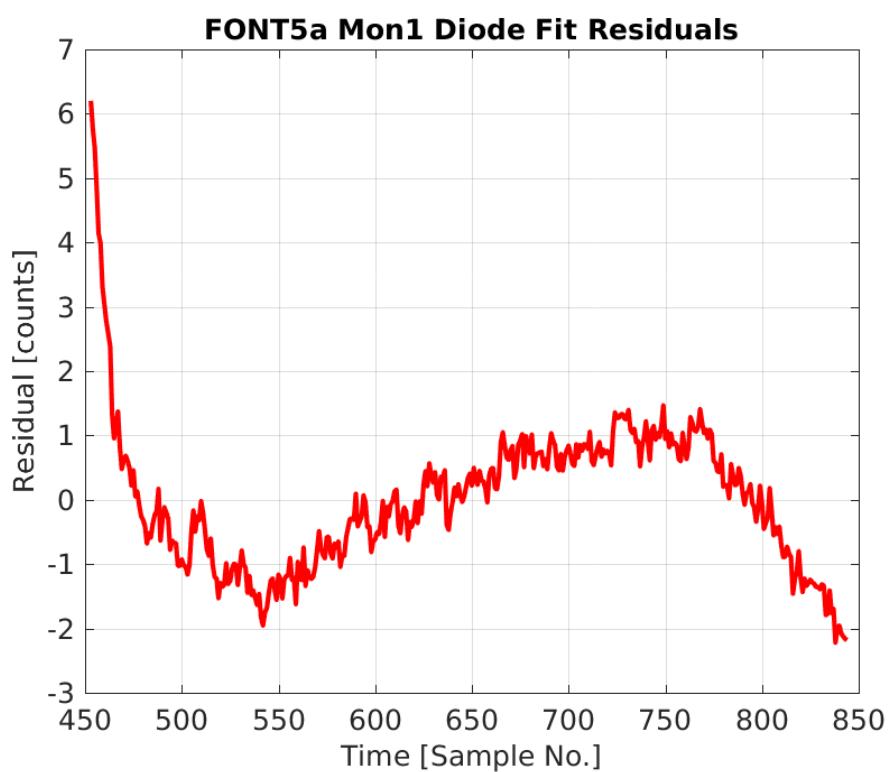


Figure 5.7: Residuals between diode exponential fit and actual diode output.

ADC	Filter Weight
1 (Mon1 Diode)	50
2 (Mon1 Mixer)	56
3	50
4	53
5	45
6	51
7	48
8	55
9	49

Table 5.1: IIR filter weights for the FONT5a board ADCs.

weight, which can be changed in the DAQ, which is then divided by a common division factor to get the real applied value of $1/\tau$ [REF]. The optimal filter weights for each ADC in the FONT5a board currently used for PFF operation are shown in Table 5.1, these can be converted in to the true decay constant values using the fitted values for ADC1 and ADC2 quoted above.

With the IIR filters enabled on the FONT5a board the droop on the diode (ADC1) is almost perfectly removed as shown in Figures 5.5 and 5.8, although in the zoomed in figure some slight deviation from flat is visible due to the residuals around the exponential fit seen previously. The effect on the reconstructed phase using only the mixer (ADC2) is shown in Figure 5.9, with a modest adjustment to the overall shape that is most visible at the start and end of the pulse.

5.1.3 Gain Calculation

As shown in Section 4.1 the effect of the PFF system can be modelled by simply subtracting the measured upstream phase, ϕ_u , from the downstream phase, ϕ_d with a gain factor, g :

$$\phi_{PFF} = \phi_d - g\phi_u \quad (5.2)$$

Where ϕ_{PFF} is the corrected downstream phase. The gain set in the FONT5a DAQ is not directly a multiplication factor in terms of the phase, but rather a multiplication in terms of DAC counts and other constants. How the set FONT gain relates to the true gain g is derived in this section.

In the case where diode normalisation is not used, which is the nominal setup, the FONT5a DAC outputs when the PFF system is enabled are given by:

$$\begin{aligned} \text{DAC1} &= \frac{F1}{64} \text{ADC2} \\ \text{DAC2} &= \frac{F2}{64} \text{ADC2} \end{aligned} \quad (5.3)$$

Where DAC1 and DAC2 are the two DAC outputs, in DAC counts, F1 and F2 are the respective gains for each DAC output set in the DAQ, and ADC2 is the Mon1 mixer signal

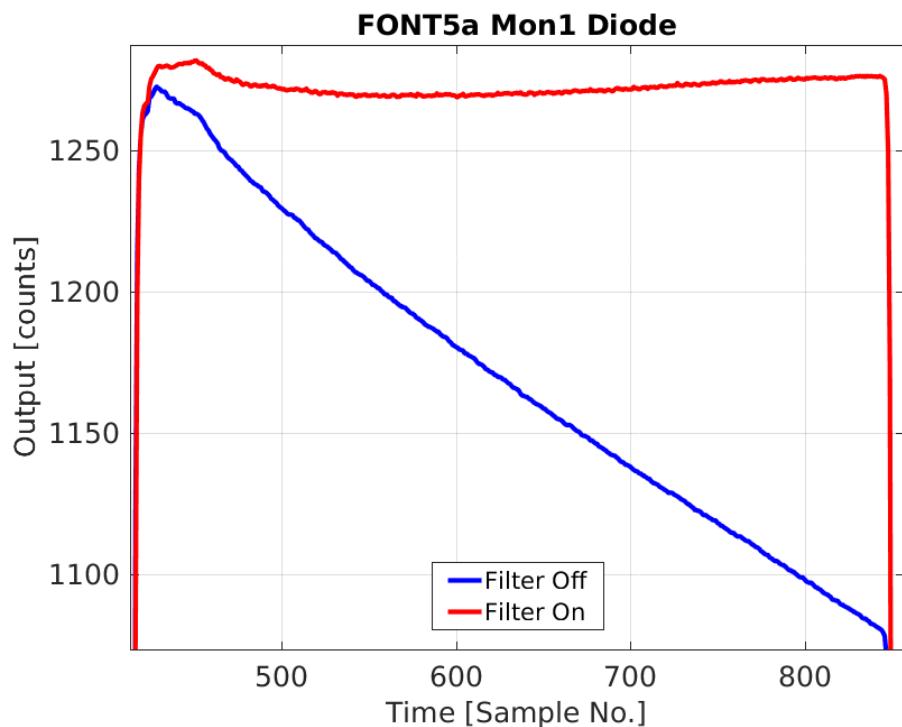


Figure 5.8: Diode output along the pulse with the IIR filter off and on. Zoomed in.

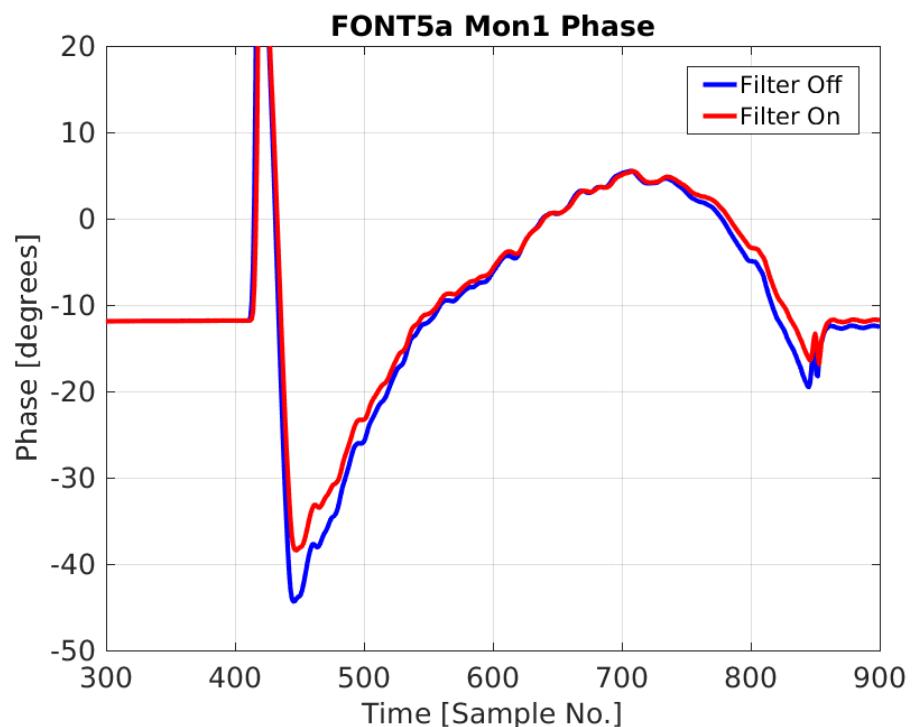


Figure 5.9: Phase along the pulse with the IIR filter off and on.

level in ADC counts, which is connected to ADC2 for normal PFF operation. The whole calculation uses 16-bit architecture, with the 3 sub-integer bits discarded when sent to the 13-bit DACs [REF]. The FONT gains F1 and F2 are 14-bit (± 8192) and the common division factor of 1/64 is set by several parameters fixed in the FONT5a firmware in order to give a sensible range of gain values based on the approximate signal levels [REF]. Typically $F2 = -F1$, so that the beam is deflected in opposite directions in each kicker to achieve orbit closure in the chicane. With a maximum FONT gain of 8192 an ADC2 input of 1 ADC count gives an output of 128 DAC counts. If diode normalisation is enabled the expressions above are multiplied by $1/\sqrt{\text{ADC}1}$ (with the Mon1 diode signal connected to ADC1 in normal operation) and different scale factors are used [REF].

To determine the optimal values for F1 and F2 given the current beam conditions it is important to know how they translate in to the real applied gain g from Equation 5.2. The real gain g is related to the DAC outputs derived above by:

$$g\phi_u = k\text{DAC} \quad (5.4)$$

Where k is the phase shift, in radians, in the chicane resulting from a DAC output of 1 count. The indices 1 and 2 are dropped from DAC1, DAC2, F1 and F2 from this point for simplicity. The upstream phase, also in radians, can then be related to the ADC2 input by:

$$\phi_u \simeq \frac{\text{ADC}2}{A} \quad (5.5)$$

Where A is the calibrated maximum Mon1 mixer amplitude, in ADC counts, determined in the same way as Section 3.7. Note that both this and the direct proportionality with ADC2 in Equations 5.3 assume the small angle approximation $\text{ADC}2 = A \sin \phi_u \simeq A\phi_u$ for small ϕ_u . This is done for latency reasons [REF] and its effect on the accuracy of the applied correction is discussed in the next section.

Combining Equations 5.3, 5.4 and 5.5 gives:

$$\begin{aligned} \frac{g}{A} \text{ADC}2 &= \frac{kF}{64} \text{ADC}2 \\ g &= \frac{kA}{64} F \end{aligned} \quad (5.6)$$

In Section 5.4.1 a phase shift of 3.5° per volt sent to the amplifier is determined. Knowing that 4096 DAC counts corresponds to 2 V sent to the amplifier this value can be converted into $k = 29.8 \mu\text{radians}/\text{DAC count}$. Typical calibrations for Mon1 on the FONT5a board give $A \simeq 3440$ ADC counts (Section 3.7.2). Overall, the real applied gain therefore relates to the set FONT gain in the DAQ via:

$$g \simeq \frac{F}{624} \quad (5.7)$$

The optimal real gain to apply (Section ??) is given by the upstream-downstream phase correlation, ρ , multiplied by downstream-upstream jitter ratio, σ_d/σ_u :

$$g_{\text{opt}} = \frac{\sigma_d}{\sigma_u} \rho \quad (5.8)$$

Therefore the optimal FONT gain to set in the DAQ dependent on the current beam conditions is:

$$F_{\text{opt}} \simeq 624 \frac{\sigma_d}{\sigma_u} \rho \quad (5.9)$$

In good conditions the correlation and jitter ratio are close to one, although the downstream phase jitter can be up to a factor two larger than the upstream jitter at CTF3 thus the applied FONT gains are typically in the range between 625 and 1250.

5.1.4 Effect of Using Small Angle Approximation

As mentioned previously the phase calculation in the PFF algorithm in the FONT5a firmware uses the small-angle approximation, thus differs from the correct full phase reconstruction method used in Chapter 3 as follows:

$$\phi_{\text{FONT}} = \frac{\text{Mixer}}{A} \quad (5.10)$$

$$\phi_{\text{FULL}} = \arcsin \left(\frac{\text{Mixer} - d}{A} \right) \quad (5.11)$$

Where ϕ_{FONT} and ϕ_{FULL} are the upstream phase with the FONT and full phase reconstruction methods respectively, Mixer is the Mon1 mixer signal, A is the Mon1 calibration amplitude and d is the calibration offset needed to take in to account asymmetry in the mixer minimum and maximum output due to cross-talk from the diode (Section ??).

The difference between the measured phase using these two reconstruction methods is plotted in Figure 5.10, versus the phase offset between the beam phase and the electronics LO phase. For small incoming phases (close to zero Mixer output) the only difference between the two methods is a static offset of $-d/A$ in the measured phase, and this holds up to $\pm 10^\circ$ (and is a good approximation up to $\pm 20^\circ$). In this case the corrected downstream phase will contain this constant offset, but the corrected phase jitter, which has been the only focus of the PFF prototype to date, will still be optimal.

With larger offsets between the beam phase and the LO (large Mixer output) the small angle approximation is no longer valid and the difference between the two methods rises to up to 35° . Most importantly the measured phase, and therefore the correction output, is no longer linearly dependent on the incoming beam phase when there is a large offset between the beam and LO phases. In the most extreme case, the measured phase difference between two pulses with a beam-LO phase offset of 80° and 90° would only be 1° with the FONT algorithm instead of the expected 10° , for example. In turn, the difference in the correction applied to the two pulses would only be 1° instead of 10° , degrading the achievable corrected phase jitter.

Figure 5.11 shows a simulation of the theoretical downstream jitter that can be achieved with both phase reconstruction methods with 0.8° initial upstream and downstream phase jitter and an upstream-downstream phase correlation of 97%. These values are chosen to represent the beam conditions that will need to be achieved in order to demonstrate 0.2° phase stability at CTF3. In the $\pm 10^\circ$ region where the small approximation holds there is

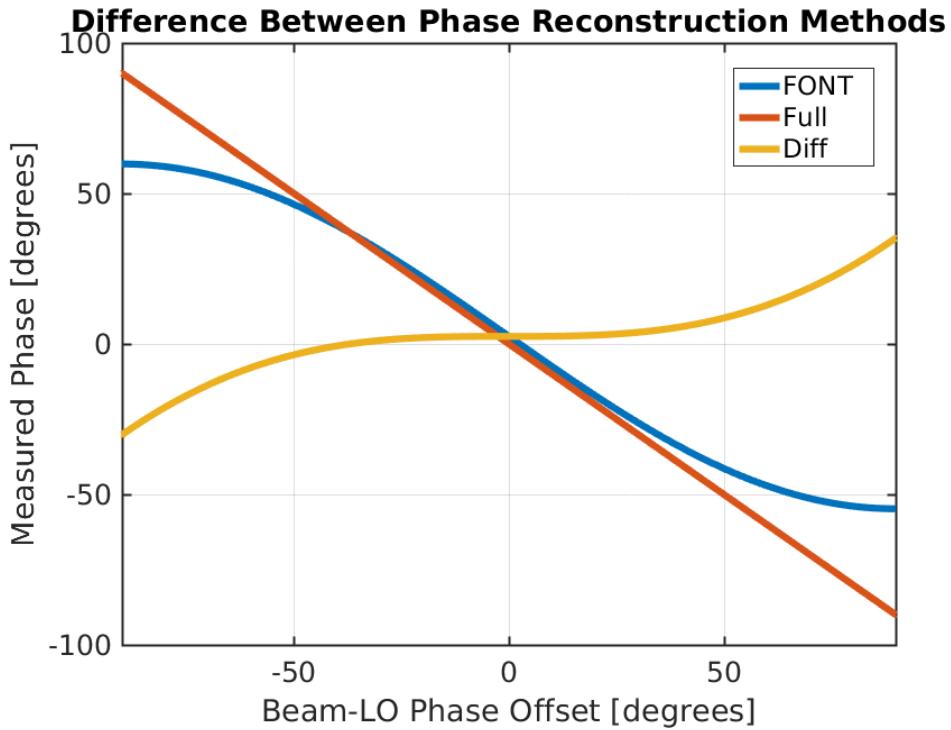


Figure 5.10: Difference between the phase reconstruction method used in the PFF algorithm on the FONT5a board (with the small angle approximation) and the full reconstruction used with data acquired from the SiS digitisers.

no degradation in the achievable corrected phase jitter using the simplified method in the PFF algorithm. Outside this region the jitter is degraded, increasing from below 0.2° to 0.3° for a beam-LO phase offset of 50° .

As the correction range of the PFF system is less than $\pm 10^\circ$ (Section 5.4.1) there should be no degradation in the PFF performance resulting from the use of the small angle approximation in the PFF algorithm, providing the LO phase shifter in the phase monitor electronics has been set correctly. This highlights the importance to adjust the phase shifters to zero the Mixer output not only in order to maximise the phase monitor resolution (Section ??) but also to ensure the correction calculation itself is valid. However, it is perhaps interesting to note that it would not be possible to correct the full CTF3 pulse length to 0.2° jitter with this implementation of the PFF algorithm even if the correction range were large enough to encompass the 40° phase sag.

5.2 Amplifier

The amplifier takes the two DAC signals from the FONT5a board and produces four high voltage outputs from them that are connected to the downstream ends of the kicker strips (two kickers and two strips per kicker gives four connections at the downstream ends in total), creating the potential difference between the kicker strips that deflects the beam in order to correct the phase. The returning signals from the upstream ends of the kicker




Figure 5.11: Achievable PFF jitter versus phase offset for full phase reconstruction and with the small angle approximation.

strips are then terminated back at the amplifier. As the amplifier is installed together with the phase monitor electronics and FONT5a board in the klystron gallery the cables between the amplifier and kickers are long and represent the single largest contribution to the overall system latency, as discussed in Section 5.5.1. This section discusses the design and performance aspects of the amplifier that are relevant to PFF operation.

5.2.1 Design

The amplifier is purpose built for the PFF prototype, also by the FONT group at Oxford University, with further details of its design available in [REF]. An annotated picture of its front panel is shown in Figure 5.12 and a simplified diagram showing the flow of signals between the FONT5a board, amplifier and kickers is shown in Figure 5.13. The amplifier is installed in a standard 3U rack and has a modular design. It consists of five individual modules split between two sides, labelled left and right. The left side of the amplifier, which uses the DAC 1 output, powers the first kicker in the chicane, and the right side of the amplifier, which uses the DAC 2 output, powers the second kicker. Each side of the amplifier contains its own “drive module” and “terminator module”. Finally there is a central “control module” that is common to both sides of the amplifier.

All the outputs from the FONT5a board are connected to the control module, this includes the DAC 1 and DAC 2 outputs as well as the trigger for the amplifier. The required 24 V, 1.1 A power supply is also connected to the control module. The control module then deals with the distribution of power, timing and input signals (derived from the DAC inputs)

to the two drive modules. The control panel can therefore include signal processing designed to protect the drive modules. It ensures the maximum signal level sent to the drive modules is at a safe level and limits the rate of large changes in the input drive signal (limits the slew-rate), as well as preventing triggering of the drive modules if a problem is detected, for example.

The amplification of the signals to create the high voltage outputs occurs in the drive modules. Each drive module takes one input signal but creates two high voltage outputs (A and B) which have equal magnitude but opposite polarity. These are connected to the downstream end of the left and right strips of one of the kickers, to create a potential difference across the kicker strips that is double the voltage of each individual drive module output. The signal amplification consists of low voltage Si FETs (Directed Energy Inc DE150-201N09A [REF]) driving high voltage SiC FETs (Cree C2M01600120D [REF]) and a final output transformer, giving a peak output of around ± 700 V (seen in more detail in Section 5.2.2). The output has a bandwidth of 47 MHz for small signal variations up to 20% of maximum output (around 140 V). The bandwidth for larger variations is slew rate limited, and around 25 MHz for variations up to 50% of the maximum output, for example [REF].

Each drive module requires a ramp up time of $1\ \mu s$, which defines when the trigger must arrive from the FONT5a board. After this period the drive modules are powered for $1.6\ \mu s$ with full output across $1.3\ \mu s$ (slightly longer than needed to correct the full $1.1\ \mu s$ CTF3 beam pulse length). The control module includes a monitoring output (labelled “ON”) that can be digitised to observe the time period within which the drive modules are on and ready to receive input.

The high voltage signals leaving the upstream ends of the kicker strips are terminated back at the amplifier on the terminator modules. Each module consists of two $50\ \omega$ terminators designed to be able to take the maximum 20 kW output from the drive modules [REF]. Each terminator has a monitoring signal, which gives output at around -40 dB or approximately 1/115 of the input voltage [REF]. These can be used to verify the applied output from the drive modules is as expected, and form the basis of most of the measurements in the remainder of this section as well as the timing checks in Section 5.5.2, for example.

The description of the amplifier here represents the version of the amplifier that was used for the latest PFF tests, and for almost all results in this thesis. A first version of the amplifier, with around half the output voltage at up to ± 350 V was used for early PFF tests. Additional FETs were added in the second version described above to boost the output voltage to ± 700 V. A third version of the amplifier was foreseen but ultimately deemed unnecessary for the application at CTF3. This would have provided a voltage of up to ± 1400 V by combining the output of four individual drive modules [REF].

5.2.2 Linearity

Figure 5.14 shows the amplifier output, as measured by the monitoring signals, at different constant input voltages sent from the FONT5a board between the minimum of -2V (-4096




Figure 5.12: Front panel of the amplifier. Inputs to the amplifier are highlighted in green, drive outputs in red and monitoring outputs in blue. Inputs to the control module are the trigger (T) and the two DAC outputs from the FONT5a board used to determine the drive to the left (L) and right (R) sides. Each side of the amplifier has two pairs of drive outputs and terminators, A and B. The signal returning to each of the terminators can be observed on their corresponding monitoring outputs. The monitoring output on the control module (ON) shows the $1.4 \mu\text{s}$ time during which the amplifiers are able to provide their output. Not highlighted in the figure are the 24 V, 1.1 A power connector at the top of the control module, and also the three custom connectors used to communicate between the control module and drive modules.




Figure 5.13: Simplified flow diagram showing the connections between the FONT5a board, the amplifier and the kickers. The kickers are inserted between the drive and terminator modules for the purposes of the diagram, but in reality the terminator modules neighbour the drive modules in the same unit as seen in Figure 5.12. Note that the A and B outputs on each side always have opposite polarity, as needed to create a large potential difference between the kicker strips.

DAC counts) and maximum of 2V (+4096 DAC counts). The output voltage from the monitoring signals is converted in to the real amplifier output Voltage using the approximate conversion factor of 115. All four amplifier outputs are shown (one for each strip of the two kickers). The plotted values are means taken across a 480 ns central part of the whole 1400 ns output pulse.

The relative polarity of the four outputs is equivalent to what would be sent to the kickers during PFF operation, with opposite polarity of the L and R amplifier outputs sent to each kicker, so that the beam is kicked in opposite directions by each kicker with the second kicker then closing the orbit bump created by the first. Within each side of the amplifier the A and B outputs (sent to each side of the kicker) also have opposite polarity, necessary to create the potential difference across the strips within each kicker that creates the deflecting field for the beam. The relative polarity of the A and B outputs is fixed in the amplifier design and cannot be controlled via the FONT5a board.

The response of the amplifier is highly linear in the region between ± 1.2 V sent to the amplifier. Outside this range the amplifier clearly begins to enter saturation, in particular above input voltages of ± 1.7 V. The linear fits shown include only the points between ± 1.2 V, excluding the first and last three points in the scan of input voltages, in order to not be biased by the effects of saturation.

Figure 5.15 shows the residuals between the linear fit and the real amplifier output across the full range of input voltages. By looking at the residuals a slight deviation from linearity in the ± 1.2 V range is also visible, although the maximum difference is only 10 V or a 3% relative error. At the maximum input voltage of ± 2 V the difference between the real output and the amplitude expected if the response was linear across the full range rises above 150 V, or a relative error of more than 25%. For example, the RB output at an input voltage of +2 V is 605 V but the fitted response gives 769 V, a difference of 164 V or 27%.

The effects of amplifier saturation are not taken in to account in the PFF algorithm on the FONT5a board, in which the DAC output is linearly dependent on the input phase (voltage from the phase monitor mixer signal) across the full range. The applied correction to the downstream phase will therefore be non-optimal when the DAC output calculated by the PFF algorithm is above an absolute value of 2500 counts (1.2 V sent to the amplifier). To date the non-linearity of the amplifier as it begins to enter saturation has also not been included in the PFF simulations presented in the following chapters. This may partially explain the small discrepancies seen between the simulated and real results in some datasets, so including the effect will be pursued in the future. In addition, it could be foreseen to incorporate the saturation characteristics in to the PFF algorithm on the FONT5a board, so that calculated outputs above 2500 counts are boosted slightly to compensate for the lower than expected amplifier output.

Discrepancies between the four amplifier outputs are also visible in Figure 5.15 and Table 5.2, both in terms of gradient and peak output. This can be partially but not completely explained by errors of up to a few percent in the precise calibration of the four monitoring outputs, which do not output exactly 1/115 of the real input voltage [TODO: Ask Colin about errors]. Differences between the A and B outputs sent to each kicker are not an issue




Figure 5.14: Amplifier output vs. input.

Amplifier Port	Output at +1 V Input
LA	+416 ± 3 V
LB	-453 ± 3 V
RA	-426 ± 3 V
RB	+409 ± 3 V

Table 5.2: Feedforward results using combined data from 20th November 2015.

for the PFF performance as both are linear (in the ± 1.2 V range) and the kick experienced by the beam in each kicker is proportional to the difference of the two. Therefore, only the calibration between the output from the FONT5a board sent to the amplifier and the resulting phase shift in the TL2 chicane is affected. However, disparity between the potential difference across each kicker (LA-LB and RA-RB), so that the deflection of the beam in each kicker is different, leads to the orbit bump created by the PFF system not being closed in the chicane, degrading the horizontal beam stability downstream. The fitted potential difference at 1 V input is 869 V for the left amplifier (LA-LB, sent to the first kicker) and 835 V for the right amplifier (RA-RB, sent to the second kicker), a difference of 4%. This can be compensated in the PFF setup on the FONT5a board by using a different gain for each correction output, so that the voltage sent to the right amplifier is higher but the resulting output voltage sent to both kickers is the same. Orbit closure is discussed further in Section 5.4.4.




Figure 5.15: Residual between amplifier output and linear fit.

5.2.3 Shape

In the previous section the linearity of the mean output was considered but the performance of the PFF correction is clearly also sensitive to any variations in output voltage along the amplifier output pulse. Figures 5.16 and 5.17 show the full $1.4 \mu\text{s}$ amplifier output pulse at a constant $+1 \text{ V}$ input sent to the left amplifier and a constant -1 V input sent to the right amplifier respectively. Spikes in the signal just prior to 2000 ns and after 3000 ns on the time axis as seen in the plots are beam pickup induced by the beam passing through the kickers. These are therefore not a property of the amplifier performance and are excluded from the analysis in this section. However, the beam pickup is used later in Section 5.5.2 for the purposes of optimising the correction timing.

For each side of the amplifier both the A and B outputs are plotted as well as the difference of the two, which is the relevant quantity in terms of the kick received by the beam as it traverses the kickers. In the ideal case the potential difference should be flat along the full pulse length. However, for both the left and right side variations in the difference are visible, with an initial increase in output across the first 500 ns of the pulse followed by a droop in response across the second half of the pulse. Although not shown here, the shape of the variations along the pulse is consistent across the full range of output voltages, and scale in magnitude with the output voltage. Figure 5.18 shows the peak-to-peak and mean deviation of the output voltage along the pulse across the full range of input voltages. The peak-to-peak deviation refers to the difference between the minimum and maximum output along the pulse, whilst the mean deviation is the average absolute difference between the mean output and the output at each sample point. For a constant input voltage the output voltage

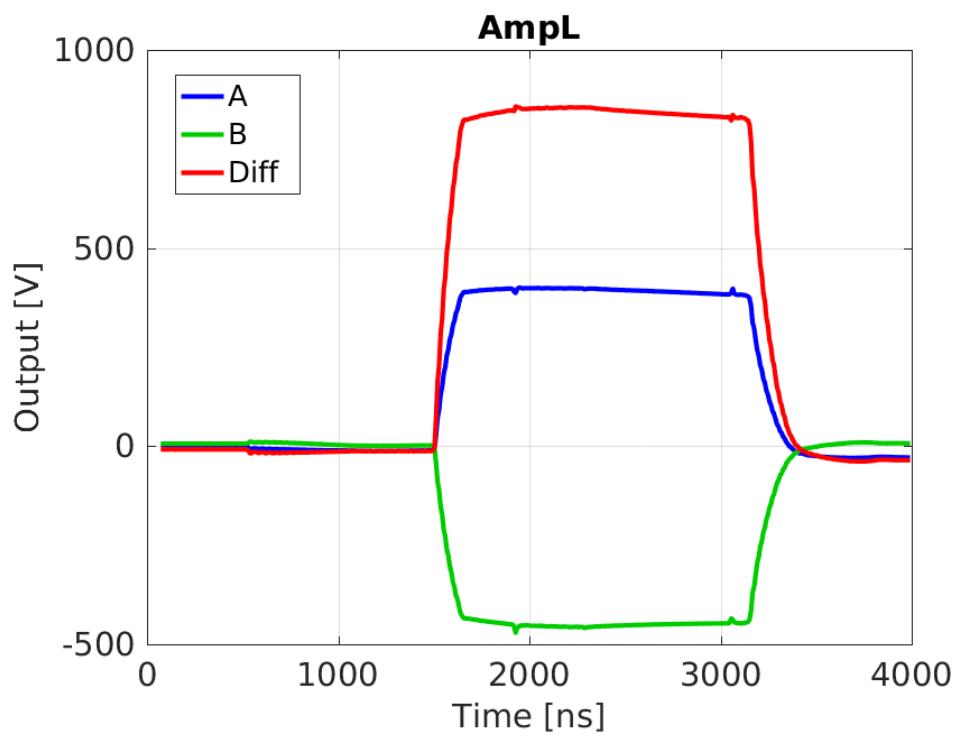


Figure 5.16: Amp L along pulse at 1 V input

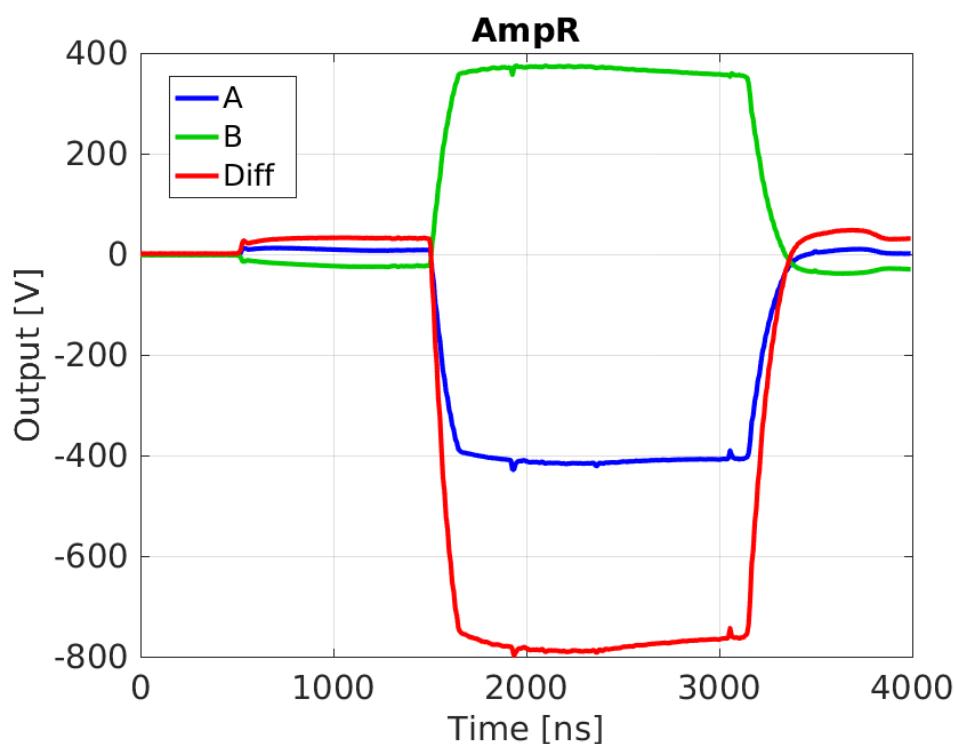


Figure 5.17: Amp R along pulse at 1 V input

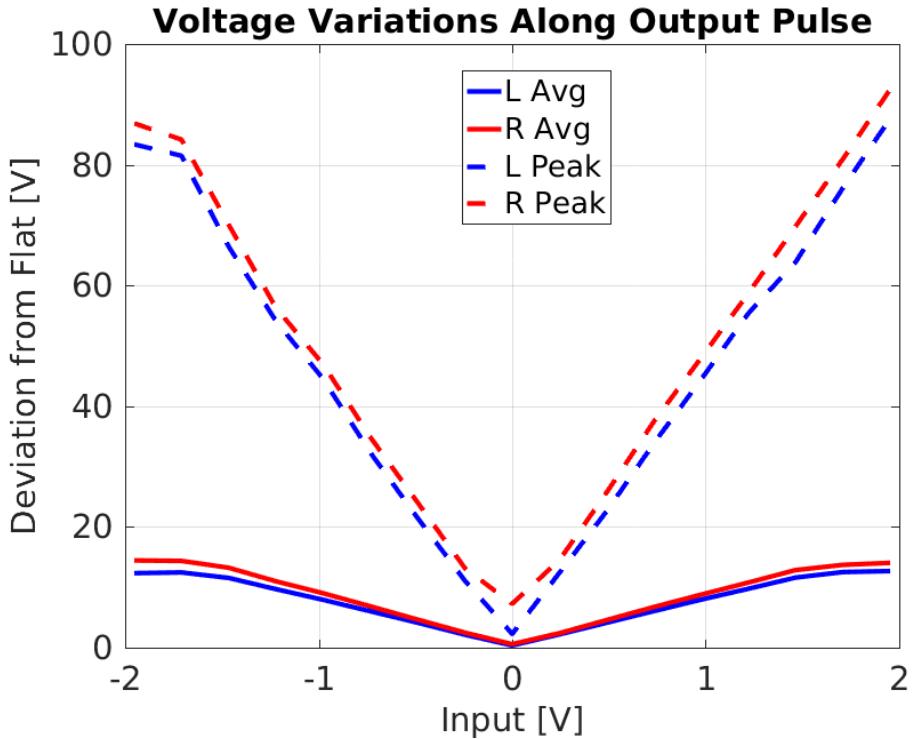


Figure 5.18: Flatness of potential difference sent to kickers.

along the pulse varies by up to 88 V peak-to-peak (mean 12 V) for the left amplifier or 93 V peak-to-peak (mean 14 V) for the right amplifier. As a relative difference, this corresponds to approximately a 6 % peak-to-peak, or 1 % mean, variation along the pulse.

The PFF algorithm on the FONT5a board uses a single gain value across the whole pulse length for each correction output, thus making the approximation that the amplifier response is flat along the pulse. The variations along the amplifier pulse therefore directly translate into discrepancies between the intended phase shift as calculated and the real phase shift experienced by the beam. As the region of interest for the correction is a few hundred nanoseconds about the central part of the pulse, as opposed to the full pulse length, the 1 % mean variation is more indicative of the resulting error than the 6 % peak-to-peak variation. With a correction range (Section 5.4.1) of $\pm 6^\circ$, the effects of the non-flat amplifier output should be below 0.06° and not measurable considering the phase monitor resolution of 0.14° . Nevertheless, it could be foreseen to implement a droop correction in the PFF algorithm on the FONT5a board, taking the variations in the amplifier output along the pulse into account.

As for the mean output voltage, the second way variations in the amplifier output along the pulse can impact the PFF performance is via the orbit closure in the chicane. For this the relevant quantity is the sum of the potential difference sent to each kicker, or $(LA - LB) + (RA - RB)$. To ensure orbit closure this quantity, named the residual kick here, should be zero across the whole pulse length for all input voltages. Figure 5.19 shows the residual kick along the pulse for all the input voltages in the scan. Clearly they are not all centred around zero, but this is expected due to the differences in the mean output voltage of the four amplifier outputs seen in the previous section. As already stated, the

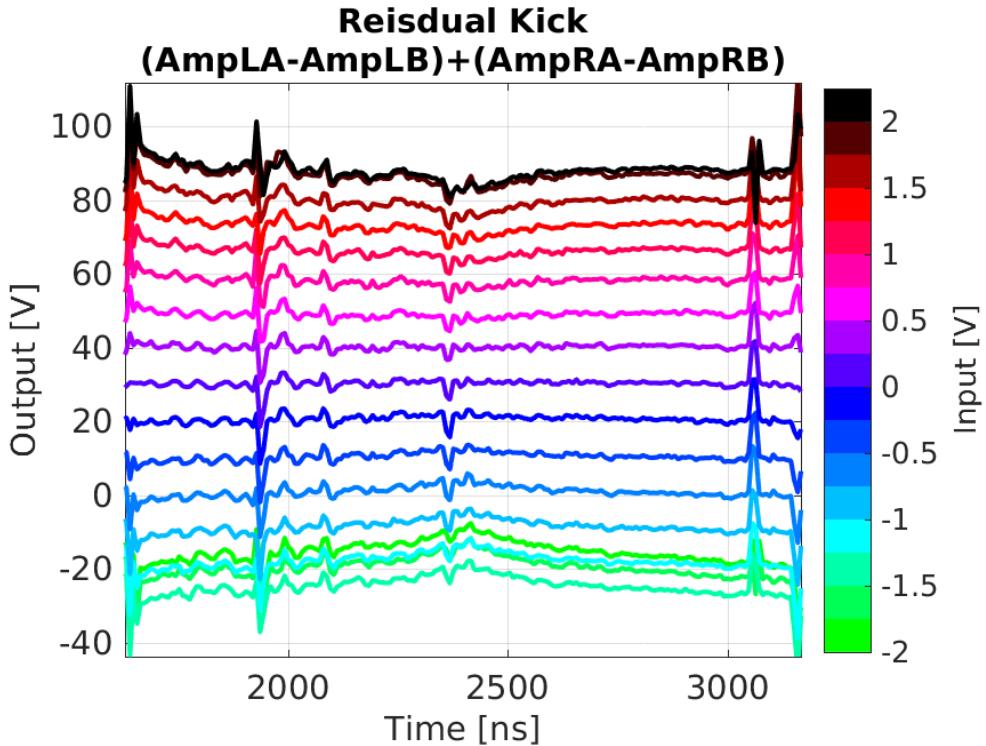


Figure 5.19: Residual kick along pulse.

overall mean offset can be removed by using a different gain for the two correction outputs. However, any variations along the pulse cannot be compensated for in the PFF algorithm. The magnitude of the effect is summarised in Figure 5.20, showing the peak-to-peak and average deviation of the residual kick from flat. The overall residual kick is very flat and effect is smaller than any of those previously shown — only up to 2 V on the mean or 21 V peak-to-peak. Whether this has any measurable effect on the orbit closure is discussed in Section 5.4.4.

5.3 Data Acquisition

5.3.1 Acquisition Tools

5.3.2 Monitoring Tools

Online display

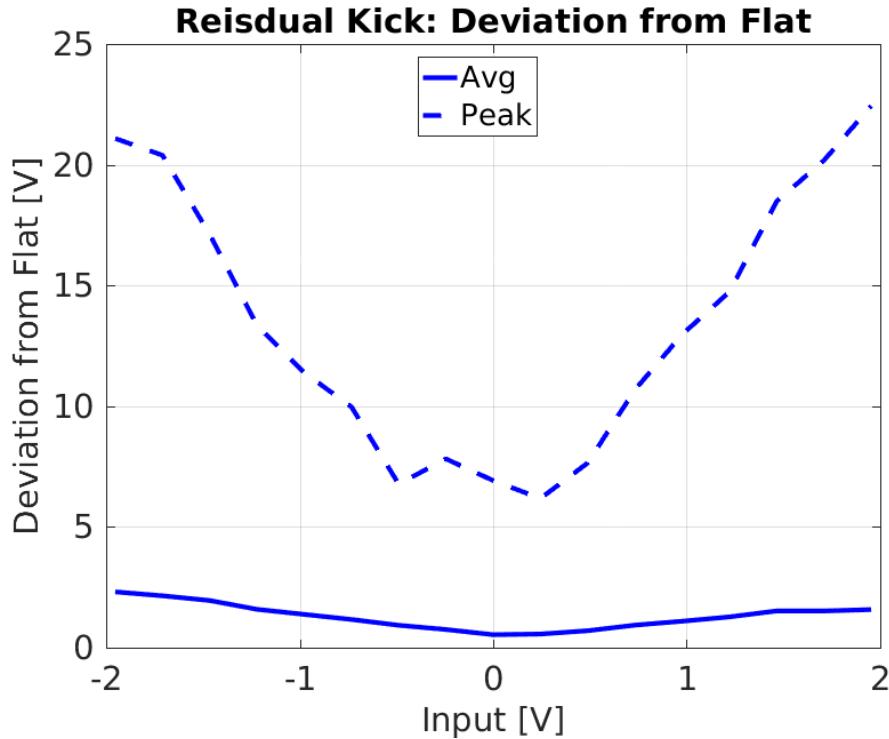


Figure 5.20: Residual kick along pulse: deviation from flat.

5.4 Kicker and Optics Performance Verification

5.4.1 Correction Range

Knowledge of the correction range of the PFF system, or more specifically the relationship between the voltage sent to the amplifier and the phase shift in the chicane, is critical for the PFF setup. The first checks of the ability to shift the phase in the TL2 chicane using the new phase feedforward optics were performed with magnetic correctors prior to the PFF amplifier being available (these correctors can be used to implement a secondary “Slow Correction” to complement the PFF system, as discussed in Section ??). Aside from their use for the PFF correction, these tests and the clear variation with beam phase versus voltage sent to the PFF kickers shown in this section are already a significant achievement and a verification of the extensive work to improve the MADX model of TL2 presented in Chapter 2.

Figure 5.21 shows the mean phase shift after the chicane (in the downstream phase monitor) across the full ± 2 V input range of the amplifier. Constant DAC outputs from the FONT5a board were sent to the amplifier in 17 steps between -4096 counts (-2 V) and +4096 counts (+2 V). In order to reduce the sensitivity to any drifts in the beam phase between data points the scan was taken in interleaved mode, alternating between pulses with no drive sent to the amplifier and a constant non-zero DAC output. The phase plotted in Figure 5.21 is therefore the difference between 50 kicked beam pulses and 50 “nominal” pulses taken at the same time for each amplifier input voltage.

At the maximum amplifier input voltage of 2 V the phase after the chicane is shifted

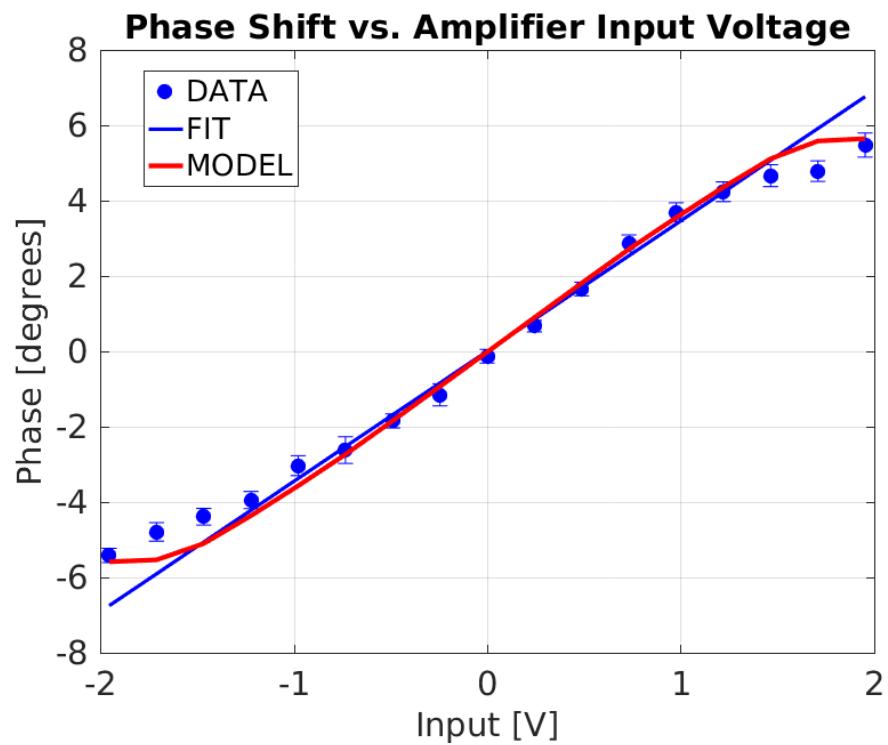


Figure 5.21: Phase shift versus amplifier input voltage.

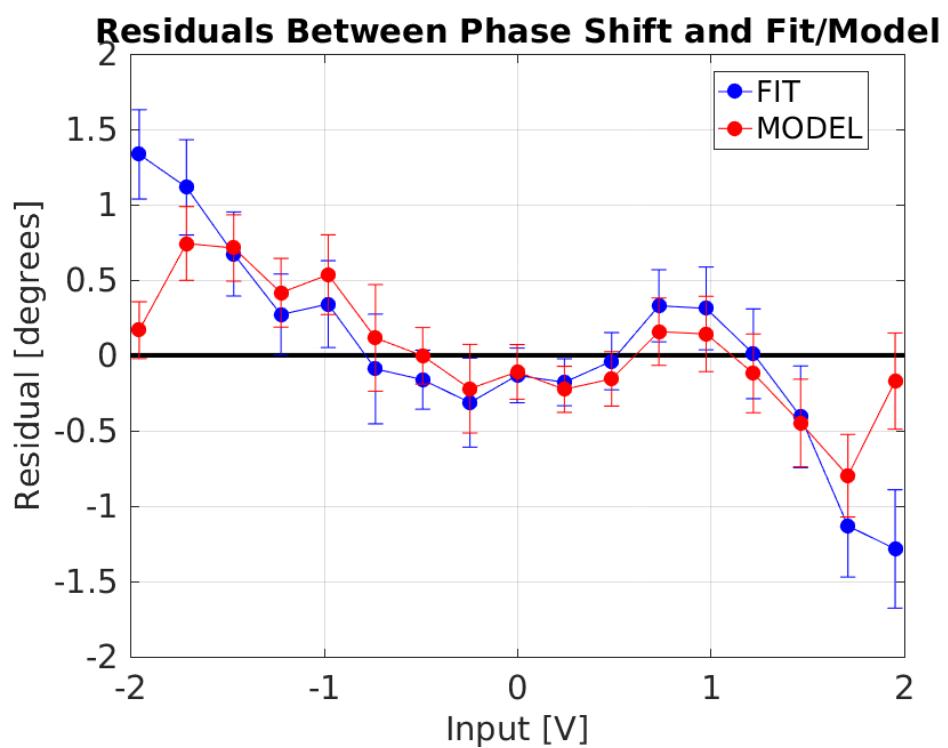


Figure 5.22: Phase shift versus amplifier input voltage.

	Phase Shift at +1 V Input	Max Phase Shift
Data	$3.5 \pm 0.1^\circ$	$5.5 \pm 0.3^\circ$
Model	3.6°	5.6°

Table 5.3: Phase shift at +1 volt input to the amplifier.

by $5.5 \pm 0.3^\circ$. The fitted phase shift per Volt sent to the amplifier is $3.5 \pm 0.1^\circ$ in the ± 1.2 V linear range of the amplifier (excluding the first and last three points, blue “FIT” line in Figure 5.21). This fitted gradient is required and was previously introduced for the conversion between the PFF gain in the units on the FONT5a board and the real applied gain in Section 5.1.1. In Section 5.1.1 it was quoted in terms of the phase shift in radians per DAC count output from the FONT5a board, rather than degrees per Volt as shown here. The value of $30\mu\text{rad}/\text{count}$ is easily derived using the conversion factors between degrees and radians and knowing that a DAC output of 4096 counts corresponds to 2 V sent to the amplifier. [TODO: Calculated factor is 29.827 microradians/count here. One I actually used for gain conversion, simulations etc. was 26.18 (used full range rather than linear range)]

Given knowledge of the amplifier output characteristics (Section 5.2.2), the kicker specifications (Section ??) and the chicane optics (Section ??) the real phase shift seen in the scan can be compared to the expected phase shift based on the system parameters. The predicted phase shift, $\Delta\phi$, in degrees is given by:

$$\Delta\phi = V_{amp}[V_{font}] \cdot K \cdot R_{52} \cdot \frac{360}{\lambda_{12\text{GHz}}} \quad (5.12)$$

Where $V_{amp}[V_{font}]$ is the amplifier output Voltage at an input voltage of V_{font} sent from the FONT5a board, K is the angular deflection of the beam per Volt applied to each kicker strip, R_{52} is the R_{52} value between the kickers in the PFF optics and $\frac{360}{\lambda_{12\text{GHz}}}$ converts the calculated orbit length difference in to an equivalent 12 GHz phase using the 12 GHz wavelength $\lambda_{12\text{GHz}}$. The value of most of these parameters has already been derived in the sections previously mentioned. They are:

$$V_{amp}[1 \text{ V}] = 435 \text{ V}$$

$$K = 0.8 \mu\text{rad}/\text{V}$$

$$R_{52} = -0.7 \text{ m}$$

$$\lambda_{12\text{GHz}} = 2.5 \text{ cm}$$

The value of $V_{amp}[1 \text{ V}]$ is given as a representative value in the linear range of the amplifier but the real amplifier output at all input voltages is used in the predictions to include the effects of saturation in the calculated phase shift values. Also, the output sent to the first kicker (from the left side of the amplifier) is used as this is most relevant for the phase shift in the chicane (the orbit should be closed after the second kicker with no further phase shift in the chicane after that point). The value of K is derived from the kicker design, in which 1.4 kV applied to each strip gives a 1 mrad kick for a 150 MeV beam [TODO: REF]. The actual CTF3 beam energy at this time was approximately 135 MeV (calculated based on the dipole currents used in the machine setup), so the value of K above is scaled by a factor 150/135.

In Figure 5.21 the red line “MODEL” shows the predicted phase shifts using Equation 5.4.1. Table 5.3 compares the fitted gradients and maximum phase shift for the model and real data. The overall agreement between the two is good, with the residuals between both the model and the data, as well as the linear fit to the data and the data, generally consistent with zero within error bars in the ± 1.2 V linear range of the amplifier as shown in Figure 5.22.

Outside the linear range some discrepancies appear, although at the maximum ± 2 V output the agreement is good so the effect is largest where the amplifier is entering saturation but before hard saturation is reached. However, most amplifier effects can be excluded as the analysis in this section uses the same dataset that was used to characterise the amplifier performance in Section 5.2. This could hint at possible remaining higher order errors in the TL2 chicane optics, or unexpected behaviour from the kickers or amplifier. Although subtracting alternating, interleaved pulses should remove the sensitivity to drifts in the machine it is possible that some residual effect remains. To determine whether the discrepancies are reproducible further scans of this type will need to be completed in the future. The residuals between between the data and the linear fit between ± 1.2 V would also be of significance for the PFF correction should they not converge to zero with additional measurements, as they are of similar magnitude to the 0.2° downstream jitter target.

However, the overall conclusion is as expected — the phase shift in the chicane linearly depends on the amplifier input in the ± 1.2 V (± 2500 DAC counts) region thus a close to optimal correction can be applied in this range, corresponding to a $\pm 4.2 \pm 0.1^\circ$ phase shift. However, when the calculated optimal correction is between an absolute input voltage of 1.2 V and 2.0 V, 2500 to 4096 DAC counts, or $\pm 4.2 \pm 0.1^\circ$ to $\pm 7.0 \pm 0.2^\circ$, the actual phase shift in the chicane is lower, only up to $\pm 5.5 \pm 0.3^\circ$, due to the amplifier entering saturation (and possibly other effects to be determined). Any calculated correction outside $\pm 5.5 \pm 0.3^\circ$ receives a static phase shift of $\pm 5.5 \pm 0.3^\circ$ in the chicane. In the limit where all pulses are outside this range the PFF system can only induce a static shift in the mean phase and makes no improvement to the phase jitter. Understanding the impact of the limited correction range on the PFF results was particularly critical for interpreting the early correction attempts with the first version of the amplifier, giving approximately half the ranges shown in this section. This is discussed using simulations of the PFF system in Chapters ?? and 6.

<http://accelconf.web.cern.ch/accelconf/ipac2011/papers/tupc007.pdf> 1.4 kV to each strip
 $= 1$ mrad kick at 150 MeV 1.26 kV to each strip $= 1$ mrad kick at 135 MeV

5.4.2 Variations Along Pulse

[TODO: In this section I intended to check the stability of the constant kick along the pulse, as I did for the ”Shape” section with the amplifier above. The results from the constant kick data do not look good, though - the errors are quite large and there are some nasty oscillations along the pulse, such as in the example Figure 5.23. With the phase stability downstream a much longer scan would probably be needed to draw conclusions here, though even on longer time scales I think it needs to be checked that the difference between odd and even pulses converges to zero. As it is I would probably choose to skip this section, or

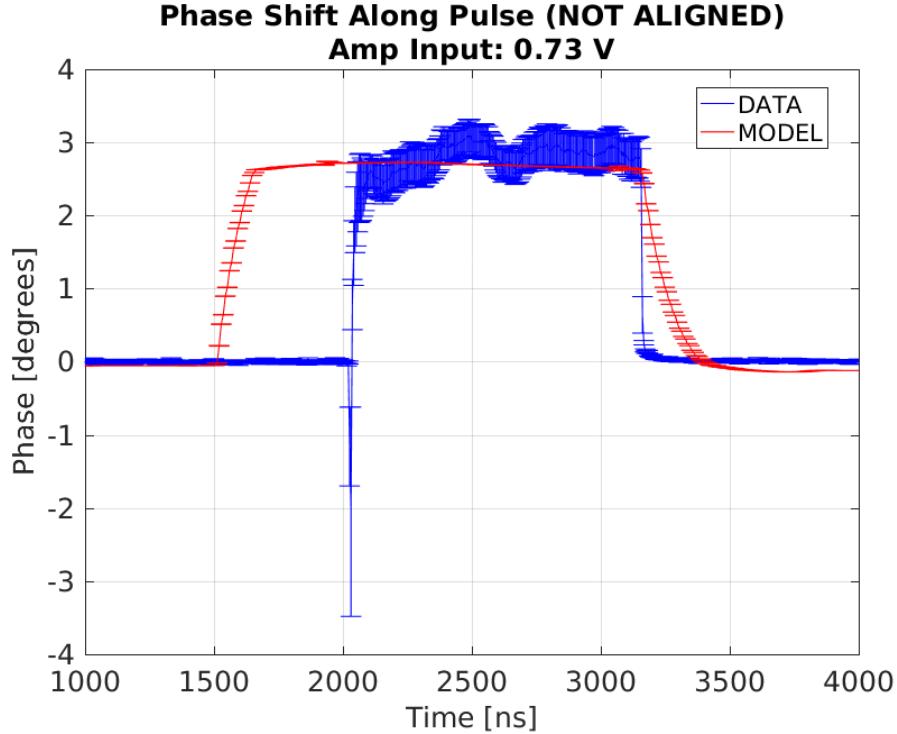


Figure 5.23: Traces relative timing scan.

possibly just show one of the better plots.]

5.4.3 Shape

[TODO: The purpose of this section is to compare the shape of the given kick to the upstream phase - i.e. compare the upstream phase to the amplifier monitoring output and the difference in the downstream phase with PFF off and on. Again here my first attempts produced results that were not as I'd like, without particularly good agreement between the difference in the downstream phase and the upstream phase shape. Looks kind of like the two are rotated with respect to each other. More analysis needed here to check more thoroughly. I think this section should definitely be included, but I need either an explanation for the differences or to find datasets where the agreement is much better.]

5.4.4 Orbit Closure

At CLIC the PFF system must not degrade the transverse beam stability. This means for any voltage sent to the kickers the horizontal beam orbit after the PFF chicane must be unchanged, or closed, despite the different orbits inside the chicane. As such, the PFF optics for the TL2 chicane at CTF3 is also designed to give a closed kick, as presented in Section ???. However, up until now the main focus during PFF operation has been the primary goal of reducing the downstream phase jitter and ensuring good beam transmission to the downstream phase monitor. As a result orbit closure after the TL2 chicane has not yet




Figure 5.24: Traces relative timing scan.

been strictly enforced during PFF operation as will be seen in this section, but the current status is shown here as an additional cross-check of the PFF optics and to highlight where improvements are needed for future tests.

Using the same constant kick data as Section 5.4.1 Figure 5.25 shows mean the horizontal orbit before, inside and after the TL2 chicane across the full ± 2 V range of inputs sent to the amplifier. The vertical black lines on the plot mark the approximate location of the entry to the chicane (index CC.500) and the exit of the chicane (index CC.800), with the two kickers being located at CC.480 and CC.780. Two BPMs before and after the chicane, as well as the four inside the chicane, are included. The plotted positions are the difference between the kicked and nominal (non-kicked) orbit at each BPM, thus removing any misalignment in the BPM centres. Before the chicane and the first kickers there is no significant effect on the orbit as expected. Inside the chicane the PFF system induces an orbit offset of up to 1.4 ± 0.1 mm. After the chicane, in BPMs CC.845 and CC.930, the orbit should return to zero in the ideal case. However, a clear residual offset can be seen, up to 0.5 ± 0.1 mm in CC.930.

During this scan the input sent to both sides of the amplifier was the same magnitude. However, in Section 5.2.2 it was seen that the right side of the amplifier, sent to the second kicker, gave 4% lower output than the left side. This could explain why the orbit after the chicane was not closed during this scan, and during PFF operation which has typically used equal gain for both correction outputs to date. Figure 5.26 shows the expected orbit in the TL2 chicane in the case where both kickers are driven with the same voltage ("nominal

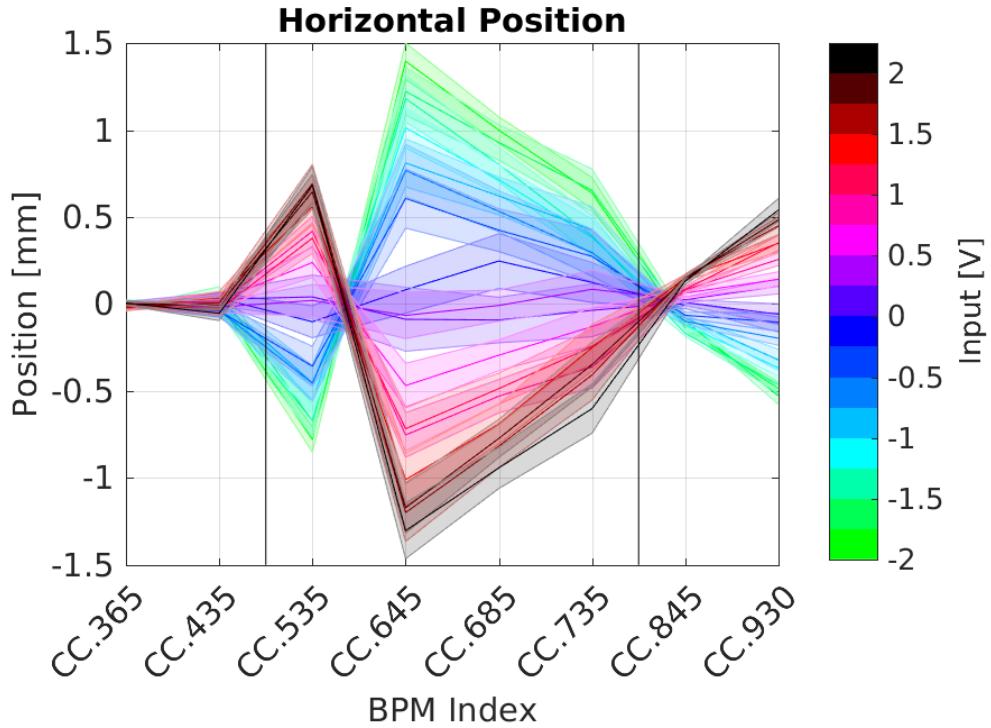


Figure 5.25: Horizontal orbit offset in and around the TL2 chicane at different input voltages sent to the amplifier.

optics") and with the 2nd driven with a 4% lower voltage ("real amplifier ratio").¹ The full MADX orbit propagated through all elements is shown, with the eight real measured BPM offsets also included at their respective positions. Each BPM point represents the gradient of a linear fit using the variation with input voltage seen previously in Figure 5.25. Before and inside the chicane the agreement between the BPM data and the model is excellent. As seen before, the BPM orbit is not closed after the chicane, unlike the nominal optics. If the actual kicker voltage ratio is taken in to account the MADX orbit after the chicane is not closed either, but the offset has opposite polarity to the real data in the last BPM.

The true explanation for the non-closed orbit is that the quadrupole strengths used in the machine setup are not exactly nominal. Although every effort has been made to keep the TL2 optics as close to nominal as possible, particularly inside the chicane, it is an extremely sensitive area for the setup of CTF3 and beam transport in to the CLEX area downstream of TL2 (including the location of the downstream phase monitor in the TBL line) is always difficult. Minor modifications have therefore been necessary in order to achieve full beam transmission to the downstream area, both for PFF and other experiments at CTF3. The largest changes have been made to the four quadrupoles following the chicane but one quadrupole inside the chicane, CC.IQFL0730 (just prior to the 2nd kicker) has a set value 10% lower than the nominal optics, as well as differences up to 2% in the other quadrupoles. Using the real quadrupole strengths used in the machine gives the result shown in Figure 5.27. In this case the agreement between the model and the data is also extremely good

¹The term "nominal optics" is used in this section to refer to the nominal PFF optics, not the nominal optics created in Chapter 2 to use when the PFF system is not under operation.

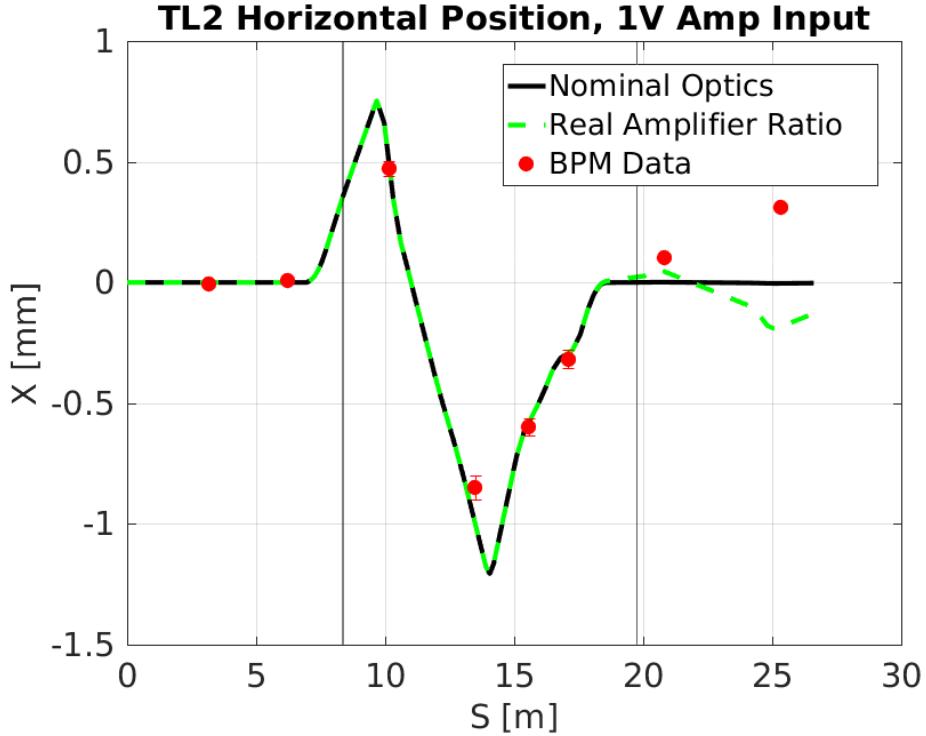


Figure 5.26: Orbit in the TL2 chicane at 1 V amplifier input for the BPM data, nominal model and model taking in to account the difference in amplifier output voltage to each kicker.

after the chicane. It may still be possible to compensate for these differences by outputting different voltages to each kicker and this will be investigated. Alternatively, a completely nominal optics can be set in the chicane purely for the purposes of verifying orbit closure, reverting back to non-nominal optics and closure if needed to achieve good beam transmission for normal PFF operation.

5.5 Correction Output Timing

All the results based on the amplifier outputs and kicked beam presented so far have used a constant output voltage sent from the FONT5a board across the full $1.4 \mu\text{s}$ time window that the amplifier is powered for. As the $1.4 \mu\text{s}$ amplifier output sent to the kickers is much longer than the $1.1 \mu\text{s}$ CTF3 beam pulse it is easy to ensure that the full length of the pulse experiences the constant kick with this setup. However, for operation of the PFF system precise control of the correction output timing becomes critical. In order to remove phase variations along the pulse with the PFF system the output correction signal, shaped by the upstream phase, must arrive at the kickers exactly in sync with the beam. Any timing misalignment between the beam and correction signal arrival will result in residual oscillations along the pulse in the downstream phase, even in the case where the upstream-downstream phase propagation is perfect. Approximating the phase sag along the pulse to be quadratic, $\phi_u(t) \sim \phi_d(t) \sim t^2$, a misaligned correction would yield a corrected downstream phase with a linear increase in phase along the pulse with time, $\phi_{PFF}(t) \sim \phi_d(t) - \phi_u(t + \delta) \sim -2\delta t - \delta^2$,

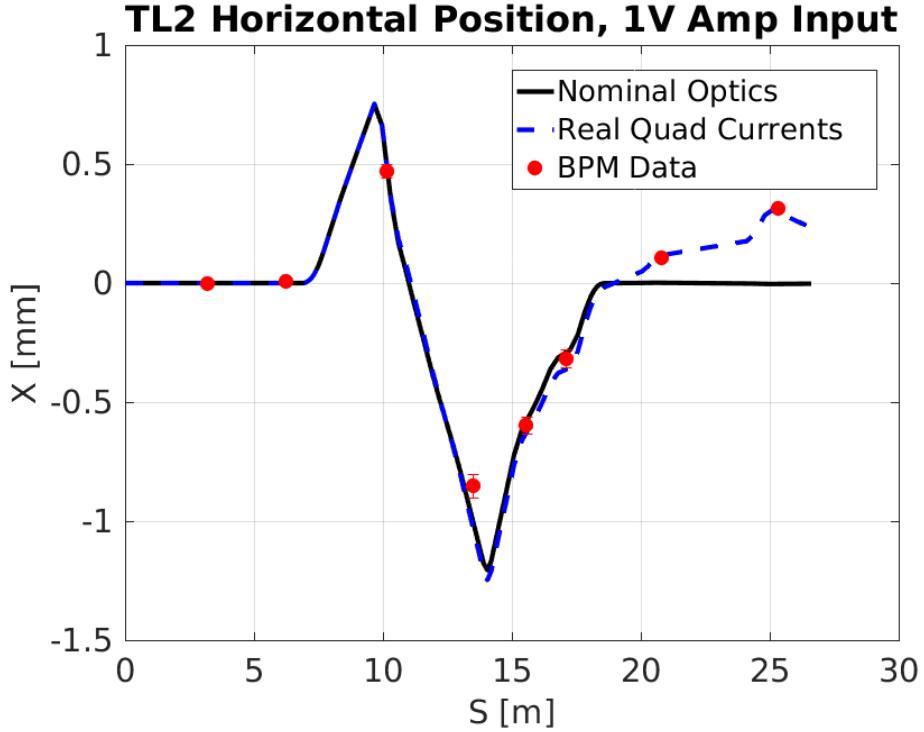


Figure 5.27: Orbit in the TL2 chicane at 1 V amplifier input for the BPM data, nominal model and model taking in to account the quadrupole currents in the real machine setup.

for example, where ϕ_{PFF} is the corrected phase, ϕ_d the uncorrected downstream phase, ϕ_u the upstream phase, t the time and δ the time misalignment in the applied correction. Also, the effect is particularly significant for any higher frequency variations in phase along the pulse. If a 40 ns oscillation is present in the upstream phase but the correction is applied with a 50 ns delay, for example, a second 40 ns oscillation with opposite sign would be introduced to the “corrected” downstream phase at a later time with no change to the initial oscillation. Although the effects are most visible along the pulse, any timing delay will also degrade the achievable mean phase jitter. This section gives an overview of the main methods that have been used to ensure that the correction output to the two kickers arrives in time with the beam.

5.5.1 Kicker Cable Lengths

The cables carrying the correction signal between the PFF amplifier and the kickers in the TL2 chicane are the single largest contributor to the overall system latency. They must be routed from the PFF electronics racks (in the klystron gallery, one floor and directly above the location of the upstream phase monitors), down in to the machine hall and across the width of the CTF3 facility to the TL2 chicane. The initial kicker cables installation used pre-existing cable trays and gave a signal transit time of 260 ns, with a signal speed of 0.66 c and approximate lengths of 50 m. This is more than two thirds the overall PFF latency budget, which must be lower than the 380 ns beam time of flight between the upstream phase monitor and the first kicker. Considering the latencies of the various pieces

of hardware in the PFF system chain as well as the cables between the upstream phase monitors and the PFF electronics (Section ??), the overall PFF system latency would have been in excess of the 380 ns beam time of flight with this setup. By re-routing the cables on to a dedicated pathway and trimming any remaining slack it was possible to reduce the cable lengths by up to 90 ns, bringing the system within the latency budget as will be seen in the following sections. Precise measurements of the cable lengths with this setup are presented in this section, as well as their significance beyond ensuring the system is within the latency requirements.

With two kickers, two strips per kicker and two ends of each strip a total of eight cables are needed. The drive from the amplifier is sent to the downstream end of each kicker strip, traverses the kicker, and is then terminated back at the amplifier after leaving the upstream end. Drive is sent to the downstream end of each kicker (meaning it propagates through the kicker in the opposite direction to the beam) so that the electric and magnetic fields between the strips are in the same direction, as discussed in Section ???. Rather than being connected directly to the amplifier, the eight kicker cables are connected to a patch panel below the amplifier in the PFF electronics racks. Eight additional cables, around 70 cm in length, are used to connect the amplifier outputs to the patch panel. This is in order to create a tidier cabling setup in the rack as well as making changes to the amplifier cabling easier, when necessary. [TODO: Picture]. The kicker cables are of type [TODO: REF] with HN-type connectors and the patch panel cables of type [TODO: REF] with N-type connectors to match the amplifier.

The length of the eight kicker cables and eight patch panel cables has been measured using time domain reflectometry (TDR) on a network analyser [TODO: REF?]. The network analyser is used to send a short pulse down the cable, with one end of the cable connected to the network analyser and the other end disconnected so it is not correctly terminated. As the signal reaches the (non-terminated) end of the cable the discontinuity in impedance creates a reflected signal that propagates back to the network analyser. The time difference between when the signal was output and when the reflected signal arrives back at the network analyser therefore corresponds to double the one-way signal transit time in the cable. [TODO: Example TDR plot?]

Table 5.4 shows the patch panel cable lengths and Table 5.5 the kicker cable lengths that were determined with this method. Quoted errors of ± 0.05 ns are estimated based on the sampling rate of the measurement. The amplifier port, patch panel port and kicker strip that the cables are connected to are also shown in the table, as well as their corresponding CTF3 identifying number for reference. For the amplifier port labels the three letters correspond to:

- Whether the cable is connected to the **L**evel or **R**ight side of the amplifier.
- Whether the cable is connected to the amplifier **A** or **B** outputs on that side.
- Whether the cable is connected to the amplifier **D**rive or **T**erminator.

And for the kicker strip labels the three letters correspond to:

Label	Length	Amplifier Port	Patch Panel Port
2907701B	2.99 ± 0.05 ns	LAT	1
2907703B	3.03 ± 0.05 ns	LBT	2
2907700B	3.03 ± 0.05 ns	LAD	3
2907702B	3.01 ± 0.05 ns	LBD	4
2907838B	3.03 ± 0.05 ns	RAD	5
2907740B	3.05 ± 0.05 ns	RBD	6
2907739B	3.03 ± 0.05 ns	RAT	7
2907741B	3.03 ± 0.05 ns	RBT	8

Table 5.4: Lengths of cables between the amplifier and the patch panel.

Label	Length	Patch Panel Port	Kicker Strip
2907701A	171.28 ± 0.05 ns	1	1UL
2907703A	171.30 ± 0.05 ns	2	1UR
2907700A	171.29 ± 0.05 ns	3	1DL
2907702A	171.30 ± 0.05 ns	4	1DR
2907838A	205.45 ± 0.05 ns	5	2DL
2907740A	205.62 ± 0.05 ns	6	2DR
2907739A	205.15 ± 0.05 ns	7	2UL
2907741A	204.49 ± 0.05 ns	8	2UR

Table 5.5: Lengths of cables between the patch panel and the kickers.

- Whether the cable is connected to the first (**1**) or second (**2**) kicker.
- Whether the cable is connected to the **Upstream** or **Downstream** end of the kicker strip.
- Whether the cable is connected to the **Left** or **Right** kicker strip, as viewed looking at the upstream end of the kicker.

Finally, the patch panel connectors are simply labelled from 1 to 8 from left to right, as viewed from the front of the rack. All of the cable connections between the amplifier, patch panel and kickers are shown in Figure 5.28.

The patch panel cables all have lengths of around 3 ns, with the lengths of each matched to within the measurement error. After the re-routing and shortening of the kicker cables the cables connected to the first kicker have a length of around 170 ns, whilst the cables for the downstream kicker are longer at around 205 ns. For the downstream kicker the latency requirements are slightly relaxed due to the additional 36 ns beam time of flight between the kickers. Rather than also shortening the downstream kicker cables as much as possible some additional slack was left so that the difference in lengths is similar to the difference in the beam time of flight between the two. This means the two correction outputs (one for each kicker) can be sent from the FONT5a board at, or close to, the same time as discussed in Section 5.5.3.

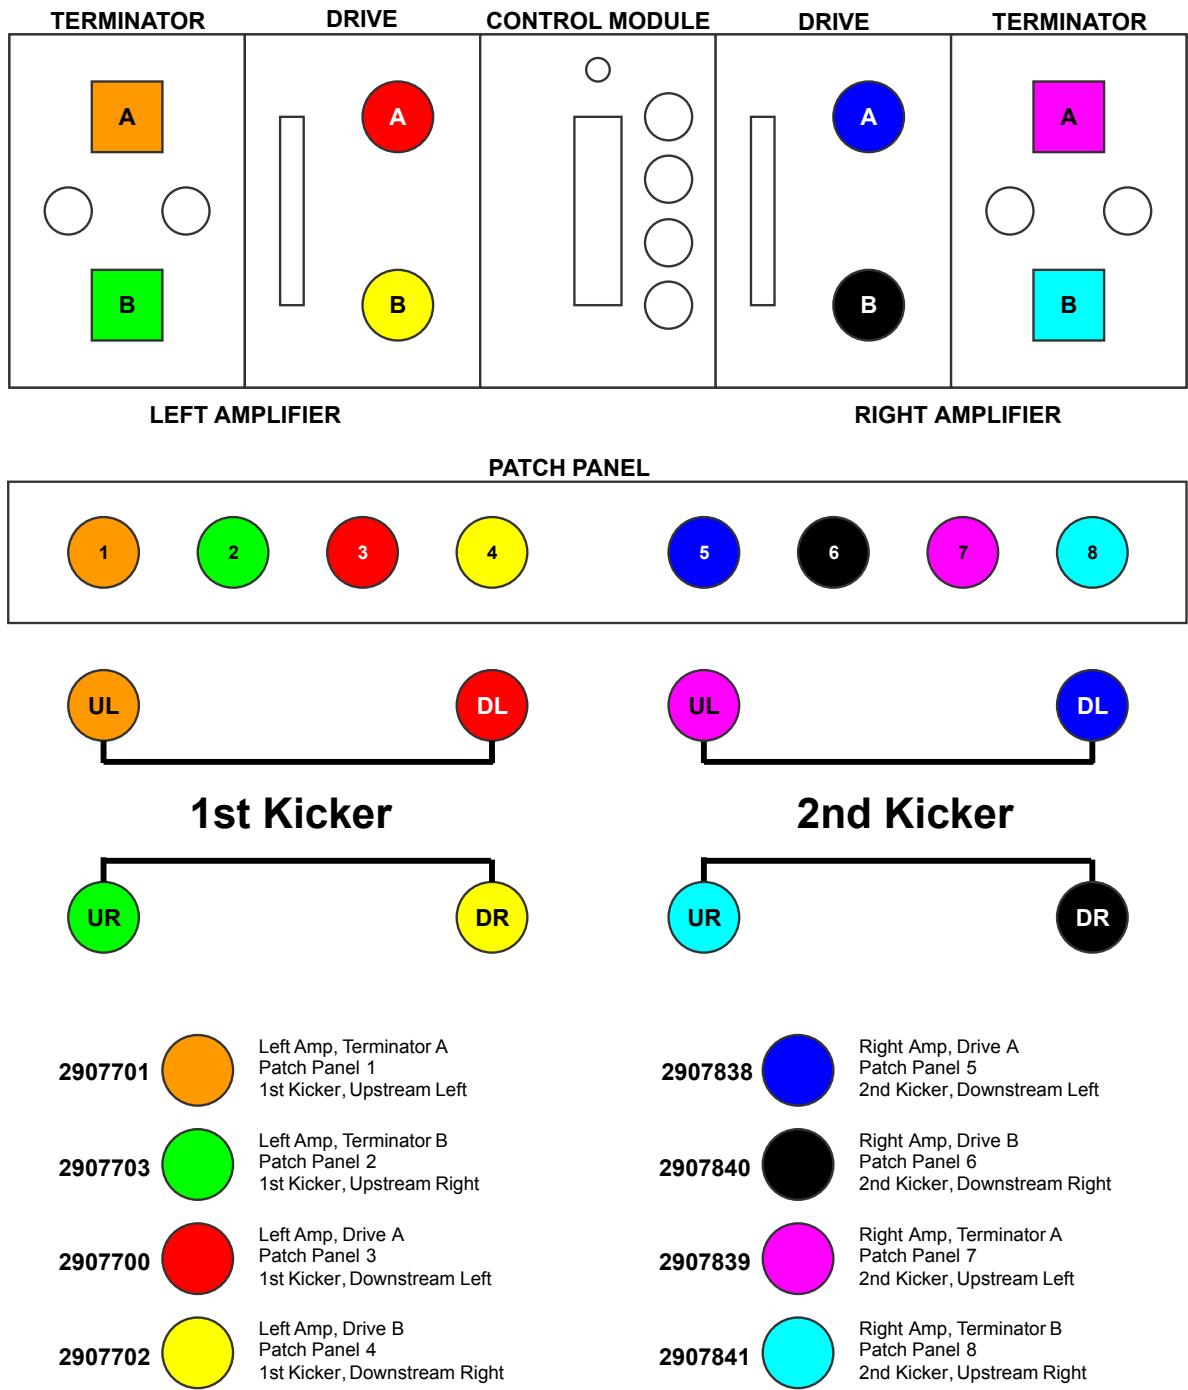


Figure 5.28: Cabling setup for cables between the amplifier and kickers.

Although all the upstream kicker cables are matched within the measurement error there are some differences in the downstream kicker cable lengths, with cable 2907741A more than 1 ns shorter than cable 2907740A, for example. If there is a difference between the lengths of the cables connected to the downstream left and downstream right strips of a kicker (the driven end) there will be a time offset in the voltage applied to each side of the kicker, which would degrade the quality of the PFF correction. However, there is no need for the cables connected to the upstream ends of the kickers to be of matched lengths, the only requirement is that they are terminated correctly at the amplifier. The shorter 2907741A cable is therefore connected to the upstream end of the second kicker, and the cables 2907838A and 2907740A, with lengths matched to within 200 ps, are used to carry the amplifier output to the downstream end of the strips.

[TODO: Tolerances for cable length matching?]

5.5.2 Absolute Timing

Using Beam Pickup

[TODO: Kicker pick-up theory. In particular, is time between pickup exactly the beam pulse length or is it longer by double the kicker length?]

Figure 5.29 shows the beam pickup at the start and end of the pulse from the PFF kickers at CTF3, as seen on one of the amplifier monitoring outputs (each of the four amplifier monitoring outputs, one for the upstream end of each strip prior to the signal being terminated at the amplifier, gives a similar response). The separation of the peaks in the beam pickup is $1.1 \mu\text{s}$, thus the same as the CTF3 pulse length as expected. By comparing the timing of these peaks with respect to the start and end of the amplifier output pulse, using the same amplifier monitoring signal, it is possible to ensure that the correction output arrives in sync with the beam.

An example of this is shown in Figure 5.30. A constant DAC output is sent from the FONT5a board to the amplifier and both this output pulse and the beam pickup, at samples 275 and 493, are visible in the figure. Importantly, the DAC output is gated using the upstream phase monitor diode signal (in other words, the constant DAC output is only sent during the time when the diode is non-zero) and this has two consequences. Firstly, the amplifier output pulse has the same length as the beam pulse in the upstream phase monitor. Secondly, the timing of the output is identical to what it would be in normal PFF operation. In the case of Figure 5.30 the drive to the amplifier is sent as quickly as possible after the arrival of the upstream diode signal at the FONT5a board. It can be seen that the amplifier pulse arrives before the beam pickup, thus with the PFF system setup this way the correction would be applied slightly early. This result therefore proves that the PFF system just meets the latency requirements, with the overall time needed to transport and process all the relevant signals a few tens of nanoseconds less than the 380 ns time of flight of the beam between the upstream phase monitor and the first kicker. However, what is also clear in the figure is that the time offset between the start of the amplifier pulse and the first beam pickup spike is much larger than the time difference between the end of the amplifier pulse

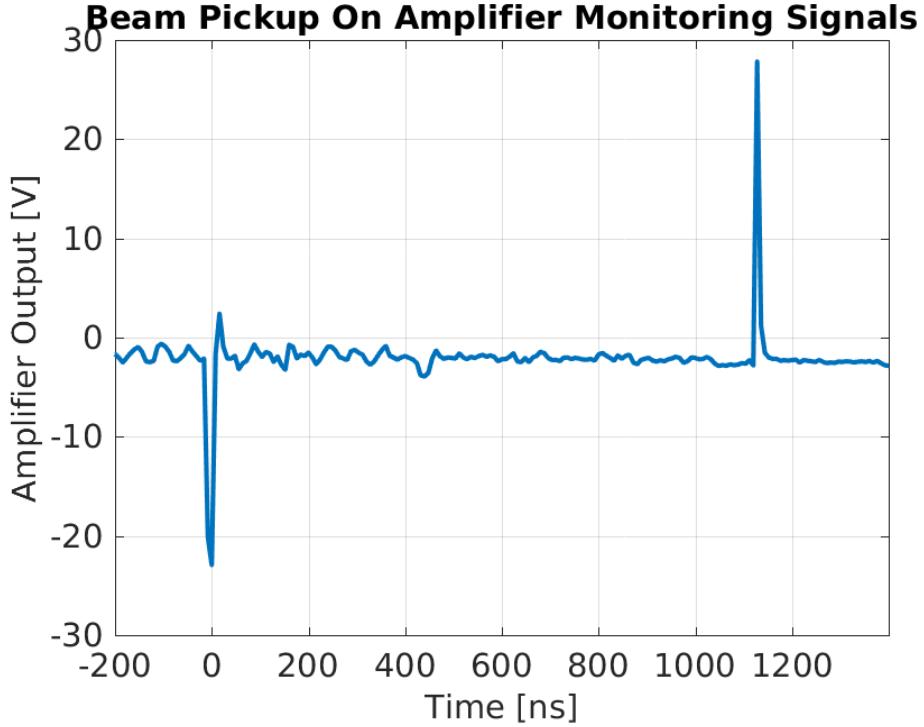


Figure 5.29: Beam pickup on kicker strips as seen on amplifier monitoring signals.

and the second beam pickup spike. This is due to the energy transient across the first 100 ns of the CTF3 beam pulse which is present in the upstream phase monitor but is then lost prior to the TL2 chicane, predominantly in TL1. As a result the downstream beam pulse is shorter than the upstream beam pulse which defines the length of the correction output. Therefore, in order to align the correction output with the beam the signals from the end of the amplifier and beam pulses must be used, not the start.

The firmware for the FONT5a board includes an output delay parameter that can be used to fine-tune the timing of the correction output sent to the amplifier. This can be done independently for each of the two correction outputs so that it can be ensured the correction arrives in sync with the beam in each kicker individually (the relative timing of the two kickers is discussed in the next section). The delay can be varied between 0 and 31 clock cycles in integer steps, with one clock cycle corresponding to one period of the 357 MHz ADC clock frequency, or 2.8 ns. A delay of up to 86.8 ns can therefore be added to the correction outputs. Figure 5.31 shows the effect of varying the output delay across the full range of possible values, zoomed in on the end of the pulse. For all output delays the beam pickup remains at sample 493, as expected. Meanwhile, the end of the amplifier pulse is moved from before the beam pickup (output too early) to after the beam pickup (output too late). To achieve the optimal correction timing the end of the amplifier pulse must be aligned with the beam pickup and this is achieved with a delay of 7 clock cycles, or 19.6 ns, as shown in Figures 5.32 and 5.33. This delay has been used for the latest PFF runs presented in Chapter 6. Due to ambiguity in which point along the falling edge of the amplifier pulse the beam pickup should be aligned to there may be a remaining error of up to 3 clock cycles in the exact alignment, and this can only be verified by beam based measurements (not using the

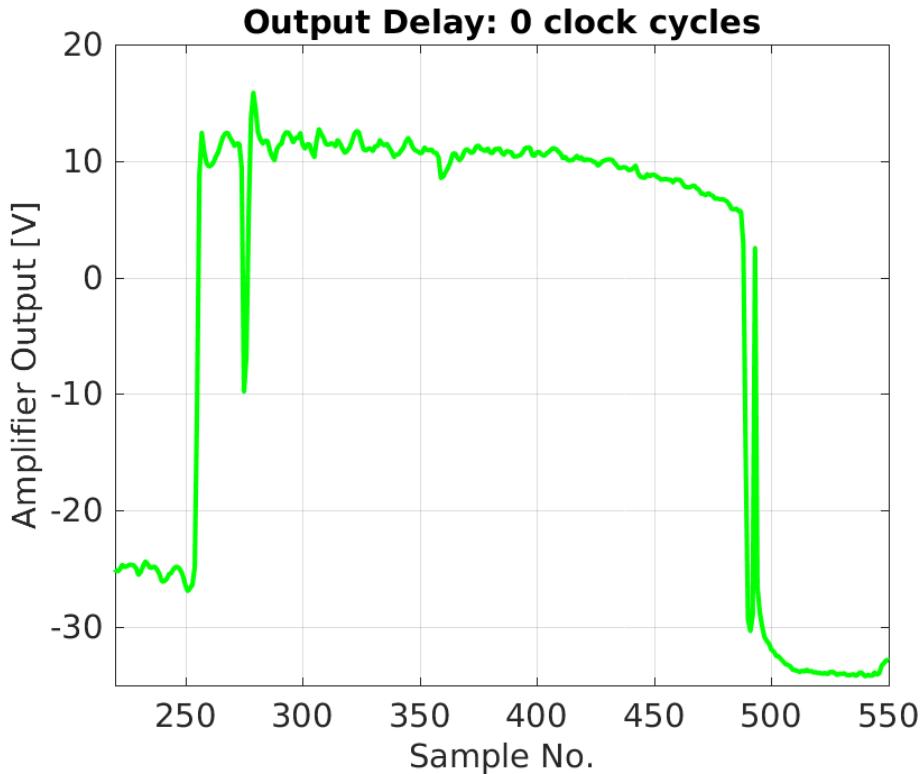


Figure 5.30: Output delay of 0 clock cycles. Full pulse.

amplifier monitoring outputs).

Using BPMs

This section presents one way in which the correction output timing can be determined using a combination of the phase monitor measurements and a BPM signal downstream of the TL2 chicane. The results shown here were performed with the first, lower power version of the amplifier and the FONT5 rather than the later FONT5a board, and because of this the optimal output delay calculated here does not agree with the value of 7 clock cycles from the beam pickup based measurement above. The newer hardware has the same latency as the previous versions, thus the difference does not come from hardware changes but rather associated changes to cabling between the phase monitor electronics, FONT5a board and amplifier. The measurement will be repeated in the future to verify that both methods give consistent results when the same hardware and cabling setup is used.

The FONT5 (and FONT5a) board firmware provides the functionality to be able to change the gain of each PFF correction output independently. This means it is possible to apply the correction to only one kicker, or to kick the beam in the same direction in each kicker (i.e. to use the same sign for the gain in each kicker, rather than gains with equal magnitude but opposite sign). In both of these cases the kicked PFF orbit in the chicane is not closed, thus the horizontal position along the beam pulse in a BPM after the chicane

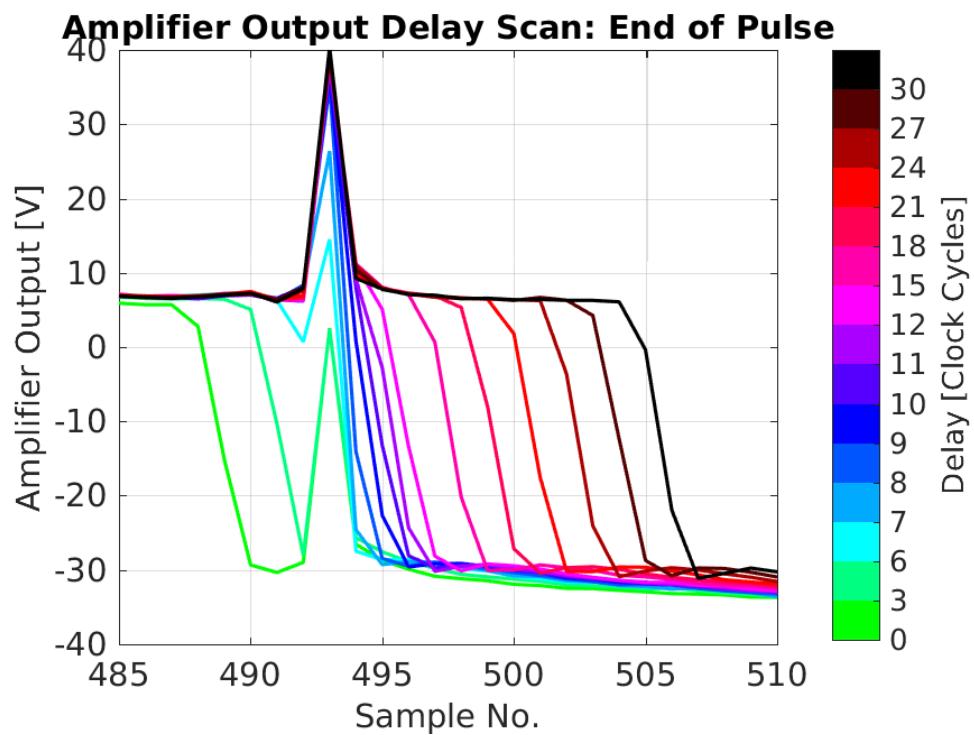


Figure 5.31: Output delay scan, end of pulse.

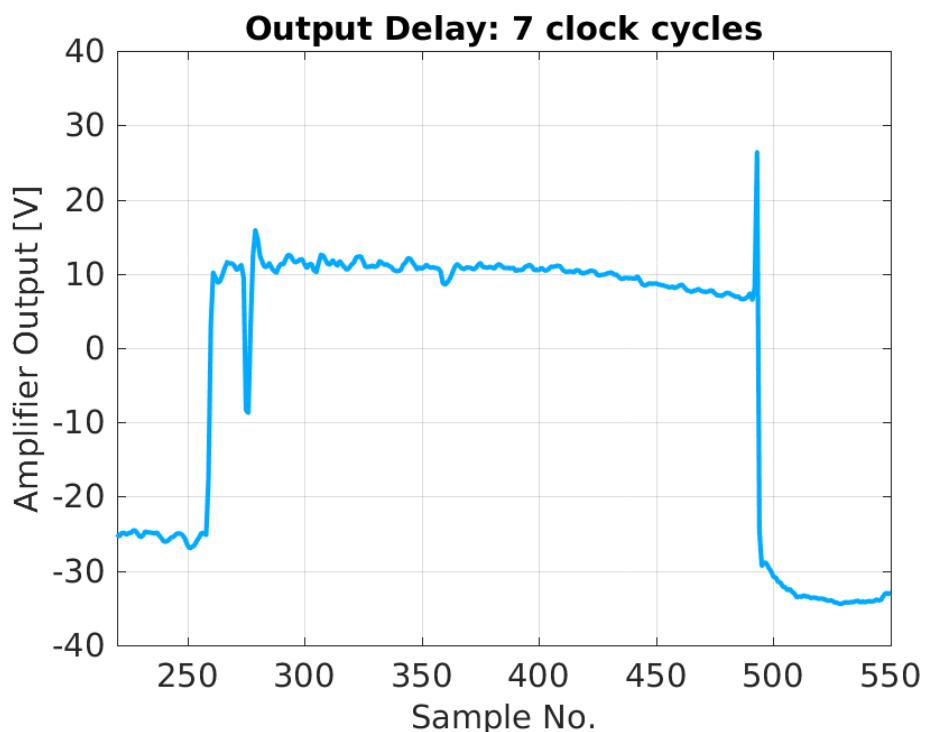


Figure 5.32: Output delay of 7 clock cycles. Full pulse.

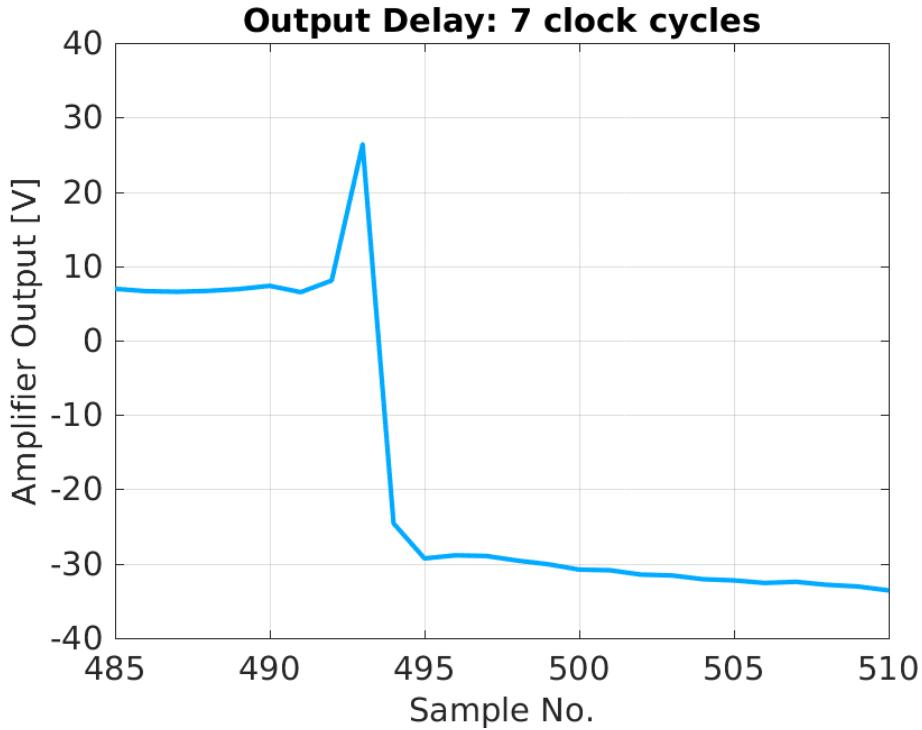


Figure 5.33: Output delay of 7 clock cycles. End of pulse.

depends on the shape and timing of the applied correction.²

Figure 5.34 compares the upstream phase, downstream phase and horizontal position (in a BPM after the TL2 chicane) along the pulse in the case where the PFF correction is applied with gains set to kick the beam in the same direction in each kicker, and with no output delay applied in the FONT5 board. The data is taken in interleaved mode, with the plotted phases shown using the PFF off data and the BPM trace being the difference between the PFF on and PFF off data. Each signal is scaled and sign flipped where necessary to give variations along the pulse with the same magnitude and sign, in arbitrary units. The BPM and phase monitor signals are acquired with the same sampling frequency of 192 MHz, with each aligned so that the end of the pulse is at the same sample number.

By taking the difference of the PFF off and PFF on data in the BPM any residual orbit variations along the pulse not related to the PFF system are removed, thus the remaining shape should match that of the PFF correction output, which in turn is linked to the upstream phase. The downstream phase should also have the same shape as the upstream phase with the PFF system off, within the limits of the upstream-downstream phase correlation achieved at this time. During this measurement many oscillations along the upstream phase were present, which usually are not desired but for this measurement are perfect points of reference to check the time alignment of the signals. As expected the overall shape of the residual horizontal position in the BPM along the pulse and the two phase signals is very similar. The largest feature in the upstream phase that is present in all three signals occurs

²In Section 5.4.4 it was shown that the corrected orbit is not perfectly closed in normal PFF operation either. However, in this case no attempt at orbit closure is made so the measured effect seen in the BPMs is much larger.

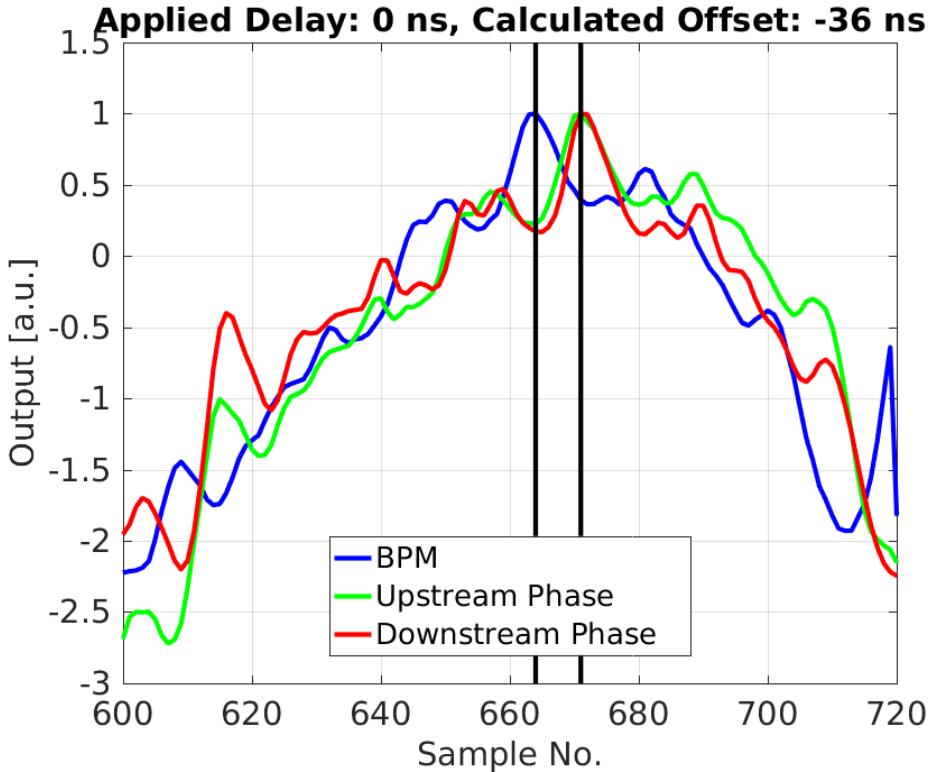


Figure 5.34: Kick output with no delay as seen on BPM and phase signals.

at sample 671 in the upstream phase, with the location of the peak of this oscillation in the phase signals and the BPM marked by vertical black lines in the figure. The peak as seen in the BPM signal is clearly before the peak in the phase monitor signals thus in this case the correction was applied early, with a measured offset of -36 ns between the peaks.

This measurement was repeated with four different correction output delays applied in the FONT5 board, at delays of 0, 10, 20 and 30 clock cycles (0 to 84 ns), which includes points where the correction is applied both early and late. Fitting the measured time offset between the peaks in the BPM and the phase in the same way as before yields an optimal correction output delay to apply of 39 ± 7 ns (Figure 5.35), or 14 ± 3 clock cycles. Applying this delay in data analysis gives the result shown in Figure 5.36, in which the similarity of the three signals becomes clear.

5.5.3 Relative Kicker Timing

For the phase correction the absolute output timing sent to the first kicker, as derived above, is the most critical as this defines the alignment of the applied phase shift in the chicane with the beam. The second kicker's main purpose is then to close the kick created by the first, ensuring the orbit after the chicane is closed (with the caveats already mentioned in Section 5.4.4). For the purposes of orbit closure it is also important to ensure that the correction arrives at the second kicker in time with the beam. As discussed in Section 5.5.1 the beam time of flight between the kickers is about 36 ns, thus the correction must arrive at




Figure 5.35: Fit time offset between kick and beam at different output delays.




Figure 5.36: Alignment between BPMs and phase signals with optimal delay applied in analysis.

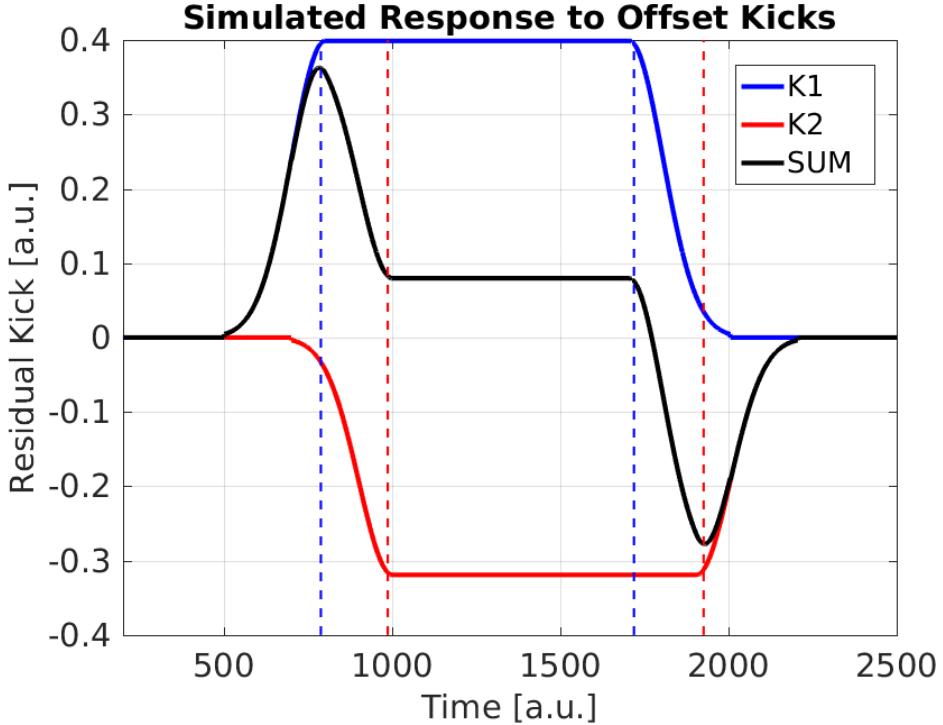


Figure 5.37: Simulated response to offset kicks.

the second kicker 36 ns later than the first kicker. Most of this difference should be accounted for by the longer cable lengths for the second kicker, but the precise relative timing is checked here. In this context the relative timing means the additional output delay that must be applied to the FONT5a correction output for the second kicker with respect to the first in order to ensure the correction is aligned in time with the beam in both kickers.

Figure 5.37 shows a simulated example of the expected effect of kicking the beam with a relative time offset in each kicker, in this case with the output to the second kicker arriving later than the first kicker (with respect to the beam pulse). The kickers are driven with opposite polarity in the same way as the PFF system, and the first kicker is shown with a larger output than the second. The total kick received in the chicane is given by the sum of the two, shown in black. In the ideal case the total/residual kick in the chicane should be zero so that the orbit is closed after the chicane. However, with a timing offset between the two kickers there are large peaks in the residual kick at the start and end of the pulse, where only one of the two kickers receives its full drive. Due to the different amplitude of the two kickers the residual kick is also non-zero in the central part of the pulse. With well-aligned timing the residual kick would be constant along the full pulse length, or zero across the full pulse length if the kicks had matched amplitudes.

By varying the relative timing of the two correction outputs on the FONT5a board (K1 and K2 delay) and using a BPM after the TL2 chicane to measure the size of the peaks at the start and end of the pulse resulting from the offset kicks (in the same way as Figure 5.37) the optimal relative delay can be determined. The optimal relative delay is the point that minimises the size of the peaks on the rising/falling edge of the pulse, with the peak magnitude approximately linearly dependent on the delay. Figure 5.38 shows the

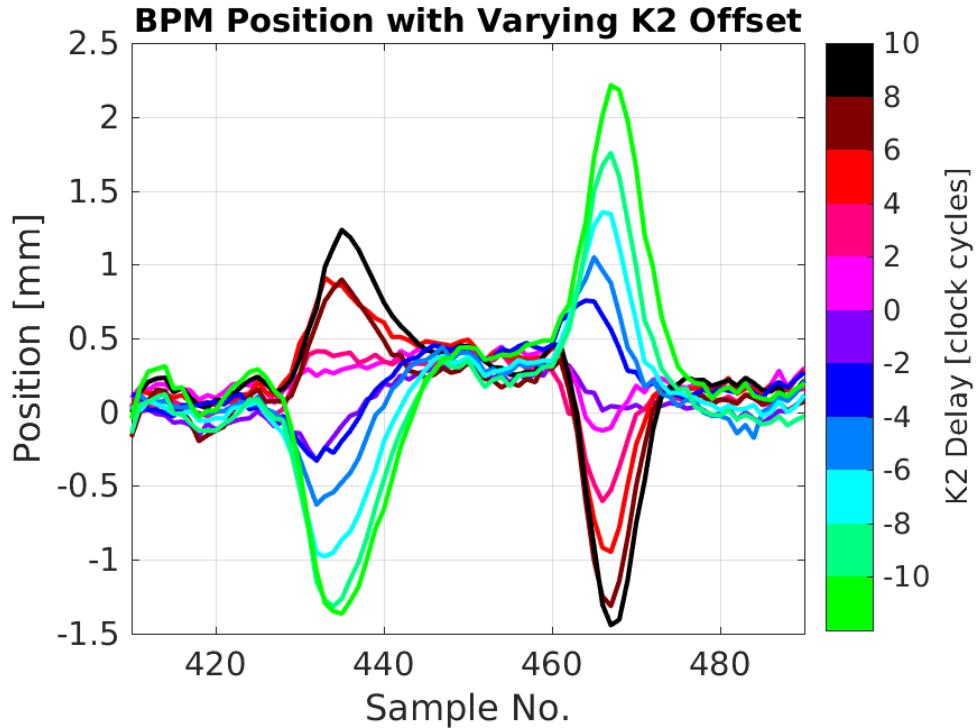


Figure 5.38: Measured BPM offset for different relative kick delays.

result of doing this, using a constant DAC output from the FONT5a board applied across a 168 ns portion of the pulse. The horizontal position in a BPM after the TL2 chicane is plotted for relative K2 delays ranging between -10 (K1 output delayed with respect to K2) to +10 (K2 output delayed with respect to K1) clock cycles. Aside from the asymmetry between the size of the peaks at the start and end of the pulse the result is as expected from the example previously discussed. Note the non-zero position offset in the central part of the pulse. Based on the orbit closure results in Section 5.4.4 this is predominantly due to optics differences leading to a non-closed orbit, rather than the small difference in amplifier output voltage to each kicker.

Figure 5.39 then shows the peak beam offset in the BPM versus the relative K2 delay using the falling edge of the pulse. The peak beam offset is defined as the difference between the maximum and minimum beam position after the chicane between sample 458 and 477 (as seen in Figure 5.38). As the K2 delay approaches the optimal value the difference in beam position in this range converges to the 0.3 mm offset in the flat central part of the kicked pulse. The point of intersection between the two linear fits shown (one for the points with a positive peak position and the other for points with a negative peak) gives the optimal relative K2 offset to be 0.1 ± 0.5 clock cycles. Repeating the procedure for the peaks at the rising edge of the pulse gives a result of 1.9 ± 2.0 clock cycles, and the two results combine to give an optimal value of 0.5 ± 0.6 clock cycles.

Relative K2 delays of both 0 and 1 clock cycles have been used during PFF operation, with no measurable difference in the PFF results between the two to date although this will have to be verified with further orbit closure tests. Adding the absolute delay of 7 clock cycles derived in Section 5.5.2, the final delays to apply in the FONT5a board are:

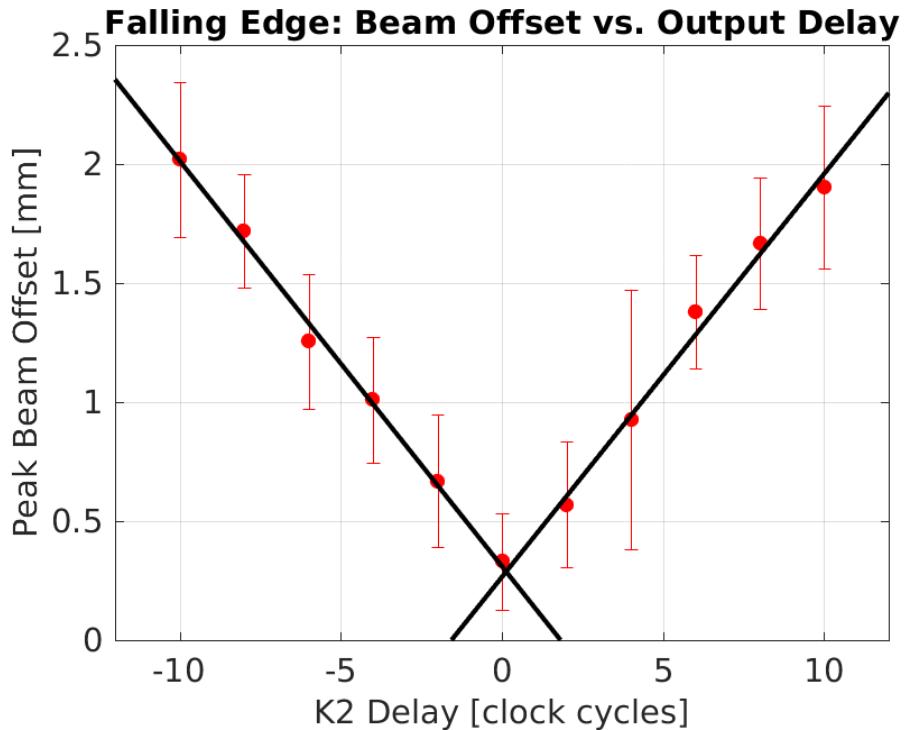


Figure 5.39: Fitted peak BPM offset vs. relative kick delay.

- K1 delay: 7 clock cycles.
- K2 delay: 7 or 8 clock cycles.

5.6 Early Phase Feedforward Attempts and Simulations

This is the introductory text.

5.6.1 Gain Scans

5.6.2 Effect of Limited Correction Range

Chapter 6

Best Achieved Phase Stabilisation

Over the course of 2014 and 2015 much experience has been gained with the PFF system and vast improvements have been made to the system setup, hardware performance and beam conditions as discussed in previous chapters. Many datasets have been taken, in particular in December 2014, July 2015 and November 2015, and in all cases a reduction in downstream phase jitter has been achieved. The purpose of this chapter is to present and discuss the downstream phase stability that was accomplished on the 20th November 2015, the day during which the best overall conditions to date for the PFF correction were obtained.

6.1 Lowest Achieved Phase Jitter

The results presented in this section show the best downstream phase jitter currently achieved at CTF3 with the PFF correction. Naturally, this dataset was taken during the best beam conditions currently achieved at CTF3 in terms of phase propagation, taken just after a series of R56 and beam energy optimisations using the same methods discussed in Chapter 4. In particular, the first attempt to smooth the upstream phase along the pulse by adjusting the waveform of the first klystron in the CTF3 injector (MKS02) as described in Section 4.5.3 yielded the highest upstream-downstream phase correlation achieved to date in normal conditions (higher correlations can be achieved by adding an additional jitter source upstream, as seen in Section ??). November 2015 also represents the first time the PFF system was operated with the improved phase monitor resolution after switching to mechanical phase shifters (Section 3.8) and with the latest version of the amplifier which doubled the correction range compared to earlier tests (Section 5.2).

The dataset was taken on the 20th November 2015 as one of a sequence of short measurements fine-tuning the gain around the optimal value. Results from the other datasets in this sequence are discussed in the following section to demonstrate the phase stability achieved on longer time scales. The individual dataset shown here comprises 150 pulses taken in interleaved mode, with the correction applied to the 75 odd indexed pulses and no correction applied to the remaining 75 even indexed pulses. The used gain in FONT5a units was 800, corresponding to an actual applied correction of 1.13 times the upstream phase

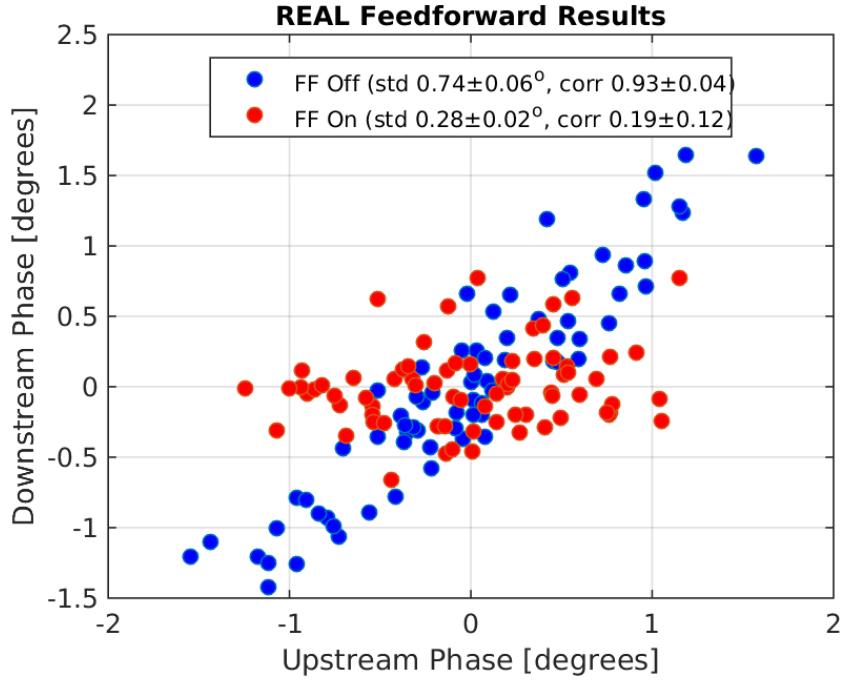


Figure 6.1: Mean phase.

using the conversion factor calculated in Section ?? [TODO: Scale based on new conversion factor].

6.1.1 Mean Phase Jitter

Initially considering the mean pulse phase, the correlation with the PFF correction off in this dataset, as shown by the blue distribution of points in Figure 6.1, is 0.93 ± 0.04 . This gives a theoretical limit of a factor 2.7 ± 0.4 reduction in the downstream jitter using Equation 4.5. The achieved uncorrected downstream mean phase jitter of $0.74 \pm 0.06^\circ$ and downstream-upstream jitter ratio of 1.1 ± 0.1 are also the lowest achieved at CTF3 to date. With this initial jitter and the theoretical reduction factor of 2.7 ± 0.4 the lowest corrected downstream jitter that could be achieved is then $0.27 \pm 0.05^\circ$. The aforementioned correlation and jitter ratio combine to give an optimal gain of 1.0 ± 0.1 (Equation 4.4). The actual system gain of 1.1 [TODO: update based on new conversion factor] and the optimal gain therefore agree within the error bars, thus the theoretical corrected downstream jitter of $0.27 \pm 0.05^\circ$ should also be possible with the actual system.

The red distribution of points in Figure 6.1 then shows the effect of the PFF correction on the phase distribution. The downstream phase jitter is reduced from $0.74 \pm 0.06^\circ$ to $0.28 \pm 0.02^\circ$, a reduction of a factor 2.6 ± 0.3 . Within the error this does indeed agree perfectly with the theoretical limit derived previously given the beam conditions in this dataset, as expected. The correction acts to remove all correlation between the upstream and downstream phase, rotating the distribution as seen in the plot. The correlation is reduced from 0.93 ± 0.04 to 0.19 ± 0.12 .

Correction Status	Upstream Jitter	Downstream Jitter	Correlation
FF Off	$0.69 \pm 0.06^\circ$	$0.74 \pm 0.06^\circ$	0.93 ± 0.04
FF On	$0.57 \pm 0.05^\circ$	$0.28 \pm 0.02^\circ$	0.19 ± 0.12
FF Simulated	$0.69 \pm 0.06^\circ$	$0.27 \pm 0.02^\circ$	0.06 ± 0.12

Table 6.1: Best PFF results.

In terms of the achieved downstream phase jitter it should be noted, however, that the measured upstream jitter of $0.57 \pm 0.05^\circ$ across the pulses with the PFF correction on in this dataset is lower than the $0.69 \pm 0.06^\circ$ measured with the PFF system off (Table 6.1). This is assumed to be a statistical fluctuation rather than being a systematic difference between the odd and even pulses at CTF3 or an effect of the correction (which can only influence the downstream phase) [TODO: calculate how likely this fluctuation is/check there is no systematic difference between odd/even]. Assuming the PFF on upstream jitter propagated downstream with the same ratio as the PFF off data, the true ‘natural’ downstream jitter without the correction applied would have been $0.61 \pm 0.09^\circ$ and the true factor reduction in the corrected jitter achieved with the PFF system would be decreased to 2.2 ± 0.4 . Assuming the upstream-downstream phase correlation was also not affected by this statistical fluctuation (so that the theoretical jitter reduction of a factor 2.7 ± 0.4 still holds), a corrected jitter of $0.23 \pm 0.05^\circ$ would have been theoretically possible for the PFF on pulses in this dataset.

[TODO: Distribution of points at around 0.5 degrees downstream?]

With interleaved data it is also possible to simulate the expected effect of the correction empirically, as an additional point of comparison between the achieved and expected results plus to verify that the complete behaviour of the system is understood. The distribution of simulated corrected phases is shown in green on Figure 6.2. It is derived by taking the initial distribution with the PFF system off (blue points) and subtracting the upstream phase, multiplied by a gain factor, from the downstream phase. This exactly mimics what the feedforward system would have done if it had been applied to the even pulses in this dataset, and can be directly compared to the odd pulses taken at the same time with the actual correction applied. In this example the simulation shown is the ideal case, considering a correction with infinite range and bandwidth applied with the optimal gain. As expected the simulated corrected downstream jitter of 0.27 ± 0.02 agrees perfectly with the theoretical prediction of $0.27 \pm 0.05^\circ$ previously derived. The achieved jitter of 0.28 ± 0.02 matches both the theoretical and simulated jitter predictions within the error, giving confidence that the overall PFF setup in this dataset (after all the commissioning steps discussed in Chapter 5) was close to optimal. There is perhaps some room for improvement due to the difference between the upstream jitter in the PFF on and off data, as mentioned previously, and this will be elaborated on in Section 6.2 below. Nevertheless, this result clearly demonstrates stability on the mean phase approaching the CLIC target of 0.2 degrees at 12 GHz and demonstrates that achieving this stability with a PFF system is feasible.

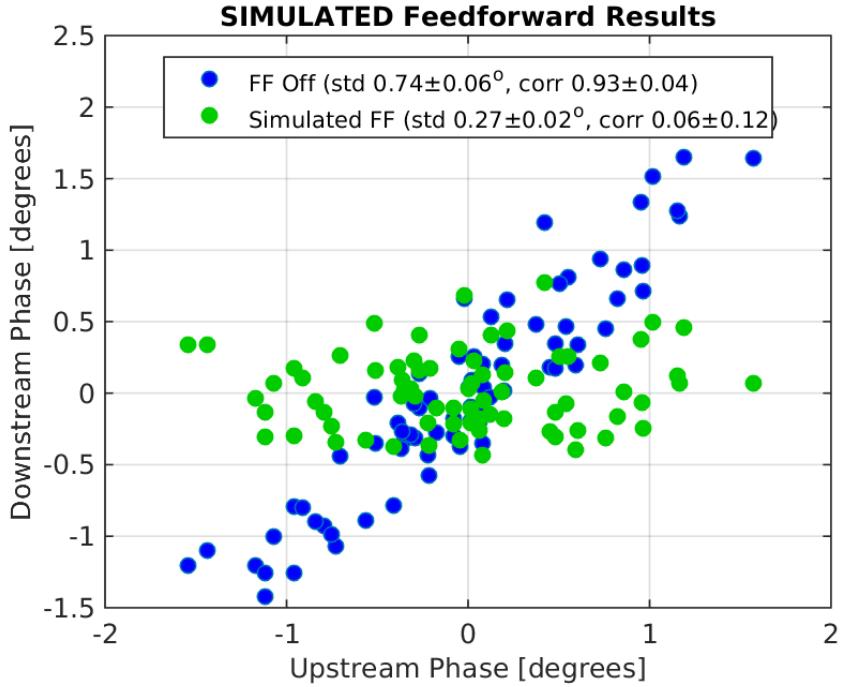


Figure 6.2: Simulated PFF.

6.1.2 Correction of Pulse Shape

Moving on to the stabilisation of the phase along the pulse, Figure 6.3 shows the mean phase along the pulse upstream, downstream with the PFF system off and downstream with the PFF system on. The vertical black lines mark the sample range that was used to calculate the mean phase results presented previously. This range is chosen to cover the maximal proportion of the pulse within which the the correction is not being saturated as a result of the phase sag (plus jitter) exceeding the $\pm 6^\circ$ correction range. It covers a total of 81 samples at 5.2 ns per sample, giving a total time span of 422 ns. The demonstration of $0.28 \pm 0.01^\circ$ mean phase stability is therefore already on a much longer pulse than is needed for CLIC, where the combined pulse length is only 240 ns. [TODO: Any significant reduction in measured jitter by using 240ns window?]

With the optimised phase propagation in place the overall shape of the upstream and (uncorrected) downstream phase, in green and blue respectively, along the pulse are very similar, although small uncorrelated variations are still visible. These uncorrelated differences are then visible in the corrected downstream phase (in red), although the overall ability of the PFF system to flatten the CTF phase sag within the sample range is strikingly clear. The original peak-to-peak variation in the mean downstream phase along the pulse of $5.76 \pm 0.14^\circ$ with the correction off is reduced to $0.65 \pm 0.07^\circ$ degrees with the correction applied within the indicated range. Outside the central region of the pulse, where the amplifier is saturated, the PFF system can no longer correct the shape of the phase along the pulse. The only effect is to shift the phase by the maximum possible correction of 5.5° .

Figure 6.4 expresses the effect of the PFF system on the phase along the pulse within the central region in terms of the distribution of ‘flatness’ values for each pulse in the data




Figure 6.3: Mean phase along.

set with PFF system off and on. For each pulse the flatness value is defined as the standard deviation of phase values about the mean across the sample range. In this case the flatness value of each pulse therefore corresponds to the standard deviation of 81 values (the length of the sample range). A pulse with a flatness value of zero would have constant phase across the whole sample range, with no small variations such as those seen in Figure 6.3. The value is also insensitive to the jitter on the overall mean pulse phase seen earlier in Figure 6.1. In Figure 6.4, the initial uncorrected downstream pulse flatness, dominated by the phase sag at CTF3, of $1.68 \pm 0.02^\circ$ is reduced to $0.26 \pm 0.01^\circ$ with the correction applied. On average, the corrected pulses are 6.5 ± 0.3 times ‘flatter’ than the uncorrected pulses.

6.1.3 Phase Jitter Along the Pulse

Finally, Figure 6.5 shows the overall phase jitter at each sample along the pulse upstream and downstream with the PFF system off and on. These jitter values contain components coming from both the jitter on the overall mean pulse phase discussed initially and from the variations along the pulse (the non-zero flatness of each pulse). These jitter values are therefore larger and taking the mean sample jitter within the sample range an initial downstream jitter of $0.72 \pm 0^\circ$ is reduced to $0.36 \pm 0^\circ$ by the correction in this case, a factor 2 reduction. There are also variations of up to a factor 2 in the jitter that was achieved at each sample point, the lowest jitter being $0.27 \pm 0^\circ$ at time 797 ns on the x-axis and the worst $0.52 \pm 0^\circ$ at time 552 ns. The achieved jitter along the pulse within the central sample range also agrees with the simulated result of $0.38 \pm 0^\circ$ using the interleaved pulses without the correction applied, as shown in Figure 6.6.[TODO: Error bars] Again, outside the sample range the actual system can not match the unlimited range simulation as the phase sag along




Figure 6.4: Flatness.

the pulse saturates the correction.

The sample by sample phase jitter along the pulse is the true figure of merit which must be reduced to 0.2° at CLIC. Although the largest component of phase jitter at CTF3 is on the pulse mean, effects such as energy variations along the pulse cause differences in the jitter and upstream-downstream phase correlation at each sample point [TODO: sec ref]. This leads to the variations in the achievable corrected downstream jitter along the pulse seen here, which can only be improved by further fine-tuning of the CTF3 injector stability and optics. [TODO: More details. Maybe correlation plot?]

6.2 Correction on Longer Time Scales

The remainder of this chapter focuses on remaining operational issues for the PFF system largely resulting from drifts in the CTF3 beam conditions. This section therefore discusses the status of the correction across longer time scales, presenting both the level of corrected phase jitter that can currently be achieved routinely and to highlight areas where improvements are still needed both in the PFF setup itself and the beam conditions. Being able to regularly demonstrate and maintain corrected downstream phase jitters at the level achieved in the best dataset shown previously (below 0.3°) on the mean phase, is one of the key goals for the PFF prototype in 2016. To be concise this section focuses on the mean phase jitter, though exactly the same arguments can be applied to the correction of the jitter along the pulse and the pulse shape.

The data used is from around 15:25 to 18:05 on the 20th November 2015, the same day as the record result previously shown which was taken during this period at 15:38. The

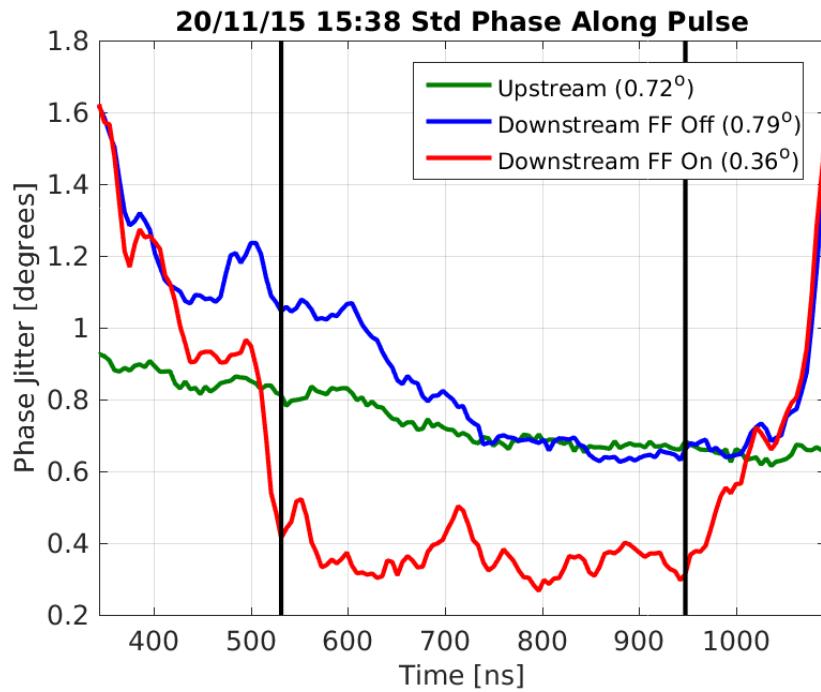


Figure 6.5: Std phase along.

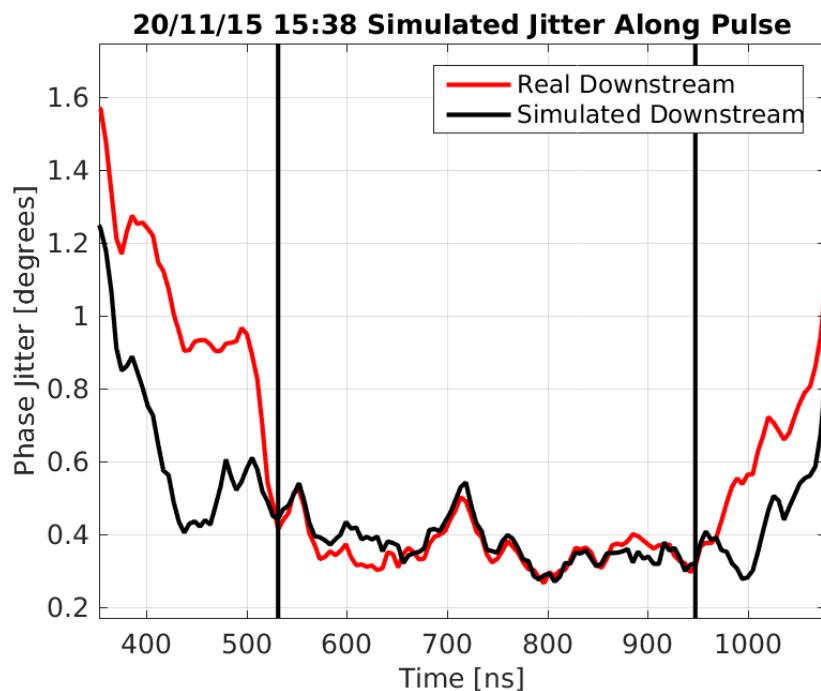


Figure 6.6: Std phase along.

PFF system was not operated continuously throughout this two and a half hour window but 15 individual datasets of a few hundred pulses each were taken and these results have been combined to create a large sample of 3083 interleaved pulses (1541 with the correction on and 1542 with the correction off). The raw history of the mean phases upstream and downstream with the correction on and off in the combined data are shown in Figure 6.7. The time span of each individual dataset is marked by vertical black lines and the times displayed on the plot represent the start time of each dataset. [TODO: Pulse no. from 1-3083 rather than 1-1500 and offset odd/even by one]. Note that the large jump in the downstream phase between the 16:00 and 16:04 datasets was caused by changes made to magnetic correctors in the TL2 chicane in order to re-optimize the beam orbit and transmission to the downstream phase monitors at this time. In Figure 6.17 the mean phase is subtracted (separately for the upstream, downstream PFF off and downstream PFF on phase) from each dataset to remove this effect, making a comparison between datasets easier. It is important to emphasise that, apart from this jump in the downstream phase, the overall picture is a fair reflection of the (uncorrected) phase stability at CTF3 in optimal conditions.

6.2.1 Upstream Phase Drifts

Over the course of the data taking period the mean upstream phase, in green, varies by ten degrees peak-to-peak or $1.75 \pm 0.02^\circ$ in terms of jitter (Figure 6.7). [TODO: Source of drift, comment on feedbacks]. Small drifts of up to a few degrees in the upstream phase are not an issue for the performance of the PFF correction providing the correlation between the upstream and downstream phase is not degraded. In some cases upstream phase drifts may lead to a loss in correlation, this could be the case if the source of the drift is a variation in beam energy due to the issues discussed in Chapter 4, for example. The variation of the correlation between datasets is discussed later in this section.

Larger changes in the upstream phase such as the ten degree fluctuation seen here may also impact the PFF performance purely via the limited correction range of $\pm 6^\circ$ combined with the phase sag along the CTF pulse. Indeed the PFF prototype's main purpose is not to remove any large, slow phase drifts but rather the faster pulse-to-pulse jitter and high frequency variations along the pulse. The phase offset applied by the PFF correction at each sample along the downstream phase, $\Delta\phi_d(t)$, is given by:

$$\Delta\phi_d(t) = \begin{cases} -6^\circ, & \text{if } g\phi_u(t) \geq +6^\circ. \\ +6^\circ, & \text{if } g\phi_u(t) \leq -6^\circ. \\ -g\phi_u(t), & \text{otherwise.} \end{cases} \quad (6.1)$$

Where $\phi_u(t)$ is the upstream phase at each sample point and g is the gain factor used. As the optimal gain (Section??) for the correction is typically larger than one due to the slight amplification in the downstream phase jitter with respect to the upstream jitter the range of the PFF system in terms of the upstream phase is less than $\pm 6^\circ$ (for example $\pm 5.3^\circ$ for the 15:38 jitter record dataset with a gain of 1.13). Any point along the upstream phase with

6.2 Correction on Longer Time Scales

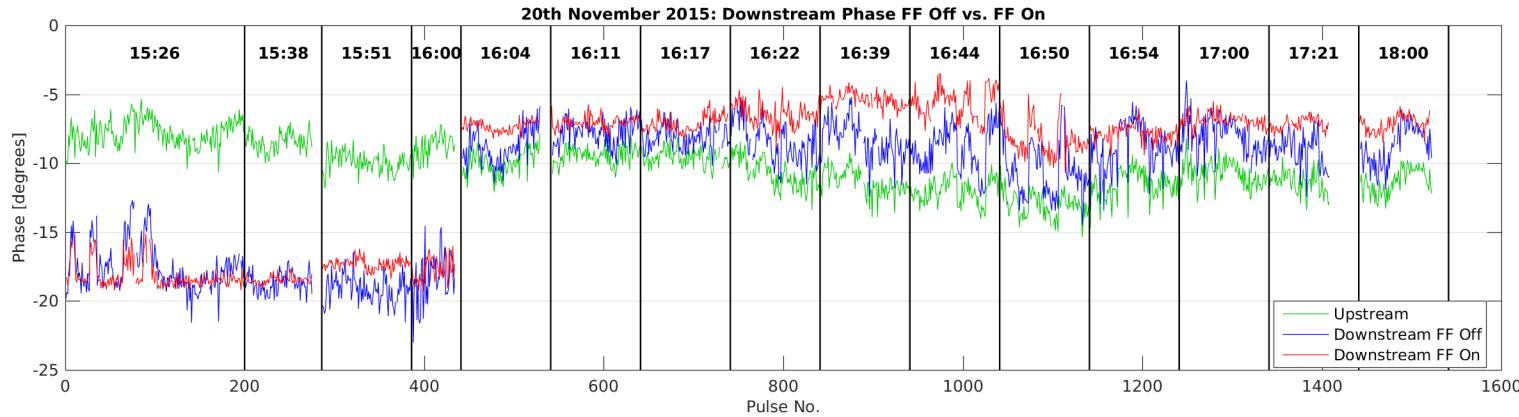


Figure 6.7: History of mean phase across datasets.

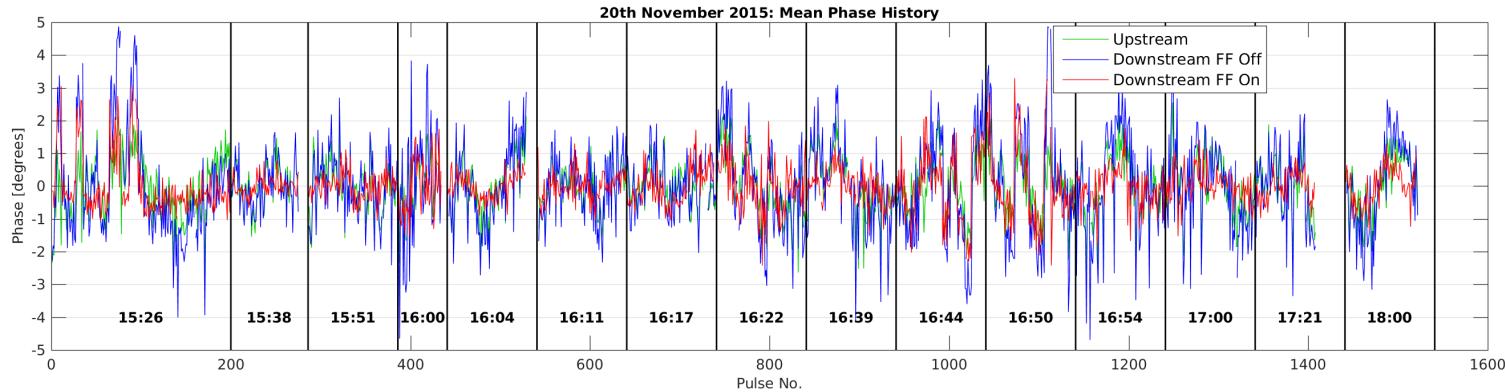


Figure 6.8: History of mean phase across datasets, with mean subtraction.

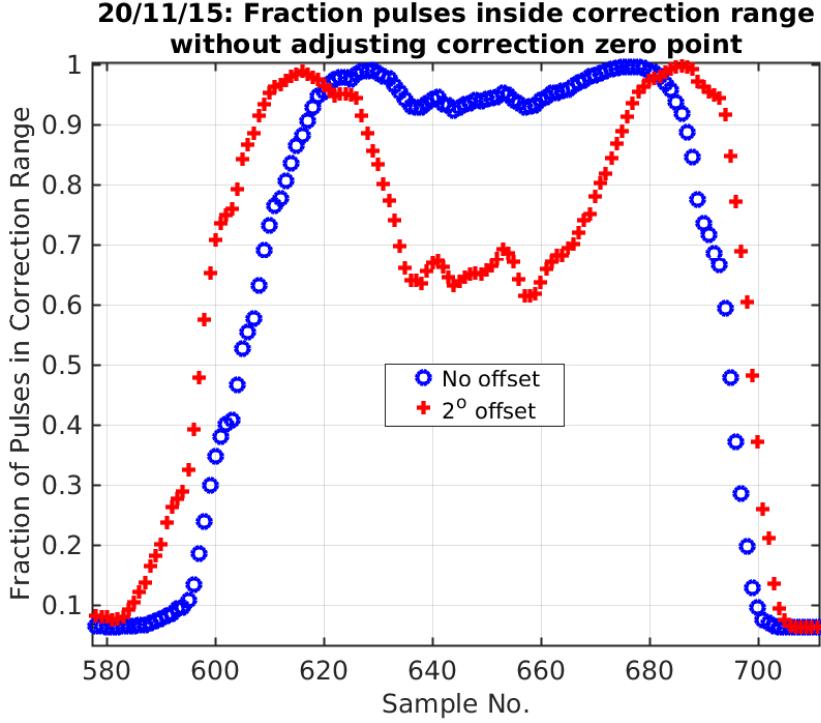


Figure 6.9: Fraction of pulses outside the correction range along the pulse. [TODO: Add line using real dataset offsets]

$|g\phi_u(t)| > 6^\circ$ receives the maximum 6° phase shift downstream but can not be corrected to zero, with this remaining residual degrading the corrected phase jitter that can be achieved. Samples with $|g\phi_u(t)| > 5^\circ$ will also receive a slightly non-optimal correction due to the effects of the amplifier entering saturation, shown in Section ??, although this effect is assumed to be small and is not yet considered in the discussion here [TODO: Calculate how significant].

Figure 6.9 shows the fraction of pulses for which the optimal correction is within the correction range in the combined dataset. During the setup of the PFF system it is necessary to choose the zero point for the correction, i.e. the incoming upstream phase at which the correction output to the kickers is 0 V. This is done in the PFF firmware on the FONT5a board by varying a channel offset applied to the raw digitiser data from the ADC to which the upstream phase monitor mixer signal is connected [TODO: ref sec]. In terms of equation 6.1 this is equivalent to adding a constant offset to ϕ_u across the full pulse length. The optimal channel offset zeroes the mean phase (ADC output) taken across the part of the pulse where the best correction is desired (usually the flatter central part of the pulse at CTF3). In this case the effects of limited correction range are minimised, as the full $\pm 6^\circ$ range can be used to remove variations about the mean phase, rather than also having to remove a static phase offset in the overall mean. In this case the ideal correction across a 310 ns portion of the pulse is within the $\pm 6^\circ$ range 96% of the time.

However, as to date this offset has been set up manually small deviations from the ideal case are possible. Figure 6.9 also shows the fraction of pulses within the correction range if there is a static two degree offset in the upstream phase. In this case as many as 39% of pulses are outside the correction range within the normally correctable central region of the

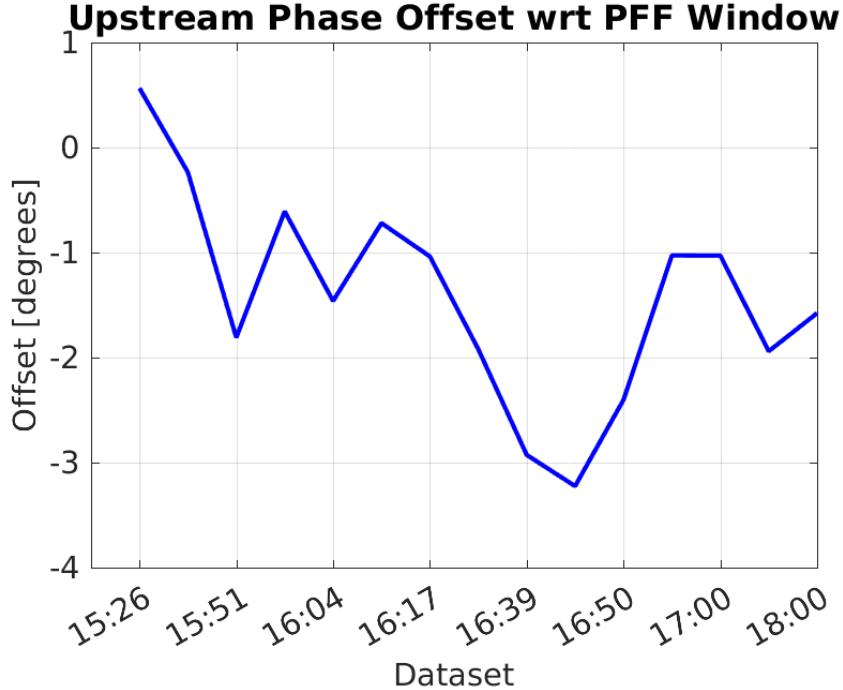


Figure 6.10: Offset between downstream phase with FF off and FF on.

pulse. To mitigate these effects and to get the largest reduction in jitter possible within each individual dataset the centring of the upstream phase in the correction range on the FONT5a board is normally adjusted between datasets. As a consequence of this differences in the upstream phase between datasets are not removed in the corrected downstream phase, as the zero point for the PFF correction is effectively moving with the phase drifts during the data taking period. These remaining slow drifts could be removed at CTF3 using a secondary “slow phase feedback”, also utilising the TL2 chicane, which is the focus of Section ??.

The accuracy to which the channel offset for the upstream phase has been set can be inferred by comparing the mean downstream phase in each dataset with the correction on (red) and off (blue) in Figure 6.7. In the ideal case the mean phase should be identical with the PFF system on and off, so that the full correction range is being used to correct jitter about the mean as mentioned previously. Although this is the case for some datasets, such as the 15:38 dataset, a clear offset between the two is often present, most visible in the datasets between 16:39 and 16:50 in which the corrected phase is clearly shifted several degrees with respect to the uncorrected phase. The offset in each dataset is plotted in Figure 6.10 [TODO: Change to table?]. In the region between 16:39 and 16:50 the offset falls below -3° . The mean offset across the combined dataset is -1.4° [TODO: Calc error, weight by dataset lengths]. In the following sections it will be shown that the effect of the non-optimal set point for the offset is small overall, although there is a noticeable degradation in the jitter that can be achieved in the datasets with the largest offsets. In any case, implementing an automatic procedure to set the zero point for the correction optimally in the FONT5a DAQ would be a useful improvement to the PFF setup procedure and this will be pursued for future PFF attempts in 2016.

6.2.2 Gain Stability

Another PFF parameter that has been set up largely empirically to date is the correction gain. Historically, the gain set point for the PFF prototype has been determined by a combination of viewing the results of gain scans (Section ??) and by observing the flatness of the corrected downstream phase in online displays of the phase monitor signals. If the applied gain is too large this can be quickly seen in the online monitors as the PFF system will act to invert the original phase sag along the pulse, for example. In this way it is relatively simple to find approximately the correct gain set point and further fine-tuning is done by varying the gain in small steps between datasets. In later PFF attempts this approach was complimented by implementing an online display of the optimal gain given the current measured upstream and downstream phase jitters and correlation (Section ??), although this only gives a representative value when the PFF system is turned off (otherwise the calculated gain value is based on the corrected jitters and correlation). However, in this section it will be shown that due to drifts in the beam conditions at CTF3 there are large variations in the optimal gain between datasets, and these variations are rarely accurately followed in the PFF setup when using this empirical approach. An automatic gain optimisation procedure is therefore another area of improvement for future PFF attempts. Particularly if the gain was automatically updated in real time during long datasets a significant reduction in jitter could be achieved, as will be seen in the remainder of this chapter. Of course, in the ideal case the stability of beam conditions at CTF3 would be improved so that the variations in optimal gain over the course of a few hours are much smaller than those shown here.

The optimal gain depends on the downstream-upstream phase jitter ratio and the downstream-upstream phase correlation (Section ??). In Figures 6.7 and 6.8 large differences in the phase stability in each dataset are clearly visible, comparing for example the large phase jumps in the 15:26 and 16:50 datasets to the comparatively calm periods at 15:38 and 16:17. This is summarised in Figure 6.11, which shows the upstream and downstream (with PFF off) phase jitter across the 5–10 minute time period of each dataset. Over the course of the data taking period the mean upstream and downstream phase jitter both vary by around a factor two — the upstream jitter between $0.6 \pm 0^\circ$ in the 16:17 dataset and $1.1 \pm 0^\circ$ at 16:22, and the downstream jitter between $0.7 \pm 0^\circ$ at 15:38 and $2.2 \pm 0^\circ$ at 16:50. [TODO: error] Given the same correlation, a factor two increase in the uncorrected downstream jitter also doubles the corrected downstream phase jitter that can be achieved with the PFF system. [TODO: explain why double jitter means double corrected jitter/refer back to prev sec].

Also of key importance for the PFF correction is that not only are there large variations in jitter between datasets but additionally in the downstream-upstream jitter ratio (dashed line in Figure 6.11). In fact, the only dataset in which the upstream and downstream jitter are comparable is the record 15:38 dataset (with a ratio of 1.1). In all other datasets the downstream jitter is more than 1.3 times larger than the upstream jitter, reaching a maximum amplification of 2.2 in the 16:50 dataset. The mean ratio across the 15 datasets is 1.5 with a standard deviation of 0.3. [TODO: errors]

As well as the jitter ratio, the upstream-downstream phase correlation also varies between datasets, as shown in Figure 6.12. The worst correlation is 0.80 ± 0.04 in the 15:26 dataset and

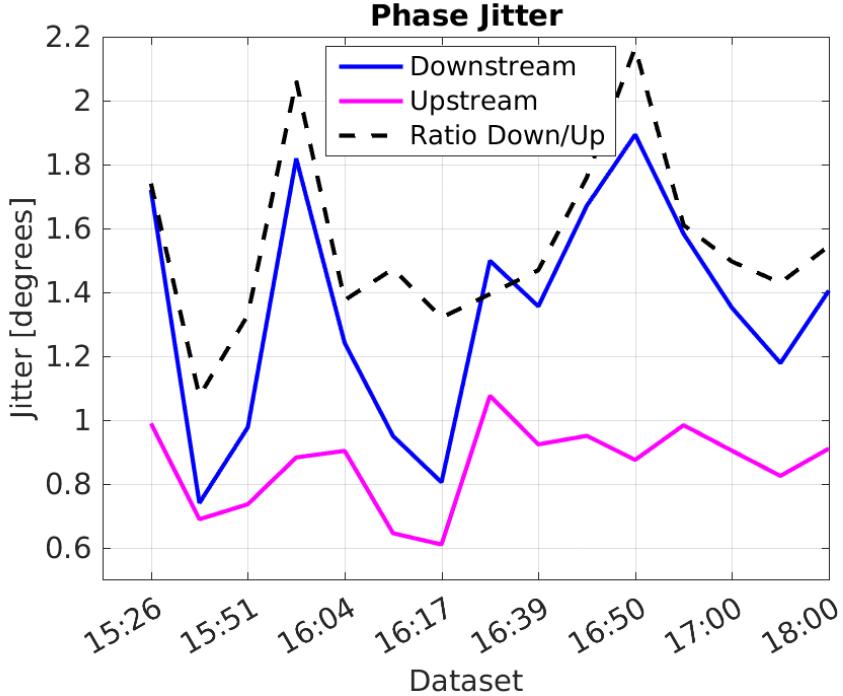


Figure 6.11: Upstream and downstream phase jitter in each data set.

the best 0.96 ± 0.03 in the 16:54 dataset. Although this has a much smaller 20% effect on the optimal gain than the factor 2 variation in jitter ratio, it has a large effect on the theoretical jitter improvement that can be achieved with the PFF system due to the dependence on the correlation squared in Equation 4.5. With 80% phase correlation only a theoretical factor 1.7 reduction in the downstream phase jitter can be achieved, whereas with 96% correlation this is increased to a factor 3.6.

There is no observed dependence of the phase jitter ratio on the phase correlation, as shown in Figure 6.13, so the effects of varying correlation and jitter ratio on the optimal gain are independent. [TODO:why?]. They combine to give the optimal gain plotted in Figure 6.14 (red line). As it is dominated by the differences in jitter ratio, the gain also varies by close to a factor two, varying from 1.00 in the 15:38 dataset to 1.95 in the 16:00 dataset. The real gain factor that was actually used in the dataset is also plotted, in blue. Although in places the empirically derived gain that was used follows the trend of the optimal gain, the changes are much smaller and it is clear that the real gain was systematically lower than the optimal gain. The smallest gain actually used was 1.06 (at 15:51) and the largest 1.34 (15:26 and 16:00), with an overall mean across the data taking period of 1.16 compared to the optimal value of 1.41. The impact of the real system generally under-correcting the downstream phase as a result of lower than optimal gain is discussed in the next sections.

6.2.3 Results

It has been shown that the frequent drifts in both phase and downstream-upstream phase jitter ratio have not been optimally taken in to account in the PFF setup in terms of the

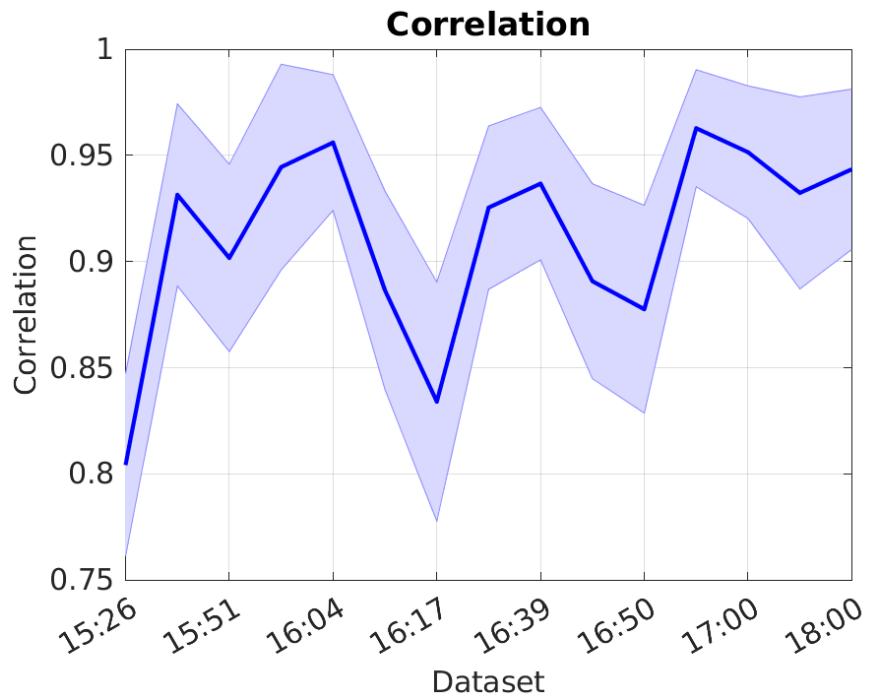


Figure 6.12: Upstream-downstream mean phase correlation in each dataset with PFF off.

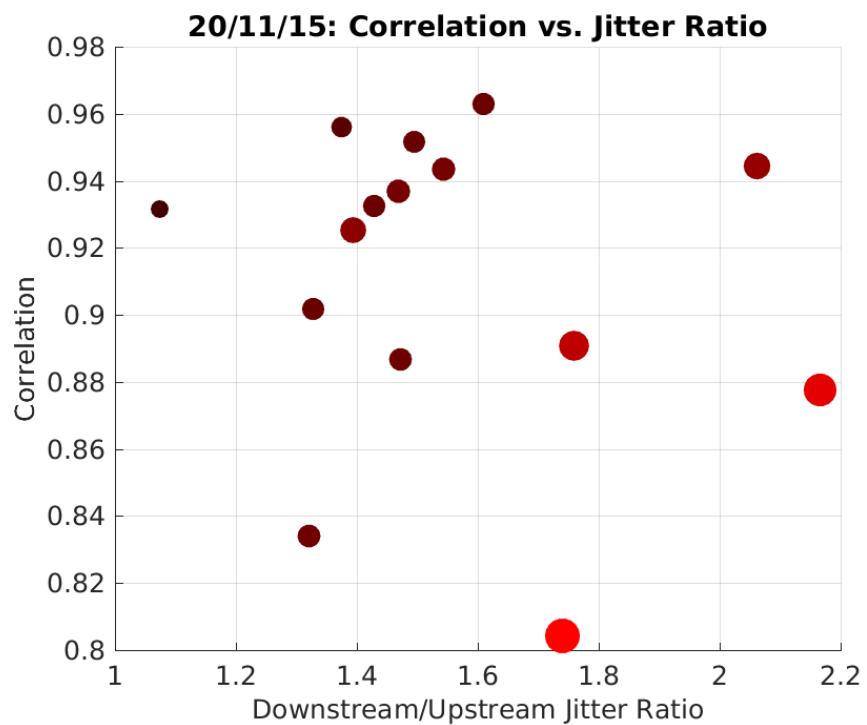


Figure 6.13: Correlation vs. phase jitter ratio.




Figure 6.14: Gain used in each dataset compared to the optimal gain.

used offset and gain. Nevertheless, even with a sub-optimal setup a large reduction in the downstream phase jitter can be achieved in all datasets. In the remainder of this section it will be shown that considering these constraints the PFF system is achieving close to peak performance, as well as highlighting the benefit that more accurate gain and offset control would have.

Firstly referring back to Figure 6.13, the size (area) and colour of the markers in the plot depend on the corrected downstream jitter that could be achieved in that dataset using the optimal gain. Small, black markers correspond to the lowest theoretical jitter and large, red markers to the largest theoretical jitter. This is to emphasise again that it is a compromise between high correlation and low initial downstream jitter (and by extension low downstream-upstream jitter ratio) that gives the best conditions for the PFF correction. There are seven datasets with an achieved correlation of above the 93% seen in the 15:38 record result, for example, but they yield worse theoretical corrections as 15:38 remains the only dataset in which a high correlation and low upstream-downstream jitter amplification have been achieved at the same time.

Figure 6.15 and Table 6.2 then show the simulated corrected downstream jitter chronologically for each dataset with five different simulation setups:

- **Unlimited:** With unlimited correction range and the optimal gain (theoretical limit).
- **6deg Range:** With $\pm 6^\circ$ correction range and the optimal gain.
- **Real Gain:** With $\pm 6^\circ$ correction range and the real gain used by the actual PFF

system.

- **Real Offset:** With $\pm 6^\circ$ correction range and the real offset in the actual PFF setup.
- **All effects:** With $\pm 6^\circ$ correction range, the real gain, and the real offset in the actual PFF setup.

[TODO: Simulations should use 5.5 degree range instead AND I used FONT gain conversion factor of 711 but correct value is 624]

By comparing the results of these five simulations it is possible to identify which PFF parameters are most critical for the correction performance. Later, by comparing the most restricted simulation, including the real offset and gain, to the phase jitter actually achieved it can be determined whether the PFF system is behaving as expected or whether there are remaining effects that need to be understood.

With the ideal PFF setup the $\pm 6^\circ$ range set by the amplifier power is sufficient to be able to optimally correct almost all the natural phase jitter, thus the difference between the unlimited and 6deg range simulation is small. The only visible effect is in datasets with the largest incoming phase jitter, with a maximal 0.05° degradation in the achievable phase jitter in the 16:00 dataset with an incoming uncorrected phase jitter of 1.8 degrees, for example. The correction range is therefore not a limiting factor for the PFF performance in normal conditions, although it may become significant if the PFF system were operated on longer time scales without updating the offset or when trying to demonstrate a factor 10 reduction in jitter (Section ??).

Depending on the dataset, the effects of using non-optimal gain and non-optimal offset are much larger. In the 15:26, 15:38, 15:51 and 16:17 datasets where both the gain and offset are close to optimal all five simulations give close to the same result, as expected. For most the other datasets the largest effect on the achievable corrected jitter comes from the non-optimal gain, with a difference of up to 0.24° in the simulated phase jitter (16:00). However, in the period between 16:22 and 16:50 where the offset in the PFF setup was largest its influence can be similar to that of the non-optimal gain, or in some cases larger, with a maximal degradation in the achievable downstream jitter of 0.16° (16:50) coming from the offset alone. With the effects of limited correction range, non-optimal offset and non-optimal gain combined the achieved corrected jitter is expected to be up to 0.29° worse than the theoretical limit (16:50), although in most datasets the effect is much smaller than this, with no significant difference between the unlimited and all effects simulations in the previously mentioned 15:26, 15:38, 15:51 and 16:17 datasets, for example.

Only the 15:38 dataset has a theoretical (and in all simulation scenarios) corrected downstream jitter of below 0.3° but in 10 out of 15 datasets below 0.5° jitter could have been achieved with an optimal PFF setup (or in 6 out of 15 with the actual setup). Further improvements not only in the peak phase propagation conditions achieved so far but also clearly on the stability of the phase propagation are therefore needed to demonstrate CLIC level phase stability both on short and long time scales at CTF3.

The achieved downstream jitter with the actual PFF system are presented in Figure

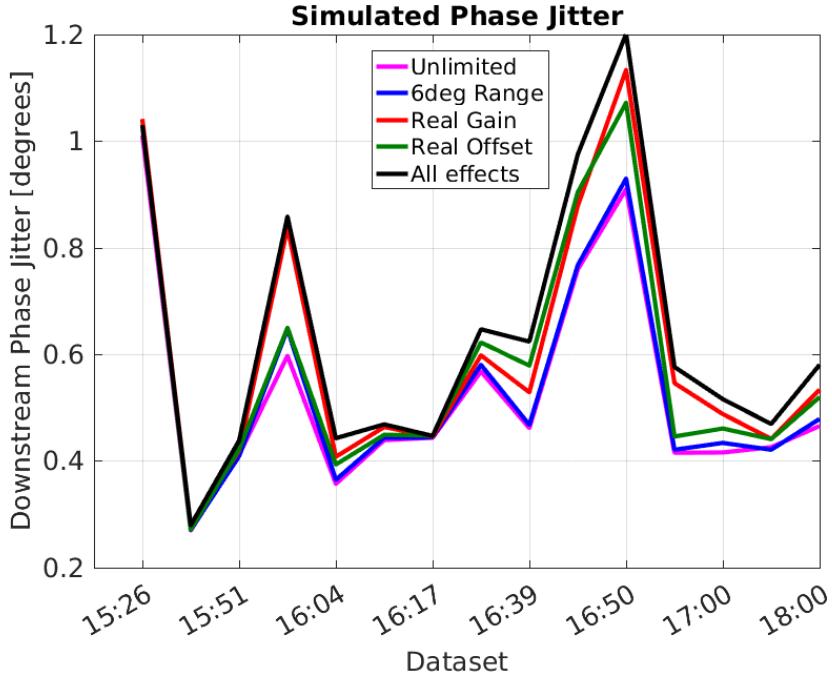


Figure 6.15: Theoretical corrected downstream jitter with optimal and used gain.

6.16 and Table 6.3, along with the uncorrected downstream and upstream jitter and the most realistic "all effects" simulation of the expected performance. Overall the agreement between the downstream jitter achieved with the actual PFF system and the simulation is very good. This gives confidence that the PFF system is behaving as expected and all the effects limiting the current performance are understood and in principle can be improved to yield lower jitter in future PFF attempts. However, there is a region between 16:17 and 16:44 where differences between the simulation and actual system can be seen. In particular, the $0.58 \pm 0.04^\circ$ and $0.82 \pm 0.06^\circ$ downstream jitter in the 16:17 and 16:22 datasets, respectively, are noticeably worse than the simulated results of 0.45 ± 0.03 and 0.65 ± 0.05 . The source of this is not yet understood and possibly hints at additional areas for improvement in the PFF setup.

Nevertheless, despite highlighting where the PFF setup and beam conditions are non-optimal during this discussion the overall benefit of the PFF system is clear - the downstream phase jitter is reduced in every dataset, with a maximum reduction factor of 3.2 in the 16:05 dataset (in which the highest correlation of 96% was achieved). Attempts to demonstrate a larger reduction factor, closer to the CLIC specification of an order of magnitude, are presented in Section ??.

Rather than showing each individual dataset, Figures 6.17—6.21 and Table 6.4 present the upstream-downstream phase distribution and overall jitter improvement with all the datasets combined. In order to yield meaningful results the mean upstream and downstream phase (both with FF on and FF off) are subtracted separately for each dataset. The effect of this can be seen by comparing Figure 6.7 (with no mean subtraction) and Figure 6.8. Without this subtraction any calculated jitter and correlation values across the combined dataset would be dominated by changes in the downstream phase resulting from changing the

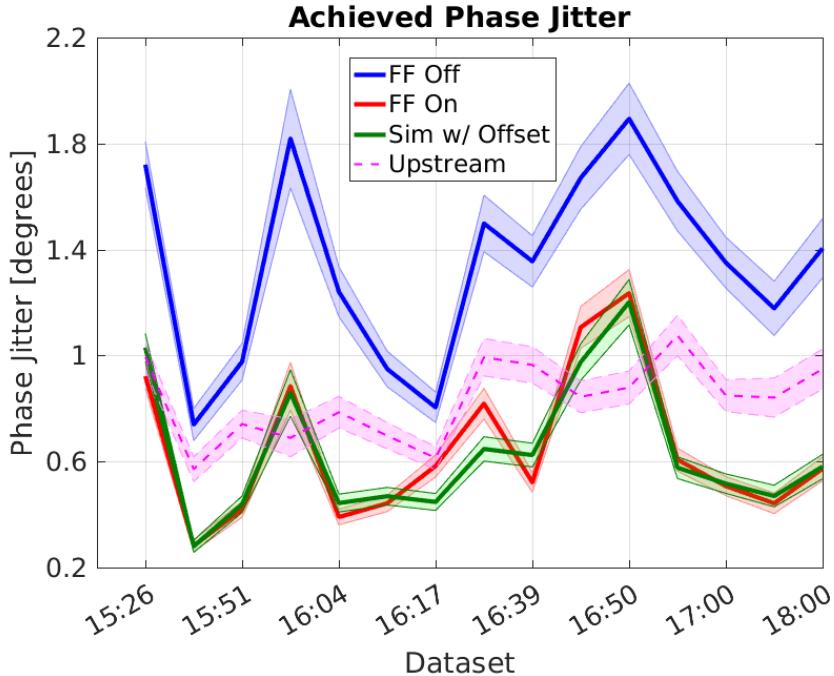


Figure 6.16: Corrected downstream jitter with the actual PFF system.

Time	Unlimited	6deg Range	Real Gain	Real Offset	All Effects
15:26	$1.01 \pm 0.05^\circ$	$1.04 \pm 0.05^\circ$	1.04 ± 0.05	$1.03 \pm 0.05^\circ$	$1.03 \pm 0.05^\circ$
15:38	$0.27 \pm 0.02^\circ$	$0.27 \pm 0.02^\circ$	0.28 ± 0.02	$0.27 \pm 0.02^\circ$	$0.28 \pm 0.02^\circ$
15:51	$0.41 \pm 0.03^\circ$	$0.41 \pm 0.03^\circ$	0.42 ± 0.03	$0.42 \pm 0.03^\circ$	$0.44 \pm 0.03^\circ$
16:00	$0.60 \pm 0.06^\circ$	$0.65 \pm 0.07^\circ$	0.84 ± 0.09	$0.65 \pm 0.07^\circ$	$0.86 \pm 0.09^\circ$
16:04	$0.36 \pm 0.03^\circ$	$0.36 \pm 0.03^\circ$	0.41 ± 0.03	$0.39 \pm 0.03^\circ$	$0.44 \pm 0.03^\circ$
16:11	$0.44 \pm 0.03^\circ$	$0.44 \pm 0.03^\circ$	0.46 ± 0.03	$0.45 \pm 0.03^\circ$	$0.47 \pm 0.03^\circ$
16:17	$0.44 \pm 0.03^\circ$	$0.44 \pm 0.03^\circ$	0.45 ± 0.03	$0.45 \pm 0.03^\circ$	$0.45 \pm 0.03^\circ$
16:22	$0.57 \pm 0.04^\circ$	$0.58 \pm 0.04^\circ$	0.60 ± 0.04	$0.62 \pm 0.04^\circ$	$0.65 \pm 0.05^\circ$
16:39	$0.46 \pm 0.03^\circ$	$0.47 \pm 0.03^\circ$	0.53 ± 0.04	$0.58 \pm 0.04^\circ$	$0.62 \pm 0.04^\circ$
16:44	$0.76 \pm 0.05^\circ$	$0.77 \pm 0.05^\circ$	0.88 ± 0.06	$0.90 \pm 0.06^\circ$	$0.98 \pm 0.07^\circ$
16:50	$0.91 \pm 0.07^\circ$	$0.93 \pm 0.07^\circ$	1.13 ± 0.08	$1.07 \pm 0.08^\circ$	$1.20 \pm 0.09^\circ$
16:54	$0.42 \pm 0.03^\circ$	$0.42 \pm 0.03^\circ$	0.55 ± 0.04	$0.45 \pm 0.03^\circ$	$0.58 \pm 0.04^\circ$
17:00	$0.42 \pm 0.03^\circ$	$0.43 \pm 0.03^\circ$	0.49 ± 0.03	$0.46 \pm 0.03^\circ$	$0.52 \pm 0.04^\circ$
17:21	$0.43 \pm 0.04^\circ$	$0.42 \pm 0.04^\circ$	0.44 ± 0.04	$0.44 \pm 0.04^\circ$	$0.47 \pm 0.04^\circ$
18:00	$0.47 \pm 0.04^\circ$	$0.48 \pm 0.04^\circ$	0.53 ± 0.04	$0.52 \pm 0.04^\circ$	$0.58 \pm 0.05^\circ$

Table 6.2: Simulated feedforward results from 20th November 2015.

Time	Up Jitter	Down Jitter FF Off	Correlation FF Off	Down Jitter FF On	Down Jitter Sim
15:26	$0.99 \pm 0.05^\circ$	$1.72 \pm 0.09^\circ$	0.80 ± 0.04	$0.92 \pm 0.05^\circ$	$1.03 \pm 0.05^\circ$
15:38	$0.57 \pm 0.05^\circ$	$0.74 \pm 0.06^\circ$	0.93 ± 0.04	$0.28 \pm 0.02^\circ$	$0.28 \pm 0.02^\circ$
15:51	$0.74 \pm 0.05^\circ$	$0.98 \pm 0.07^\circ$	0.90 ± 0.04	$0.42 \pm 0.03^\circ$	$0.44 \pm 0.03^\circ$
16:00	$0.69 \pm 0.07^\circ$	$1.82 \pm 0.19^\circ$	0.94 ± 0.05	$0.88 \pm 0.09^\circ$	$0.86 \pm 0.09^\circ$
16:04	$0.79 \pm 0.06^\circ$	$1.24 \pm 0.09^\circ$	0.96 ± 0.03	$0.39 \pm 0.03^\circ$	$0.44 \pm 0.03^\circ$
16:11	$0.70 \pm 0.05^\circ$	$0.95 \pm 0.07^\circ$	0.89 ± 0.05	$0.44 \pm 0.03^\circ$	$0.47 \pm 0.03^\circ$
16:17	$0.61 \pm 0.04^\circ$	$0.80 \pm 0.06^\circ$	0.83 ± 0.06	$0.58 \pm 0.04^\circ$	$0.45 \pm 0.03^\circ$
16:22	$0.99 \pm 0.07^\circ$	$1.50 \pm 0.11^\circ$	0.93 ± 0.04	$0.82 \pm 0.06^\circ$	$0.65 \pm 0.05^\circ$
16:39	$0.96 \pm 0.07^\circ$	$1.36 \pm 0.10^\circ$	0.94 ± 0.04	$0.52 \pm 0.04^\circ$	$0.62 \pm 0.04^\circ$
16:44	$0.84 \pm 0.06^\circ$	$1.67 \pm 0.12^\circ$	0.89 ± 0.05	$1.11 \pm 0.08^\circ$	$0.98 \pm 0.07^\circ$
16:50	$0.88 \pm 0.06^\circ$	$1.89 \pm 0.13^\circ$	0.88 ± 0.05	$1.24 \pm 0.09^\circ$	$1.20 \pm 0.09^\circ$
16:54	$1.08 \pm 0.08^\circ$	$1.58 \pm 0.11^\circ$	0.96 ± 0.03	$0.61 \pm 0.04^\circ$	$0.58 \pm 0.04^\circ$
17:00	$0.85 \pm 0.06^\circ$	$1.35 \pm 0.10^\circ$	0.95 ± 0.03	$0.51 \pm 0.04^\circ$	$0.52 \pm 0.04^\circ$
17:21	$0.84 \pm 0.07^\circ$	$1.18 \pm 0.10^\circ$	0.93 ± 0.05	$0.44 \pm 0.04^\circ$	$0.47 \pm 0.04^\circ$
18:00	$0.95 \pm 0.08^\circ$	$1.40 \pm 0.11^\circ$	0.94 ± 0.04	$0.57 \pm 0.05^\circ$	$0.58 \pm 0.05^\circ$

Table 6.3: Feedforward results from 20th November 2015.

Correction Status	Upstream Jitter	Downstream Jitter	Correlation
FF Off	$0.88 \pm 0.02^\circ$	$1.40 \pm 0.03^\circ$	0.89 ± 0.01
FF On	$0.86 \pm 0.02^\circ$	$0.72 \pm 0.01^\circ$	0.48 ± 0.02
FF Sim Opt Gain	$0.88 \pm 0.02^\circ$	$0.61 \pm 0.01^\circ$	-0.01 ± 0.03
FF Sim Real Gain	$0.88 \pm 0.02^\circ$	$0.68 \pm 0.01^\circ$	0.35 ± 0.02
FF Sim Offset	$0.88 \pm 0.02^\circ$	$0.69 \pm 0.01^\circ$	0.36 ± 0.02
FF Sim 90% Real Gain	$0.88 \pm 0.02^\circ$	$0.72 \pm 0.01^\circ$	0.46 ± 0.02

Table 6.4: Feedforward results using combined data from 20th November 2015. [TODO: this table shows results from old simulations!]

zero point (offset) for the correction between datasets, plus the large step in the downstream phase between the 16:00 and 16:04 datasets due to a beam setup change.

Overall, the actual system is able to reduce an initial downstream jitter of $1.40 \pm 0.03^\circ$ by a factor of two, down to $0.72 \pm 0.01^\circ$ (Figures 6.17, 6.18 and 6.19). Due to the non-optimal setup in some datasets as shown the PFF system does not remove all correlation between the upstream and corrected downstream phase, with the initial correlation of 0.89 ± 0.01 reduced only to 0.48 ± 0.02 . With a completely optimal setup and unlimited correction range all the correlation would be removed and the jitter could have been reduced further to $0.61 \pm 0.01^\circ$ (Figure 6.20). To achieve better than this improved beam conditions are required, including more stable and higher upstream-downstream phase correlation and lower and more stable phase jitters. Considering the constraints of the actual system and non-optimal setup, the achieved downstream jitter and residual correlation are as expected (Figure 6.21).




Figure 6.17: Histogram showing overall distribution of downstream phase with FF off and on.

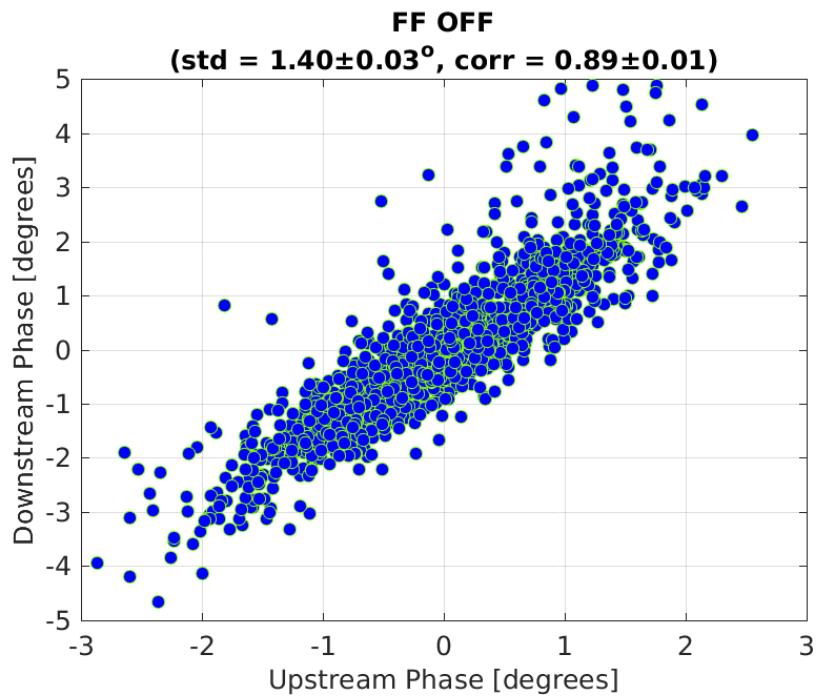


Figure 6.18: Downstream phase vs. upstream phase with FF off.

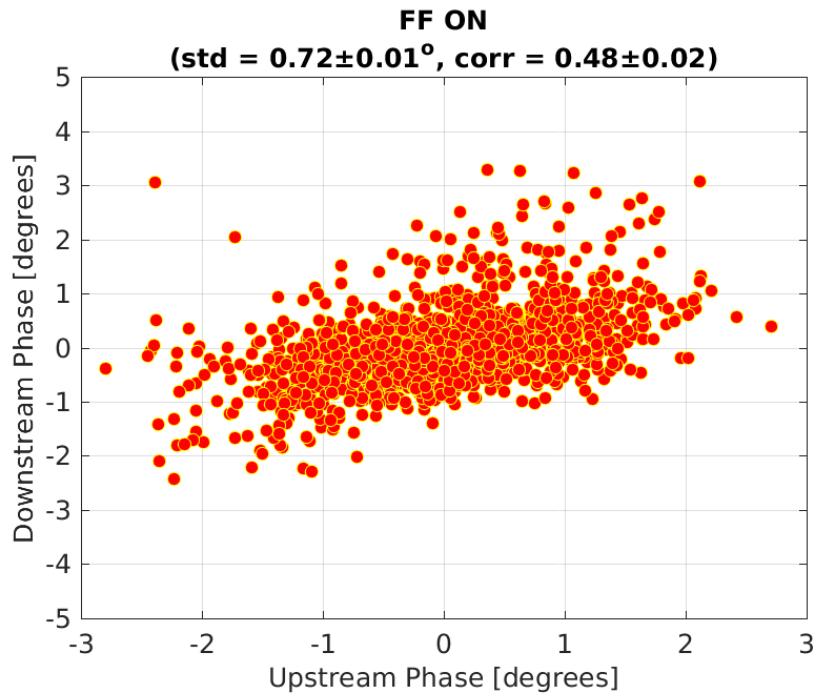


Figure 6.19: Downstream phase vs. upstream phase with FF on.

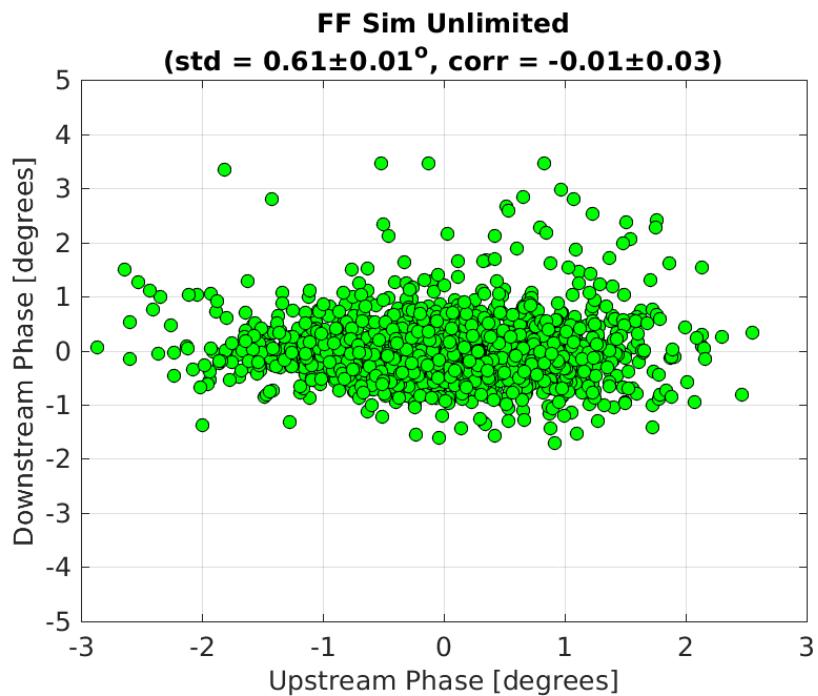


Figure 6.20: Downstream phase vs. upstream phase with FF simulated at optimal gain.

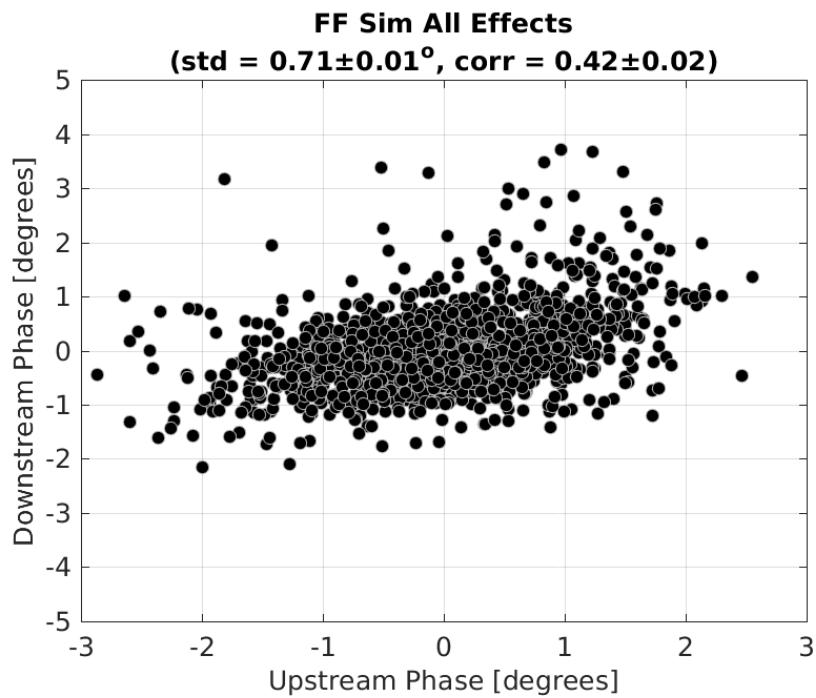


Figure 6.21: Downstream phase vs. upstream phase with FF simulated with actual gain used.

6.3 Possible Areas for Future Improvement

Chapter 7

Conclusions

This is the introductory text.

7.1 Summary

7.2 Future Work

Bibliography

- [1] Dummy One & Dummy Two. *Phys. Journal*, **1**, 1 (2002) 1–5. hep-ph/0000000. <http://some.web.address>